

# **CS3002: Artificial Intelligence**

**Images and Vision**

**Yongmin Li**

# Acknowledgement

- **Sarbjeet Singh**
- **Antonio Torralba**
- **Srinivasa Narasimhan**

# Overview – Early Vision

- **Image Fundamentals**
- **Edge Detection**
- **Corner Detection**
- **Line Detection by Hough Transform**

# 1-bit Images

- Each pixel is stored as a single bit (0 or 1), so also referred to as binary image.
- Such an image is also called a 1-bit monochrome image since it contains no colour.



# 8-bit Gray-Level Images

- Each pixel has a gray-value between 0 and 255. Each pixel is represented by a single byte; e.g., a dark pixel might have a value of 10, and a bright one might be 230.



# 24-bit Colour Images

- In a colour 24-bit image, each pixel is represented by three bytes, usually representing RGB.



# 24-bit Colour Images



(a)



(b)



(c)



(d)

**(a): Example of 24-bit colour image “forestfire.bmp”. (b, c, d): R, G, and B colour channels for this image**

# Other Colour Images

- **8-bit colour images:** using 8-bit index to 24-bit real colours.
- **32-bit colour images:** many 24-bit colour images are actually stored as 32-bit images, with the extra byte of data for each pixel used to store an *alpha* value representing special effect information (e.g., transparency).






# Edge Detection



# Edges in Images

- Edge - Area of significant change in the image intensity / contrast
- Edge Detection – Locating areas with strong intensity contrasts
- Use of Edge Detection – Extracting information about the image. E.g. location of objects present in the image, their shape, size, image sharpening and enhancement



# Types of Edges

- Variation of Intensity / Gray Level
  - Step Edge
  - Line Edge
  - Roof Edge




# Steps in Edge Detection

- Filtering – Filter image to improve performance of the Edge Detector wrt noise
- Enhancement – Emphasize pixels having significant change in local intensity
- Detection – Identify edges - thresholding
- Localization – Locate the edge accurately, estimate edge orientation


# Methods of Edge Detection

- First Order Derivative / Gradient Methods
  - Roberts Operator
  - Sobel Operator
  - Prewitt Operator
- Second Order Derivative
  - Laplacian
  - Laplacian of Gaussian
  - Difference of Gaussian
- Optimal Edge Detection
  - Canny Edge Detection

# First Derivative



- At the point of greatest slope, the first derivative has maximum value
  - E.g. For a Continuous 1-dimensional function  $f(t)$



# Gradient

- For a continuous two dimensional function Gradient is defined as

$$G[f(x, y)] = \begin{bmatrix} G_x \\ G_y \end{bmatrix} = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$$

$$|G| = \sqrt{G_x^2 + G_y^2} \approx |G_x| + |G_y|$$


$$\theta = \tan^{-1}\left(\frac{G_y}{G_x}\right)$$

# Gradient

- Approximation of Gradient for a discrete two dimensional function
- **Convolution** Mask

$$G_x \cong f[i, j+1] - f[i, j]$$

$$G_y \cong f[i, j] - f[i+1, j]$$



- Differences are computed at the interpolated points  $[i, j+1/2]$  and  $[i+1/2, j]$



# Roberts Operator

- **Provides an approximation to the gradient**

$$G[f(i, j)] = |G_x| + |G_y| = |f(i, j) - f(i+1, j+1)| + |f(i+1, j) - f(i, j+1)|$$

- **Convolution Mask**

**G<sub>x</sub>**=


|   |    |
|---|----|
| 1 | 0  |
| 0 | -1 |

**G<sub>y</sub>** =

|   |    |
|---|----|
| 0 | -1 |
| 1 | 0  |

# Roberts Operator - Example

- The output image has been scaled by a factor of 5
- Spurious dots indicate that the operator is susceptible to noise



# Sobel Operator

- The 3x3 convolution mask smoothes the image by some amount , hence it is less susceptible to noise. But it produces thicker edges. So edge localization is poor
- Convolution Mask


$$- \mathbf{G_x} = \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$

$$\mathbf{G_y} = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$


- The differences are calculated at the centre pixel of the mask.

# Sobel Operator - Example

- Compare the output of the Sobel Operator with that of the Roberts Operator:
  - The spurious edges are still present but they are relatively less intense compared to genuine lines
  - Roberts operator has missed a few edges
  - Sobel operator detects thicker edges



Sobel



Roberts

# Prewitt Operator

- It is similar to the Sobel operator but uses slightly different masks
- Convolution Mask

**P<sub>x</sub>** =

|    |   |   |
|----|---|---|
| -1 | 0 | 1 |
| -1 | 0 | 1 |
| -1 | 0 | 1 |

**P<sub>y</sub>** =

|    |    |    |
|----|----|----|
| 1  | 1  | 1  |
| 0  | 0  | 0  |
| -1 | -1 | -1 |

# First Order Derivative Methods

- Noise – simple edge detectors are affected by noise – filters can be used to reduce noise
- Edge Thickness – Edge is several pixels wide for Sobel operator– edge is not localized properly
- Roberts operator is very sensitive to noise
- Sobel operator goes for averaging and emphasizes on the pixel closer to the centre of the mask. It is less affected by noise and is one of the most popular Edge Detectors.

# Canny Edge Detector

- Step 1
  - Noise is filtered out – usually a Gaussian filter is used
  - Width is chosen carefully
- Step 2
  - Edge strength is found out by taking the gradient of the image
  - A Roberts mask or a Sobel mask can be used


$$|G| = \sqrt{Gx^2 + Gy^2} \approx |Gx| + |Gy|$$

# Canny Edge Detector

- Step 3
  - Find the edge direction

$$\theta = \tan^{-1}\left(\frac{G_y}{G_x}\right)$$

- Step 4
  - Resolve edge direction





# Canny Edge Detector


- Step 5
  - Non-maxima suppression – trace along the edge direction and suppress any pixel value not considered to be an edge. Gives a thin line for edge
- Step 6
  - Use double / hysteresis thresholding to eliminate streaking

# Canny Edge Detector


- Compare the results of Canny and Sobel



Input image



Canny




Sobel

# Corner Detection




# The Basic Idea


- We should easily localize the point by looking through a small window
- Shifting a window in *any direction* should give *a large change* in intensity




# Harris Detector: Basic Idea



"flat" region:  
no change as  
shift window in  
all directions



"edge":  
no change as shift  
window along the  
edge direction



"corner":  
significant change  
as shift window in  
all directions

# Harris Detector: Mathematics

Window-averaged change of intensity induced by shifting the image data by  $[u, v]$ :


$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Window  
function

Shifted  
intensity


Intensity

Window function  $w(x, y) =$



1 in window, 0 outside

or



Gaussian

# Taylor Series Approximation

$$\begin{aligned} E(u, v) &\approx \sum_{x, y} w(x, y) [I(x, y) + uI_x + vI_y - I(x, y)]^2 \\ &= \sum_{x, y} w(x, y) [uI_x + vI_y]^2 \\ &= \sum_{x, y} w(x, y) \begin{pmatrix} u & v \end{pmatrix} \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix} \begin{pmatrix} u \\ v \end{pmatrix} \end{aligned}$$

# Harris Detector: Mathematics

Expanding  $I(x,y)$  in a Taylor series expansion, we have, for small shifts  $[u, v]$ , a *bilinear* approximation:

$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where  $M$  is a  $2 \times 2$  matrix computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

$M$  is also called “structure tensor”




# Harris Detector: Mathematics

Intensity change in shifting window: eigenvalue analysis



$$E(u, v) \cong \begin{bmatrix} u & v \end{bmatrix} M \begin{bmatrix} u \\ v \end{bmatrix} \quad \lambda_1, \lambda_2 - \text{eigenvalues of } M$$

Ellipse  $E(u, v) = \text{const}$

Iso-intensity contour of  $E(u, v)$





# Selecting Good Features





$\lambda_1$  and  $\lambda_2$  are large

# Selecting Good Features



large  $\lambda_1$ , small  $\lambda_2$


# Selecting Good Features



small  $\lambda_1$ , small  $\lambda_2$

# Harris Detector: Mathematics

Classification of image points using eigenvalues of  $M$ :



# Harris Detector: Mathematics

Measure of corner response:

$$R = \det M - k (\text{trace } M)^2$$


This expression does not require computing the eigenvalues.

$$\begin{aligned}\det M &= \lambda_1 \lambda_2 \\ \text{trace } M &= \lambda_1 + \lambda_2\end{aligned}$$

( $k$  – empirical constant,  $k = 0.04$ - $0.06$ )

# Harris Detector: Mathematics

- $R$  depends only on eigenvalues of  $M$
- $R$  is large for a **corner**
- $R$  is negative with large magnitude for an **edge**
- $|R|$  is small for a **flat** region



# Harris Detector

- The Algorithm:
  - Find points with large corner response function  $R$  ( $R > \text{threshold}$ )
  - Take the points of local maxima of  $R$

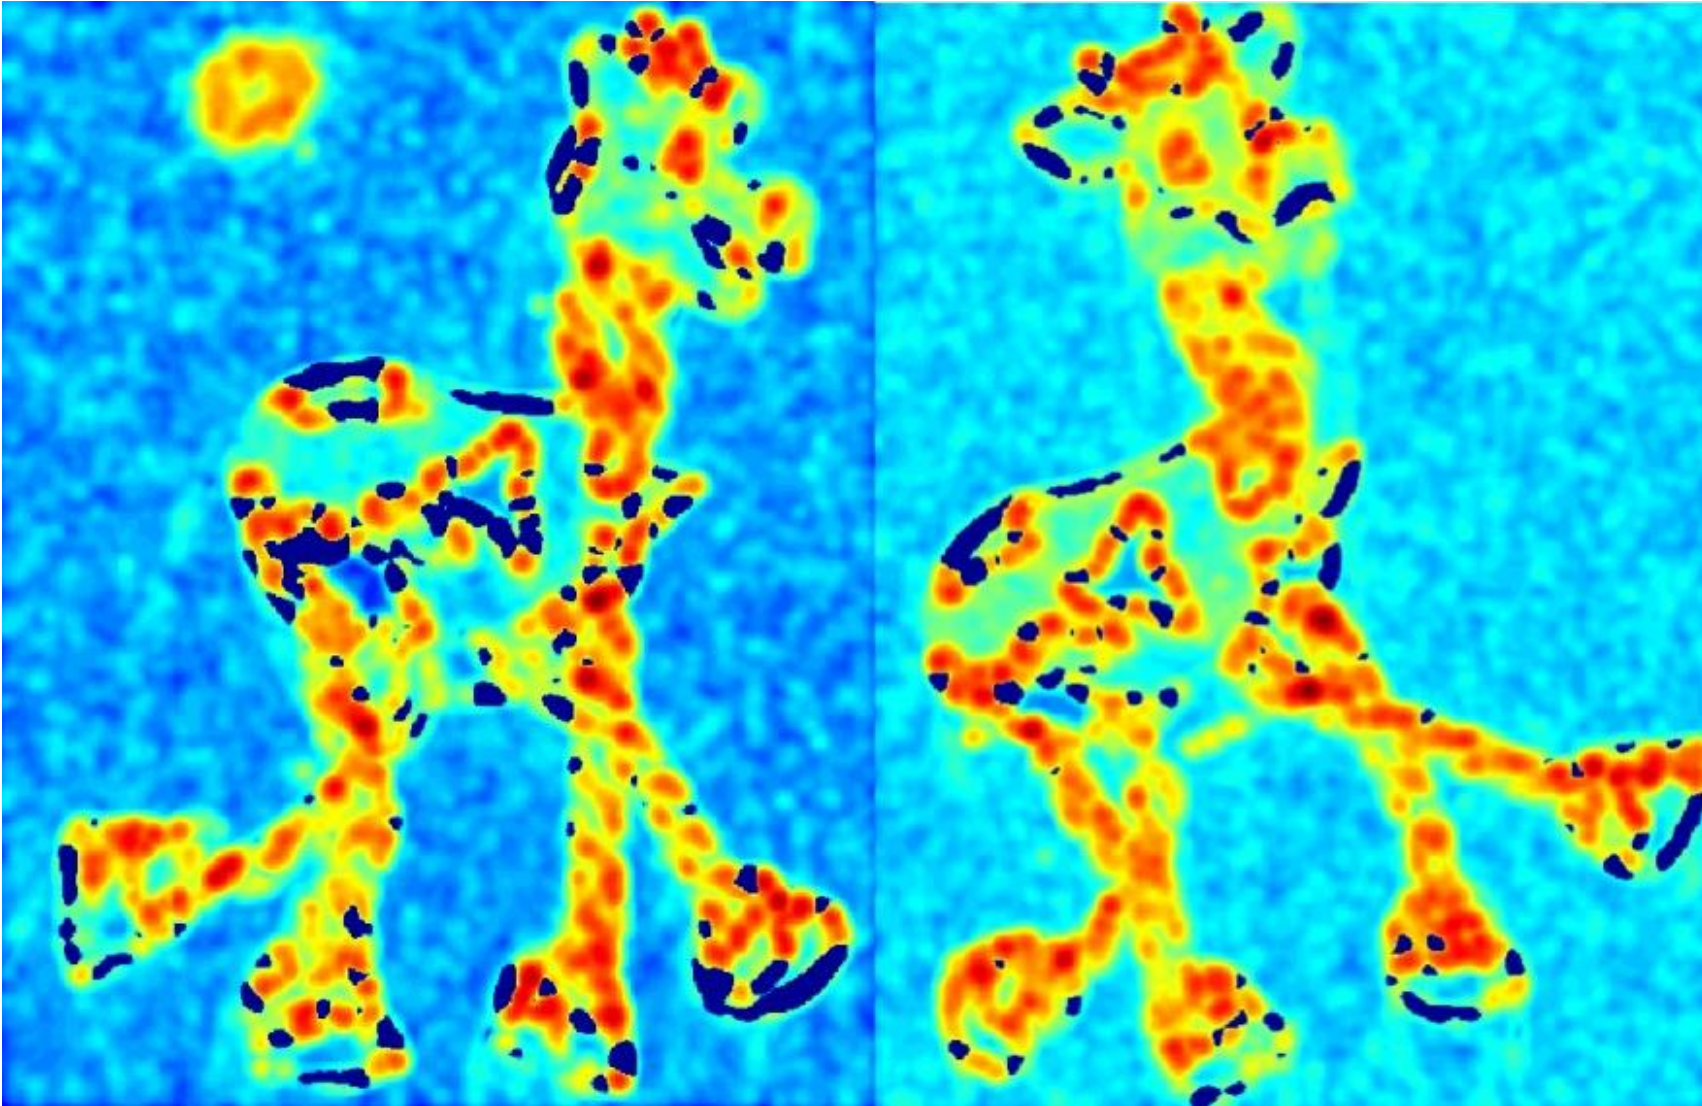


# Harris Detector: Workflow




# Harris Detector: Workflow

Compute corner response  $R$



# Harris Detector: Workflow

Find points with large corner response:  $R > \text{threshold}$



# Harris Detector: Workflow

Take only the points of local maxima of  $R$





# Harris Detector: Workflow



# Harris Detector: Summary (1)

- Estimate the structure tensor  $M$  – an indicator of intensity changes:

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$


- Calculate the *corner response*  $R$  from the eigenvalues of  $M$ :

$$R = \lambda_1 \lambda_2 - k (\lambda_1 + \lambda_2)^2$$

- A good (corner) point should have a *large intensity change* in *all directions*, i.e.  $R$  should be large positive

# Harris Detector: Summary (2)

- $R$  depends only on eigenvalues of  $M$
- $R$  is large for a **corner**
- $R$  is negative with large magnitude for an **edge**
- $|R|$  is small for a **flat** region




# Line Detection







# Preprocessing Edge Images




Image



Edge Detection  
and Thresholding





Noisy edge image  
Incomplete boundaries



Shrink and Expand



Thinning




# Edge Tracking Methods

## Adjusting a priori Boundaries:

Given: Approximate Location of Boundary

Task: Find Accurate Location of Boundary



- Search for STRONG EDGES along normals to approximate boundary.
- Fit curve (eg., polynomials) to strong edges.


# Edge Tracking Methods

## Divide and Conquer:

Given: Boundary lies between points A and B

Task: Find Boundary

- Connect A and B with Line
- Find strongest edge along line bisect
- Use edge point as break point
- Repeat



# Fitting Lines (Least Squares)

Given: Many  $(x_i, y_i)$  pairs  
Find: Parameters  $(m, c)$

Minimize: Average square distance:


$$E = \sum_i \frac{(y_i - mx_i - c)^2}{N}$$

Using:

$$\frac{\partial E}{\partial m} = 0 \quad \& \quad \frac{\partial E}{\partial c} = 0$$

Note:

$$\bar{y} = \frac{\sum_i y_i}{N} \quad \bar{x} = \frac{\sum_i x_i}{N}$$



$$c = \bar{y} - m \bar{x}$$
$$m = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sum_i (x_i - \bar{x})^2}$$

# Curve Fitting

Find Polynomial:

$$y = f(x) = ax^3 + bx^2 + cx + d$$


that best fits the given points  $(x_i, y_i)$

Minimize:


$$\frac{1}{N} \sum_i [y_i - (ax_i^3 + bx_i^2 + cx_i + d)]^2$$

Using:  $\frac{\partial E}{\partial a} = 0$  ,  $\frac{\partial E}{\partial b} = 0$  ,  $\frac{\partial E}{\partial c} = 0$  ,  $\frac{\partial E}{\partial d} = 0$


Note:  $f(x)$  is LINEAR in the parameters (a, b, c, d)



# Problem with Parameterization



# Line Grouping Problem



# This is difficult because of:

- Extraneous data: clutter or multiple models
  - We do not know what is part of the model?
  - Can we pull out models with a few parts from much larger amounts of background clutter?
- Missing data: only some parts of model are present
- Noise
- **Cost:**
  - It is not feasible to check all combinations of features by fitting a model to each possible subset



# Hough Transform

- Elegant method for direct object recognition
- Edges need not be connected
- Complete object need not be visible
- Key Idea: Edges VOTE for the possible model

# Image and Parameter Spaces



Equation of Line:  $y = mx + c$

Find:  $(m, c)$

Consider point:  $(x_i, y_i)$

$$y_i = mx_i + c \quad \text{or} \quad c = -x_i m + y_i$$

Parameter space also called Hough Space



# Line Detection by Hough Transform

### Algorithm:


- Quantize Parameter Space  $(m, c)$
- Create Accumulator Array  $A(m, c)$
- Set  $A(m, c) = 0 \quad \forall m, c$
- For each image edge  $(x_i, y_i)$  increment:

$$A(m, c) = A(m, c) + 1$$

If  $(m, c)$  lies on the line:

$$c = -x_i m + y_i$$

- Find local maxima in  $A(m, c)$



## Parameter Space

$$A(m, c)$$
[illegible]

# Better Parameterization

m: slope  
c: intercept

NOTE:  $-\infty \leq m \leq \infty$   
Large Accumulator  
More memory and computations  
(Finite Accumulator Array Size)

Improvement:


Line equation:  $\rho = x \cos \theta + y \sin \theta$

Here  $0 \leq \theta \leq 2\pi$


$0 \leq \rho \leq \rho_{\max}$

Given points  $(x_i, y_i)$  find  $(\rho, \theta)$


Hough Space Sinusoid



# Example 1




**Image space**




**Votes**

Horizontal axis is  $\theta$ ,  
vertical is  $\rho$ .

# Example 2





**Image  
space**



**votes**

# Example 3



# Application Issues

- Difficulties
  - how big should the cells be? (too big, and we merge quite different lines; too small, and noise causes lines to be missed)
- How many lines?
  - Count the peaks in the Hough array
  - Treat adjacent peaks as a single peak
- Which points belong to each line?
  - Search for points close to the line
  - Solve again for line and iterate



# Reading

- Not required, but if you really, really want...
- Try this:  
Hlavac, Sonka & Boyle (2014). Image Processing, Analysis, and Machine Vision. 4<sup>th</sup> edition. Nelson Engineering.  
**Chapter 5 (5.1, 5.2 and 5.3)**