

A gentle introduction to Deep Learning

Part 1: Computer Vision

Alina Miron

alina.miron@brunel.ac.uk

In today's lecture we'll discuss:

- What is deep learning
- Recap Neural Networks
- Image classification task
- Convolutional Neural Networks
 - Convolution
 - Pooling
 - Dropout
 - RELU activation function
- Basic deep learning software libraries
- Deep learning advantages & disadvantages

A short history of deep learning

The term of deep
learning is introduced

1986

A short history of deep learning

The term of deep learning is introduced

1986

LEARNING WHILE SEARCHING IN CONSTRAINT-SATISFACTION-PROBLEMS*

Rina Dechter

Artificial Intelligence Center
Hughes Aircraft Company, Calabasas, CA 91302
and
Cognitive Systems Laboratory, Computer Science Department
University of California, Los Angeles, CA 90024

ABSTRACT

The popular use of backtracking as a control strategy for theorem proving in PROLOG and in Truth-Maintenance-Systems (TMS) led to increased interest in various schemes for enhancing the efficiency of backtrack search. Researchers have referred to these enhancement schemes by the names

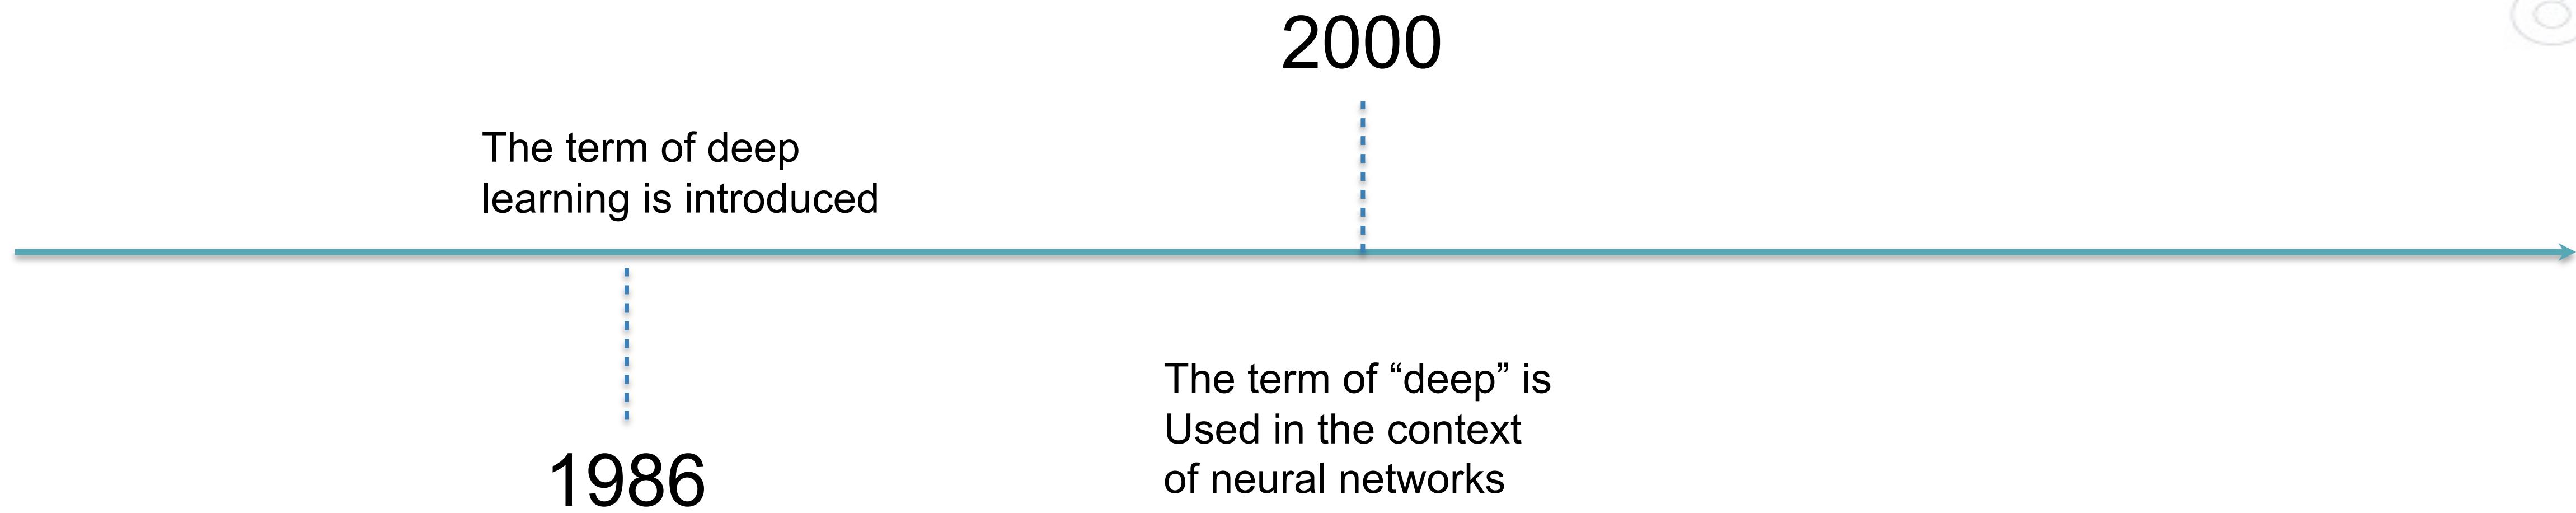
2. **Look-back schemes:** affecting the decision of where and how to go in case of a dead-end situation. Look-back schemes are centered around two fundamental ideas:
 - a. **Go-back to source of failure:** an attempt is made to detect and change previous decisions

Discovering all minimal conflict-sets amounts to acquiring all the possible information out of a dead-end. Yet, such **deep learning** may require considerable amount of work. While the number of minimal conflict-sets is less than 2^r , where r is the cardinality of the Conf-set, we can envision a worst case where all subsets of $\text{Conf}(S, X_i)$ having $\frac{r}{2}$ elements are in conflict with X_i . The number of minimal conflict-sets should then satisfy

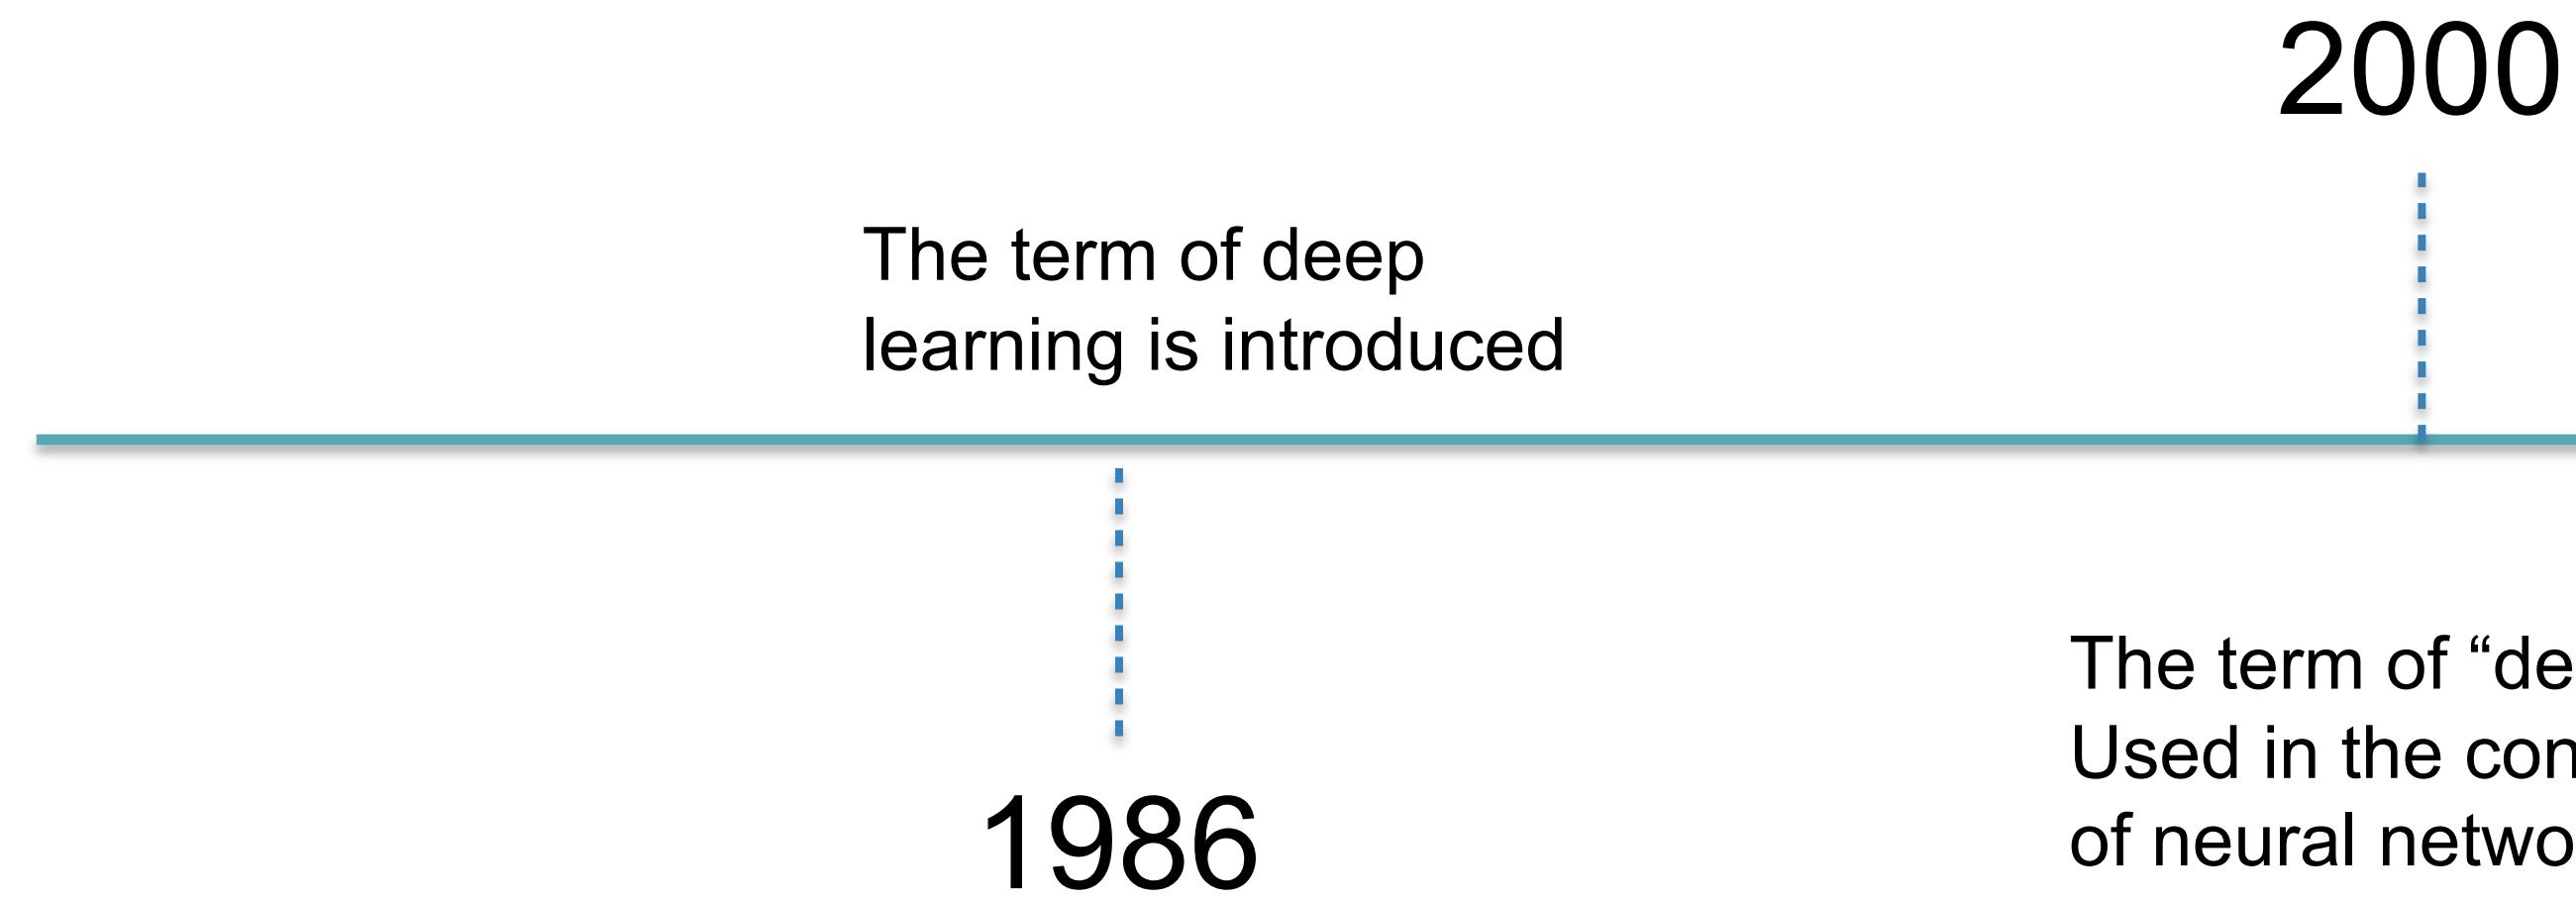
$$\#\text{min-conflict-sets} = \binom{r}{\frac{r}{2}} \approx 2^r, \quad (2)$$

which is still exponential in the size of the Conf-set. If the size of this Conf-set is small it may still be reasonable to recognize all minimal conflict-sets.

A short history of deep learning



A short history of deep learning



Co-Evolving Recurrent Neurons Learn Deep Memory POMDPs

Faustino J. Gomez¹
tino@idsia.ch

Jürgen Schmidhuber^{1,2}
juergen@idsia.ch

¹ IDSIA, Galleria 2, 6928 Lugano, Switzerland

² TU Munich, Boltzmannstr. 3, 85748 Garching, München, Germany

ABSTRACT

Recurrent neural networks are theoretically capable of learning complex temporal sequences, but training them through gradient-descent is too slow and unstable for practical use in reinforcement learning environments. Neuroevolution, the evolution of artificial neural networks using genetic algorithms, can potentially solve real-world reinforcement learning tasks that require **deep** use of memory, i.e. memory spanning hundreds or thousands of inputs, by searching the space of recurrent neural networks directly. In this paper, we introduce a new neuroevolution algorithm called Hierarchical Enforced SubPopulations that simultaneously evolves networks at two levels of granularity: full networks and network components or *neurons*. We demonstrate the method

previous network inputs. This ability to act based upon the recollection of past events is essential for truly complex behavior. Unfortunately, training RNNs with standard gradient-descent methods such as Real-Time Recurrent Learning (RTRL; [11, 14]) or Back Propagation Through Time (BPTT; [12]) is not possible when actions depend on inputs from more than as few as 10 time-steps in the past. The error signal used to adjust synaptic weights either vanishes or explodes as it is propagated back through time, making learning prohibitively slow or preventing it altogether [5, 3]. For this reason, using conventional RNN architectures to approximate a value-function or policy for reinforcement learning (i.e. Q-learning, SARSA) is not practical when the environment exhibits long temporal dependencies.

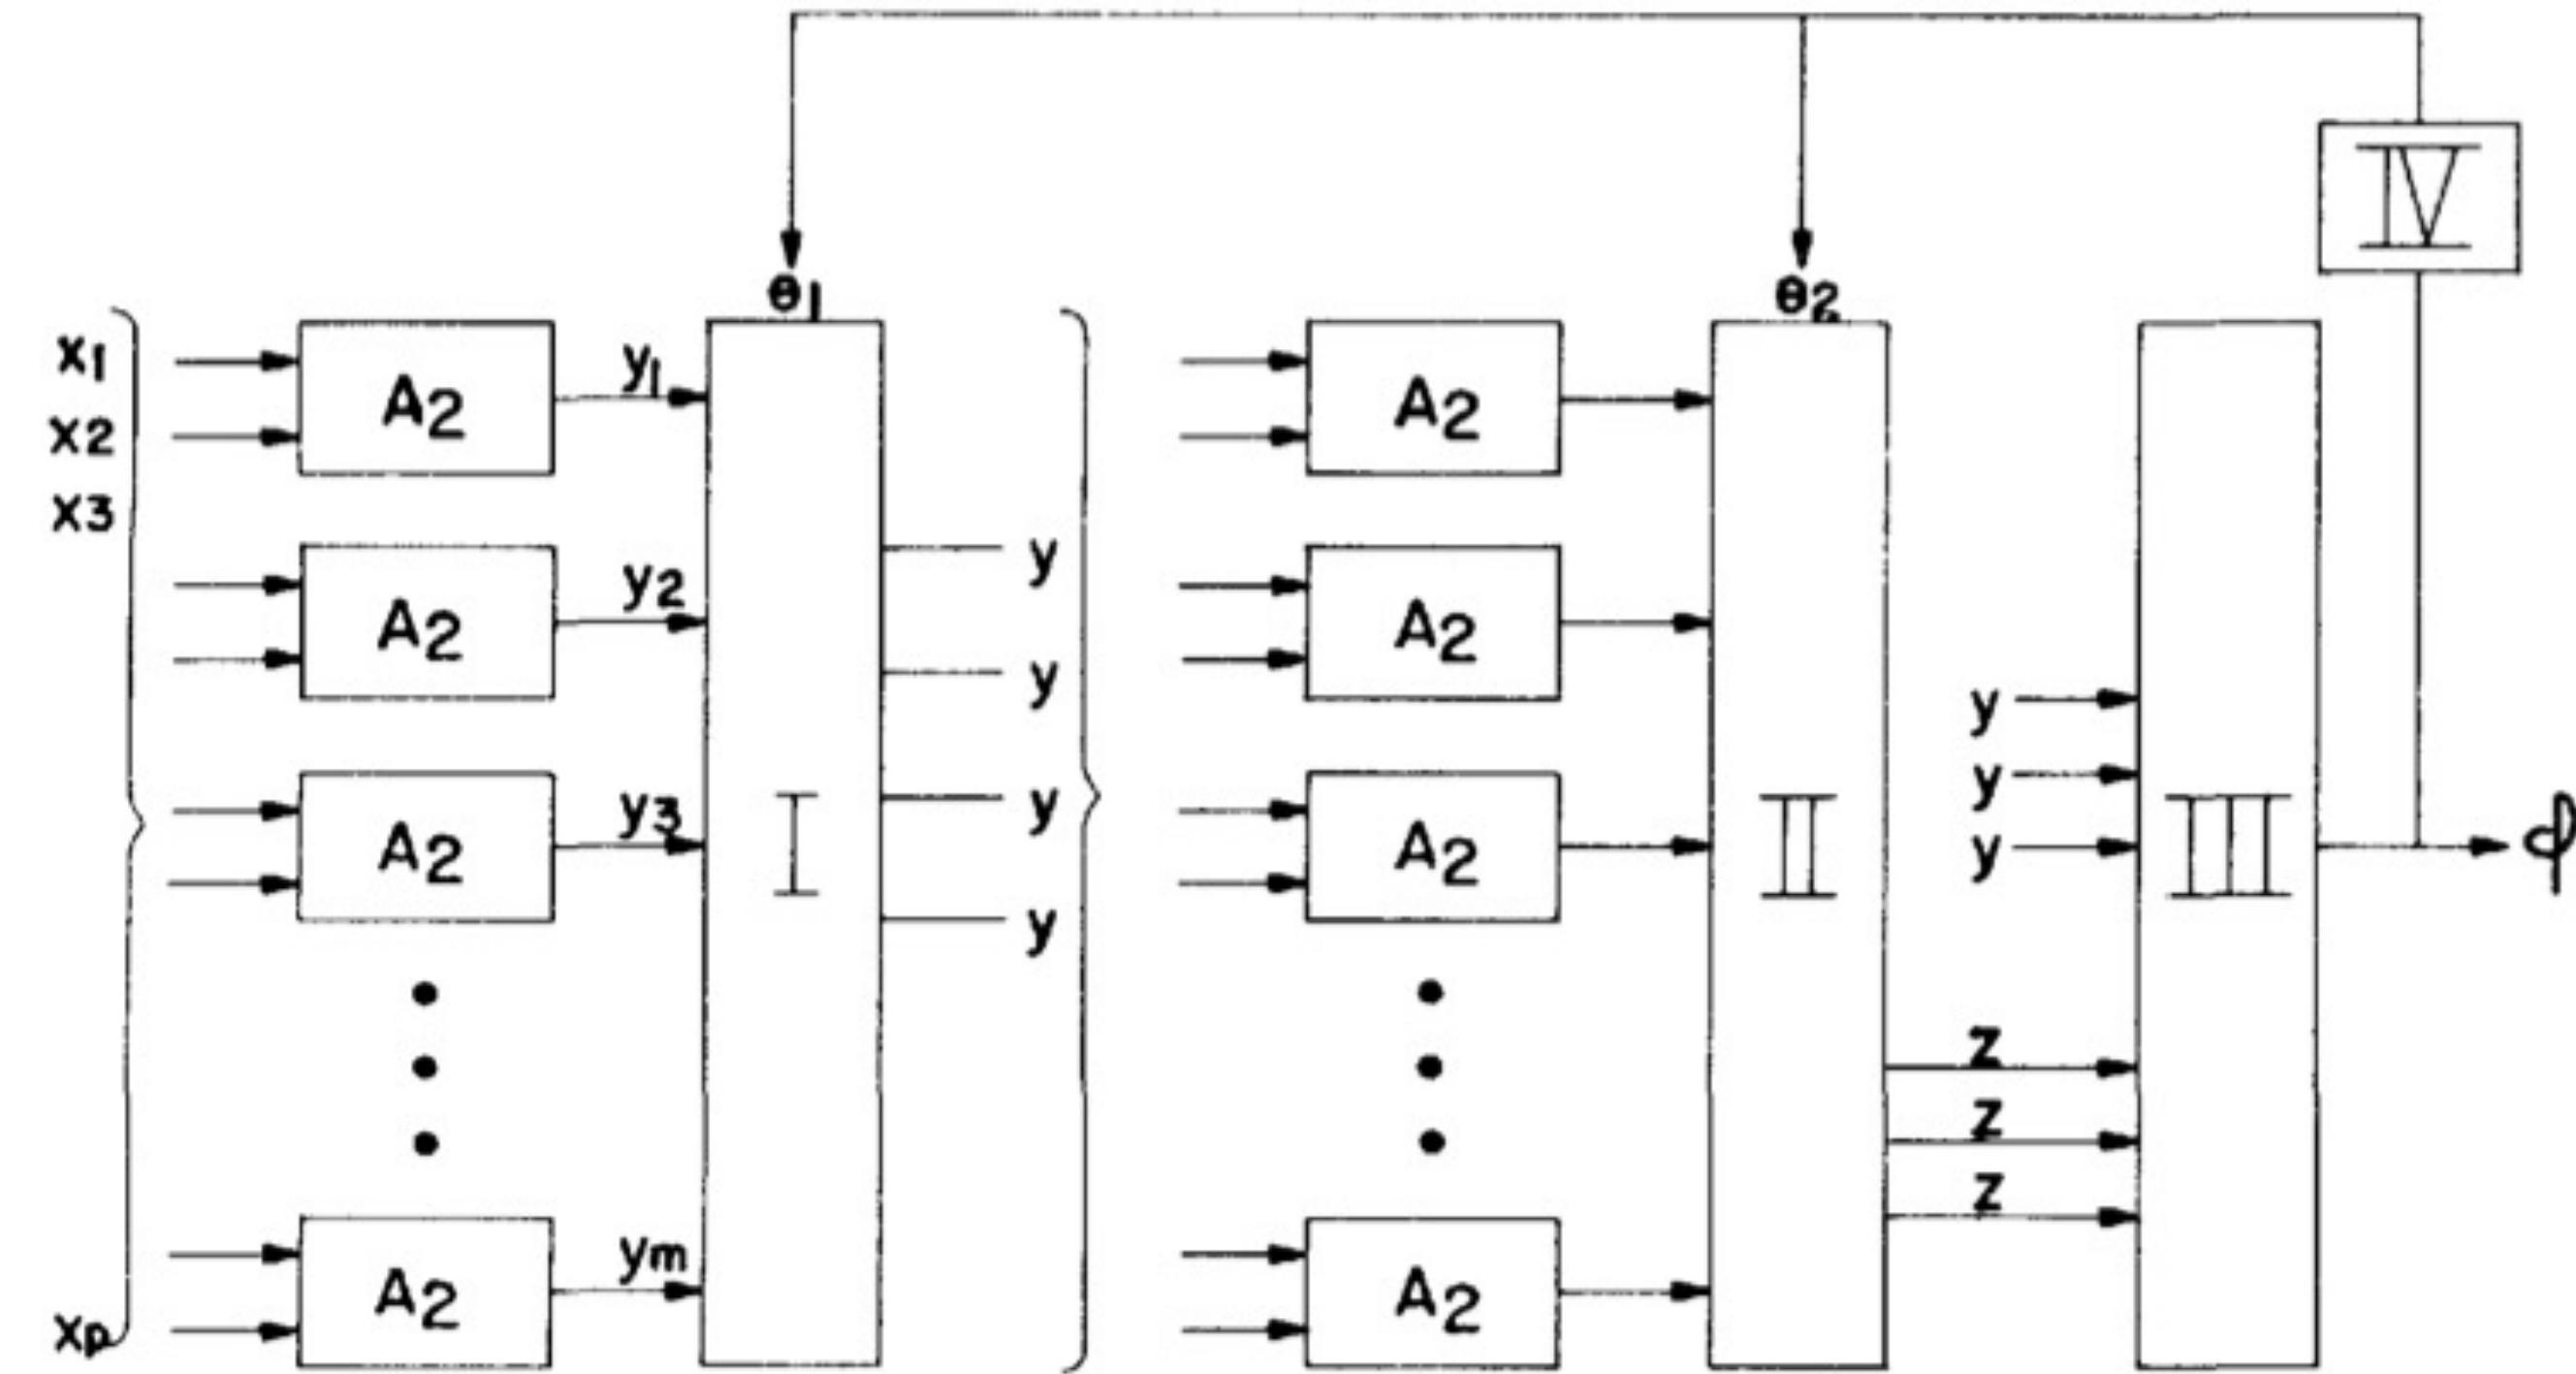
A short history of deep learning

Ivakhnenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965



The
lea



Algorithm of group method of data handling (GMDH) using second-degree polynomials. I—first threshold self-selection; II—second threshold self-selection; III—selection from all solutions; IV—threshold optimization.

<https://kpi.ua/en/ivakhnenko>

<https://www.gmdh.net/articles/history/polynomial.pdf>

A short history of deep learning

Ivakhnenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

The term of deep
learning is introduced

1986

2000

The term of “deep” is
Used in the context
of neural networks

<https://kpi.ua/en/ivakhnenko>

<https://www.gmdh.net/articles/history/polynomial.pdf>

A short history of deep learning

Ivaknenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

1980

The term of deep
learning is introduced

1986

2000

The term of “deep” is
Used in the context
of neural networks

Fukushima, Kunihiko, and Sei Miyake. "Neocognitron: A self-organizing neural network model for a mechanism of visual pattern recognition." *Competition and cooperation in neural nets*. Springer, Berlin, Heidelberg, 1982. 267-285.

A short history of deep learning

Ivakhnenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

1980

The term of deep
learning is introduced

1986

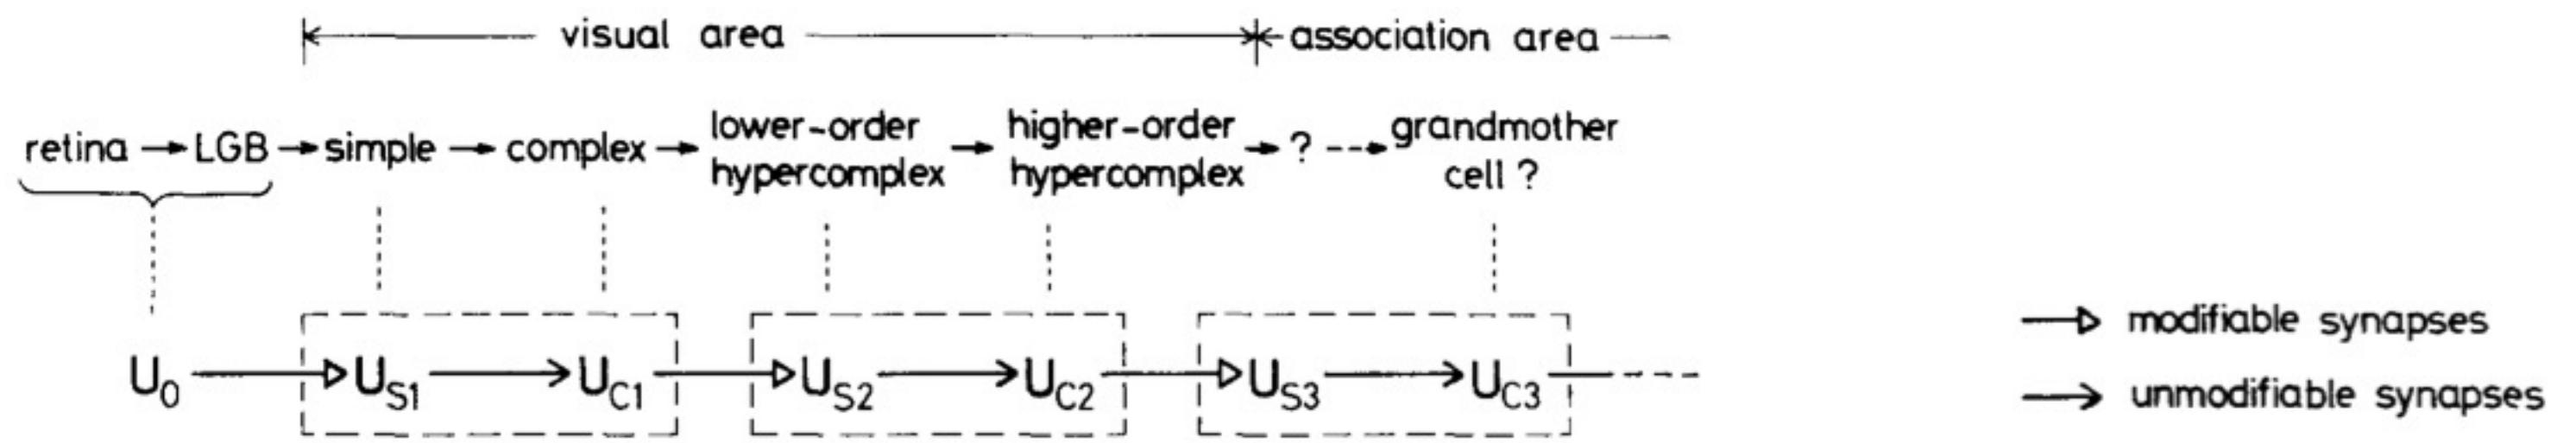


Fig. 1. Correspondence between the hierarchy model by Hubel and Wiesel, and the neural network of the neocognitron

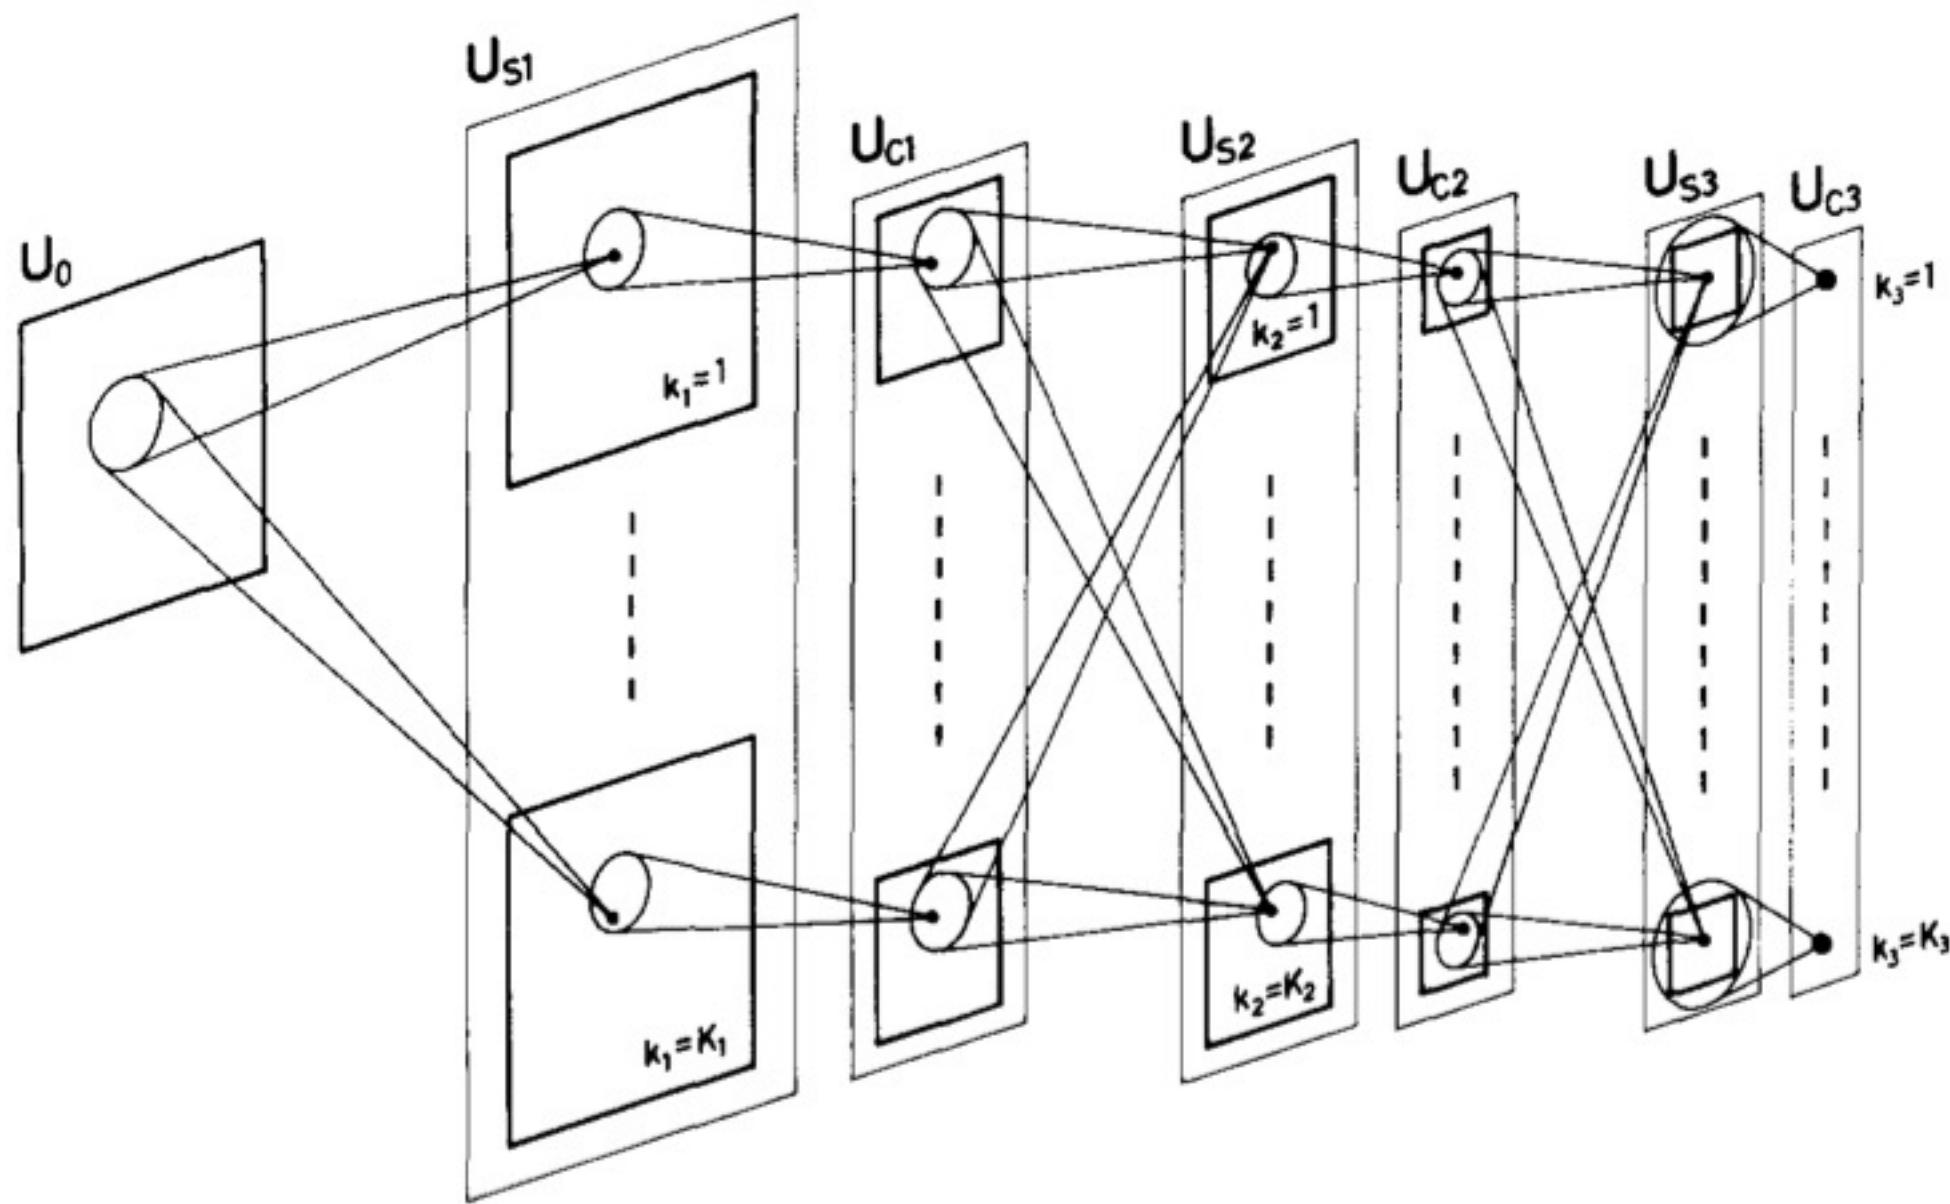


Fig. 2. Schematic diagram illustrating the interconnections between layers in the neocognitron

Fukushima, Kunihiko, and Sei Miyake. "Neocognitron: A self-organizing neural network model for a mechanism of visual pattern recognition." *Competition and cooperation in neural nets*. Springer, Berlin, Heidelberg, 1982. 267-285.

A short history of deep learning

Ivaknenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

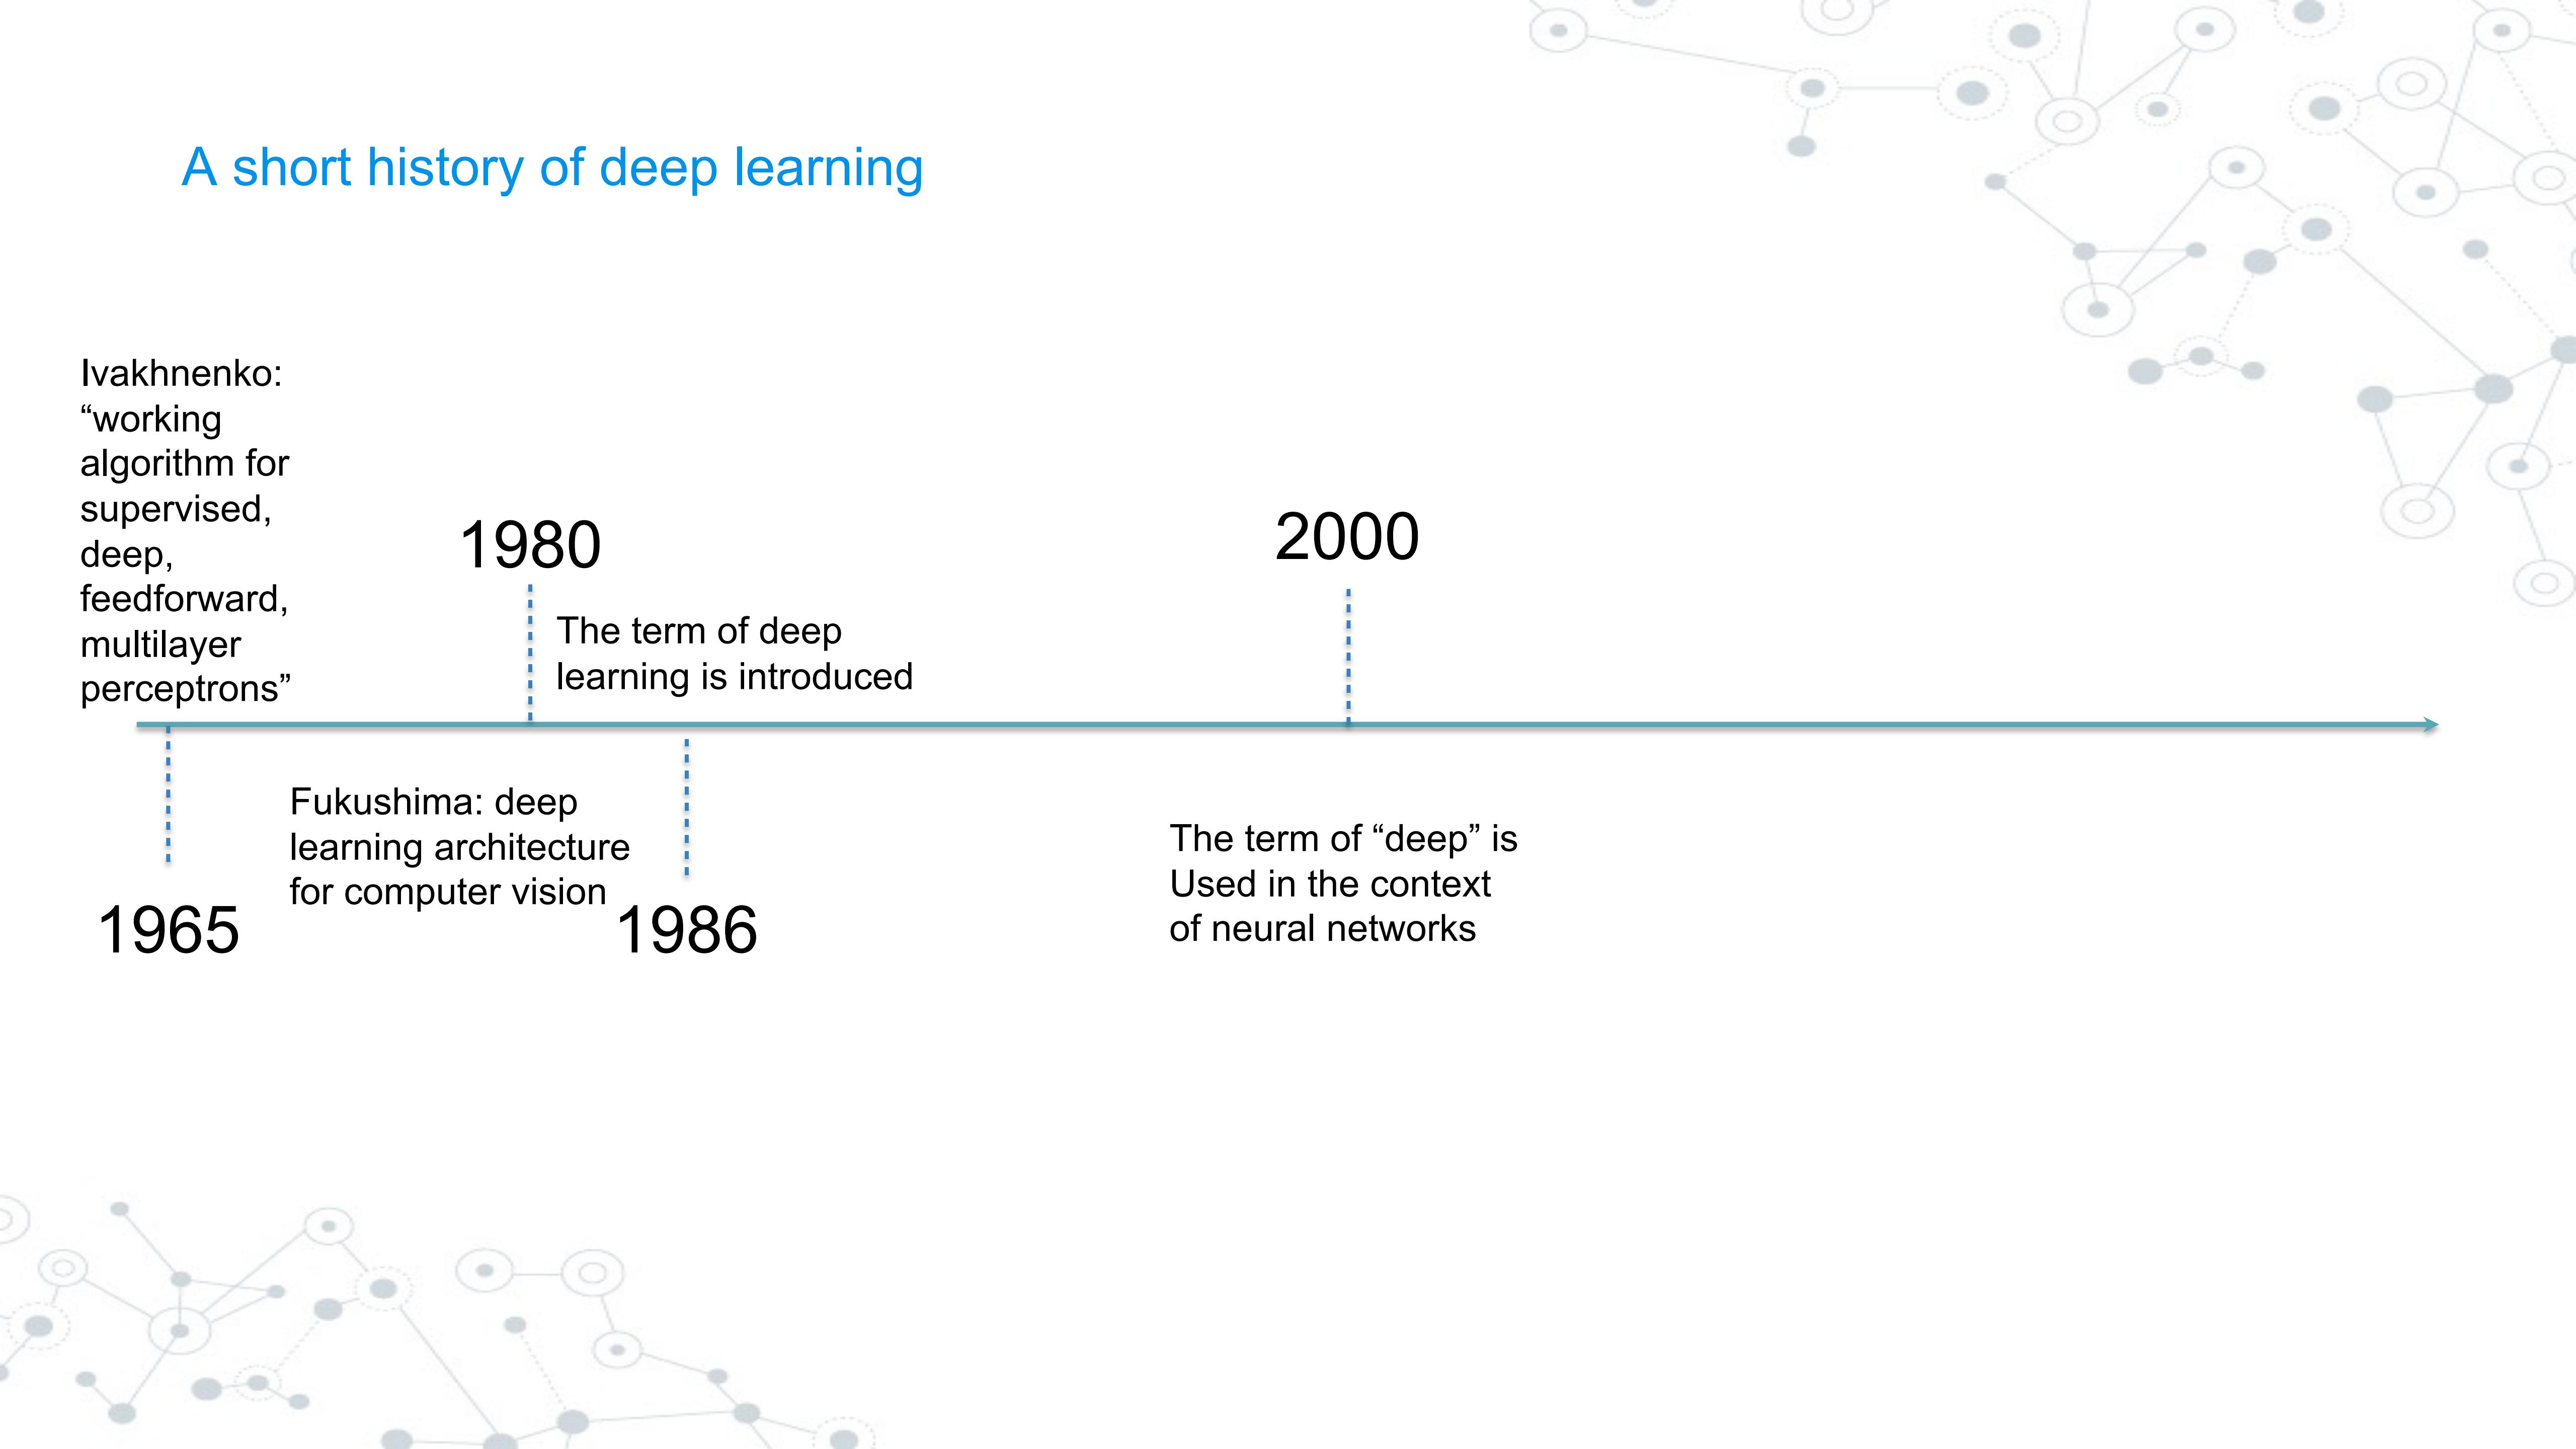
1980

The term of deep
learning is introduced

1986

2000

The term of “deep” is
Used in the context
of neural networks



A short history of deep learning

Ivaknenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

1980

The term of deep
learning is introduced

1986

LeCunn:
Applied back-
propagation
algorithm, to a
deep neural
network with the
purpose of
recognising
handwritten ZIP
codes

1989

The term of “deep” is
Used in the context
of neural networks



A short history of deep learning

Ivaknenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

1980

The term of deep
learning is introduced

1986

LeCunn:
Applied back-
propagation
algorithm, to a
deep neural
network with the
purpose of
recognising
handwritten ZIP
codes

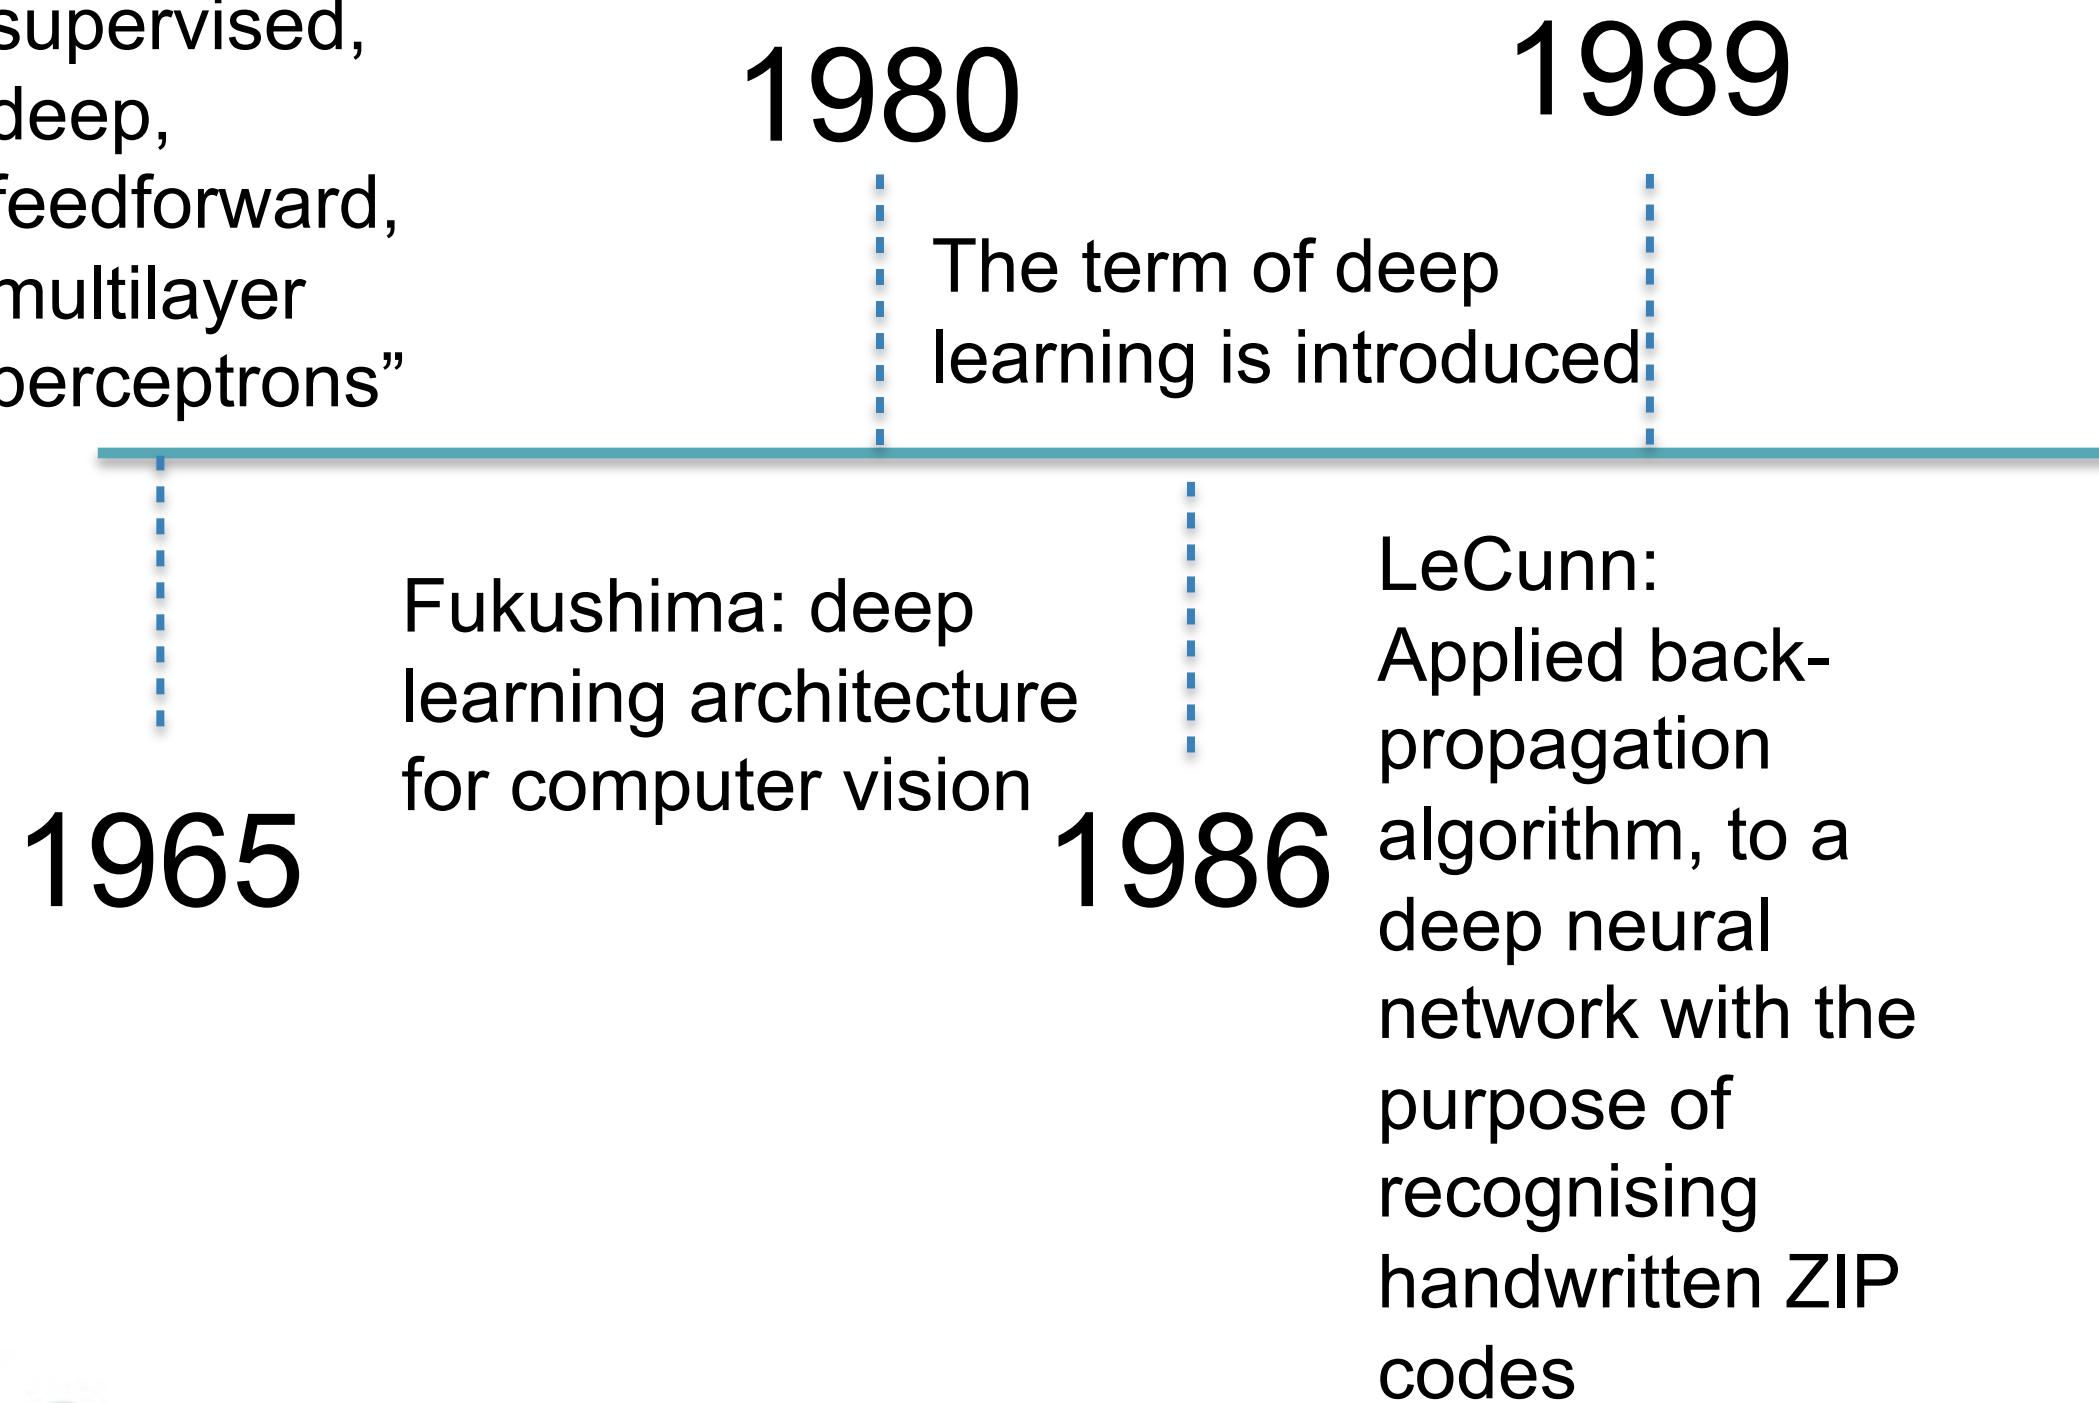
1989

The term of “deep” is
Used in the context
of neural networks

LeCun, Y., Boser, B., Denker, J.S., Henderson, D., Howard, R.E., Hubbard, W.
and Jackel, L.D., 1989. Backpropagation applied to handwritten zip code
recognition. *Neural computation*, 1(4), pp.541-551.

A short history of deep learning

Ivakhnenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”



80322 - 4129 80206

40004 14310

37878 05153

~~35502~~ 75216

35460 · 44209

LeCun, Y., Boser, B., Denker, J.S., Henderson, D., Howard, R.E., Hubbard, W. and Jackel, L.D., 1989. Backpropagation applied to handwritten zip code recognition. *Neural computation*, 1(4), pp.541-551.

A short history of deep learning

Ivakhnenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

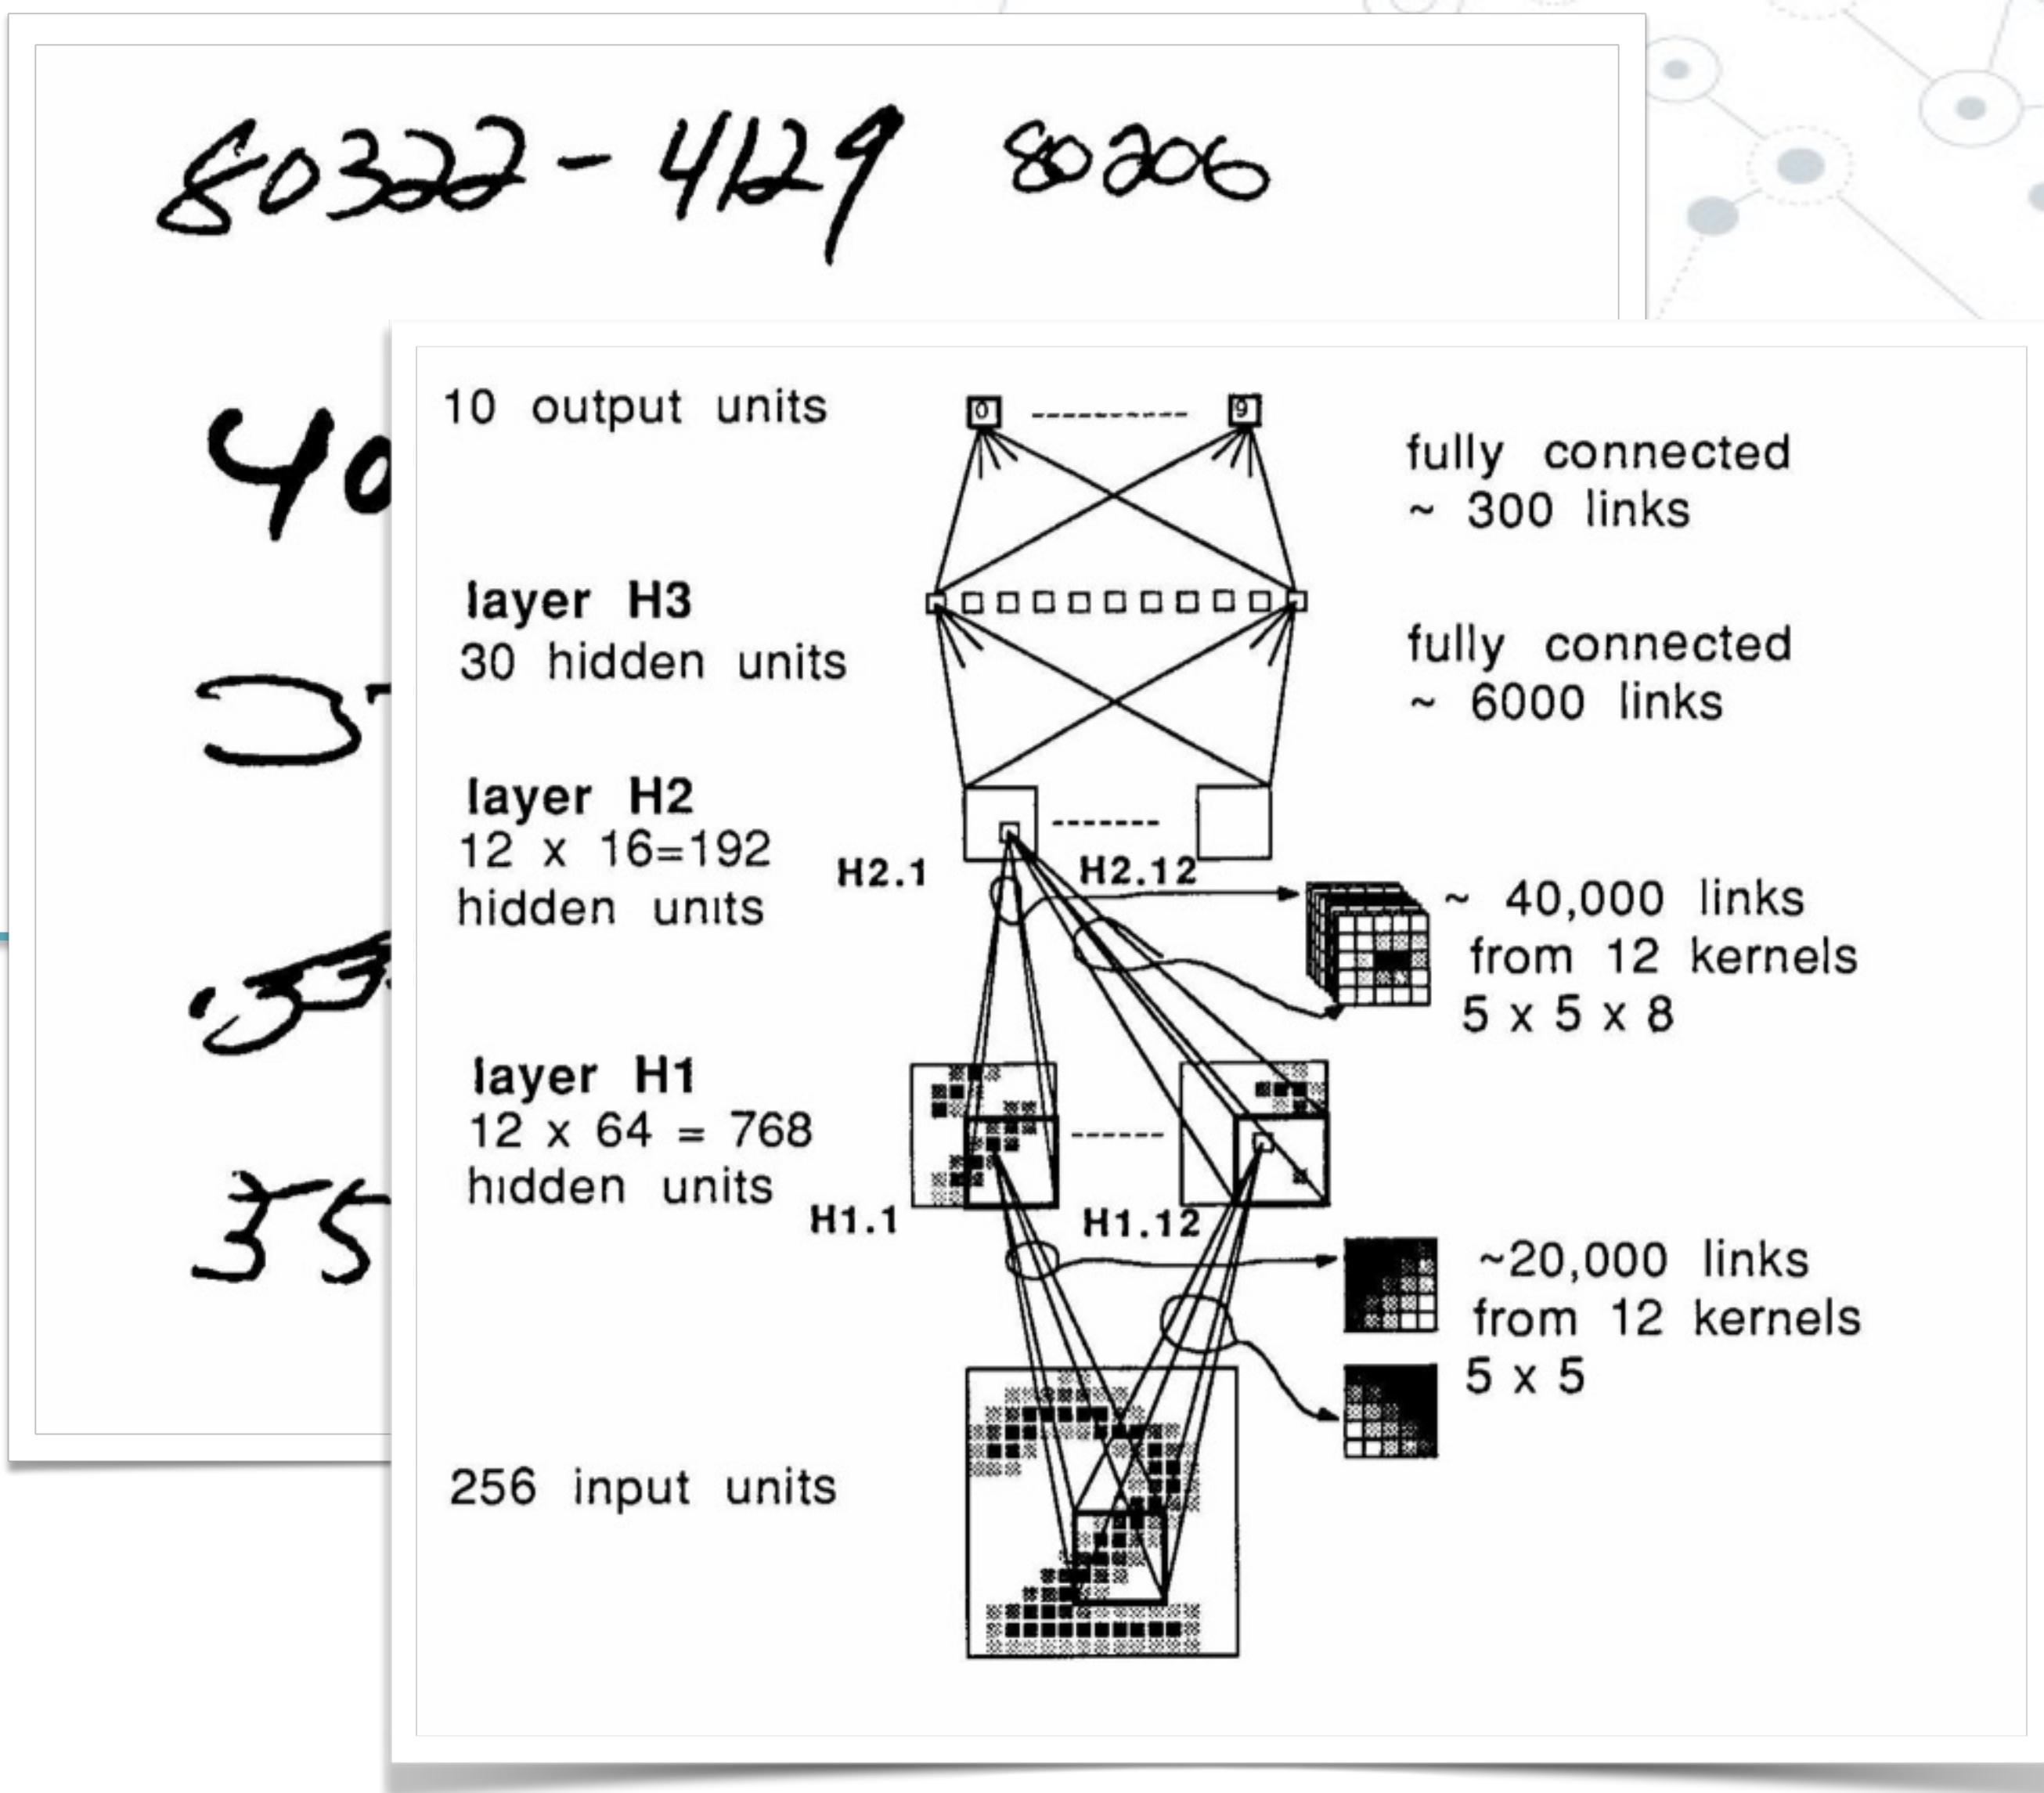
1980

The term of deep
learning is introduced

1986

LeCunn:
Applied back-
propagation
algorithm, to a
deep neural
network with the
purpose of
recognising
handwritten ZIP
codes

1989



LeCun, Y., Boser, B., Denker, J.S., Henderson, D., Howard, R.E., Hubbard, W. and Jackel, L.D., 1989. Backpropagation applied to handwritten zip code recognition. *Neural computation*, 1(4), pp.541-551.

A short history of deep learning

Ivaknenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

1980

The term of deep
learning is introduced

1986

LeCunn:
Applied back-
propagation
algorithm, to a
deep neural
network with the
purpose of
recognising
handwritten ZIP
codes

1989

The term of “deep” is
Used in the context
of neural networks

LeCun, Y., Boser, B., Denker, J.S., Henderson, D., Howard, R.E., Hubbard, W.
and Jackel, L.D., 1989. Backpropagation applied to handwritten zip code
recognition. *Neural computation*, 1(4), pp.541-551.

A short history of deep learning

Ivaknenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

1980

The term of deep
learning is introduced

1986

LeCunn:
Applied back-
propagation
algorithm, to a
deep neural
network with the
purpose of
recognising
handwritten ZIP
codes

LSTM
networks
proposed for
speech
recognition

2000

The term of “deep” is
Used in the context
of neural networks

1997

Hochreiter, S. and Schmidhuber, J., 1997. Long short-term memory. *Neural computation*, 9(8), pp.1735-1780.

A short history of deep learning

Ivakhnenco:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

1980

The term of deep
learning is introduced

1989

LSTM
networks
proposed for
speech
recognition

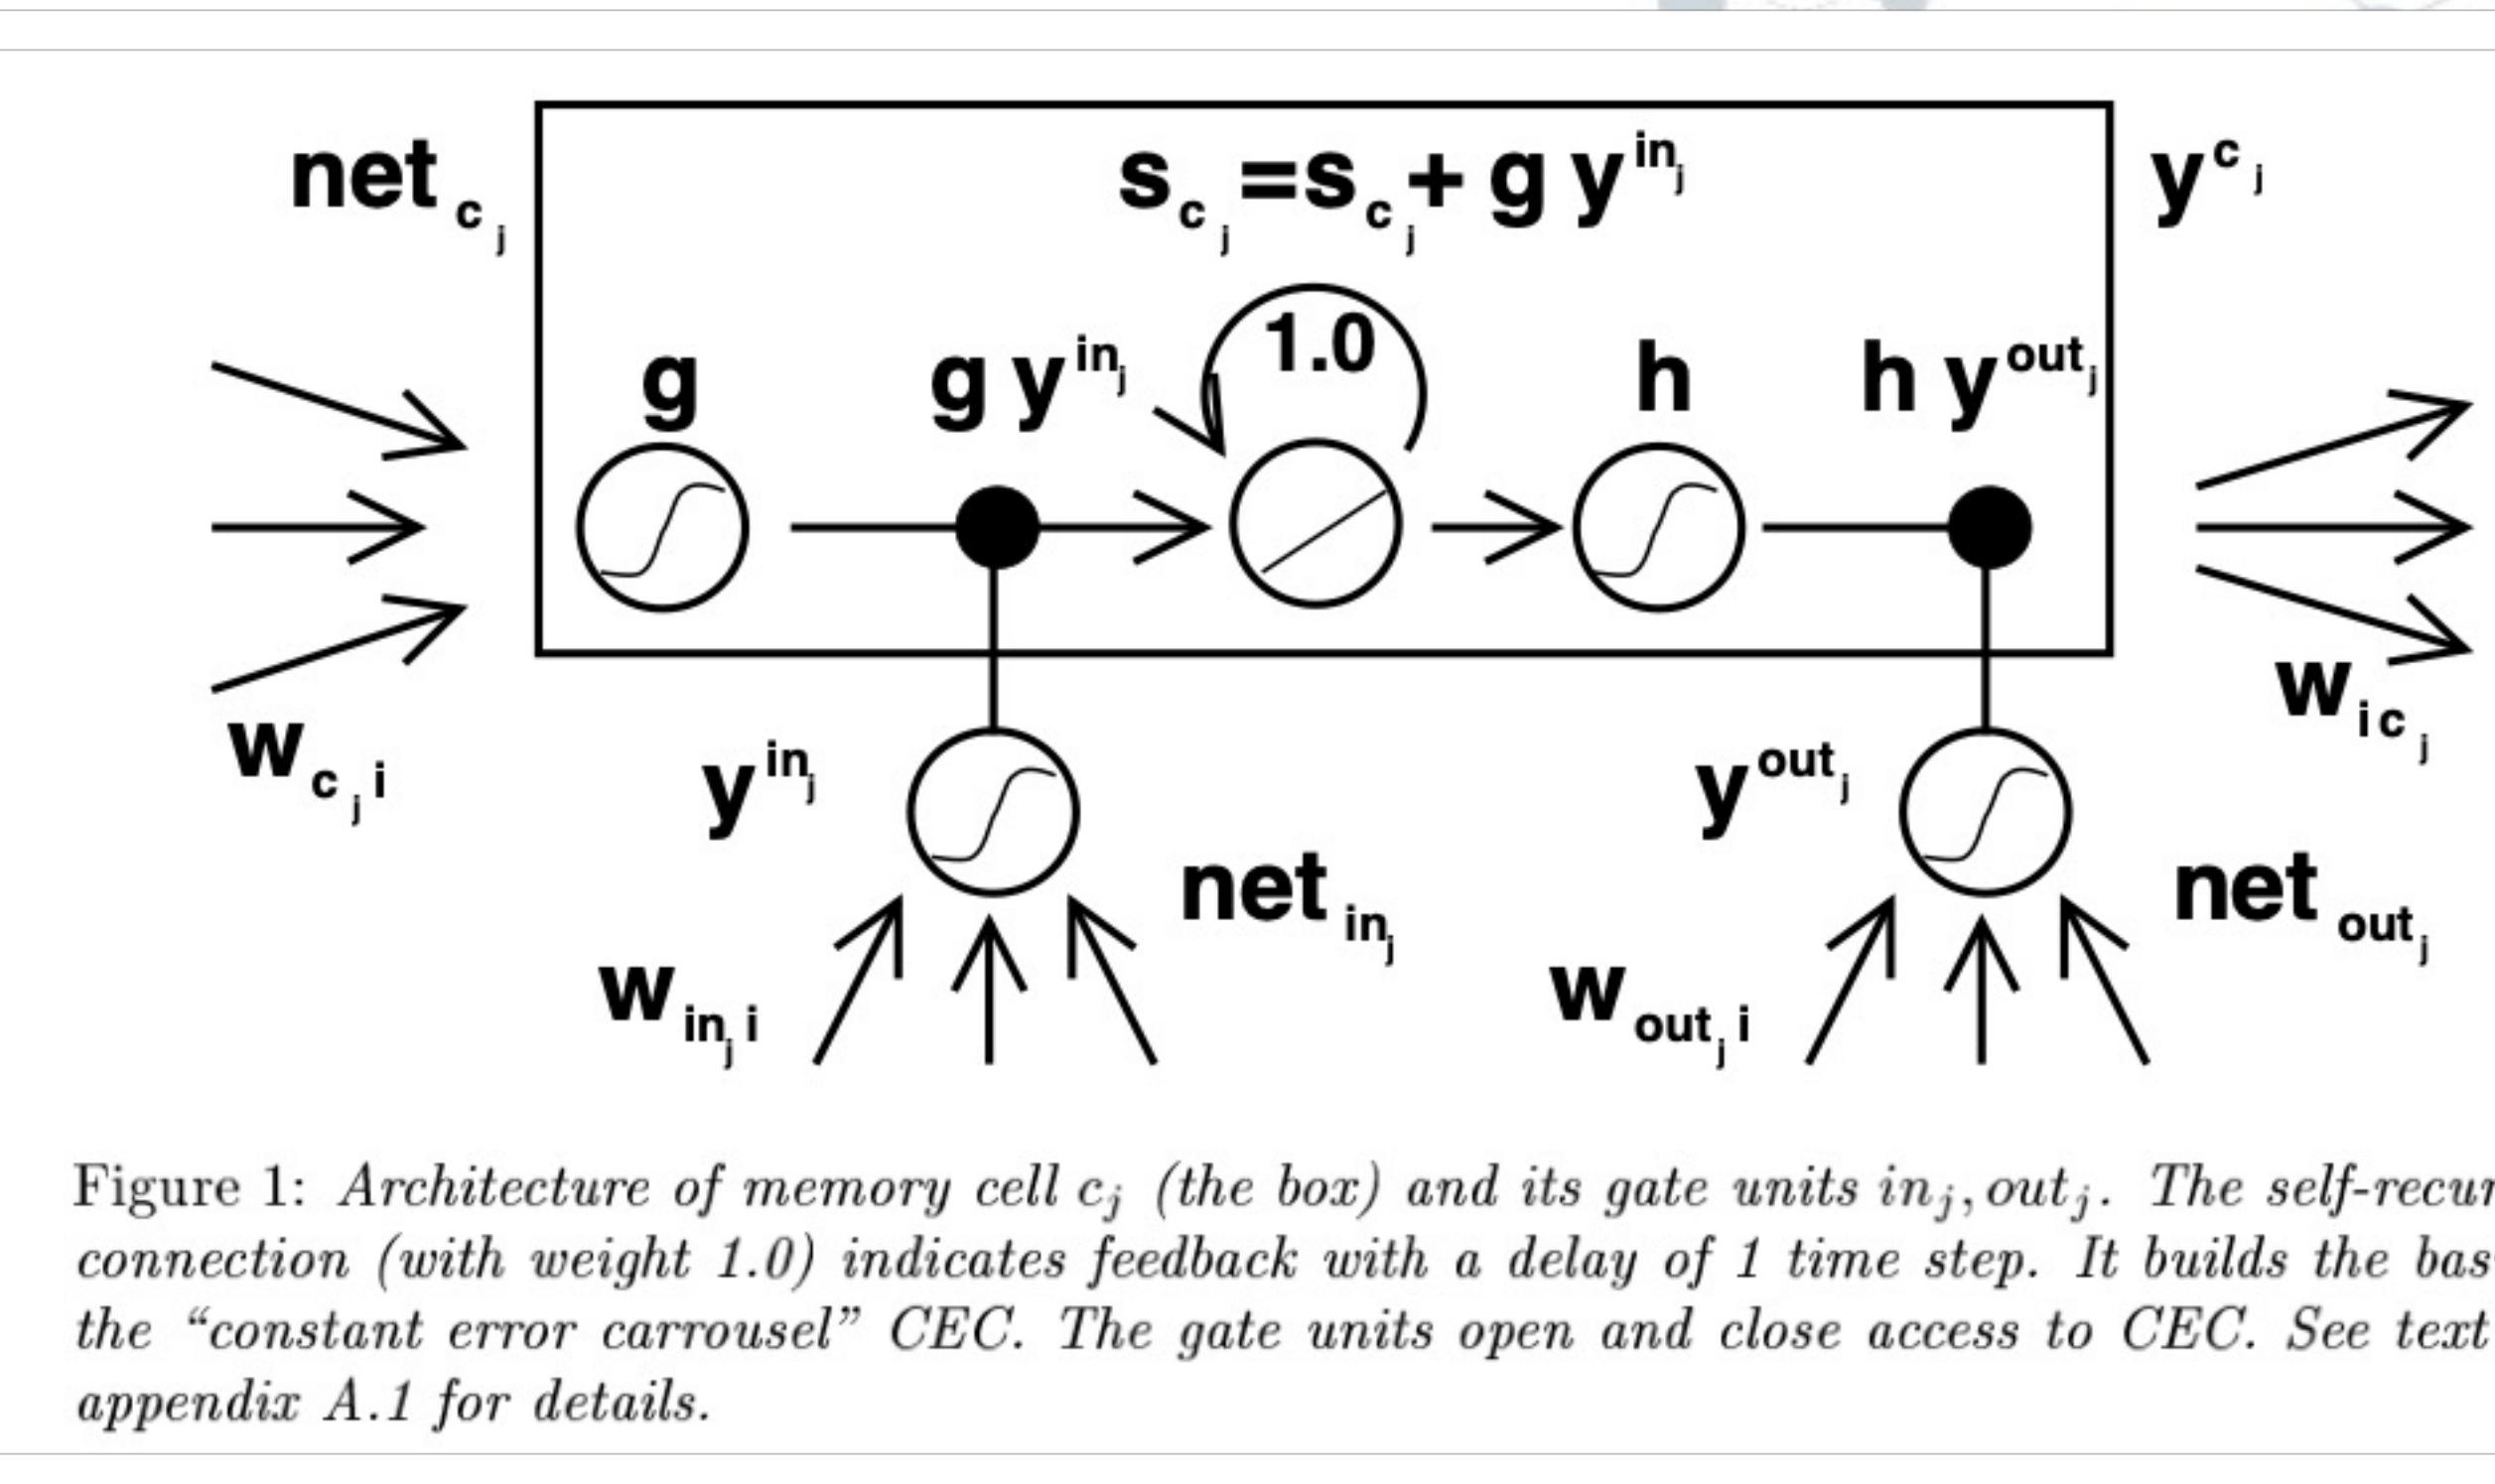
1986

LeCunn:
Applied back-
propagation
algorithm, to a
deep neural
network with the
purpose of
recognising
handwritten ZIP
codes

1997

The t
Used
of ne

Hochreiter, S. and Schmidhuber, J., 1997. Long short-term memory. *Neural computation*, 9(8), pp.1735-1780.



A short history of deep learning

Ivaknenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

1980

The term of deep
learning is introduced

1986

LeCunn:
Applied back-
propagation
algorithm, to a
deep neural
network with the
purpose of
recognising
handwritten ZIP
codes

LSTM
networks
proposed for
speech
recognition

2000

The term of “deep” is
Used in the context
of neural networks

1997

Hochreiter, S. and Schmidhuber, J., 1997. Long short-term memory. *Neural computation*, 9(8), pp.1735-1780.

A short history of deep learning

Ivaknenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

1980

The term of deep
learning is introduced

1989

LSTM
networks
proposed for
speech
recognition

1986

LeCunn:
Applied back-
propagation
algorithm, to a
deep neural
network with the
purpose of
recognising
handwritten ZIP
codes

2000

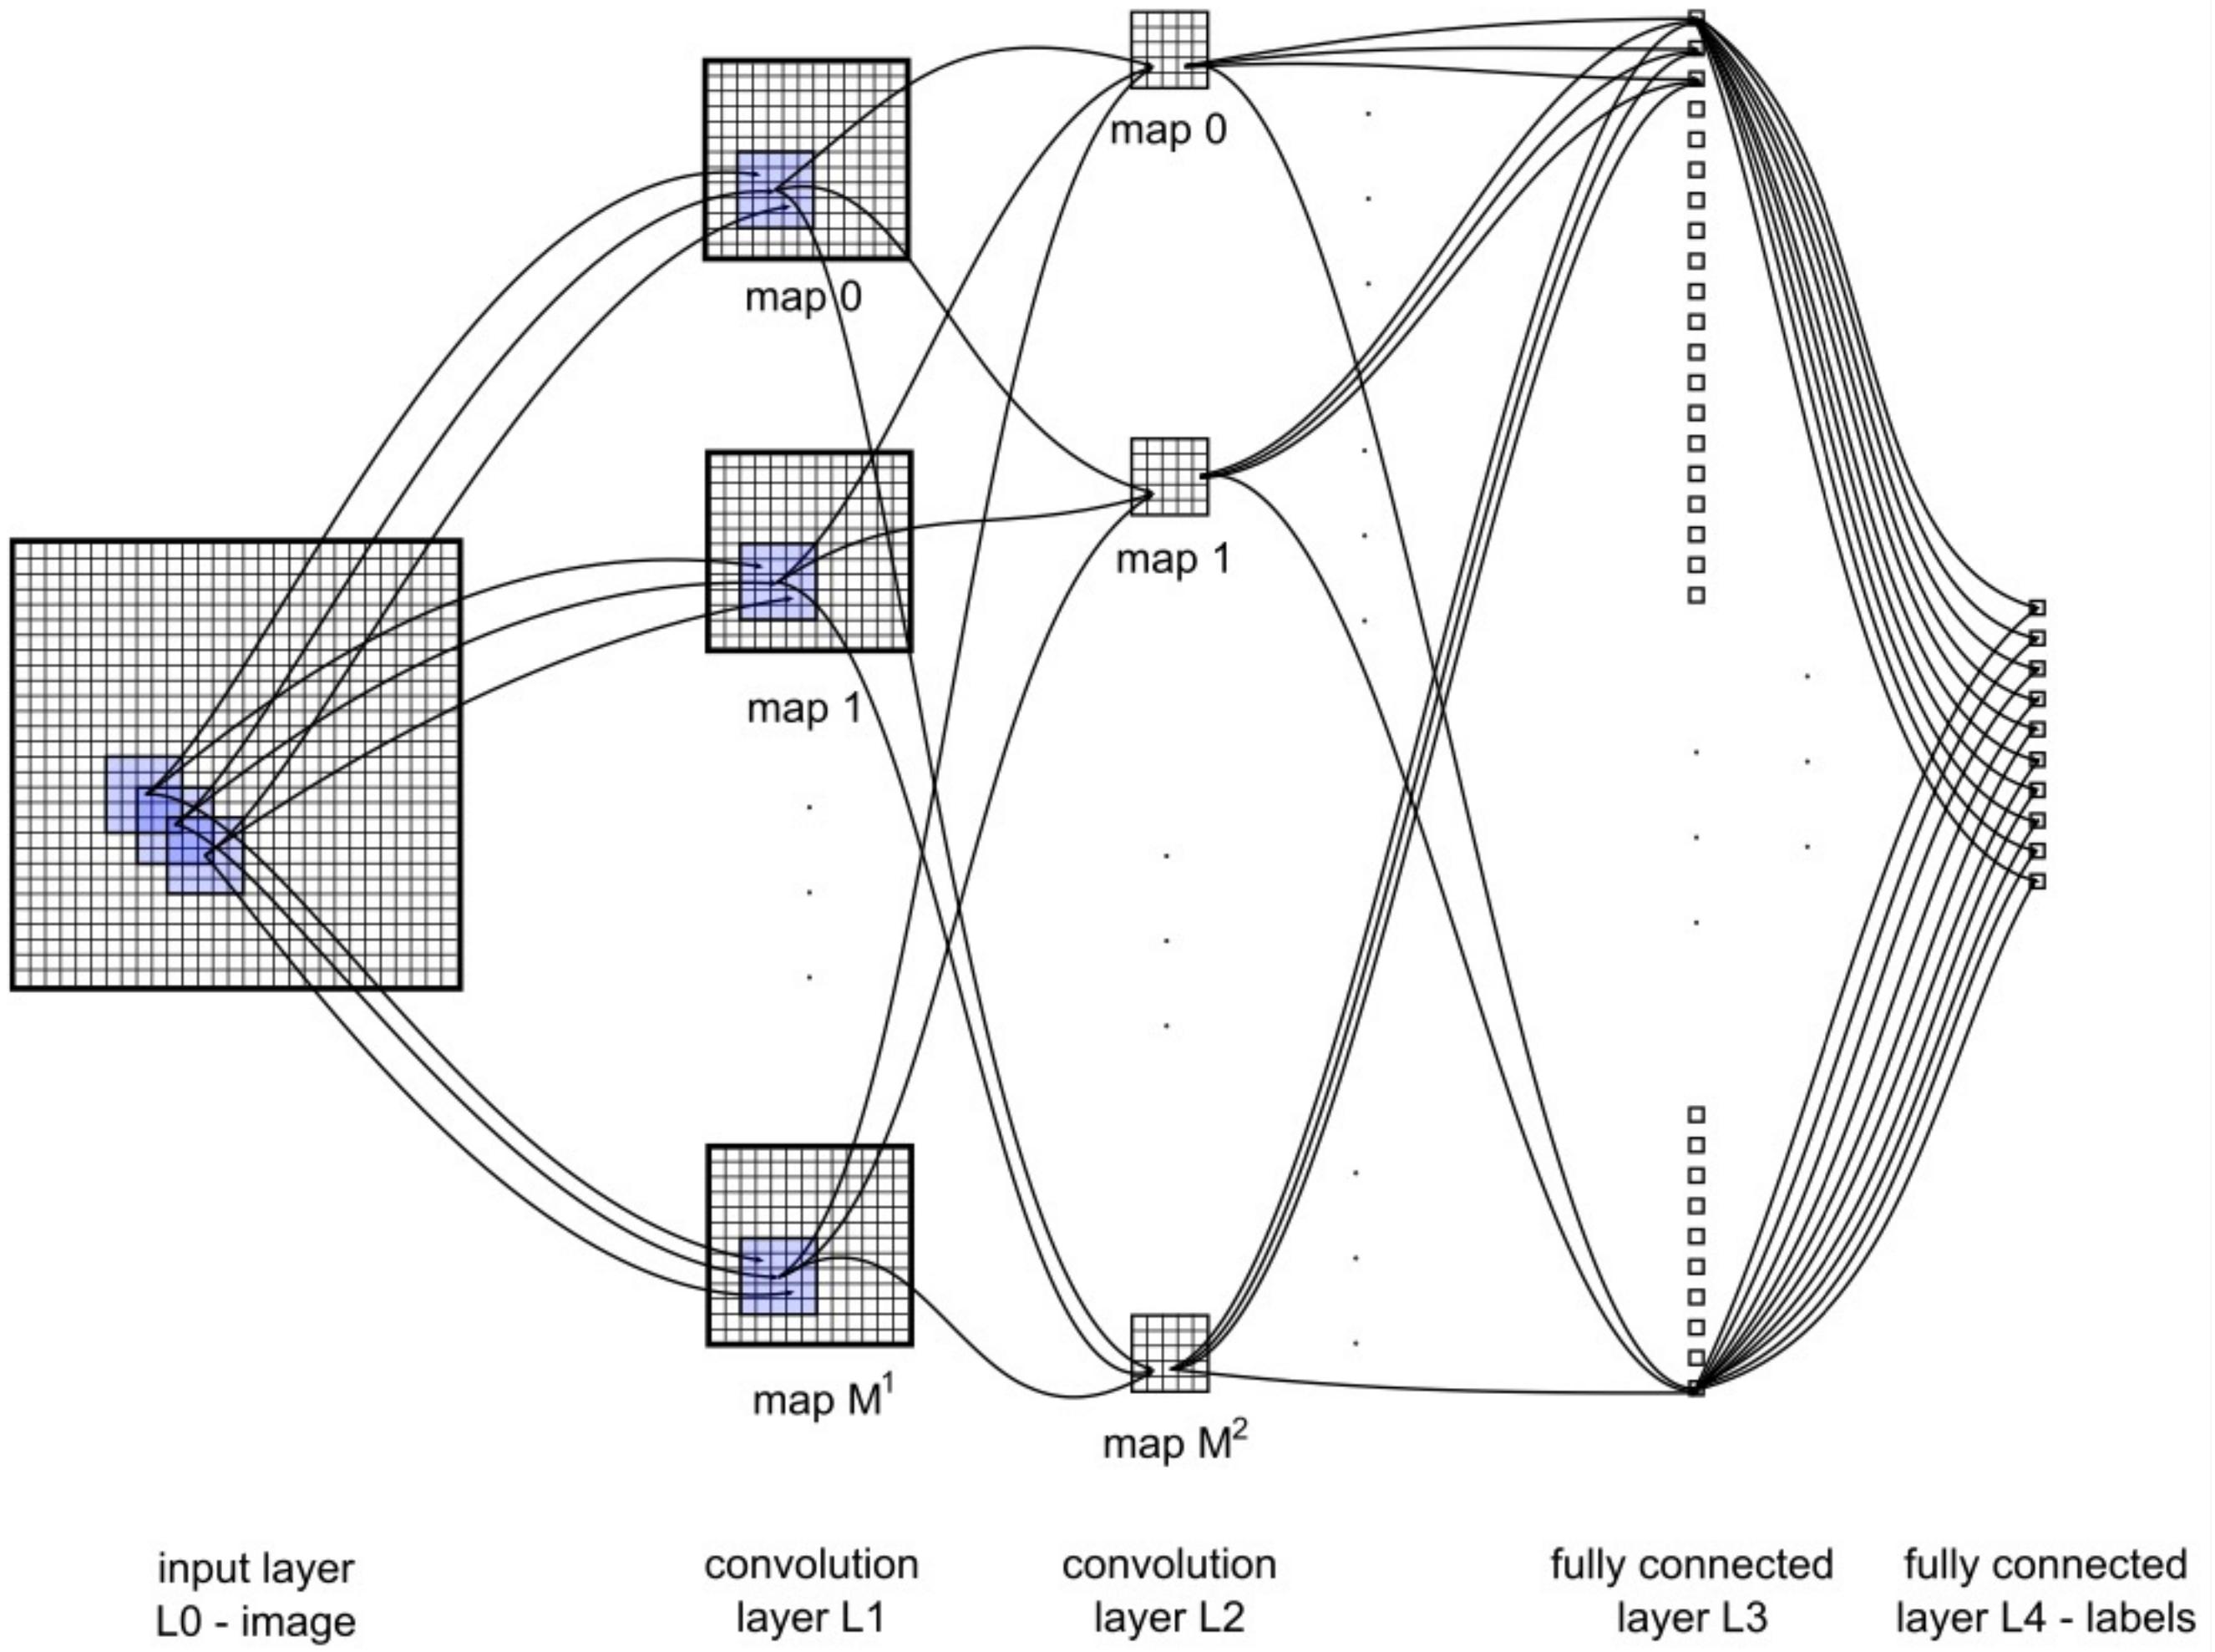
The term of “deep” is
Used in the context
of neural networks

2011

Ciresan: Proposed
CNN architecture
trained on GPU to
improve the results
of image
classification

Cireşan, D.C., Meier, U., Masci, J., Gambardella, L.M. and Schmidhuber, J.,
2011. High-performance neural networks for visual object classification. *arXiv preprint arXiv:1102.0183*.

A short history of deep learning



Ciresan: Proposed
CNN architecture
trained on GPU to
improve the results
of image
classification

is

2011

.., Masci, J., Gambardella, L.M. and Schmidhuber, J.,
2011. High-performance neural networks for visual object classification. *arXiv preprint arXiv:1102.0183*.

A short history of deep learning

Ivaknenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

1980

The term of deep
learning is introduced

1989

LSTM
networks
proposed for
speech
recognition

1986

LeCunn:
Applied back-
propagation
algorithm, to a
deep neural
network with the
purpose of
recognising
handwritten ZIP
codes

2000

The term of “deep” is
Used in the context
of neural networks

2011

Ciresan: Proposed
CNN architecture
trained on GPU to
improve the results
of image
classification

Cireşan, D.C., Meier, U., Masci, J., Gambardella, L.M. and Schmidhuber, J.,
2011. High-performance neural networks for visual object classification. *arXiv preprint arXiv:1102.0183*.

A short history of deep learning

Ivaknenko:
“working
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

1980

The term of deep
learning is introduced

1989

LSTM
networks
proposed for
speech
recognition

1986

LeCunn:
Applied back-
propagation
algorithm, to a
deep neural
network with the
purpose of
recognising
handwritten ZIP
codes

2000

The term of “deep” is
Used in the context
of neural networks

1997

Ciresan: Proposed
CNN architecture
trained on GPU to
improve the results
of image
classification

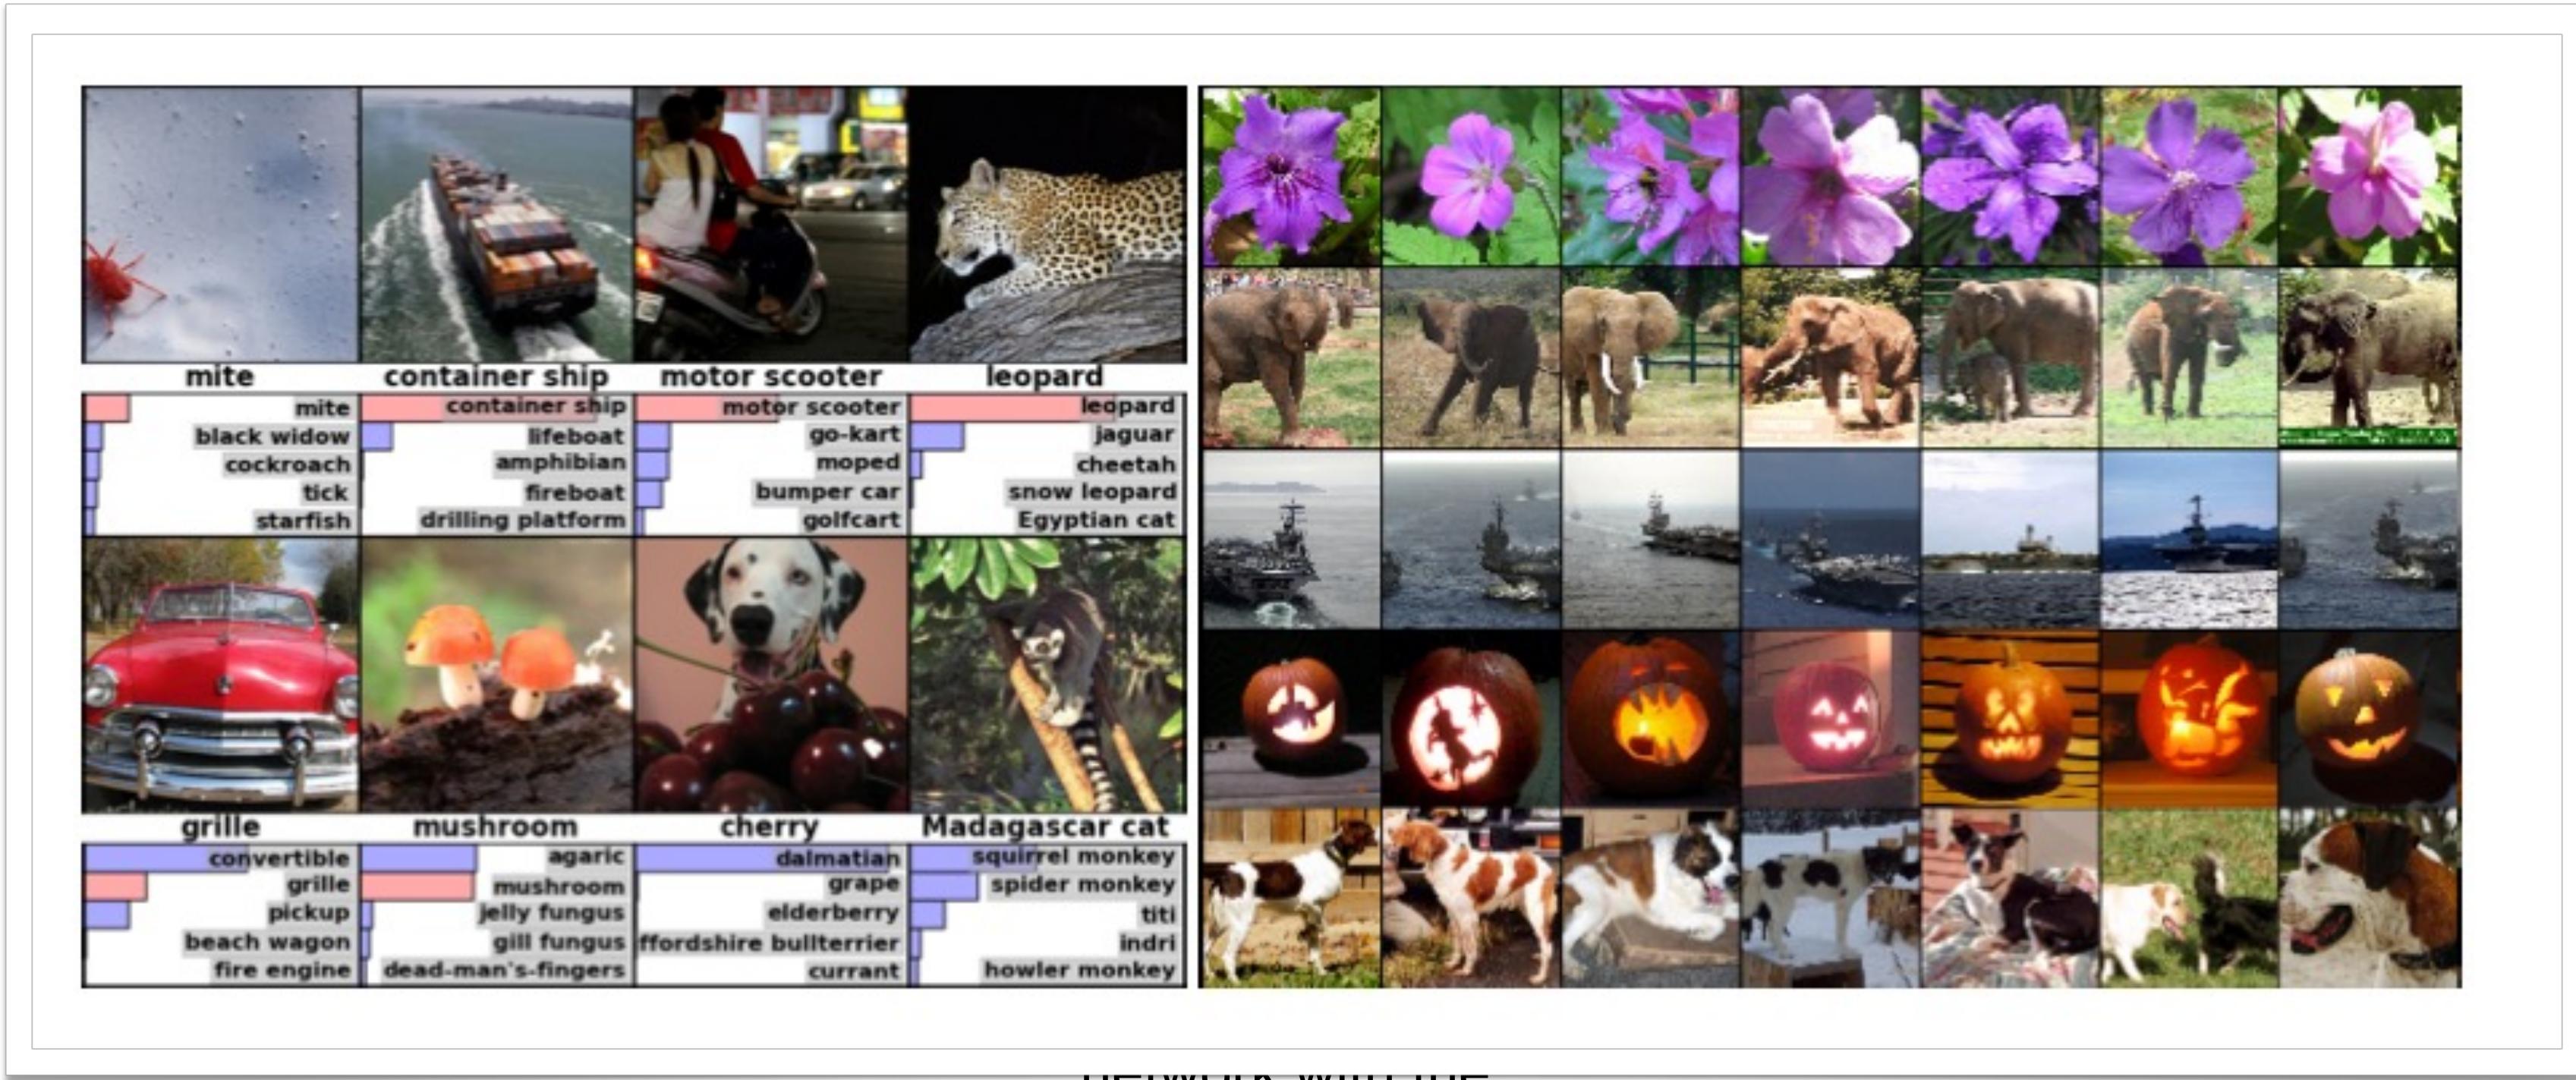
2011

AlexNet
architecture
won the
ImageNet
competition
by a large
margin

2012

Krizhevsky, A., Sutskever, I. and Hinton, G.E., 2017. Imagenet classification with deep convolutional neural networks. *Communications of the ACM*, 60(6), pp.84-90.

A short history of deep learning

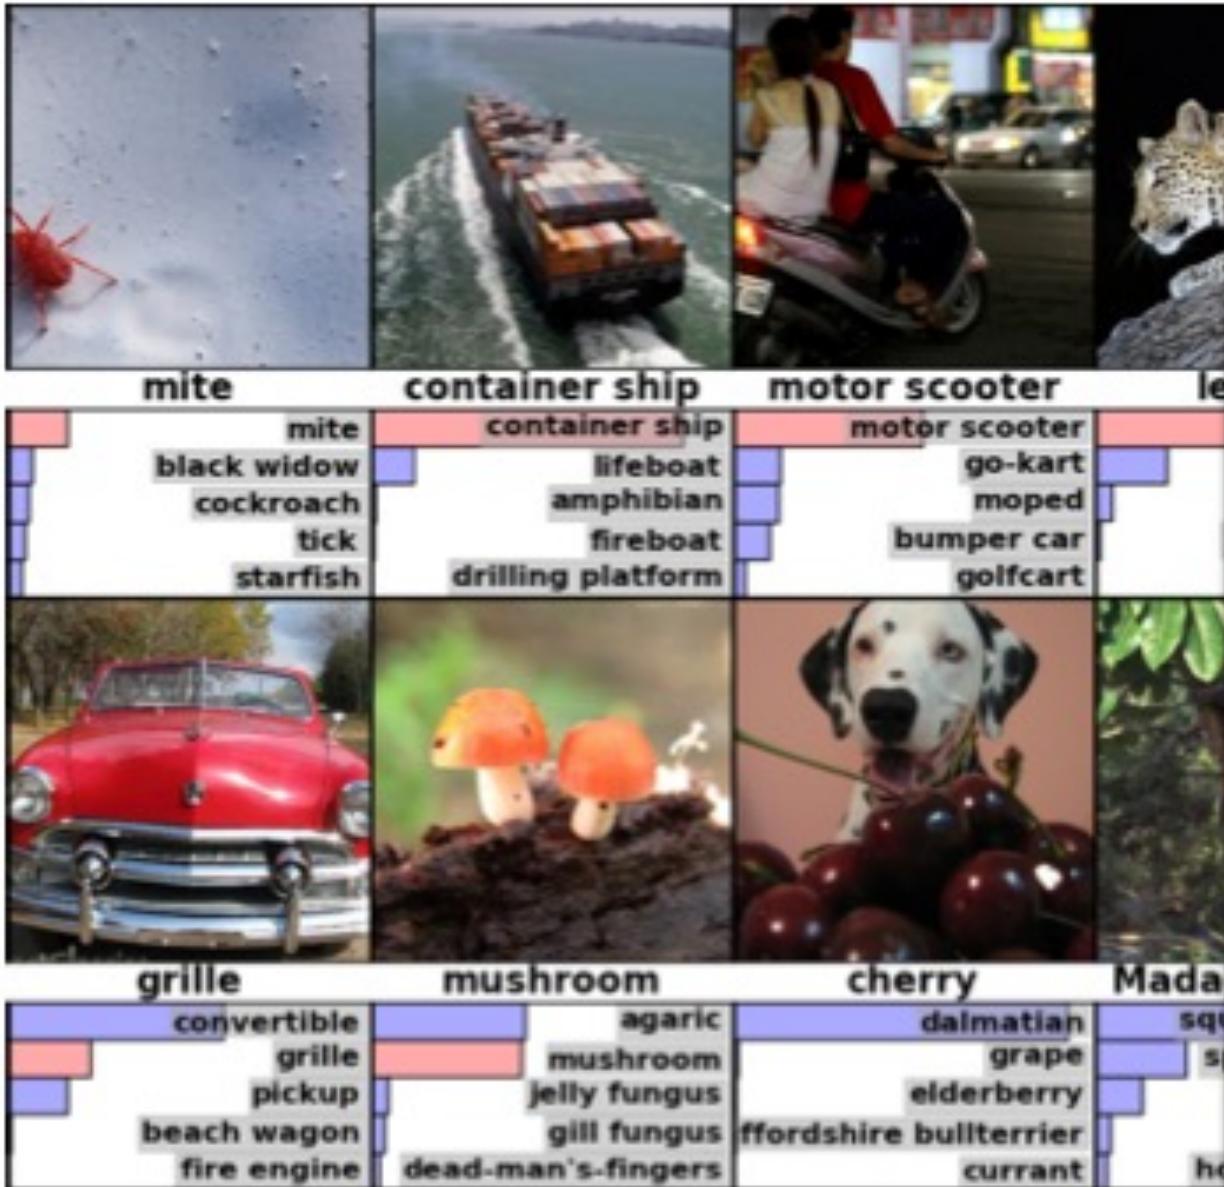


osed
ture
PU to
results
2012

AlexNet
architecture
won the
ImageNet
competition
by a large
margin

Krizhevsky, A., Sutskever, I. and Hinton, G.E., 2017. Imagenet classification with deep convolutional neural networks. *Communications of the ACM*, 60(6), pp.84-90.

A short history of deep learning



Model	Top-1 (val)	Top-5 (val)	Top-5 (test)
<i>SIFT + FVs [7]</i>	—	—	26.2%
1 CNN	40.7%	18.2%	—
5 CNNs	38.1%	16.4%	16.4%
1 CNN*	39.0%	16.6%	—
7 CNNs*	36.7%	15.4%	15.3%

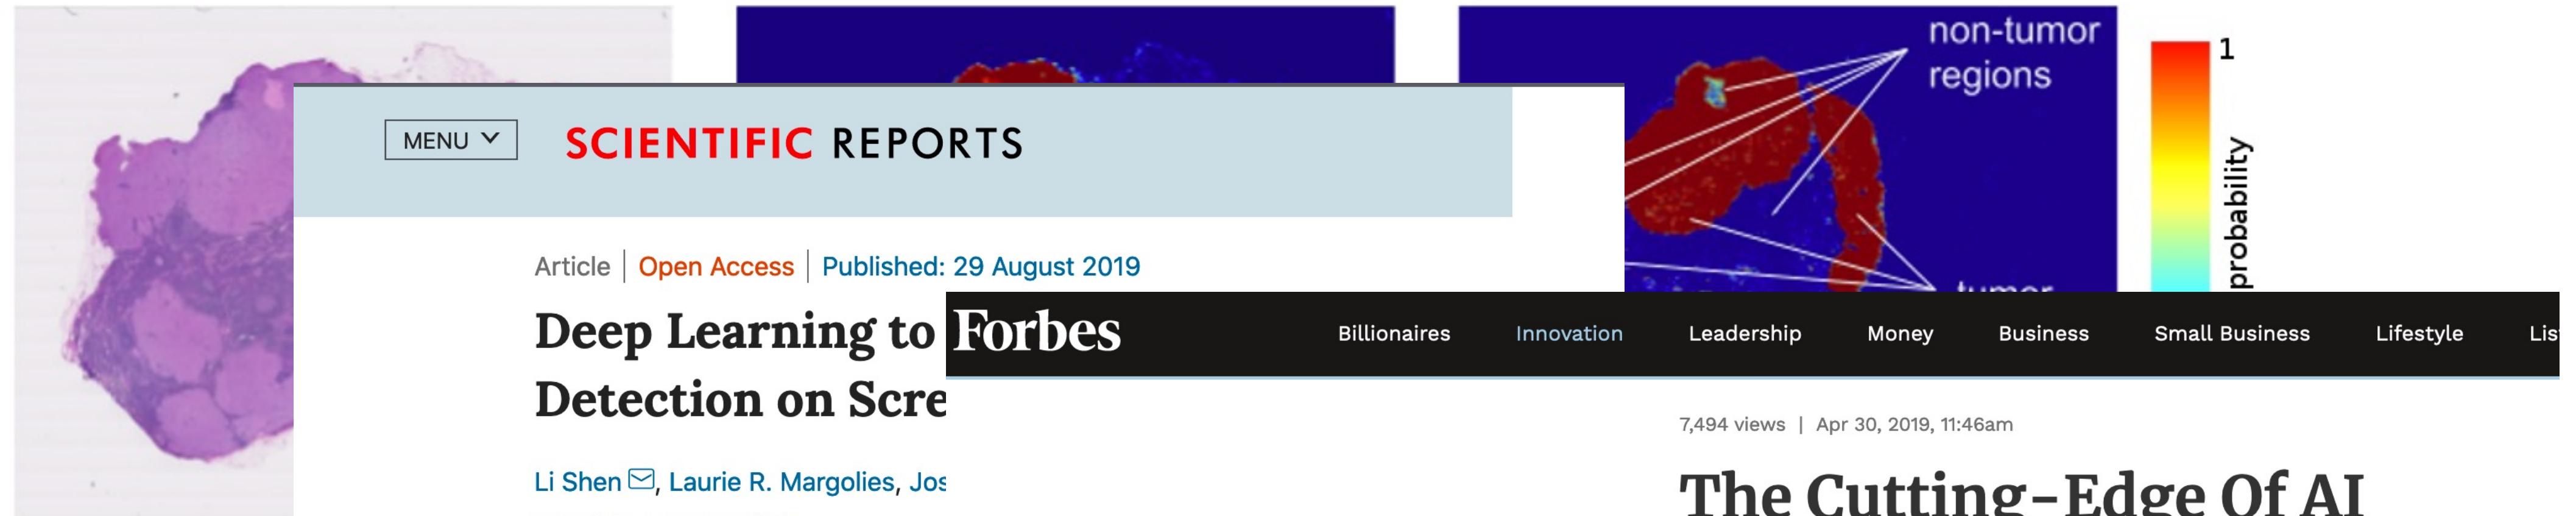
Table 2: Comparison of error rates on ILSVRC-2012 validation and test sets. In *italics* are best results achieved by others. Models with an asterisk* were “pre-trained” to classify the entire ImageNet 2011 Fall release. See Section 6 for details.

Krizhevsky, A., Sutskever, I. and Hinton, G.E., 2017. Imagenet classification with deep convolutional neural networks. *Communications of the ACM*, 60(6), pp.84-90.



15:09:49

Deep Learning to detect tumours



7,494 views | Apr 30, 2019, 11:46am

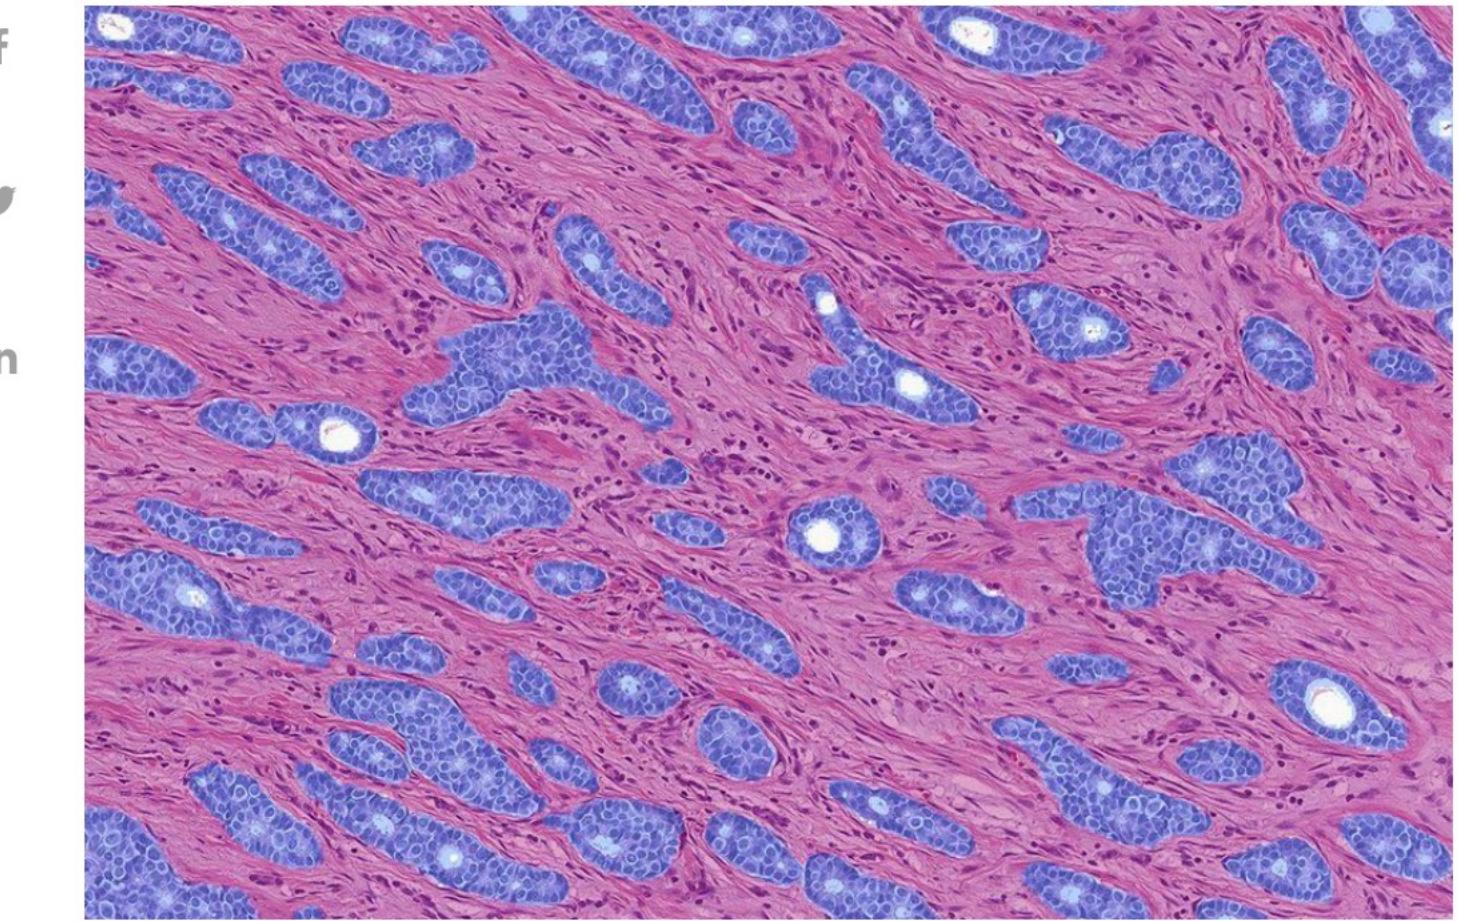
The Cutting-Edge Of AI Cancer Detection



Charles Towers-Clark Contributor

AI & Big Data

I write about AI, data, deep tech & self-management in the digital age



Detecting and treating cancer could benefit from a dose of AI, whether in pathological H&E slides or... [\[+\]](#) LUNIT.IO

1



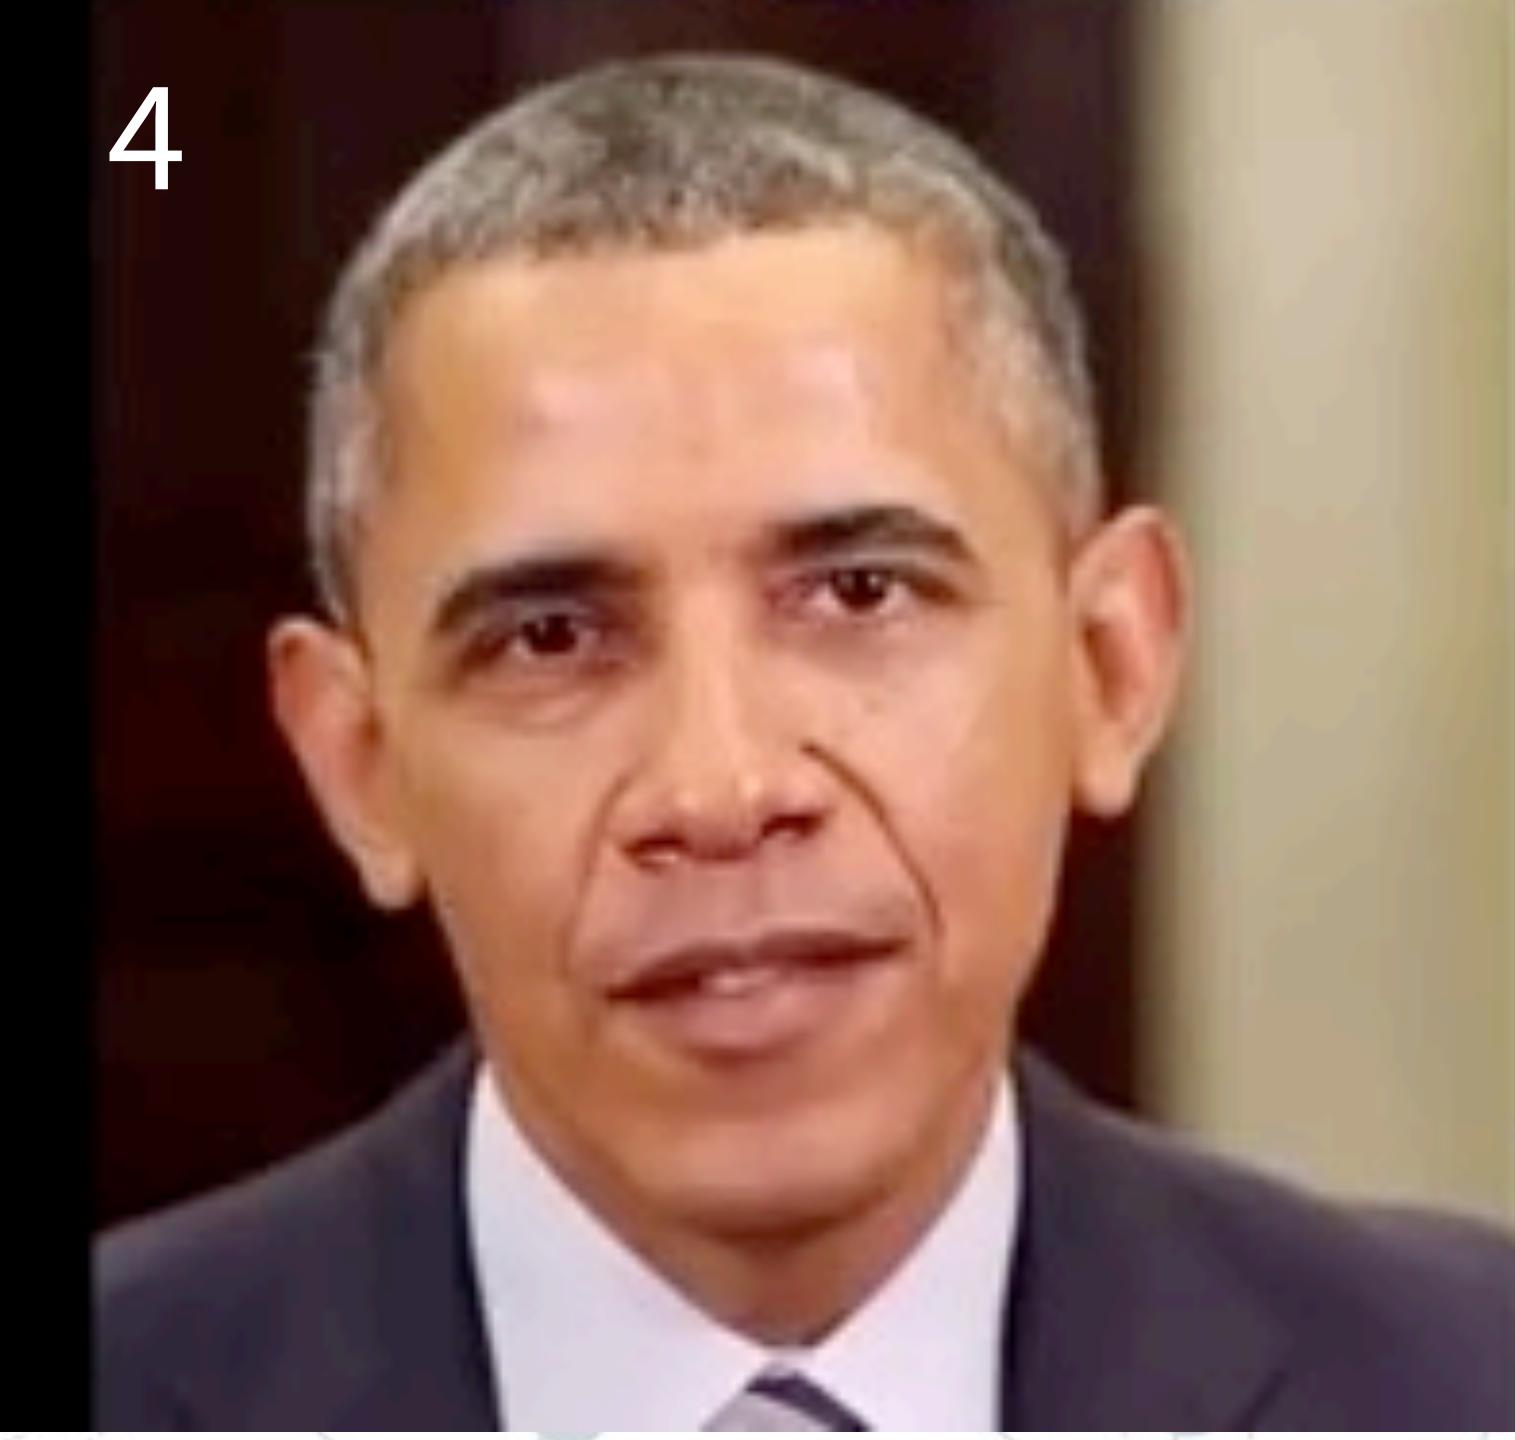
2



3



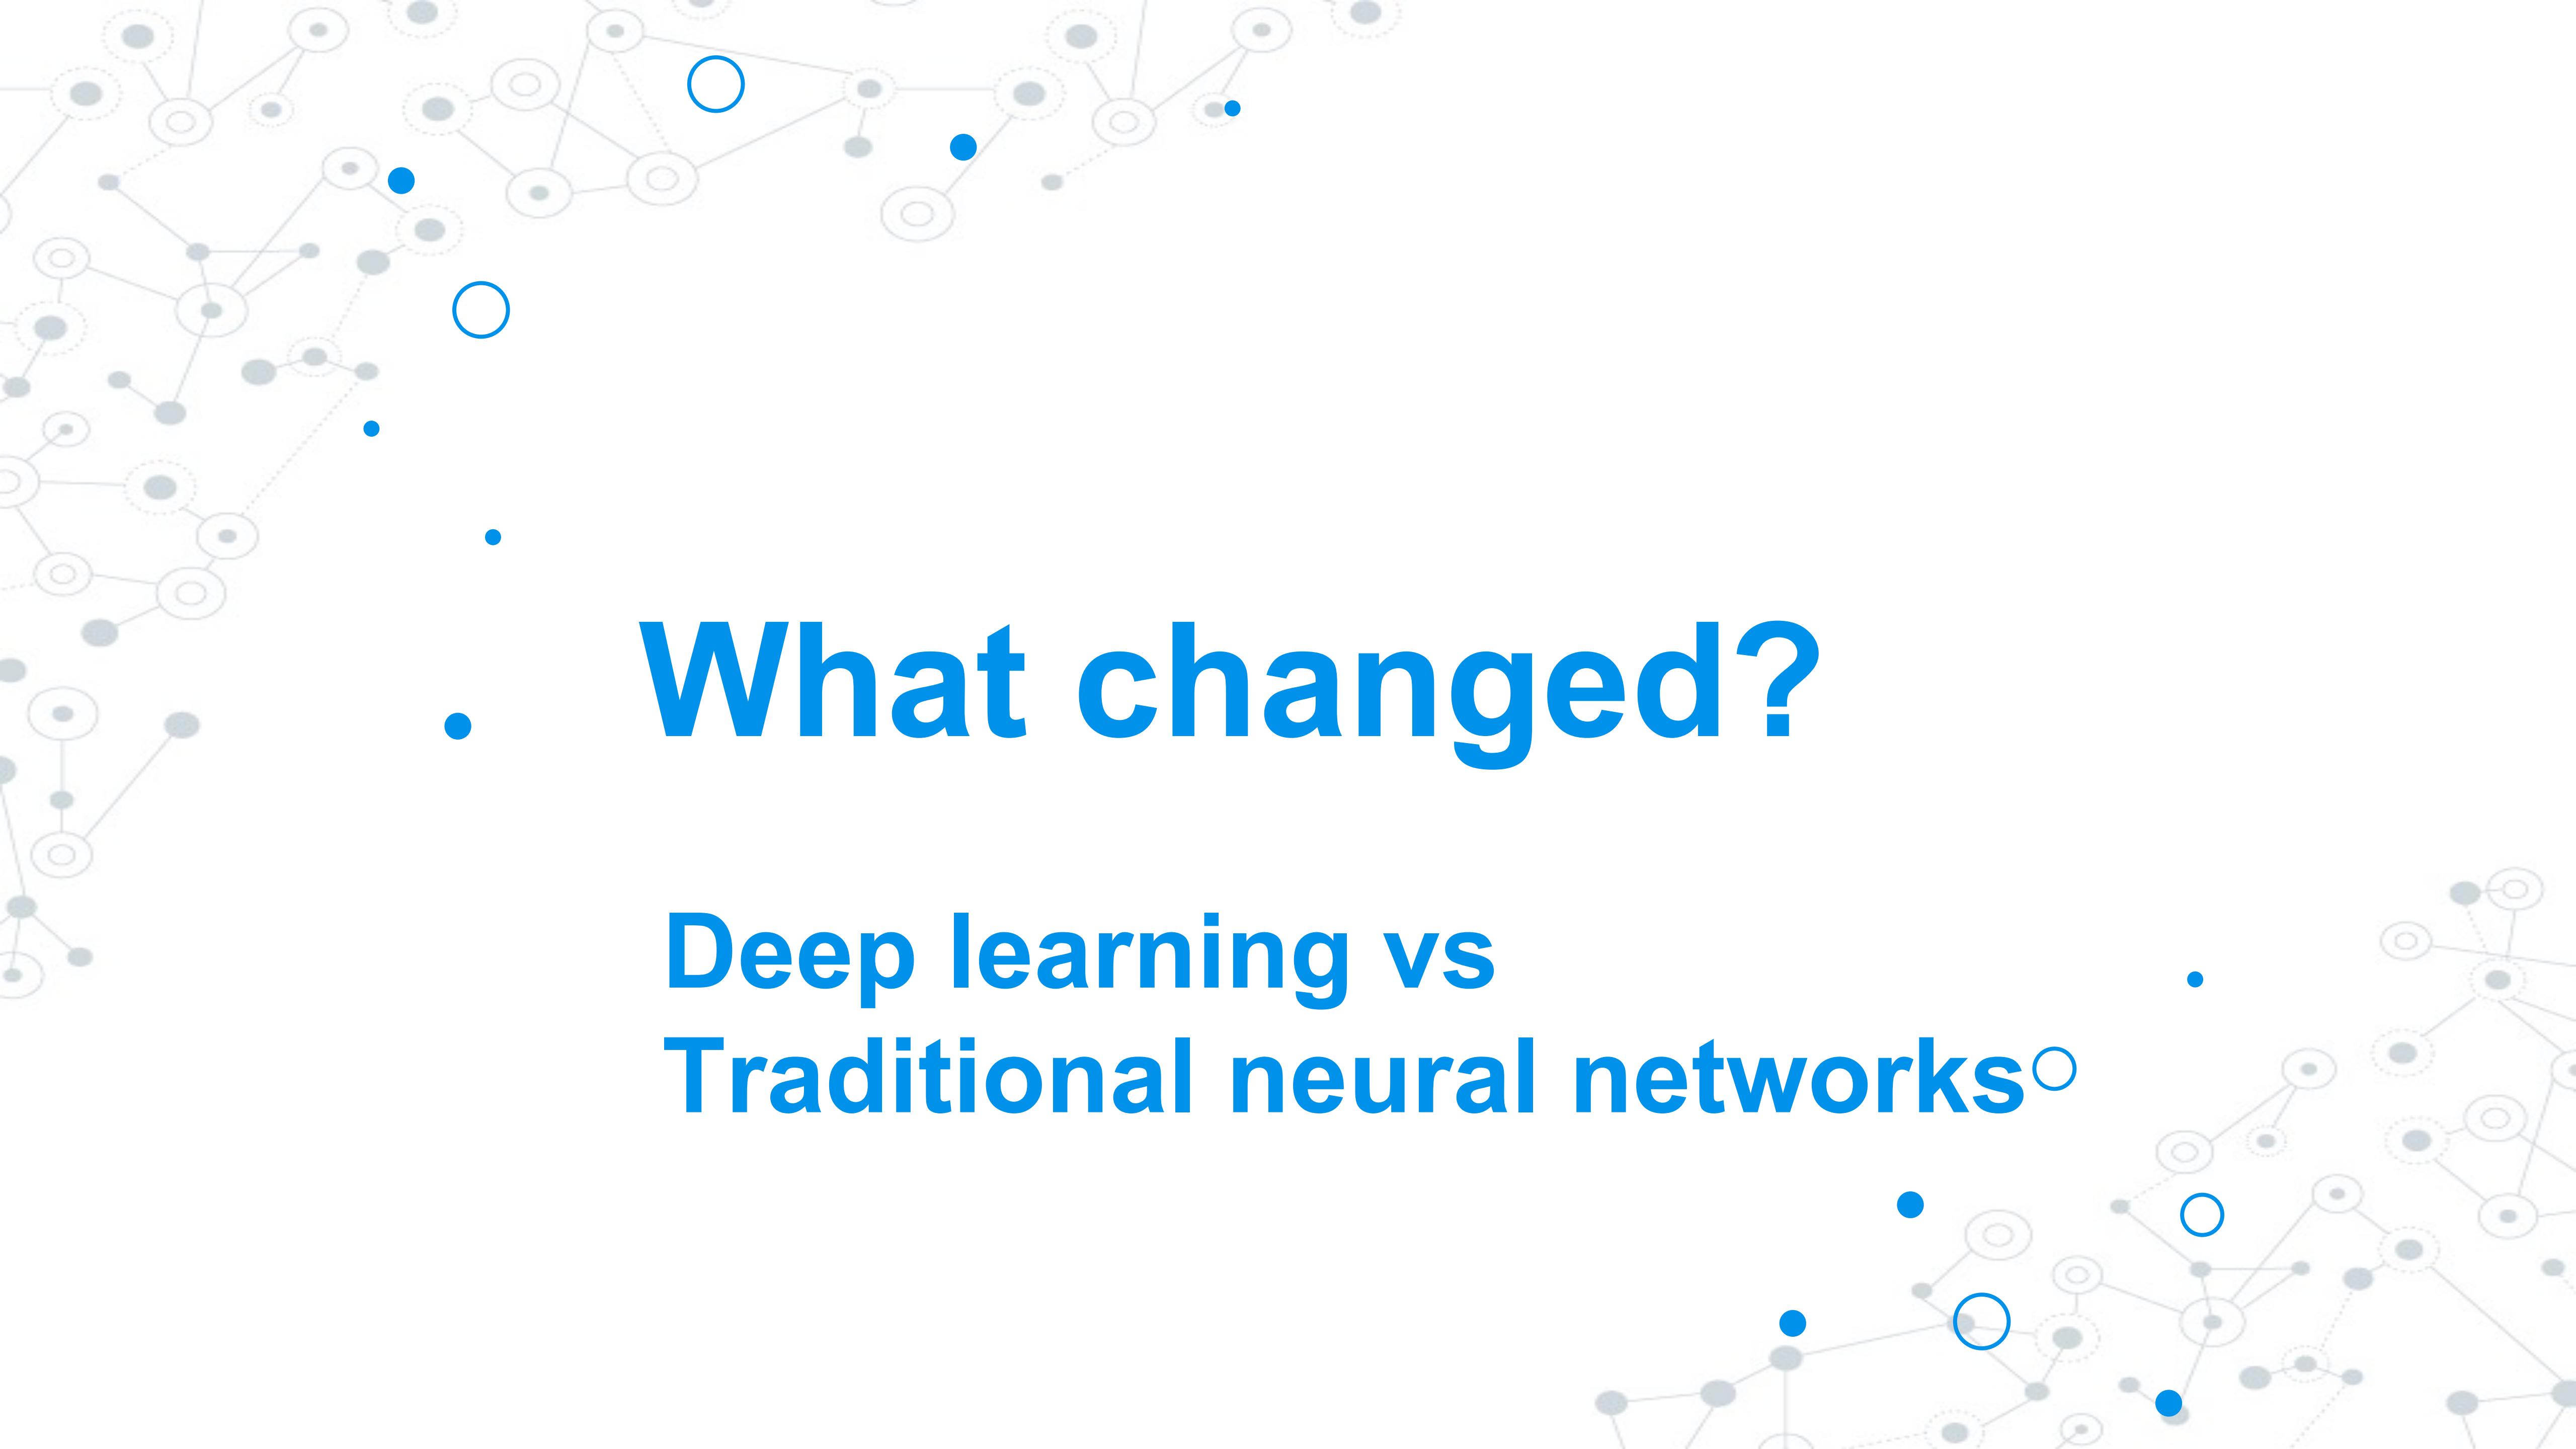
4





Check out these links:

- <https://www.nytimes.com/interactive/2020/11/21/science/artificial-intelligence-fake-people-faces.html>
- <https://www.thispersondoesnotexist.com>



What changed?

**Deep learning vs
Traditional neural networks**

● Decision Trees
Search term

+ Compare

United States ▾

Past 12 months

All categories ▾

Web Search ▾

Interest over time

Past hour

Past 4 hours

Past day

Past 7 days

Past 30 days

Past 90 days

Past 5 years

Nov 17, 2019

Mar 8, 2020

Jun 28, 2020

Oct 18, 2020

2004 - present

Custom time rang...

Interest by subregion

?



1 Massachusetts

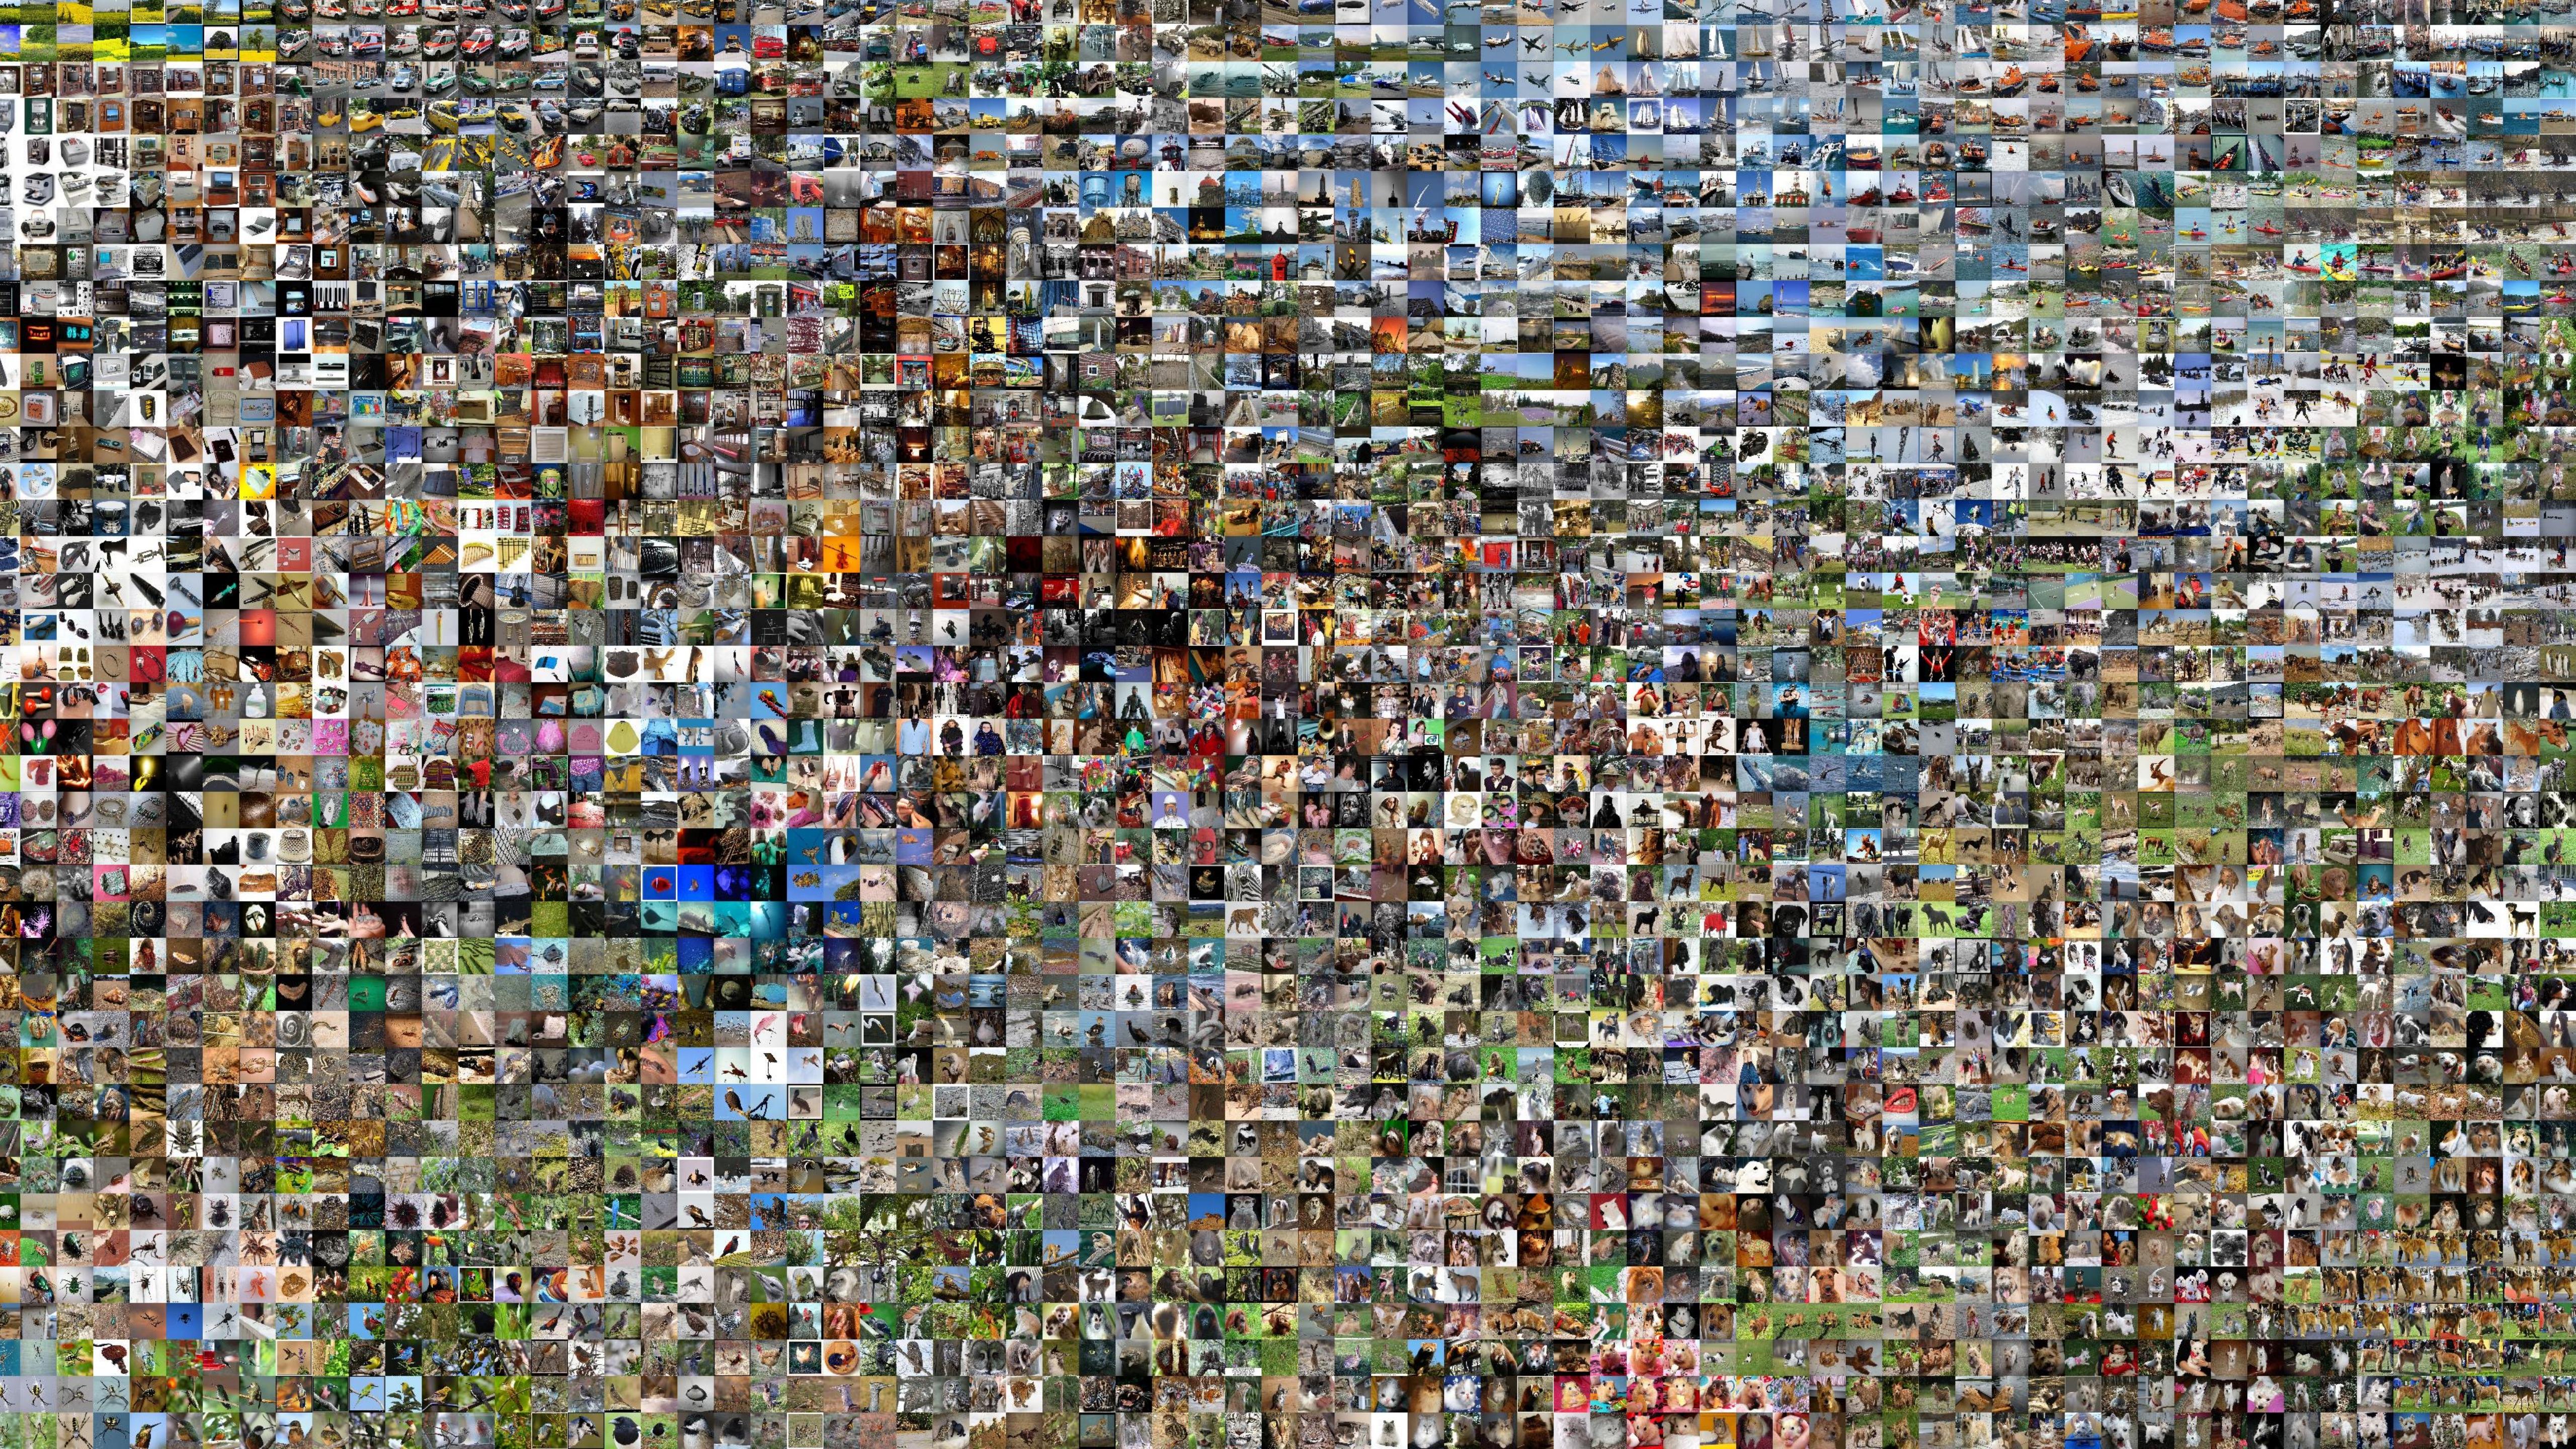
100

2 New Jersey

94

Lots and lots of

DATA



COMPUTATION POWER

GPU

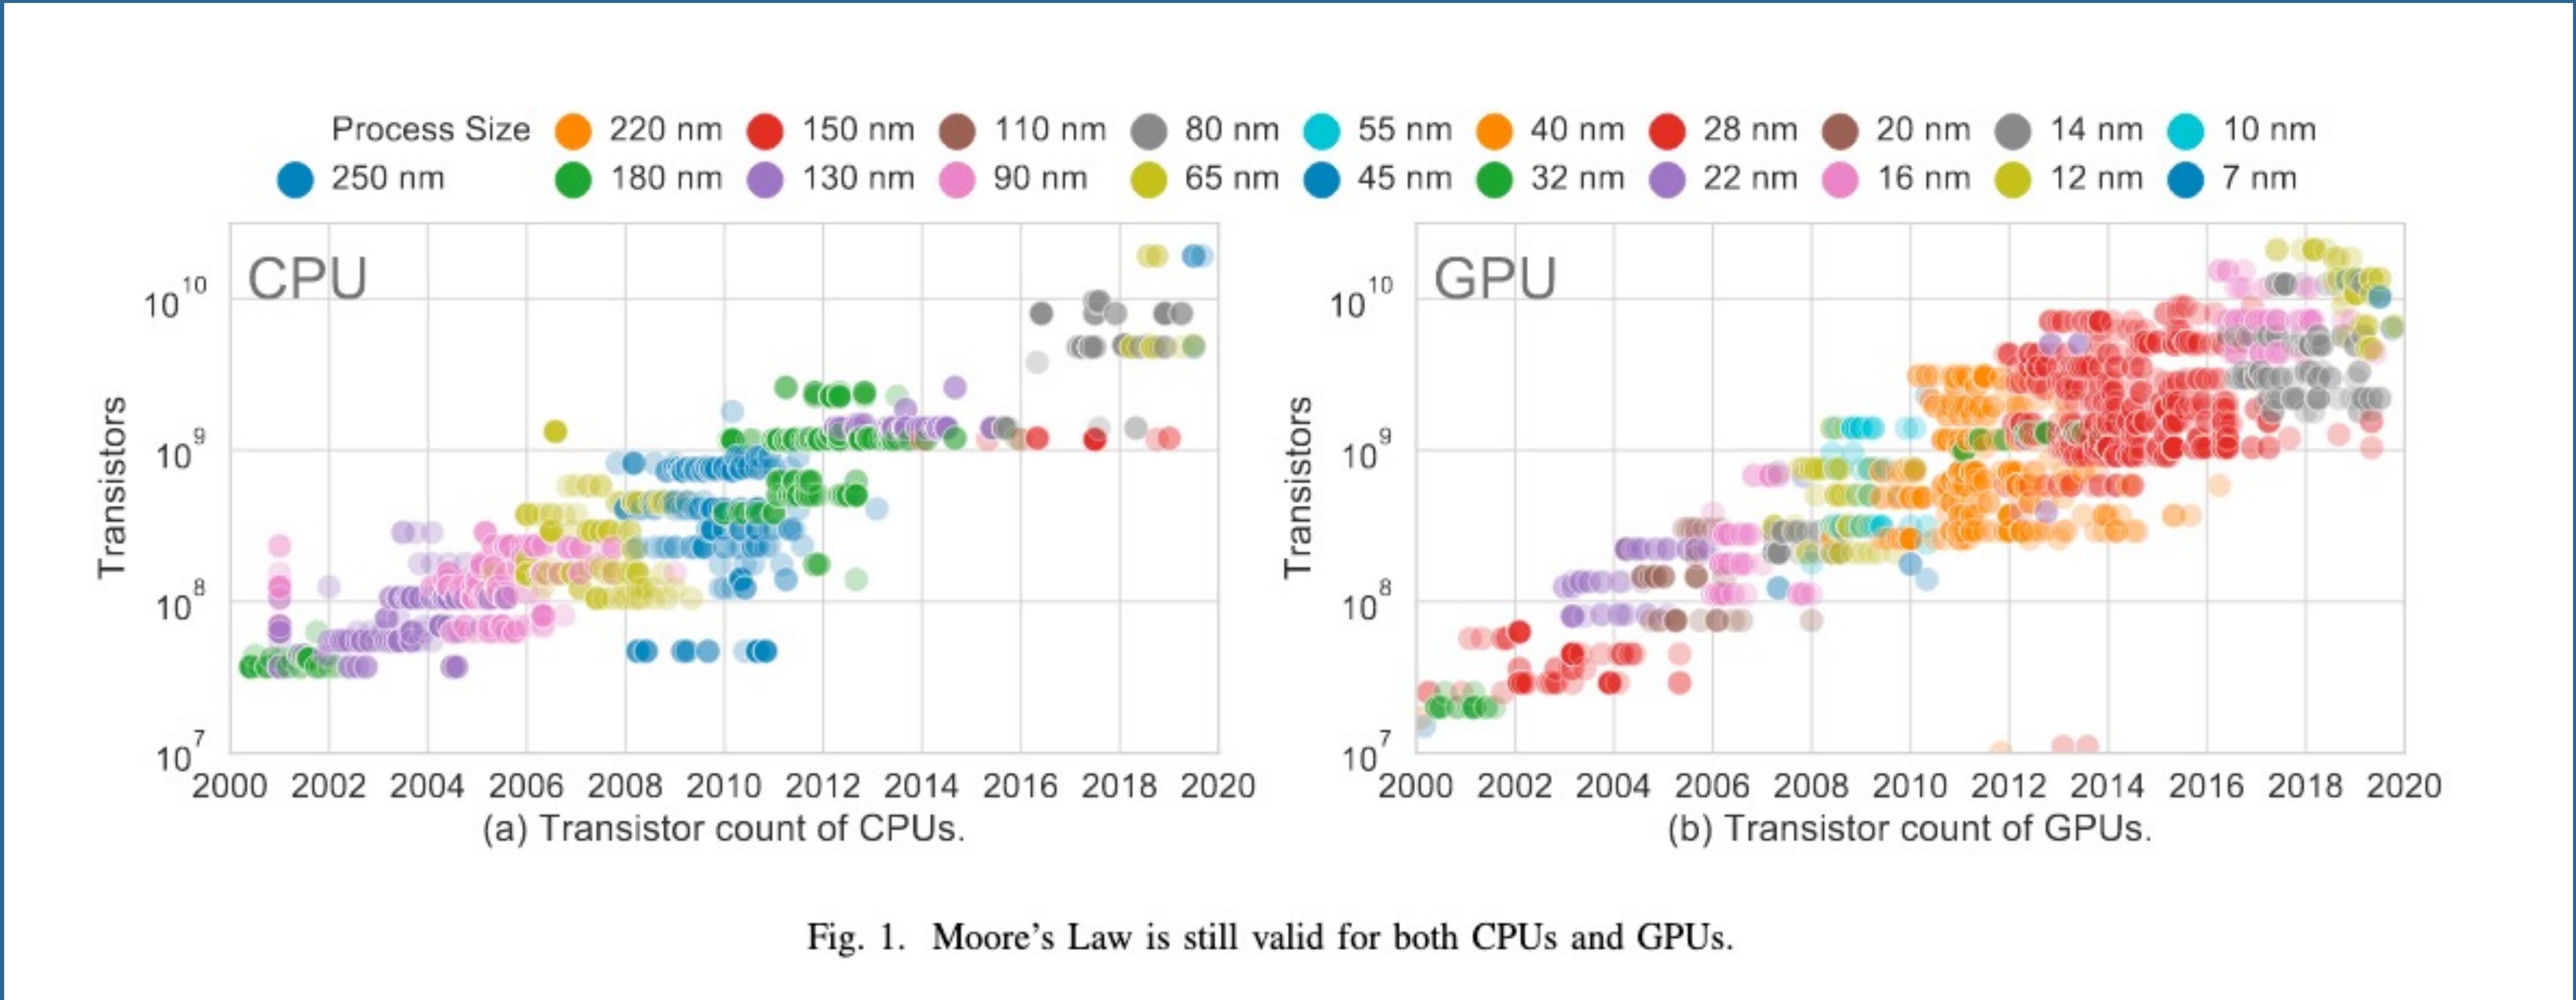
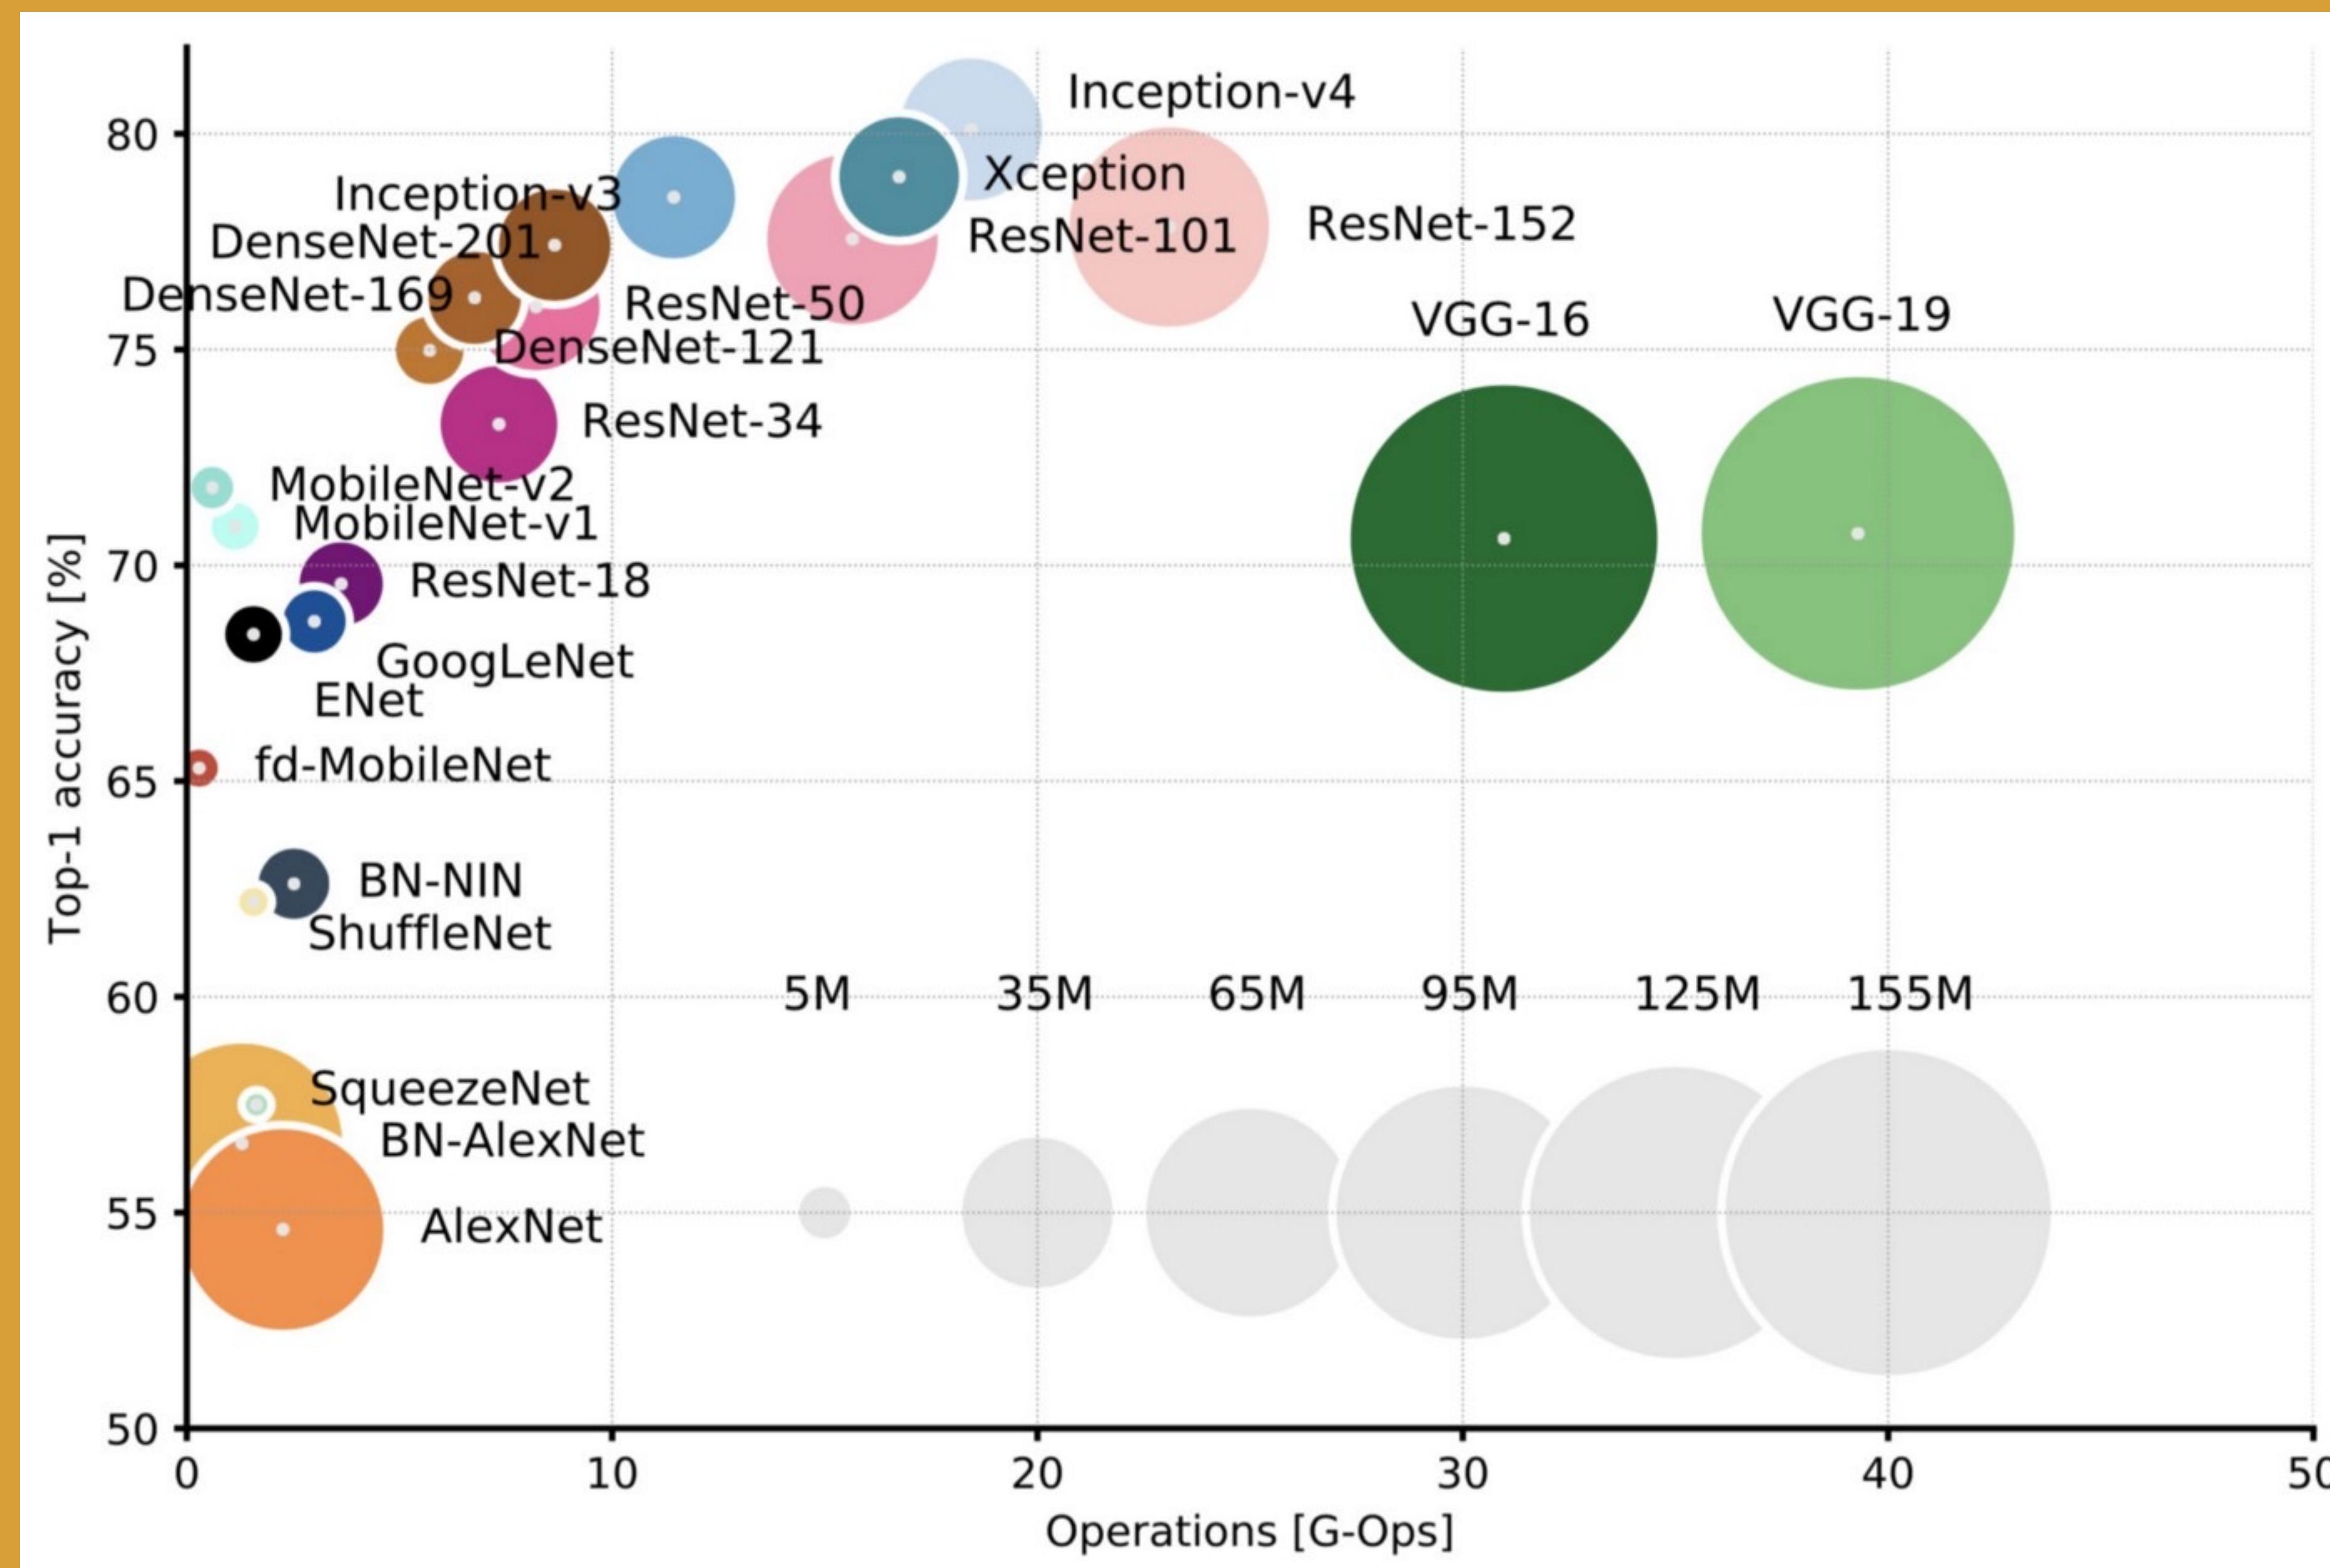


Fig. 1. Moore's Law is still valid for both CPUs and GPUs.

Sun, Y., Agostini, N.B., Dong, S. and Kaeli, D., 2019. Summarizing CPU and GPU Design Trends with Product Data. *arXiv preprint arXiv:1911.11313*.

Size of the Neural Network

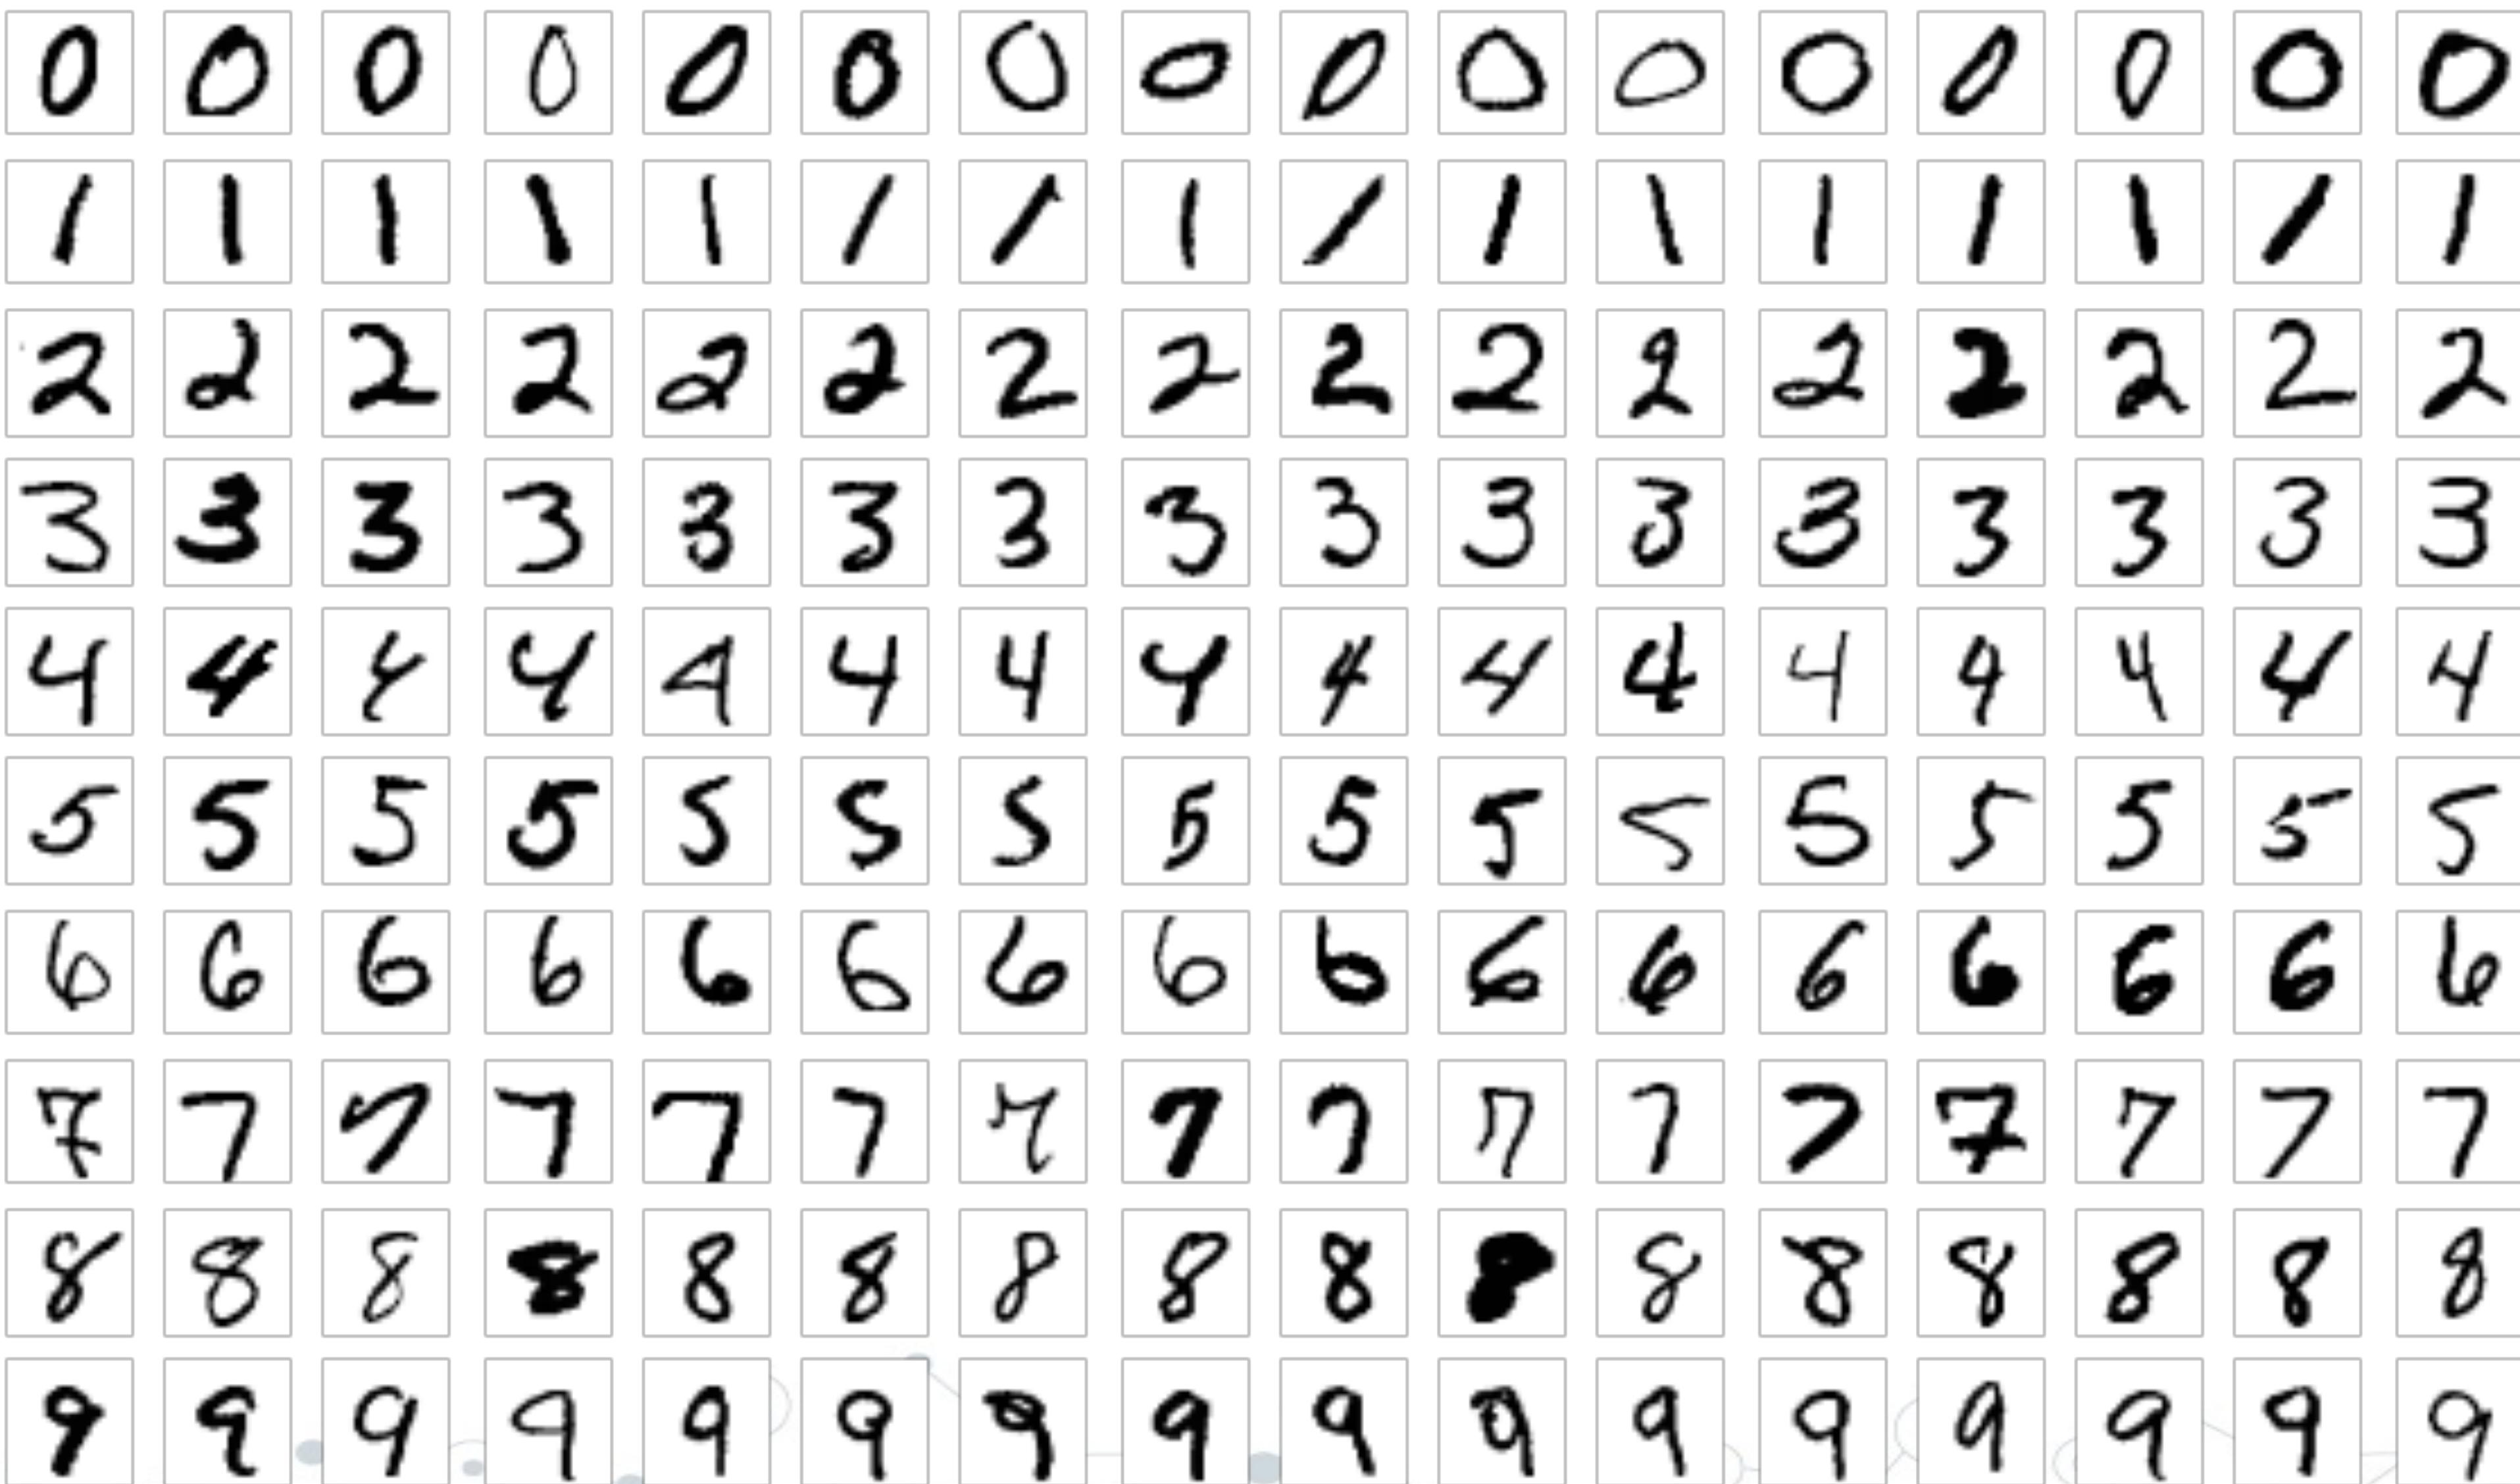


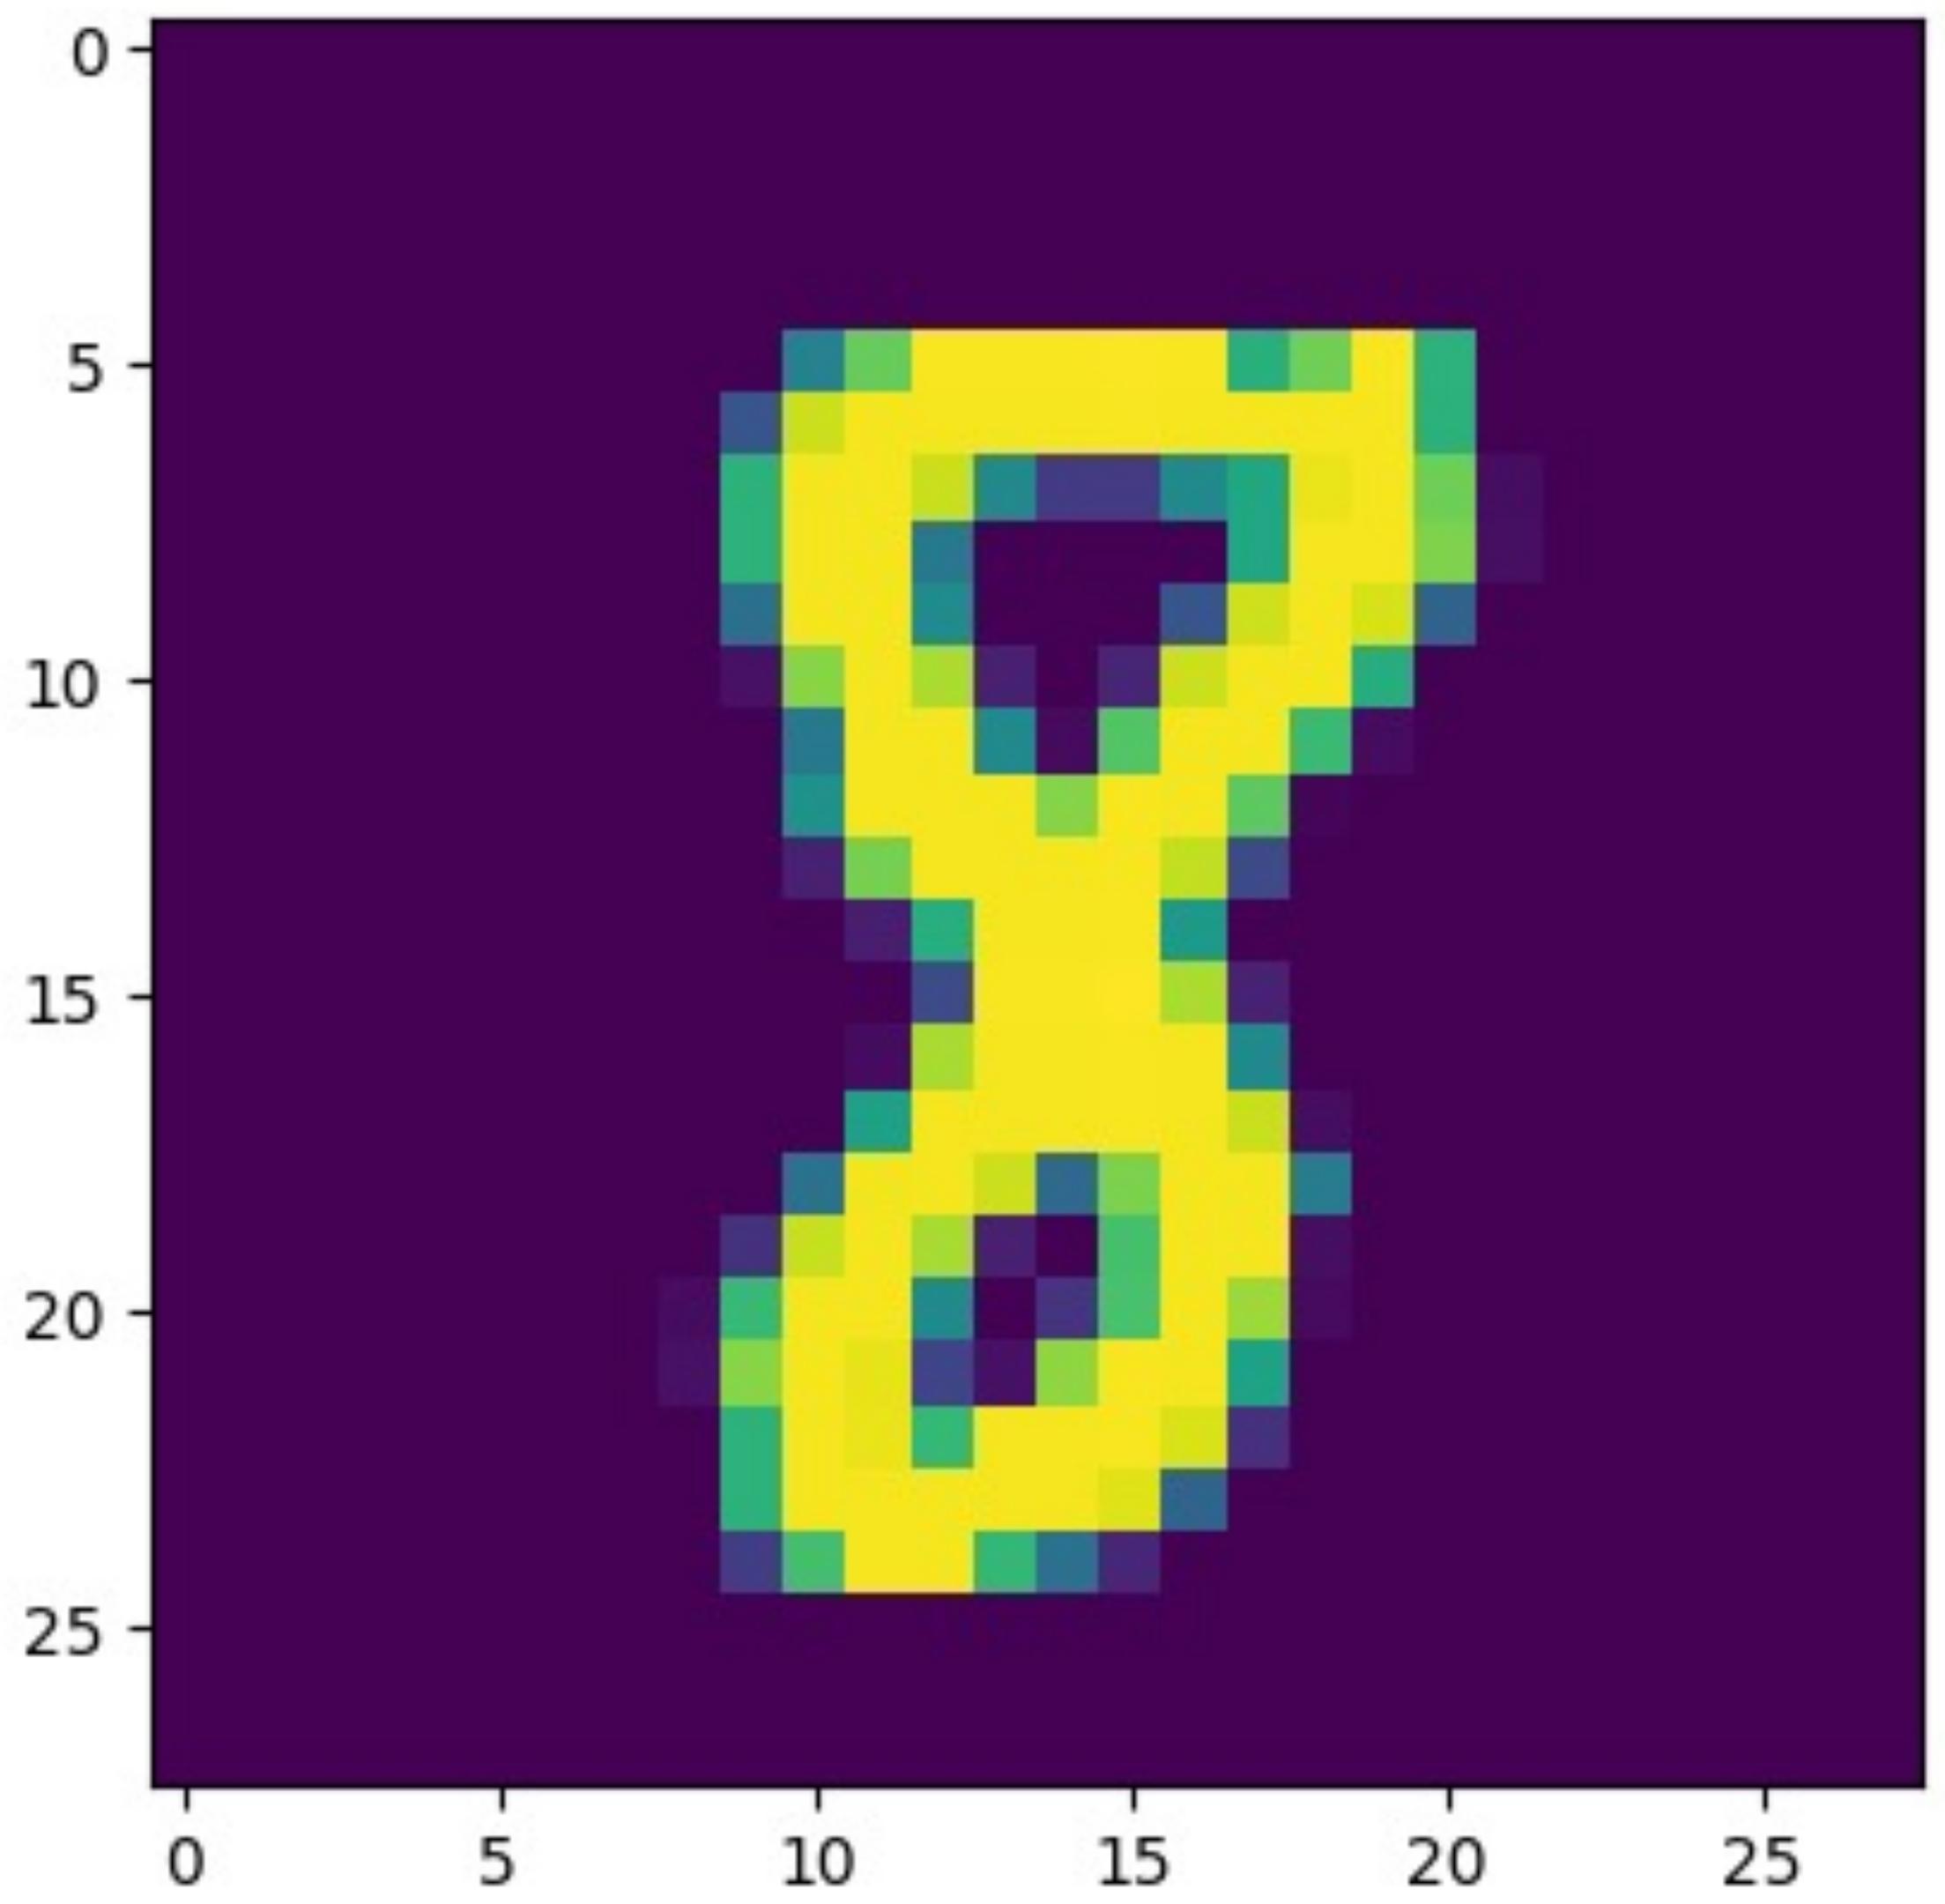
Canziani, A., Paszke, A. and Culurciello, E., 2016. An analysis of deep neural network models for practical applications. *arXiv preprint arXiv:1605.07678*.

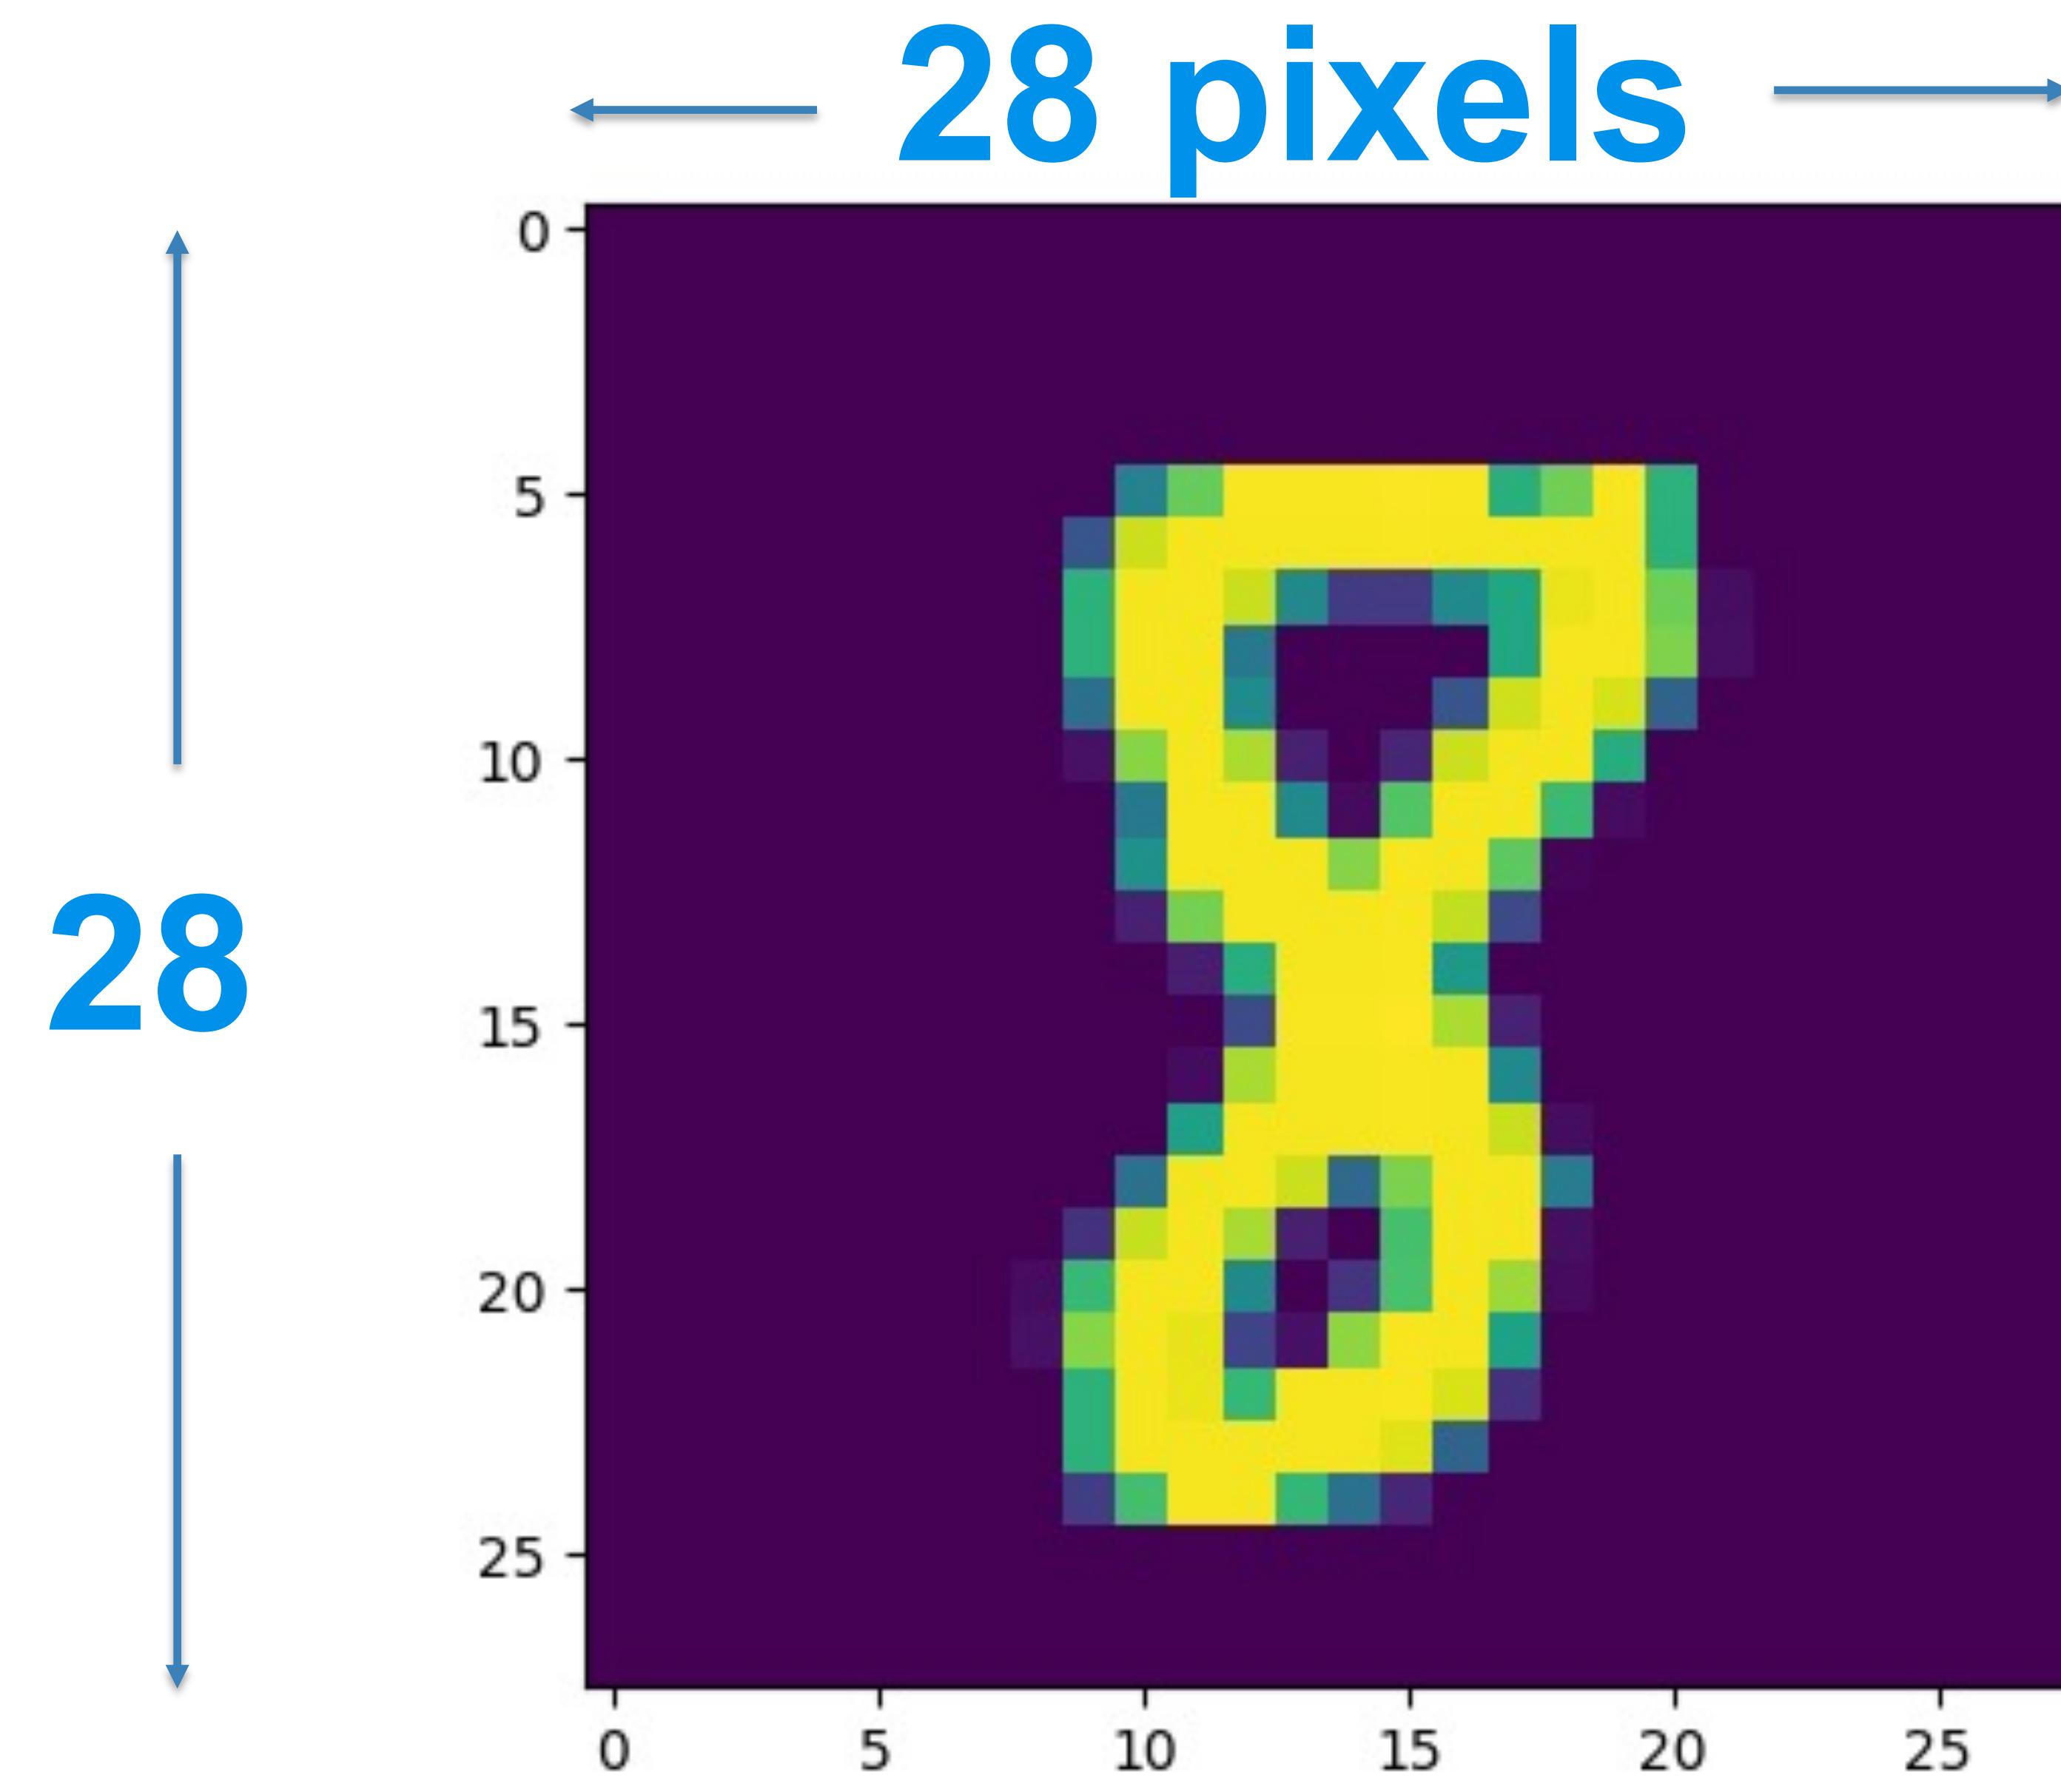
A few **tweaks** about the
neural network itself
(activation function,
weight initialisation)

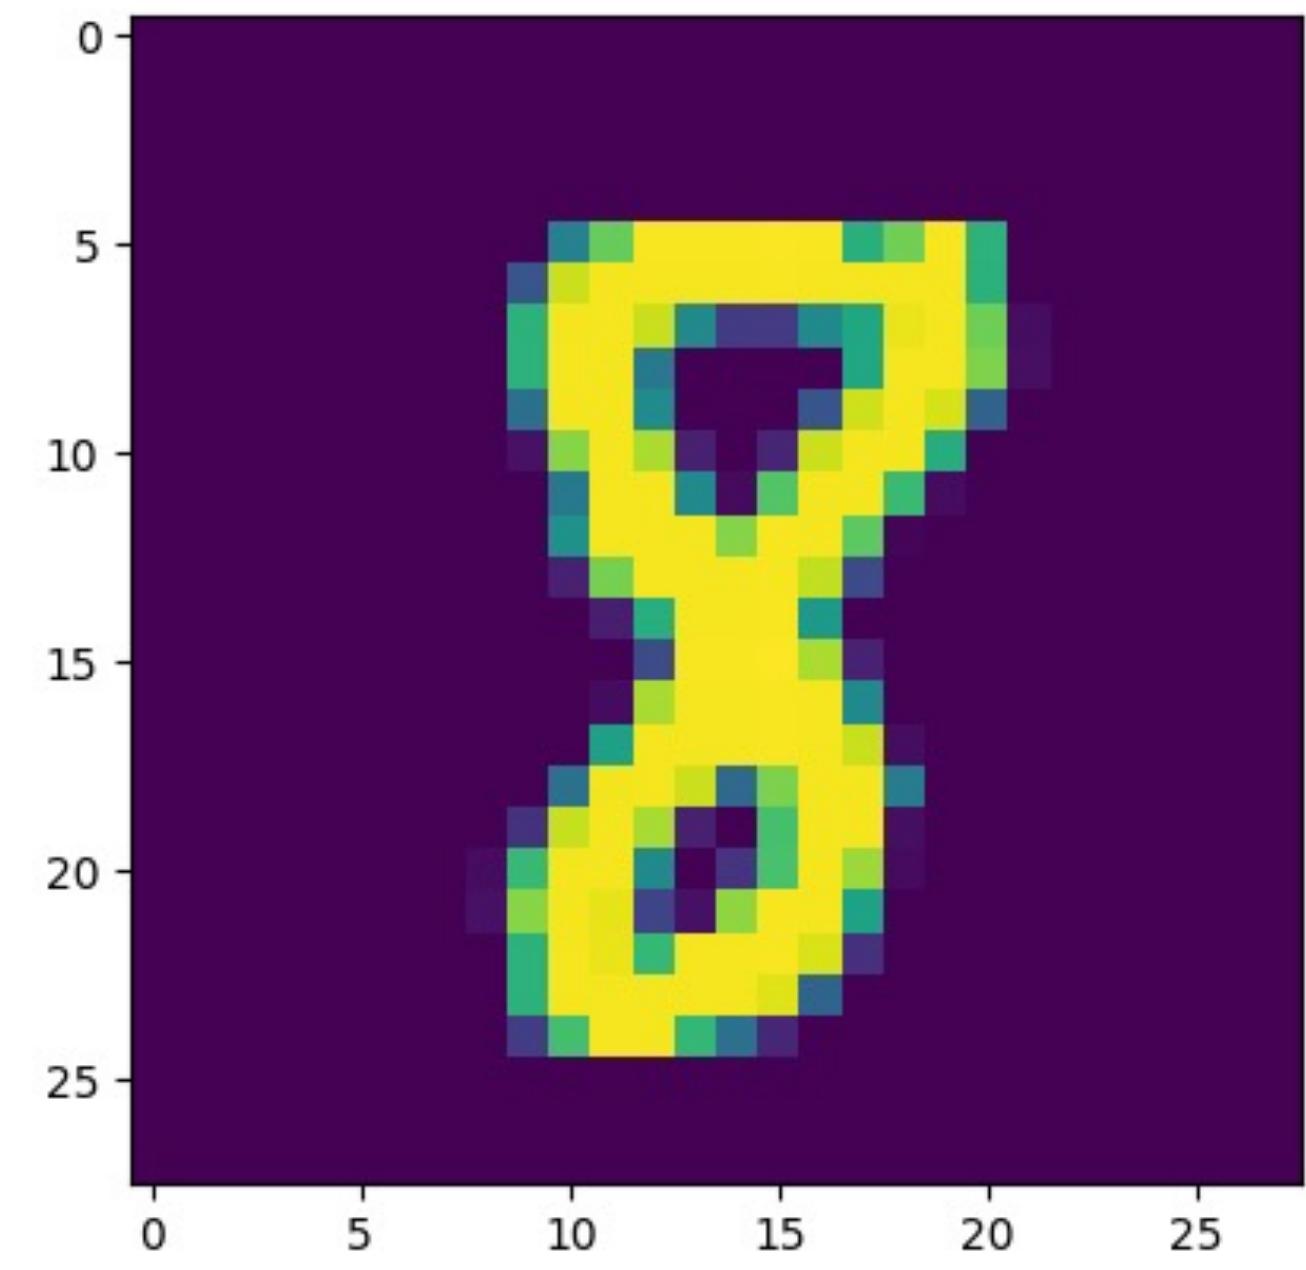
Notations

MNIST dataset - “Hello World” for image recognition

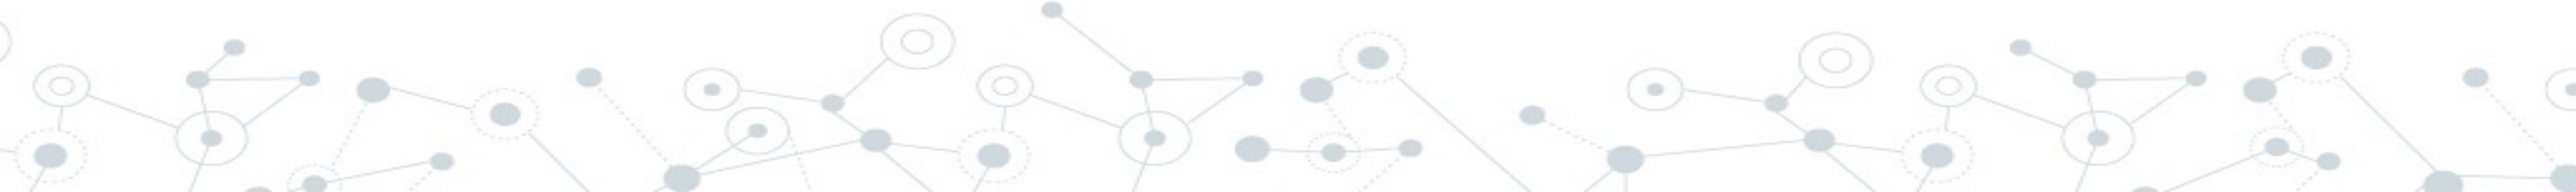
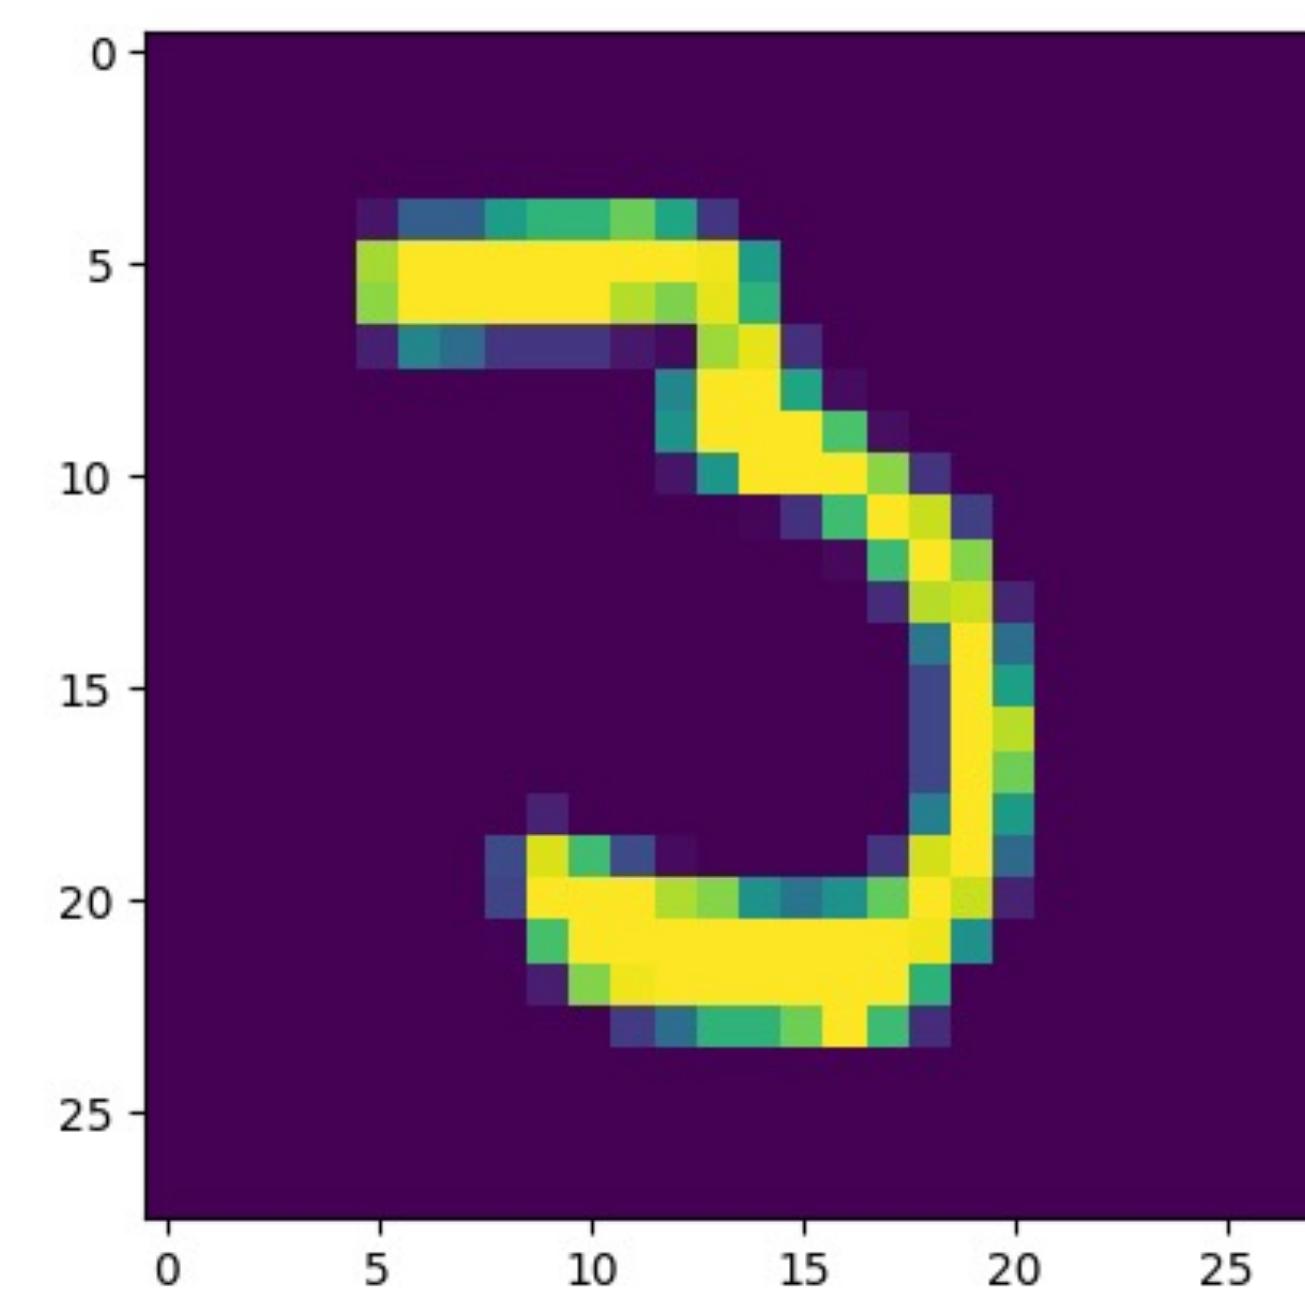








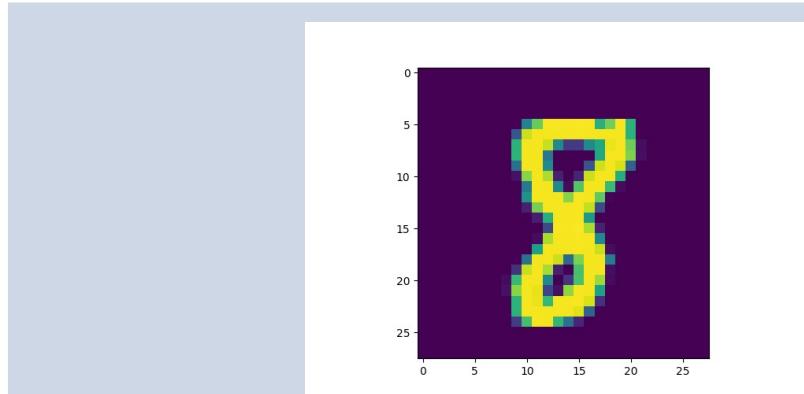
VS.



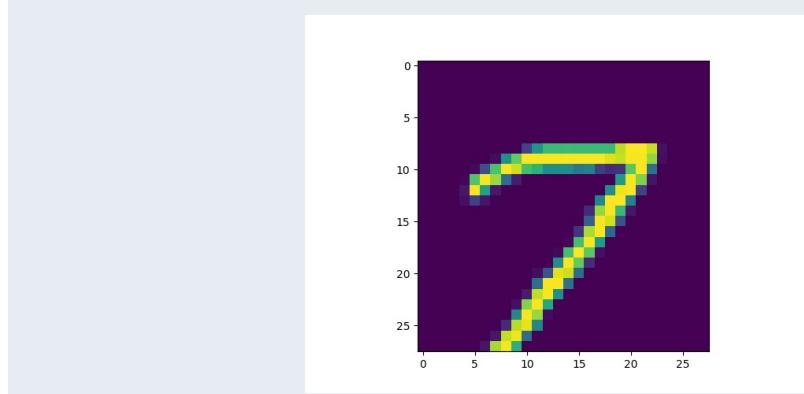


INPUT (X)

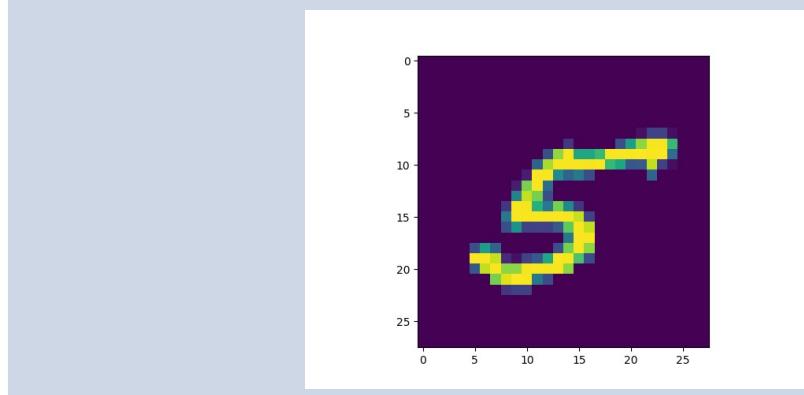
OUTPUT (Y)



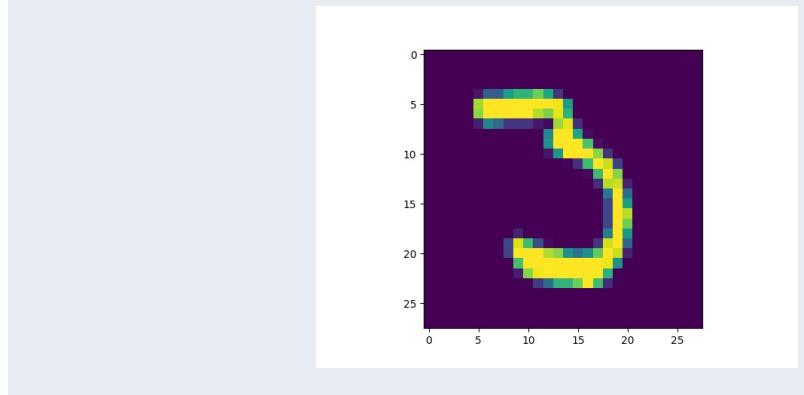
8



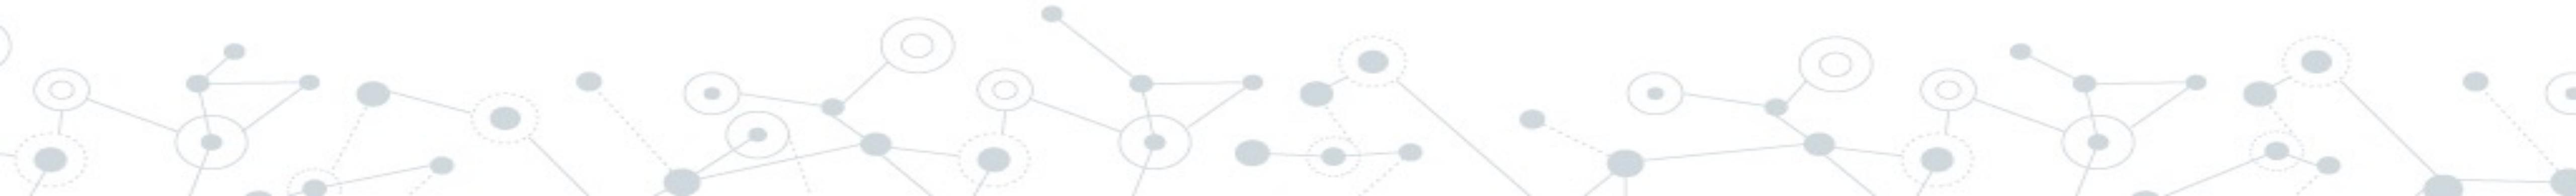
7



5



3



x - the set examples

y - the set of labels

$x^{(1)}$ - refers to the first example

x_1 - first feature from the first example

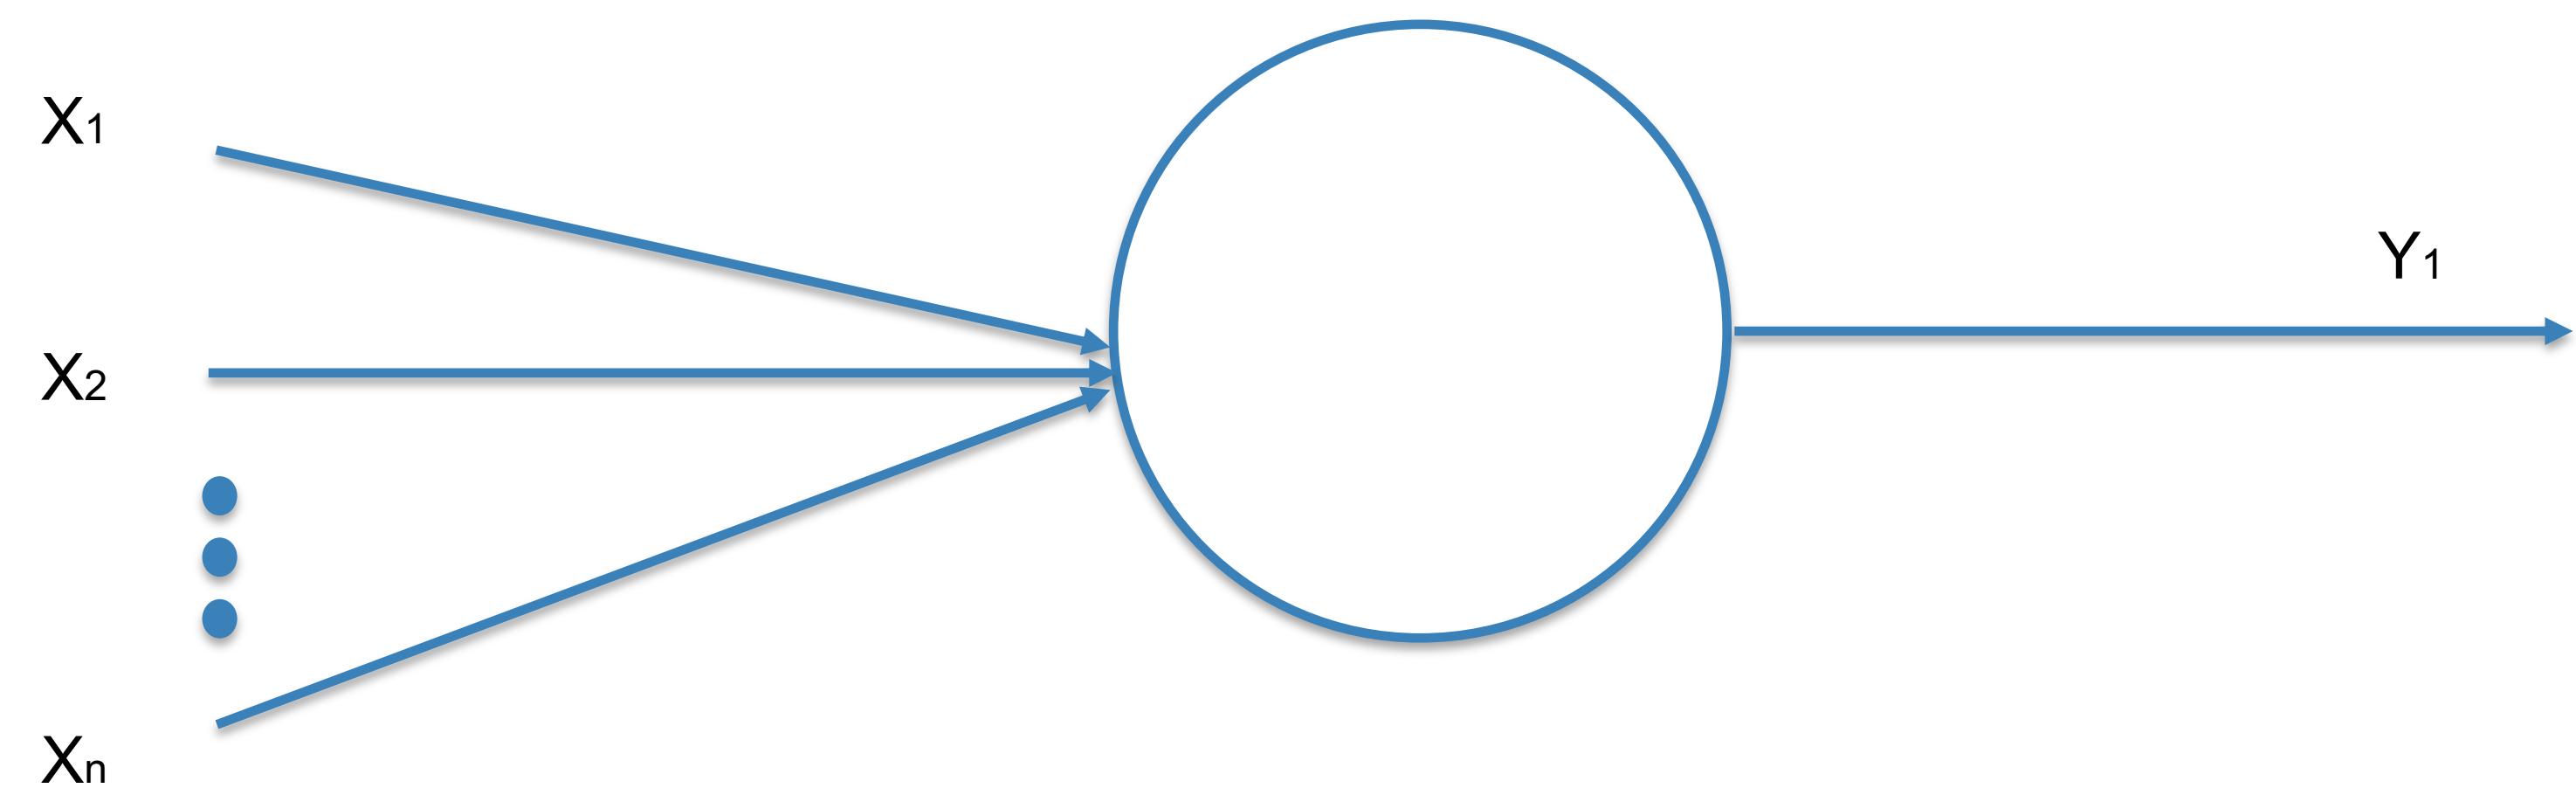
$y^{(1)}$ - label of the first example

N - number of features

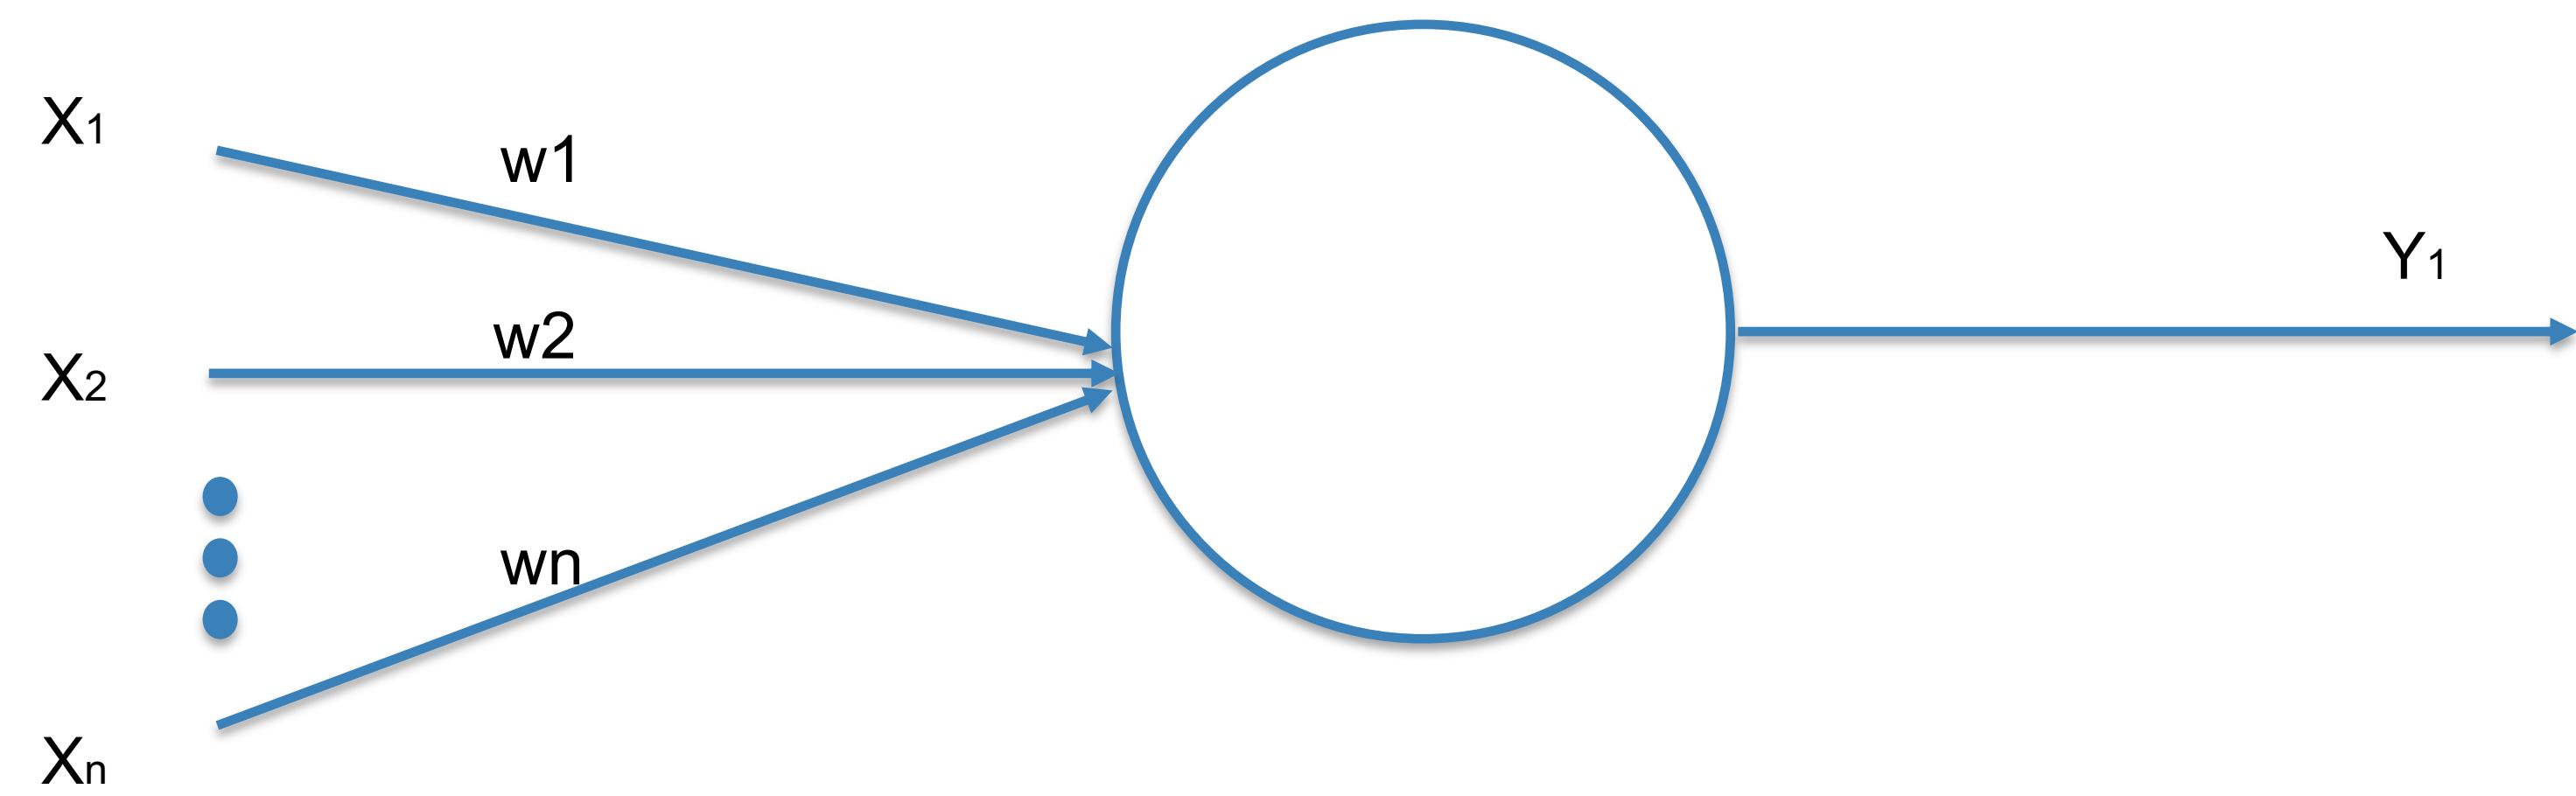
M - number of examples

Recap: Neural Networks

Neural Networks

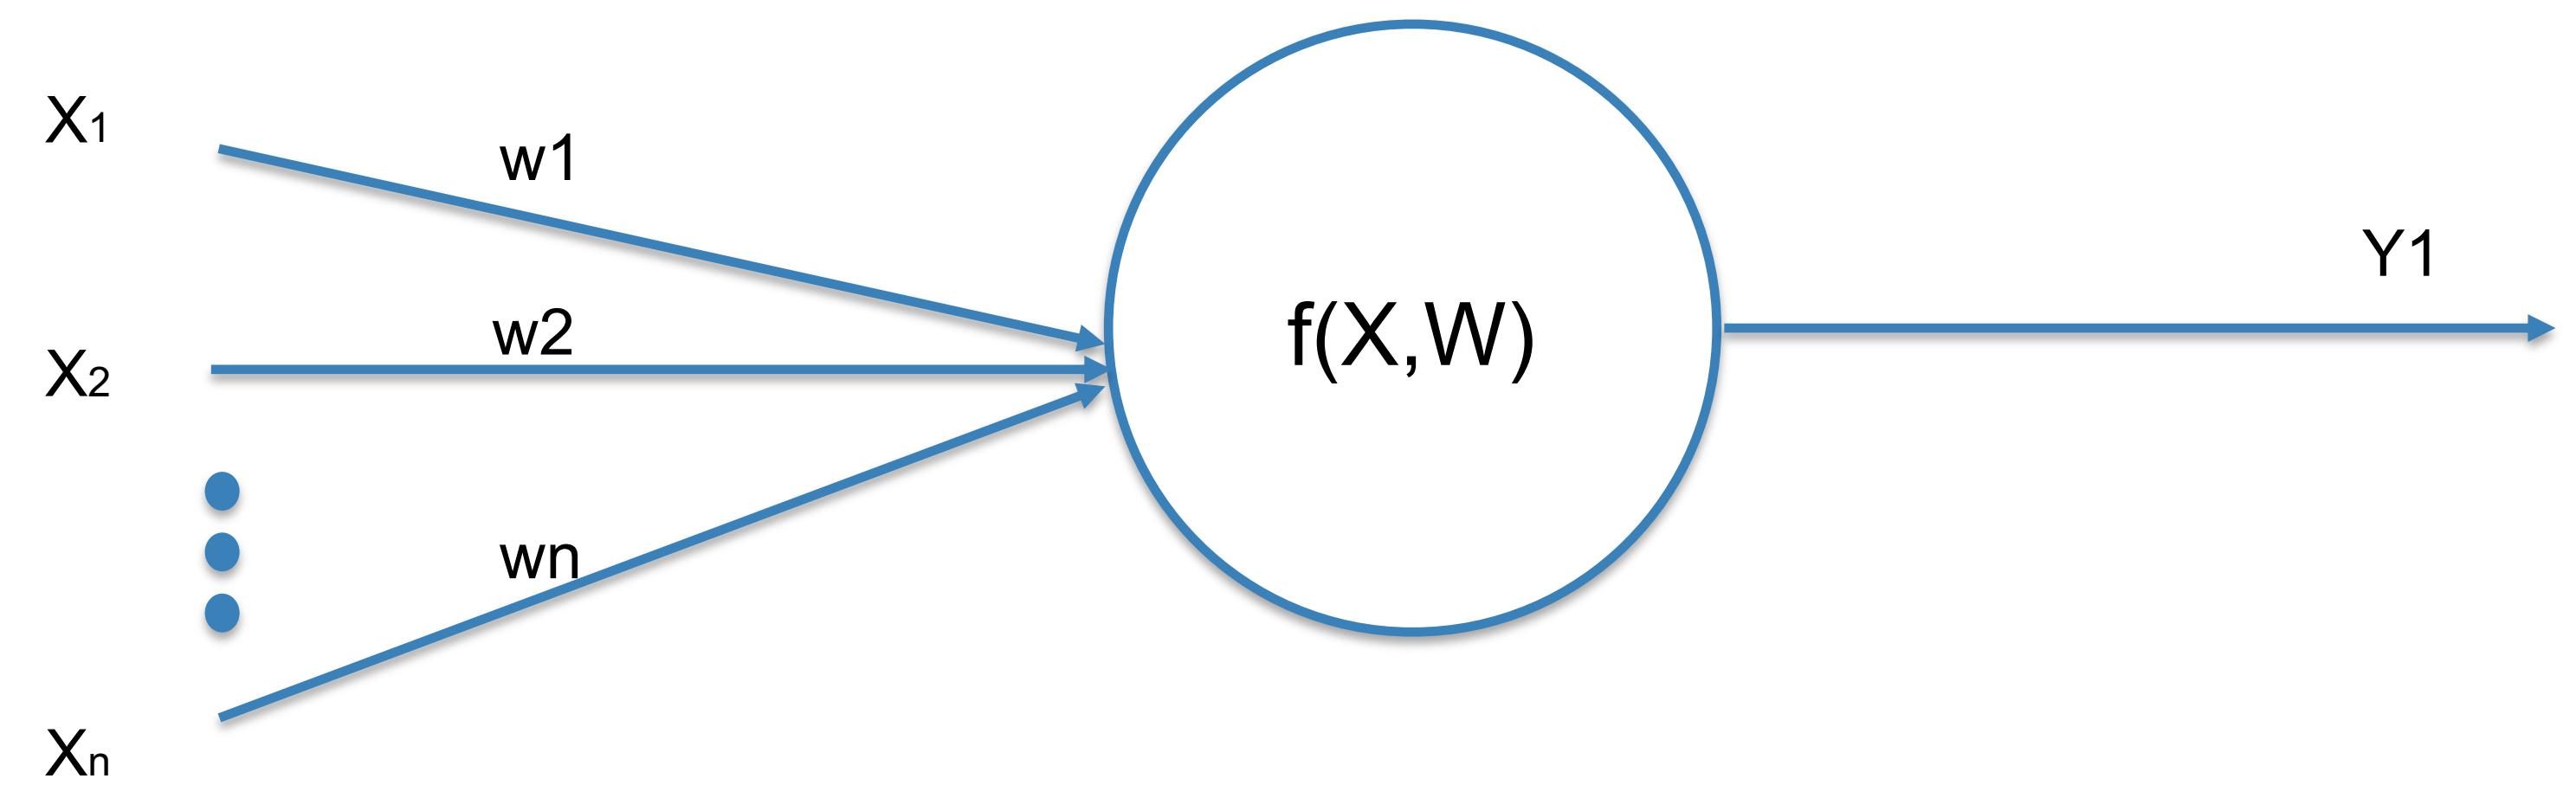


Neural Networks



The **weights** (w) give the importance of a input

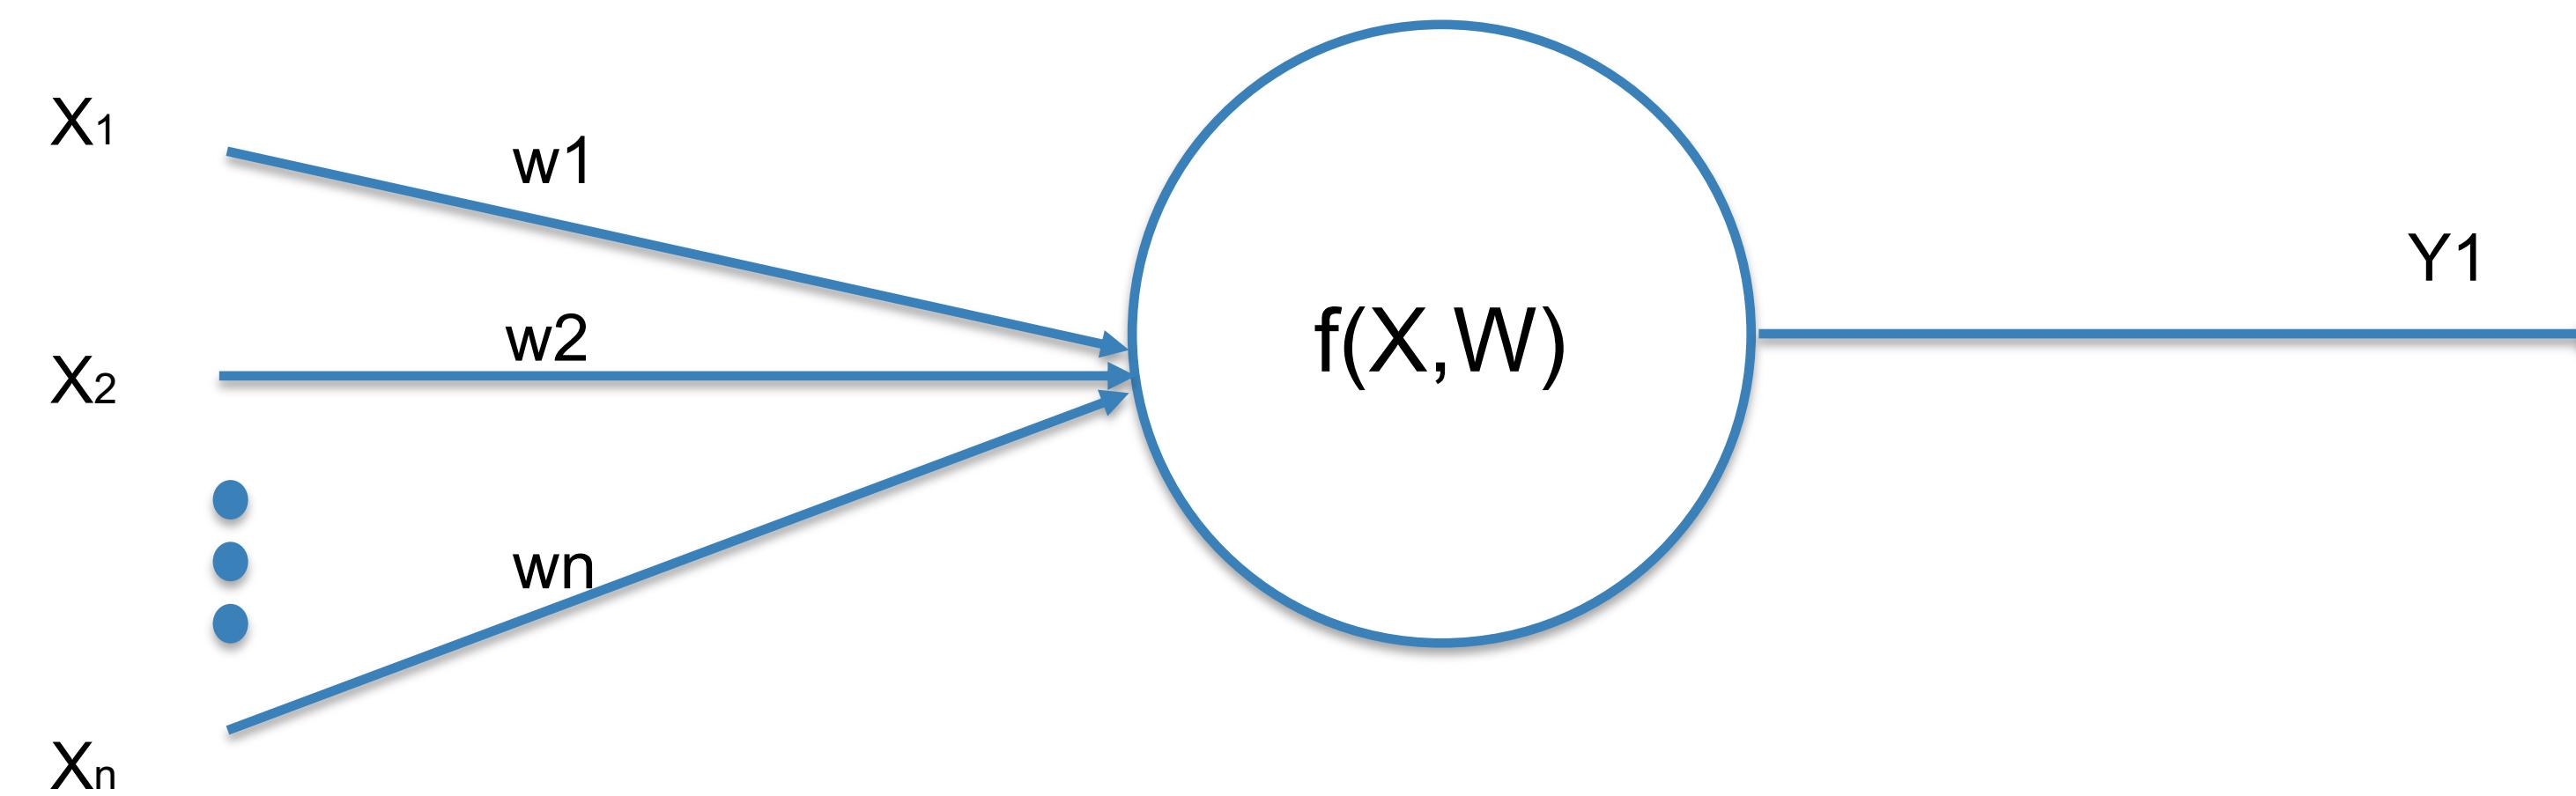
Neural Networks



A NEURON:

- Receives input (x)
- Computes output (y) based on an activation function f

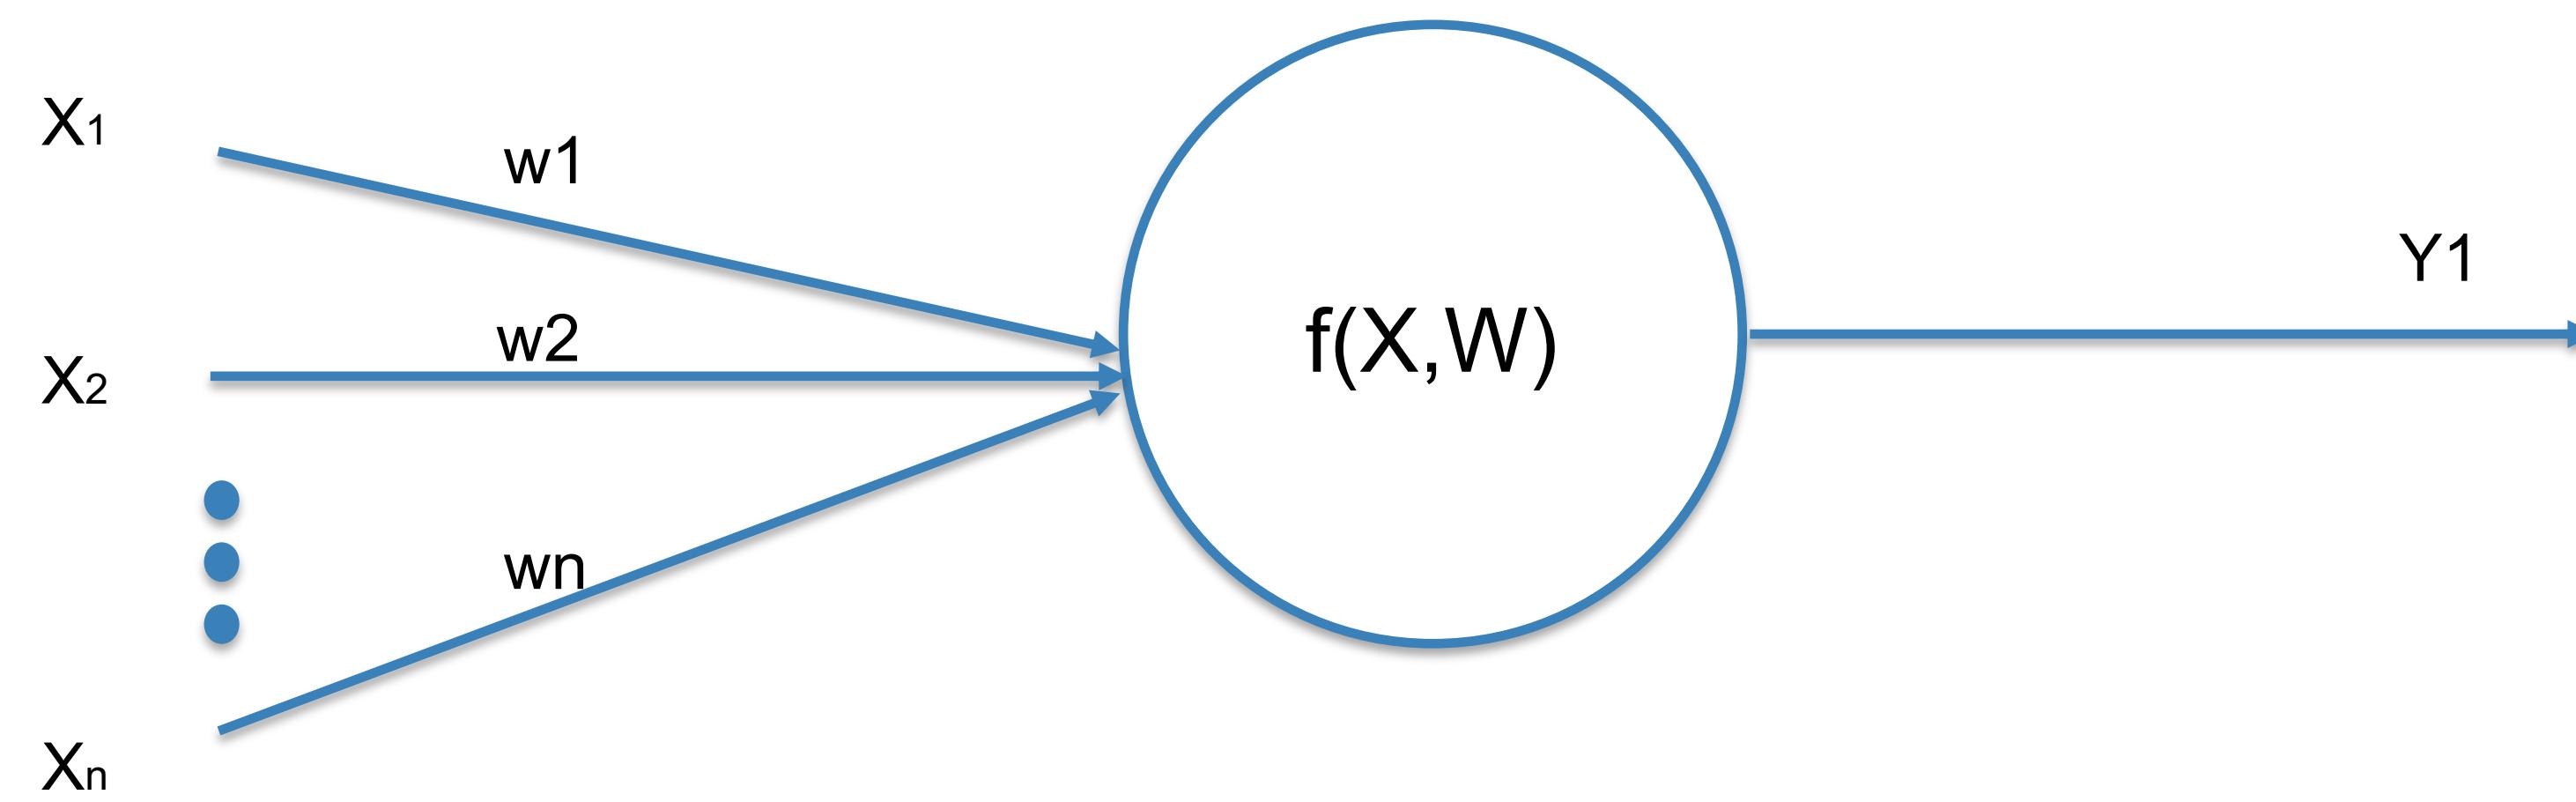
Artificial neuron



A NEURON:

- Receives input (x)
- Computes output (y) based on an activation function f

Neural Networks



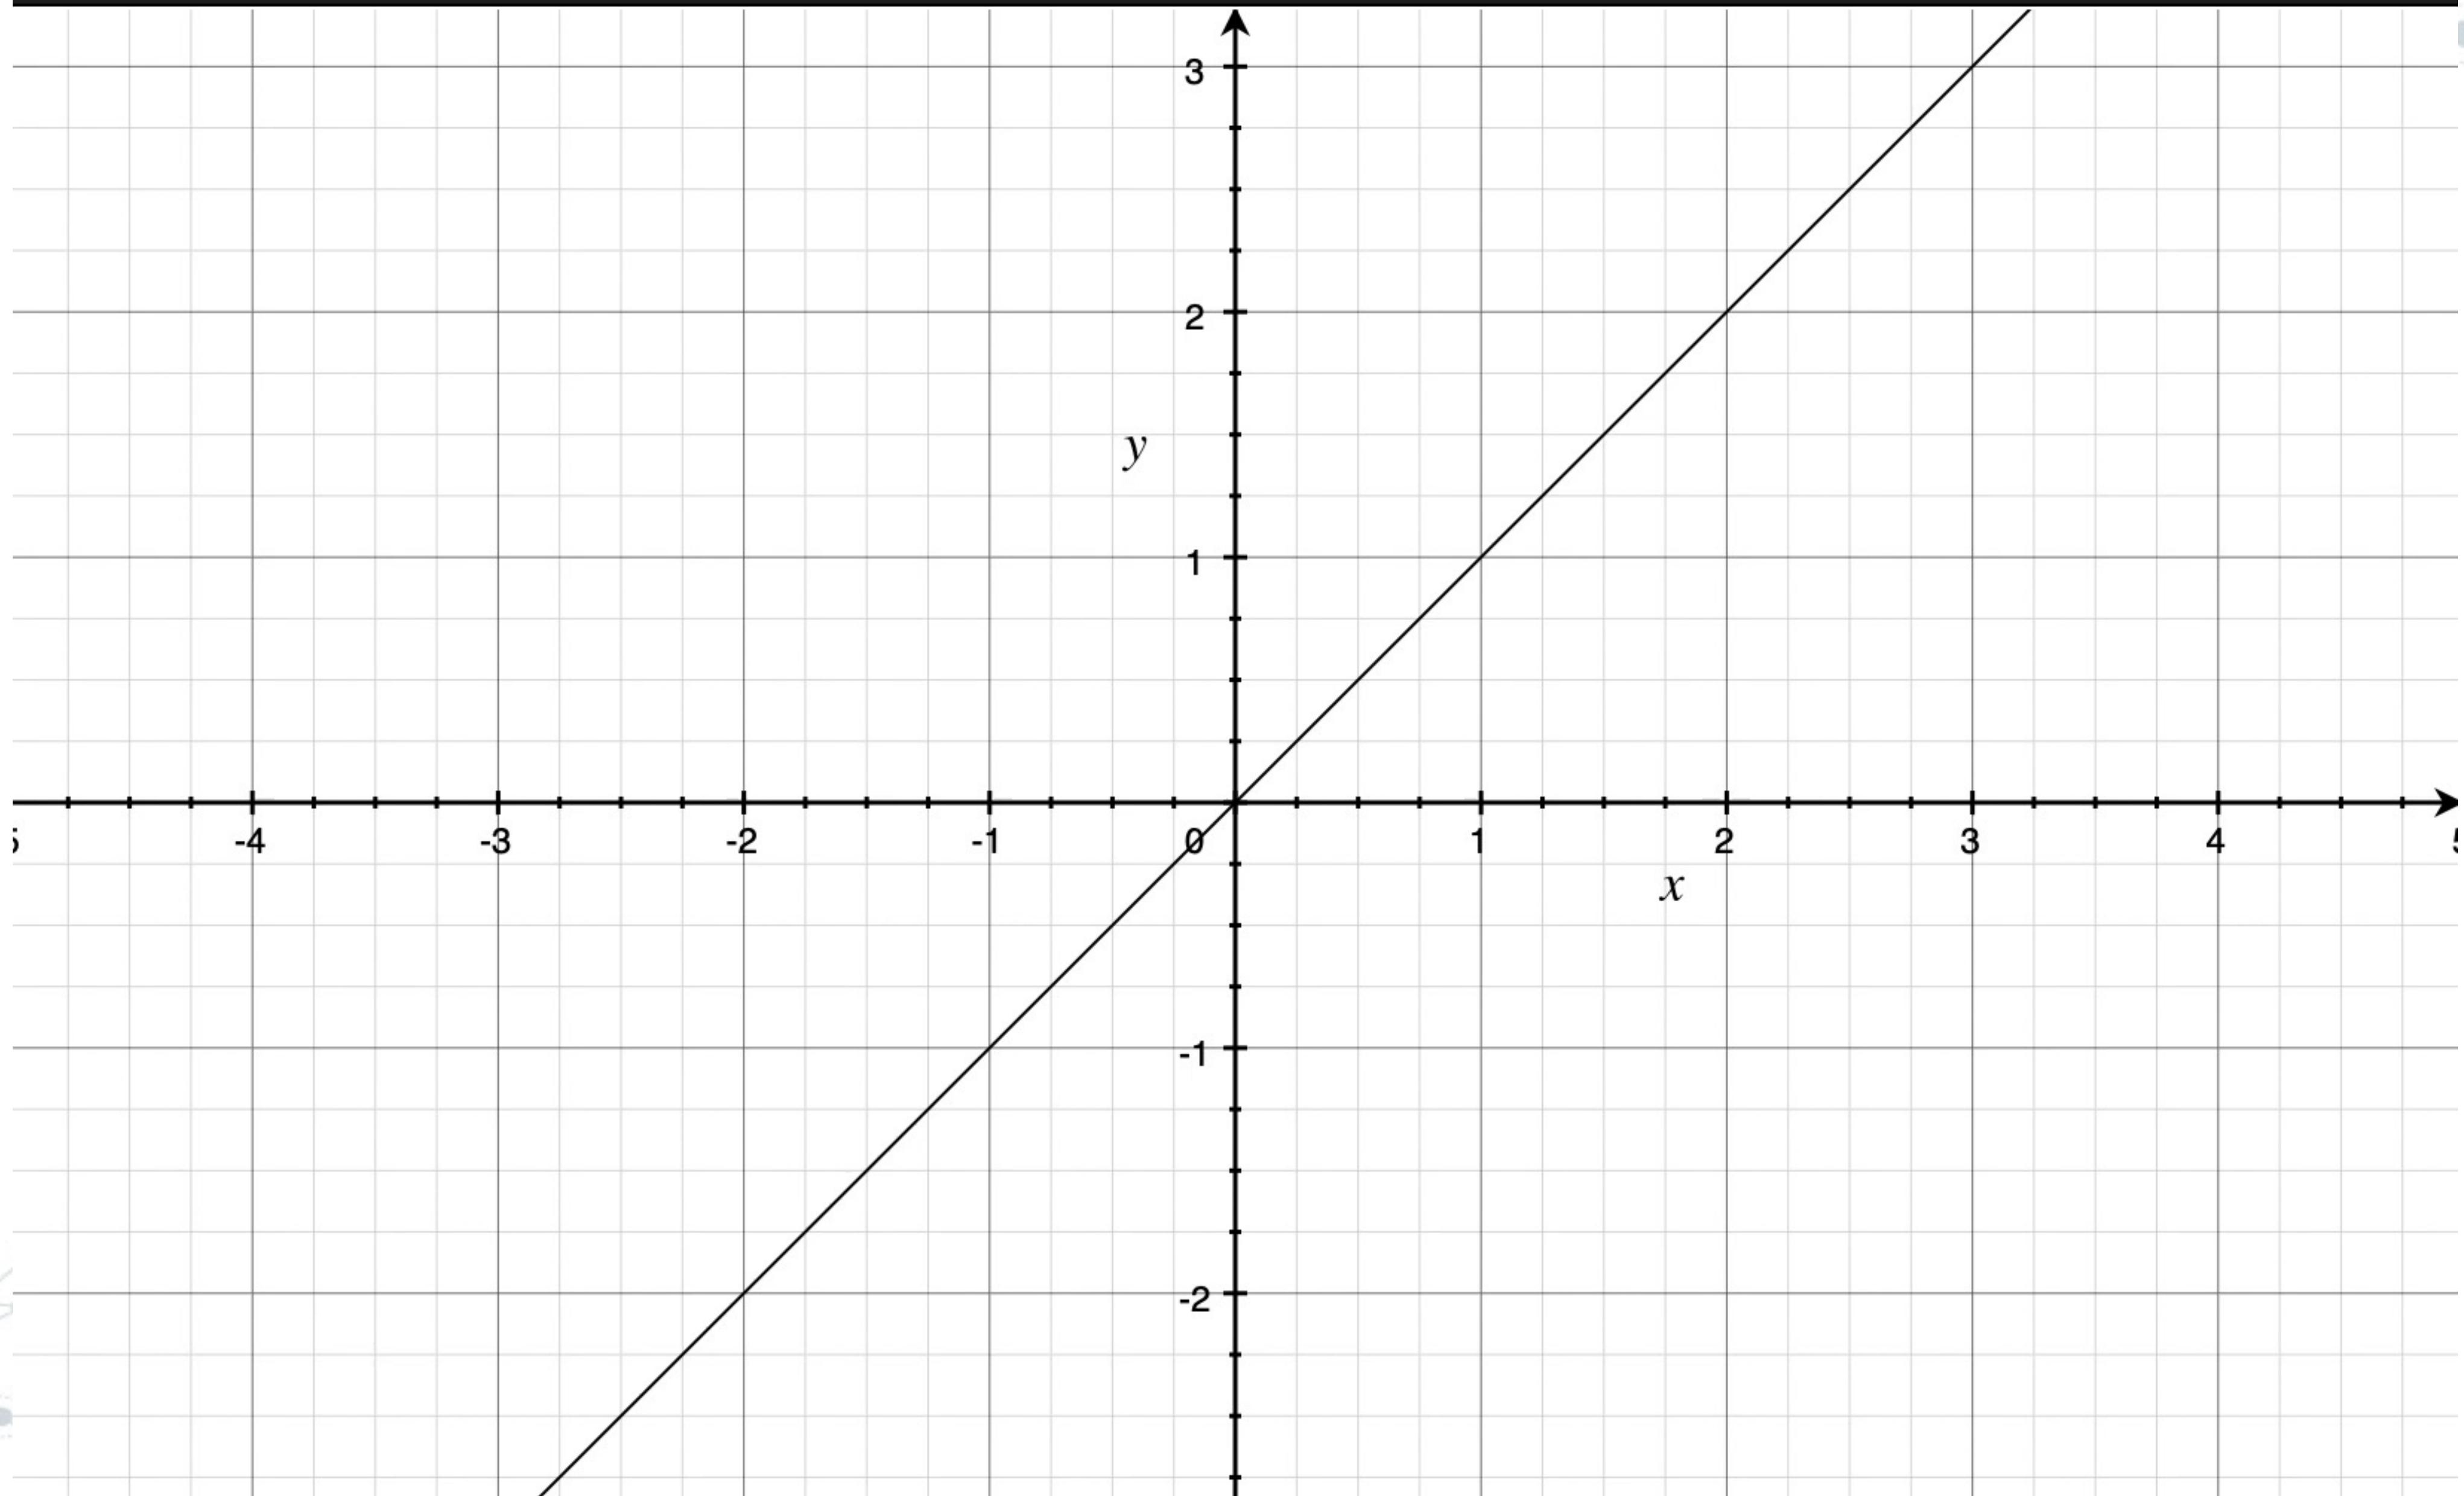
Activation function:

$$f(w_1x_1 + w_2x_2 + w_3x_3 + \dots + w_nx_n + b)$$

b is a bias term. Think of it as the it is somehow similar to the constant b of a linear function ($y = ax+b$)

$y =$

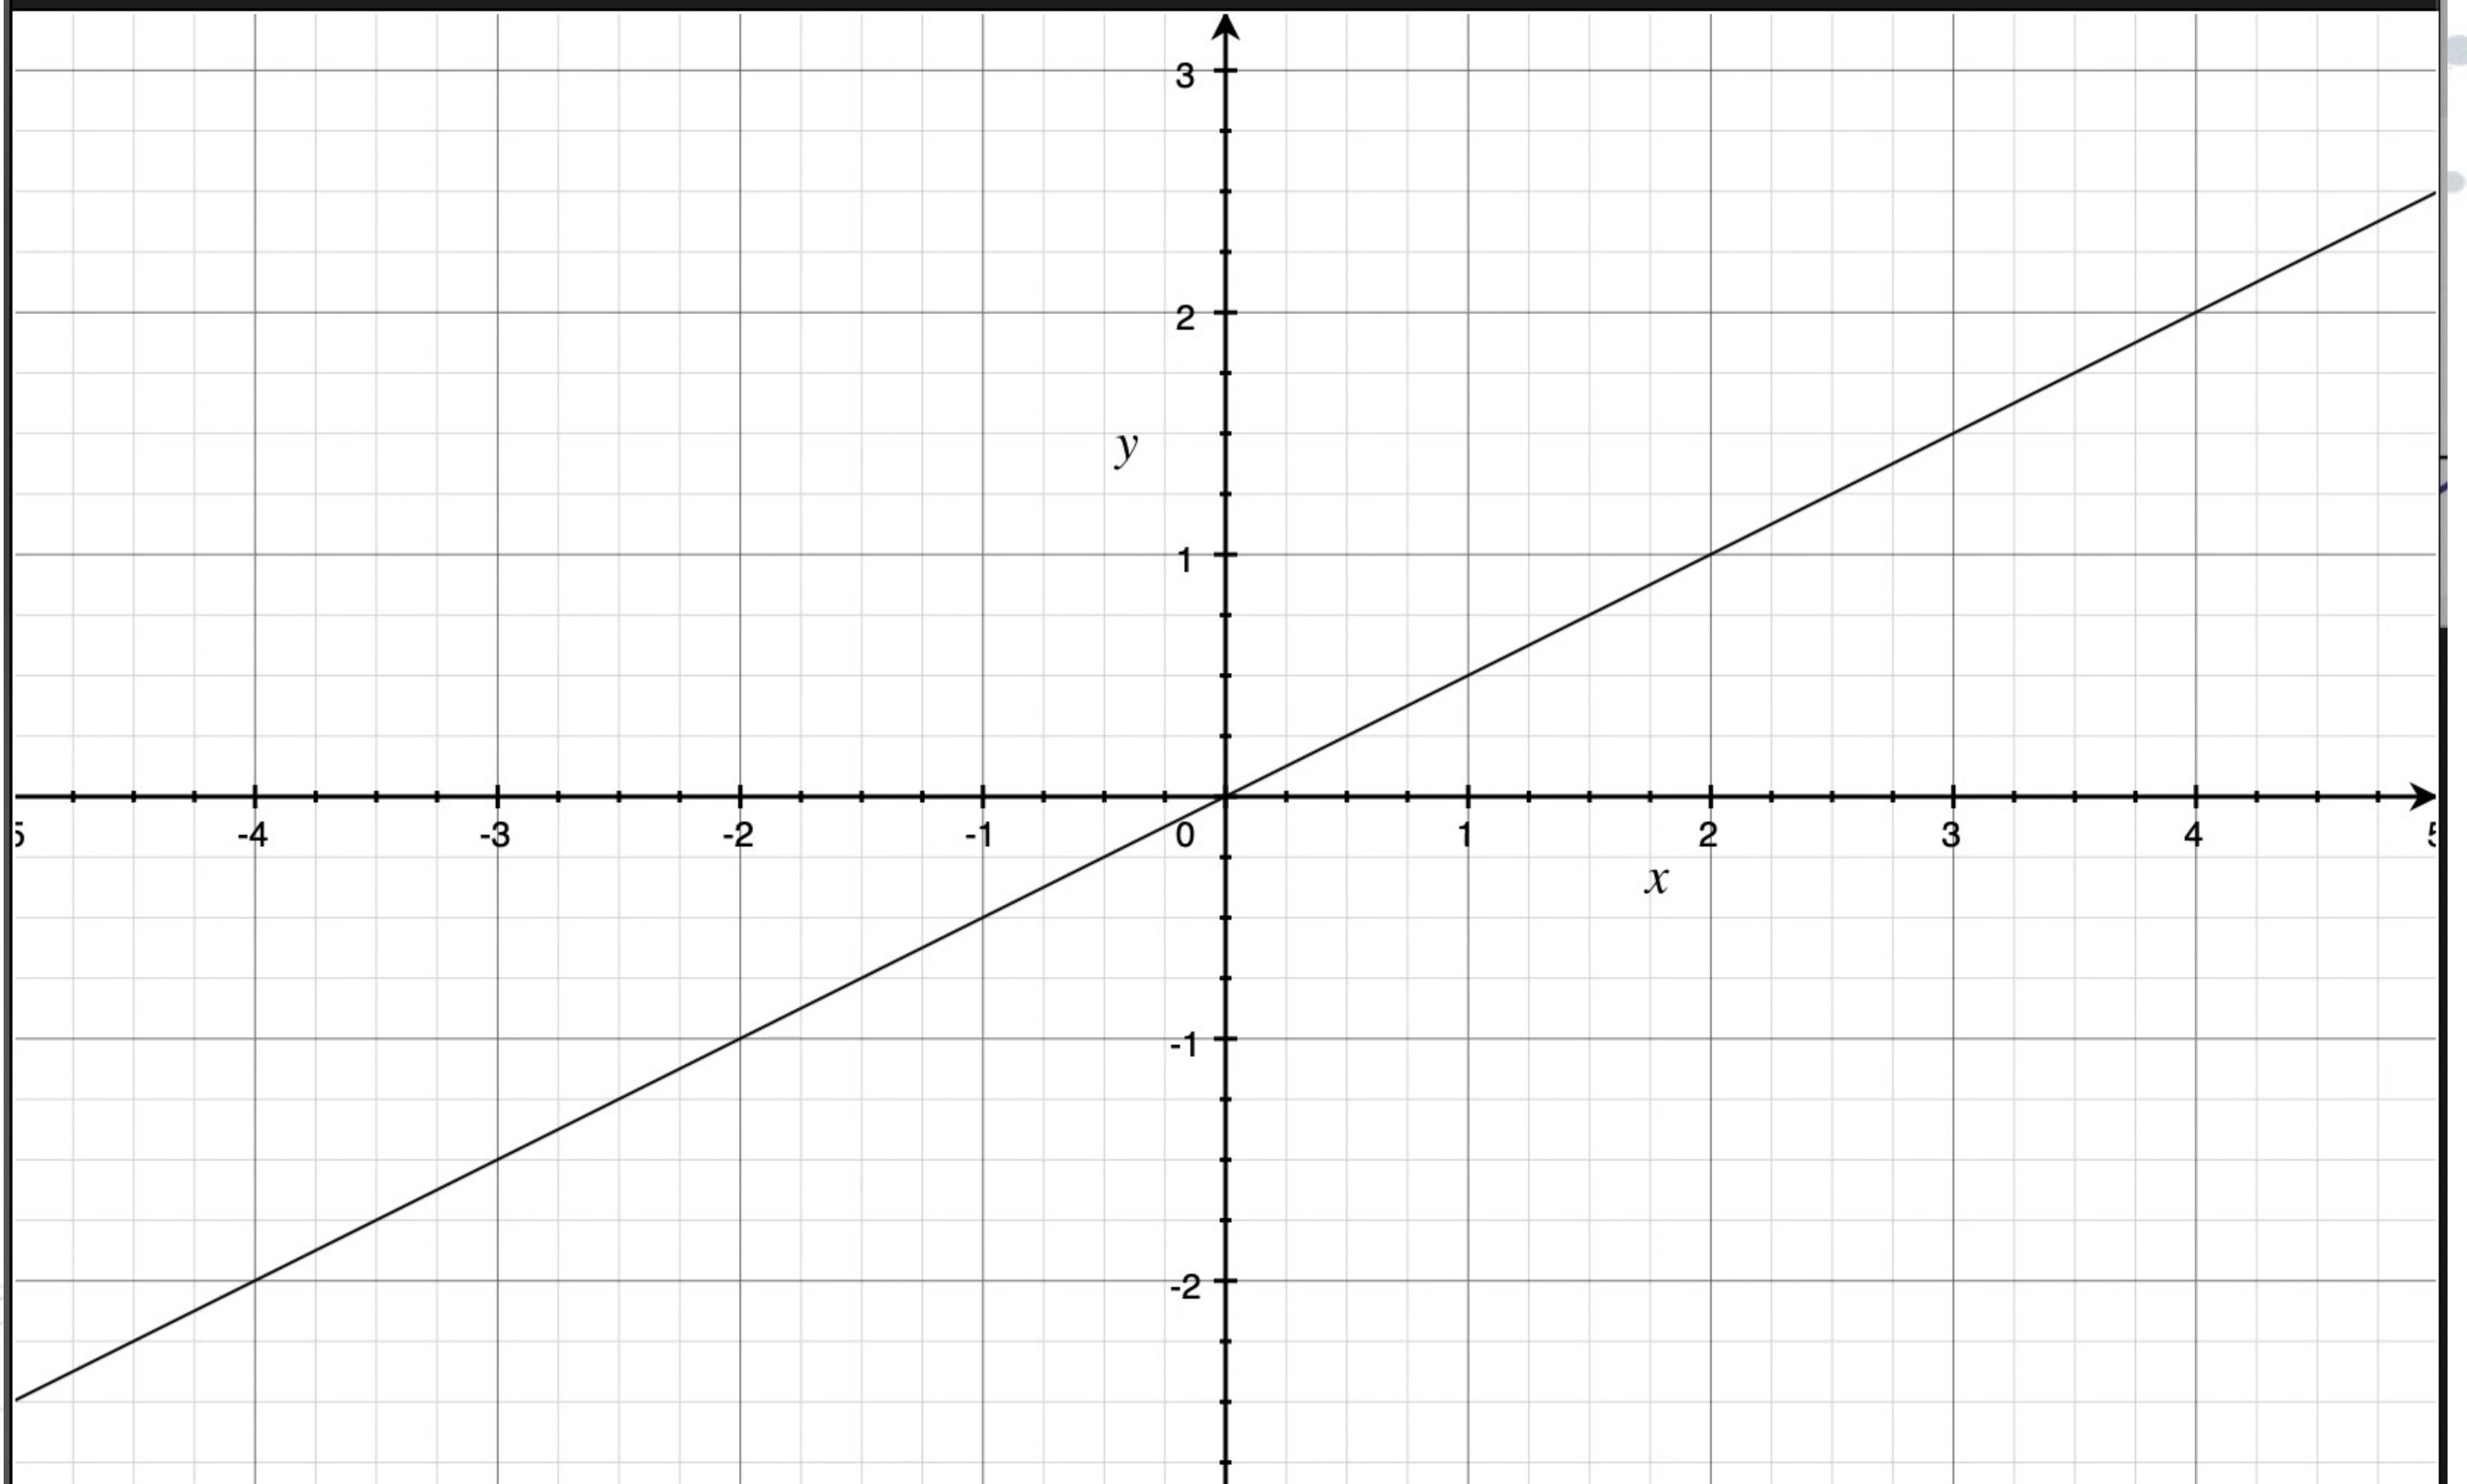
Σx^2



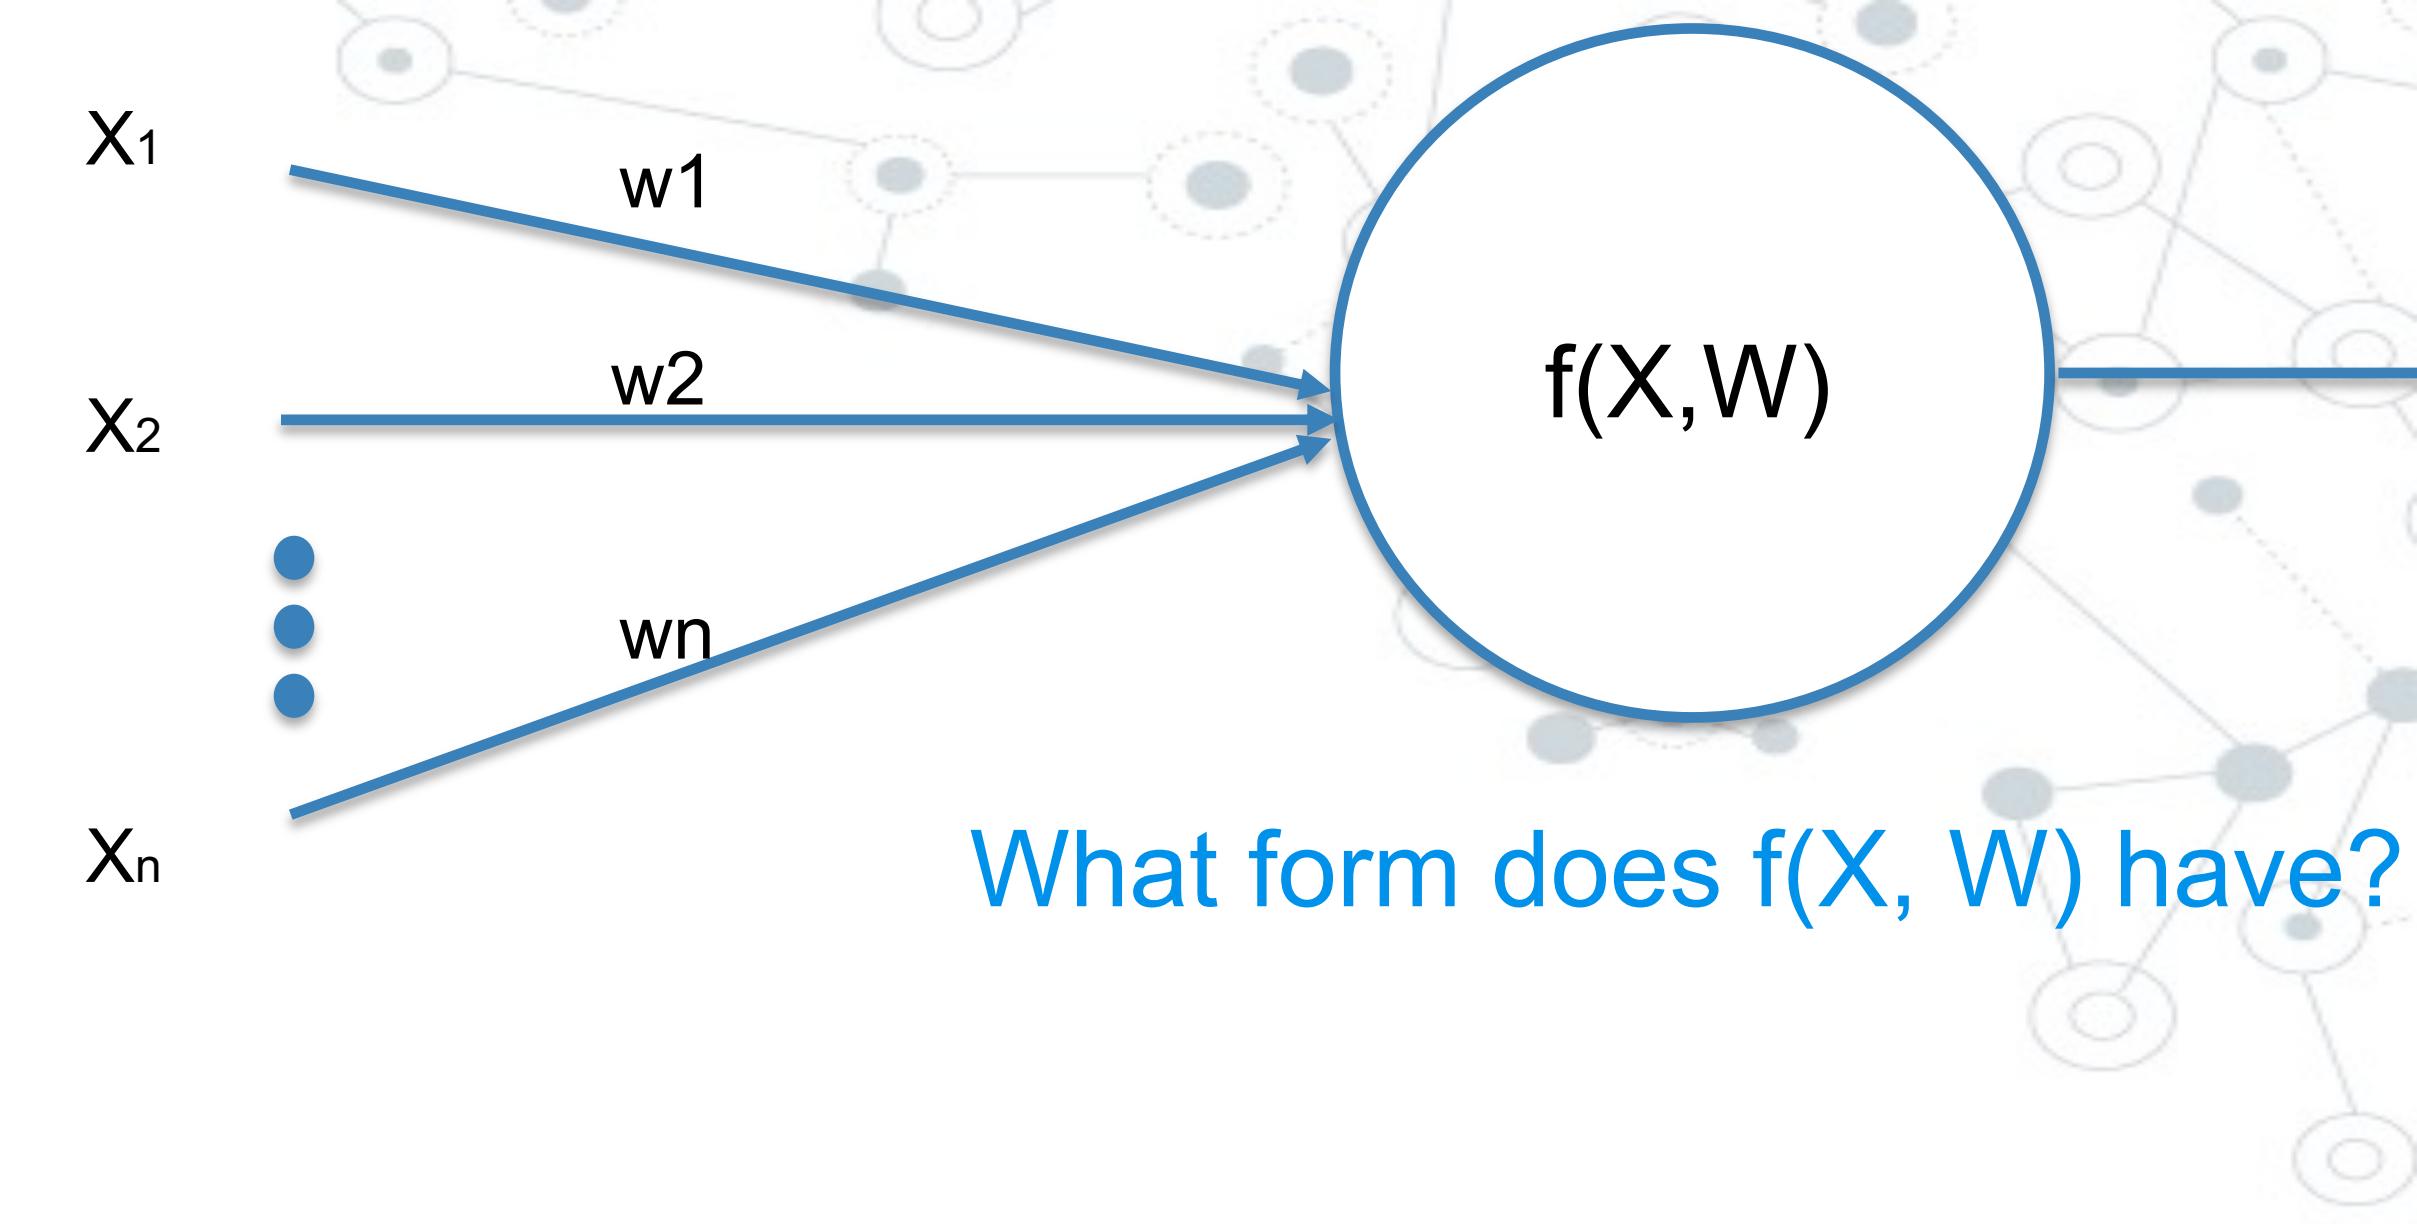
$y=0.5x$

I

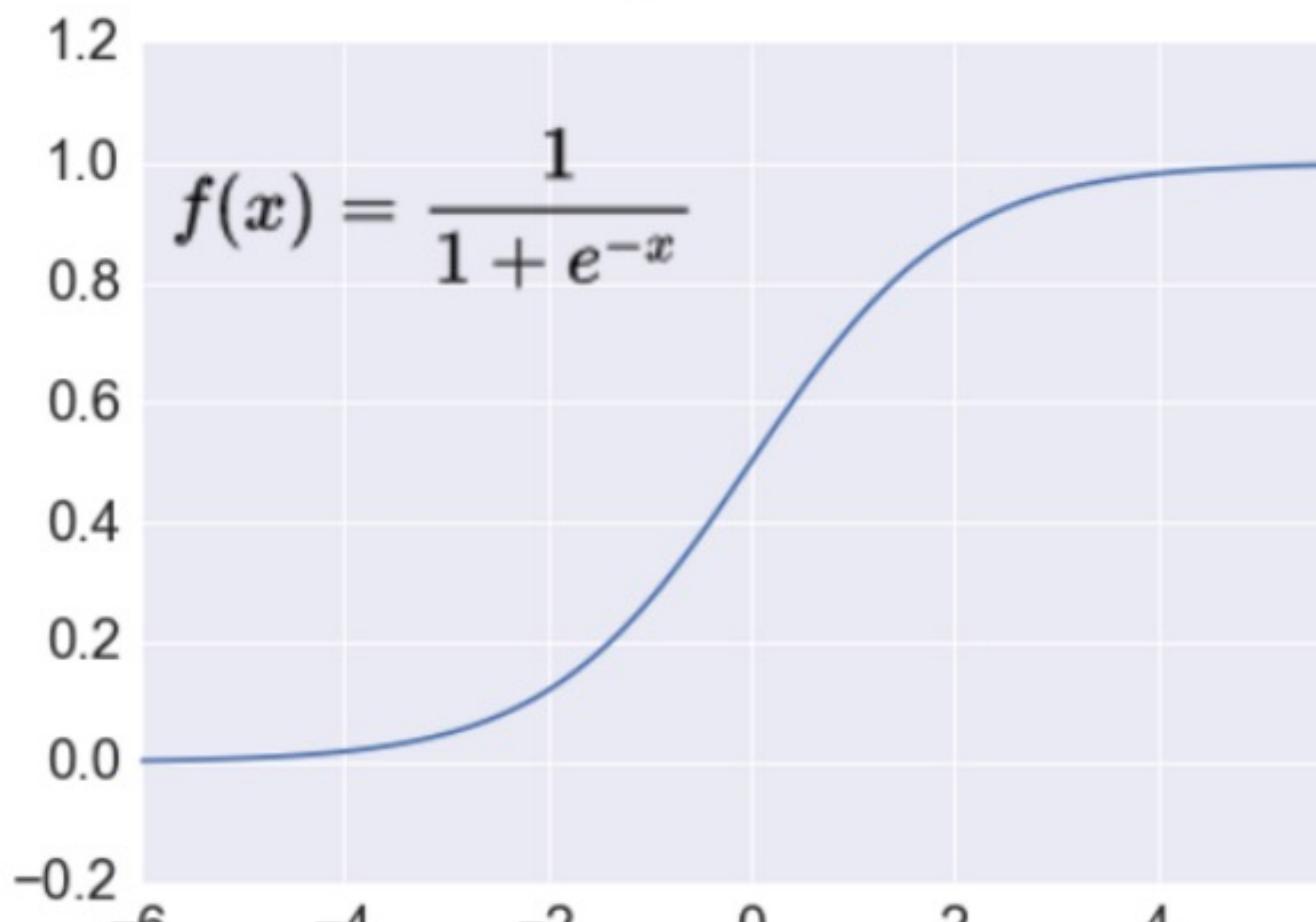
Σx^2



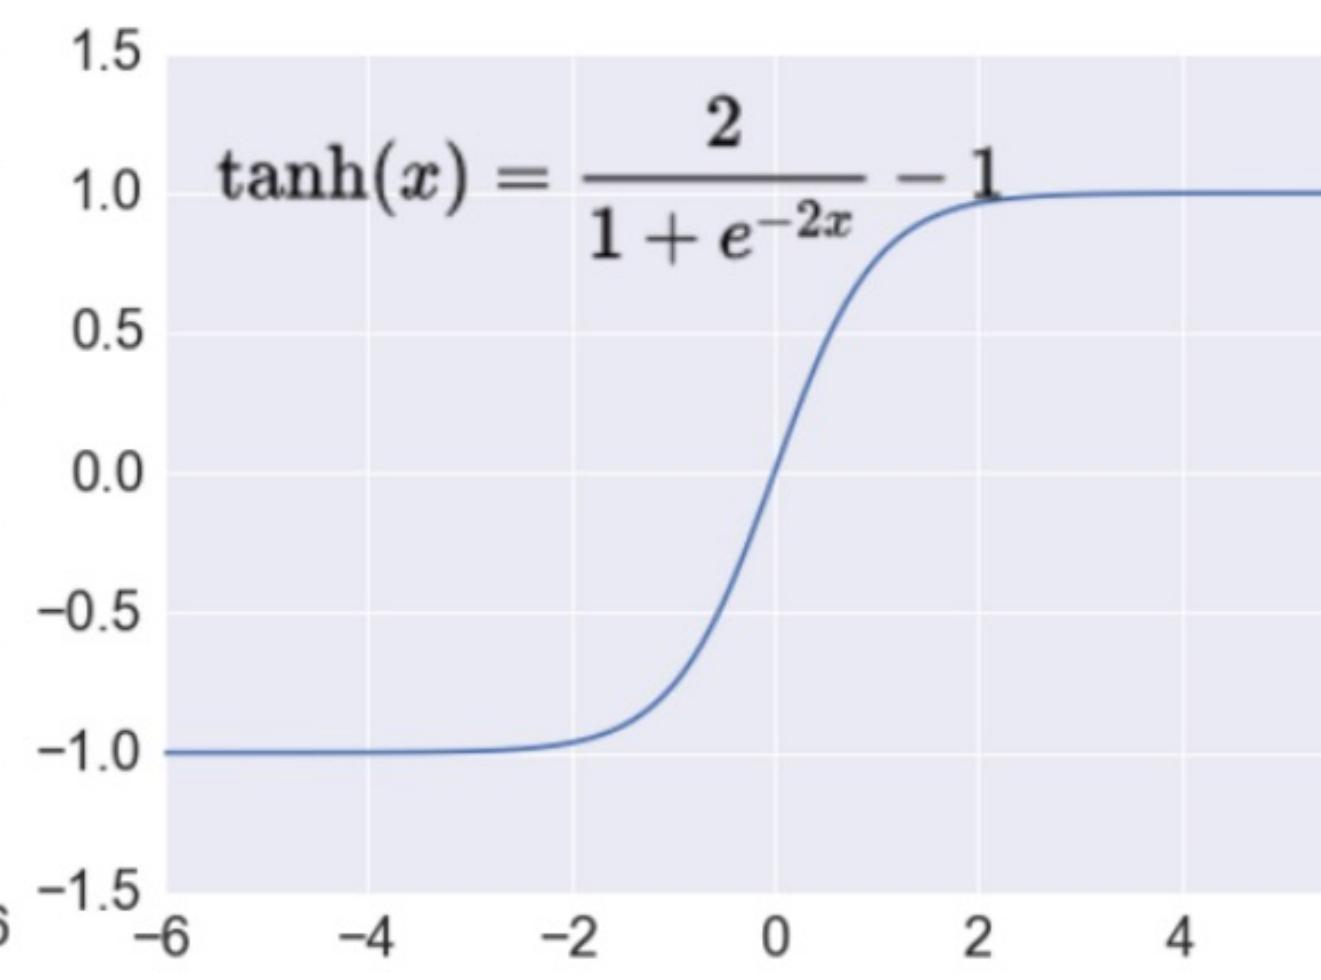
Neural Network - Activation functions



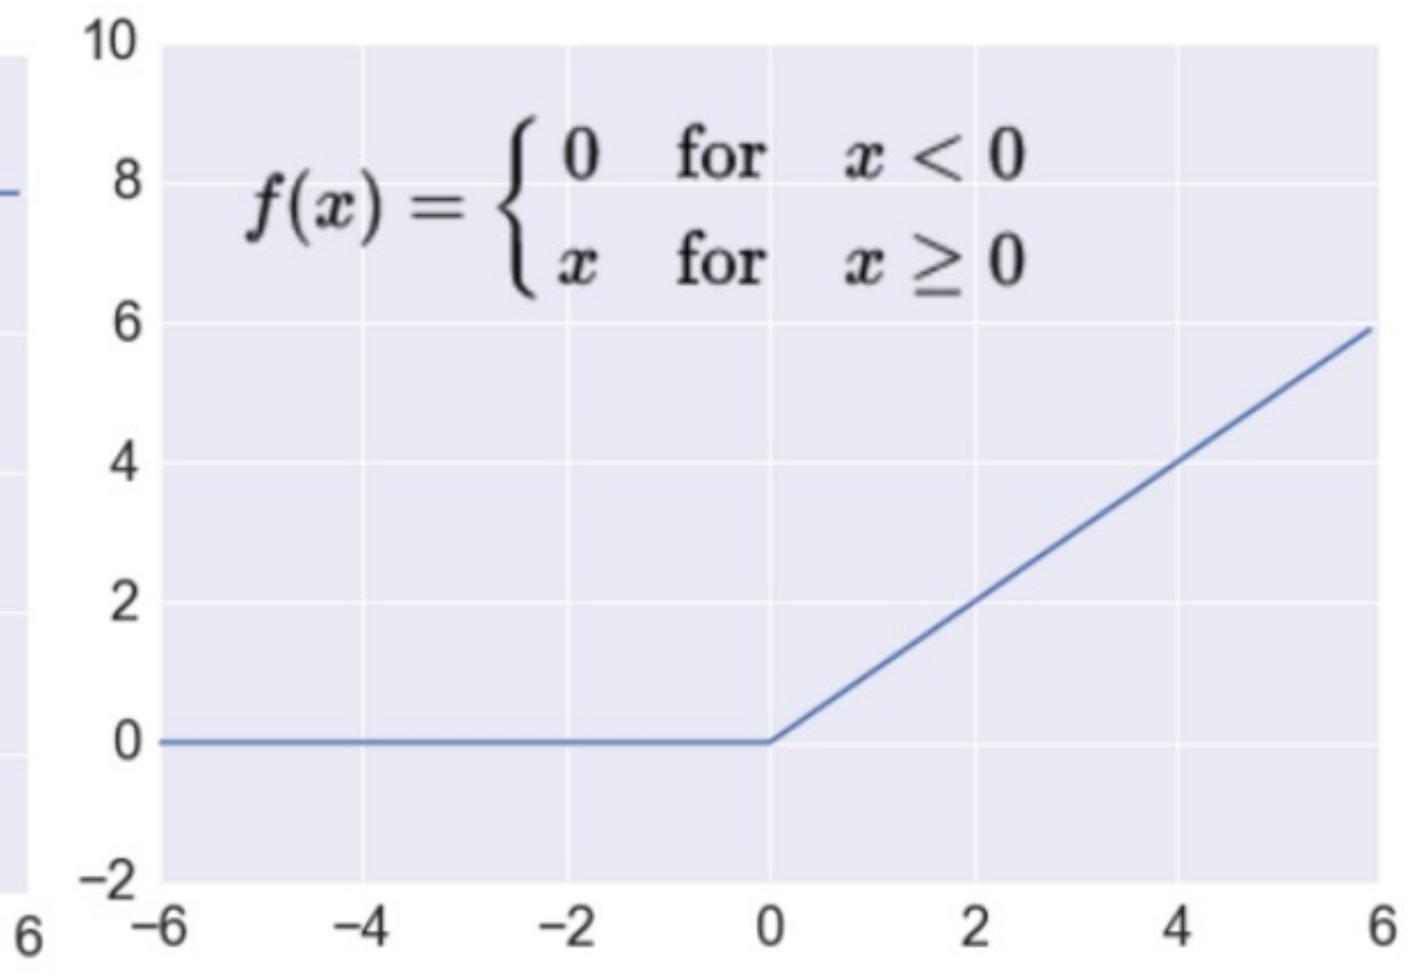
Sigmoid



TanH



ReLU

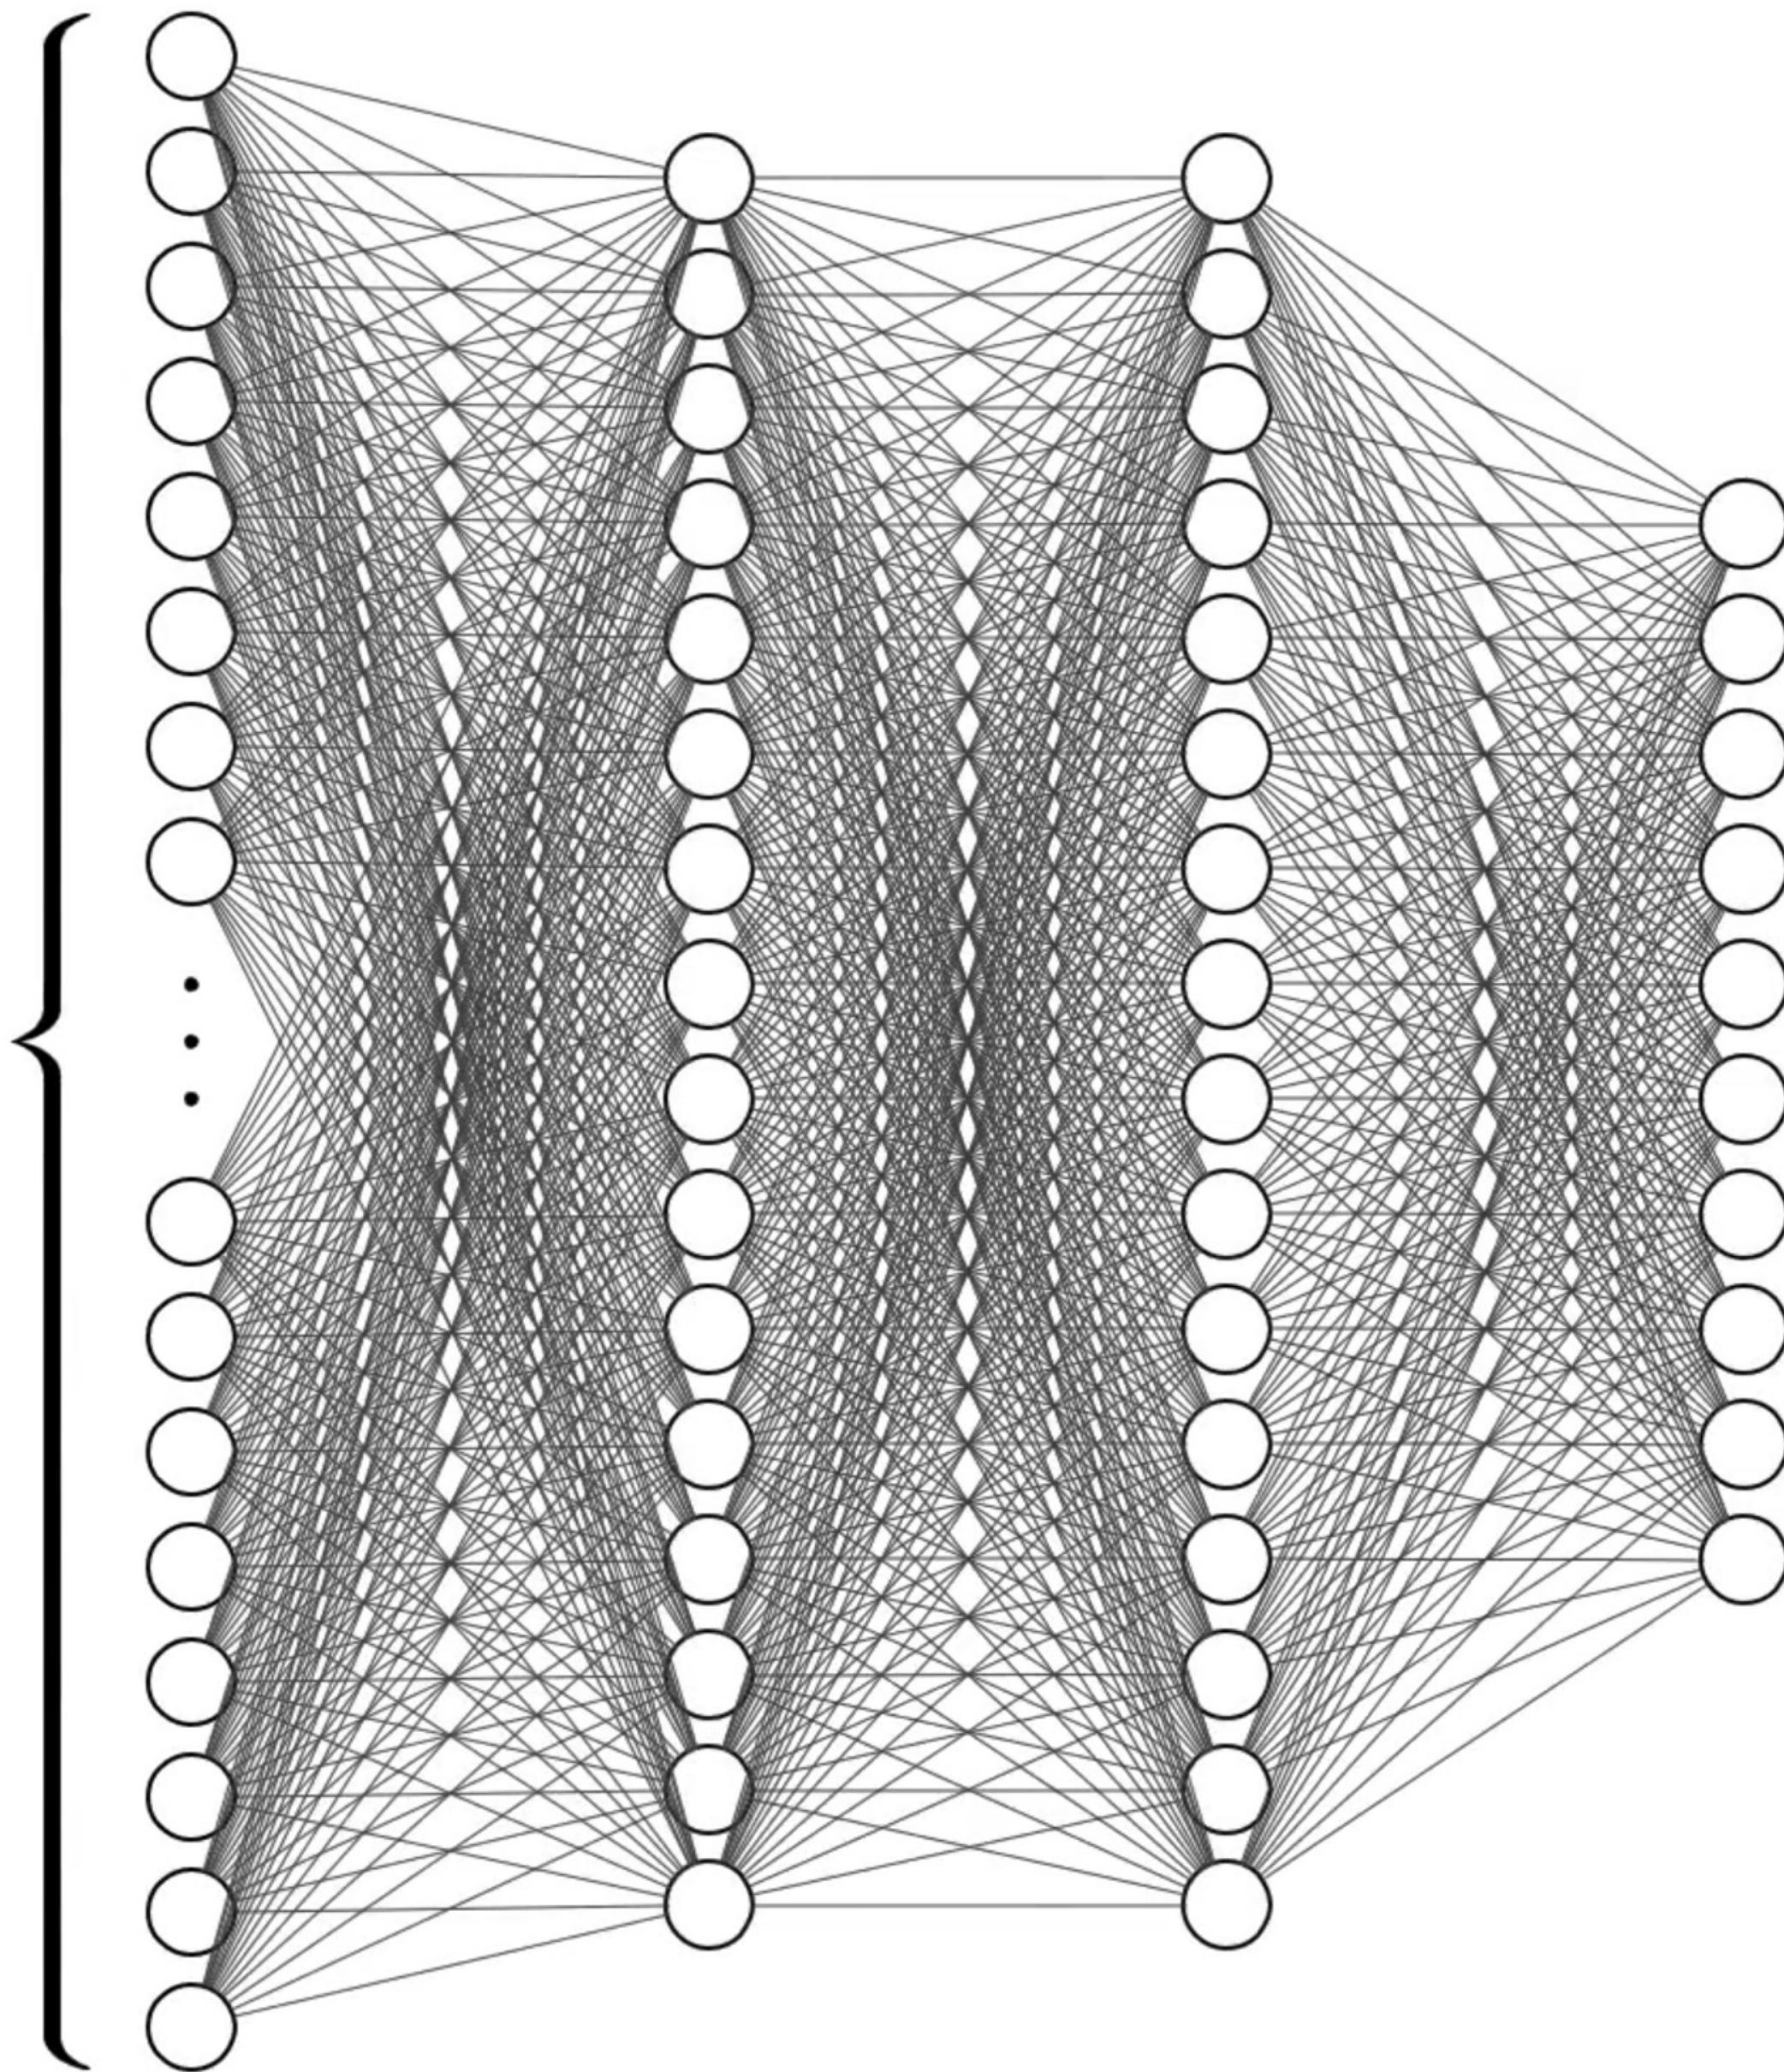


28x28 => 784 pixels

The value of the pixels are concatenated in a vector.

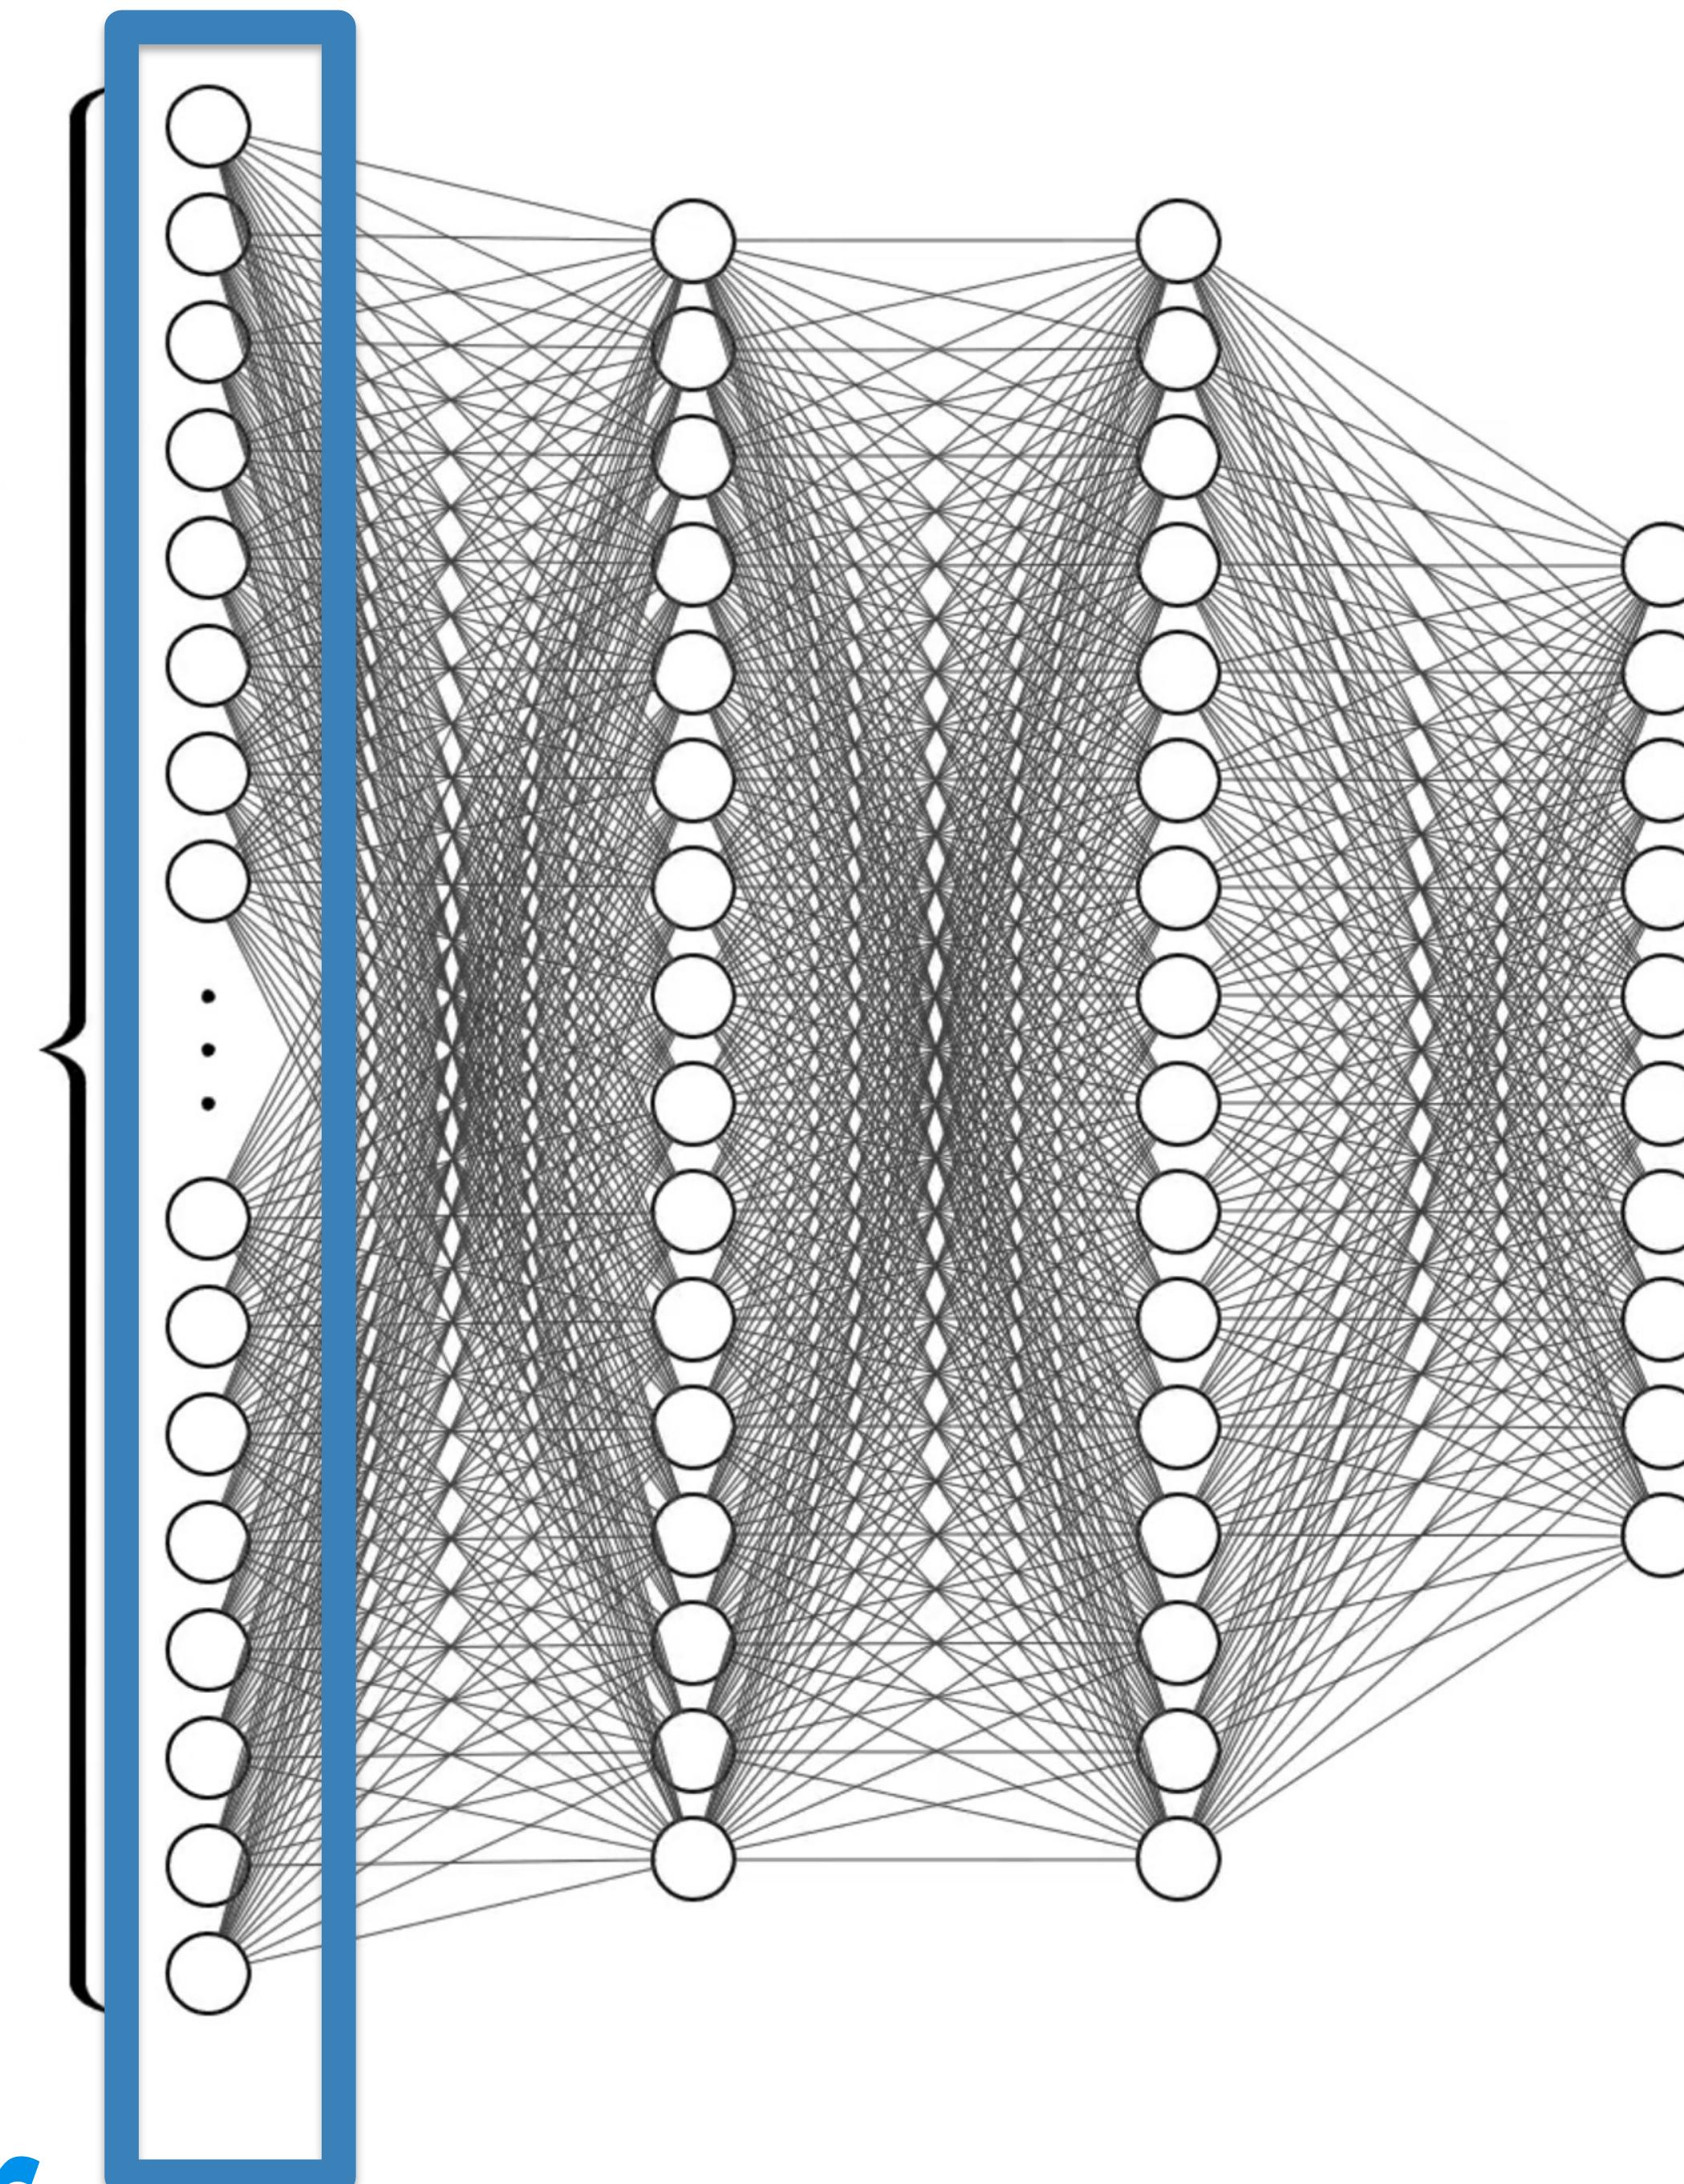
The 784 numbers represent the input layer to our network

784

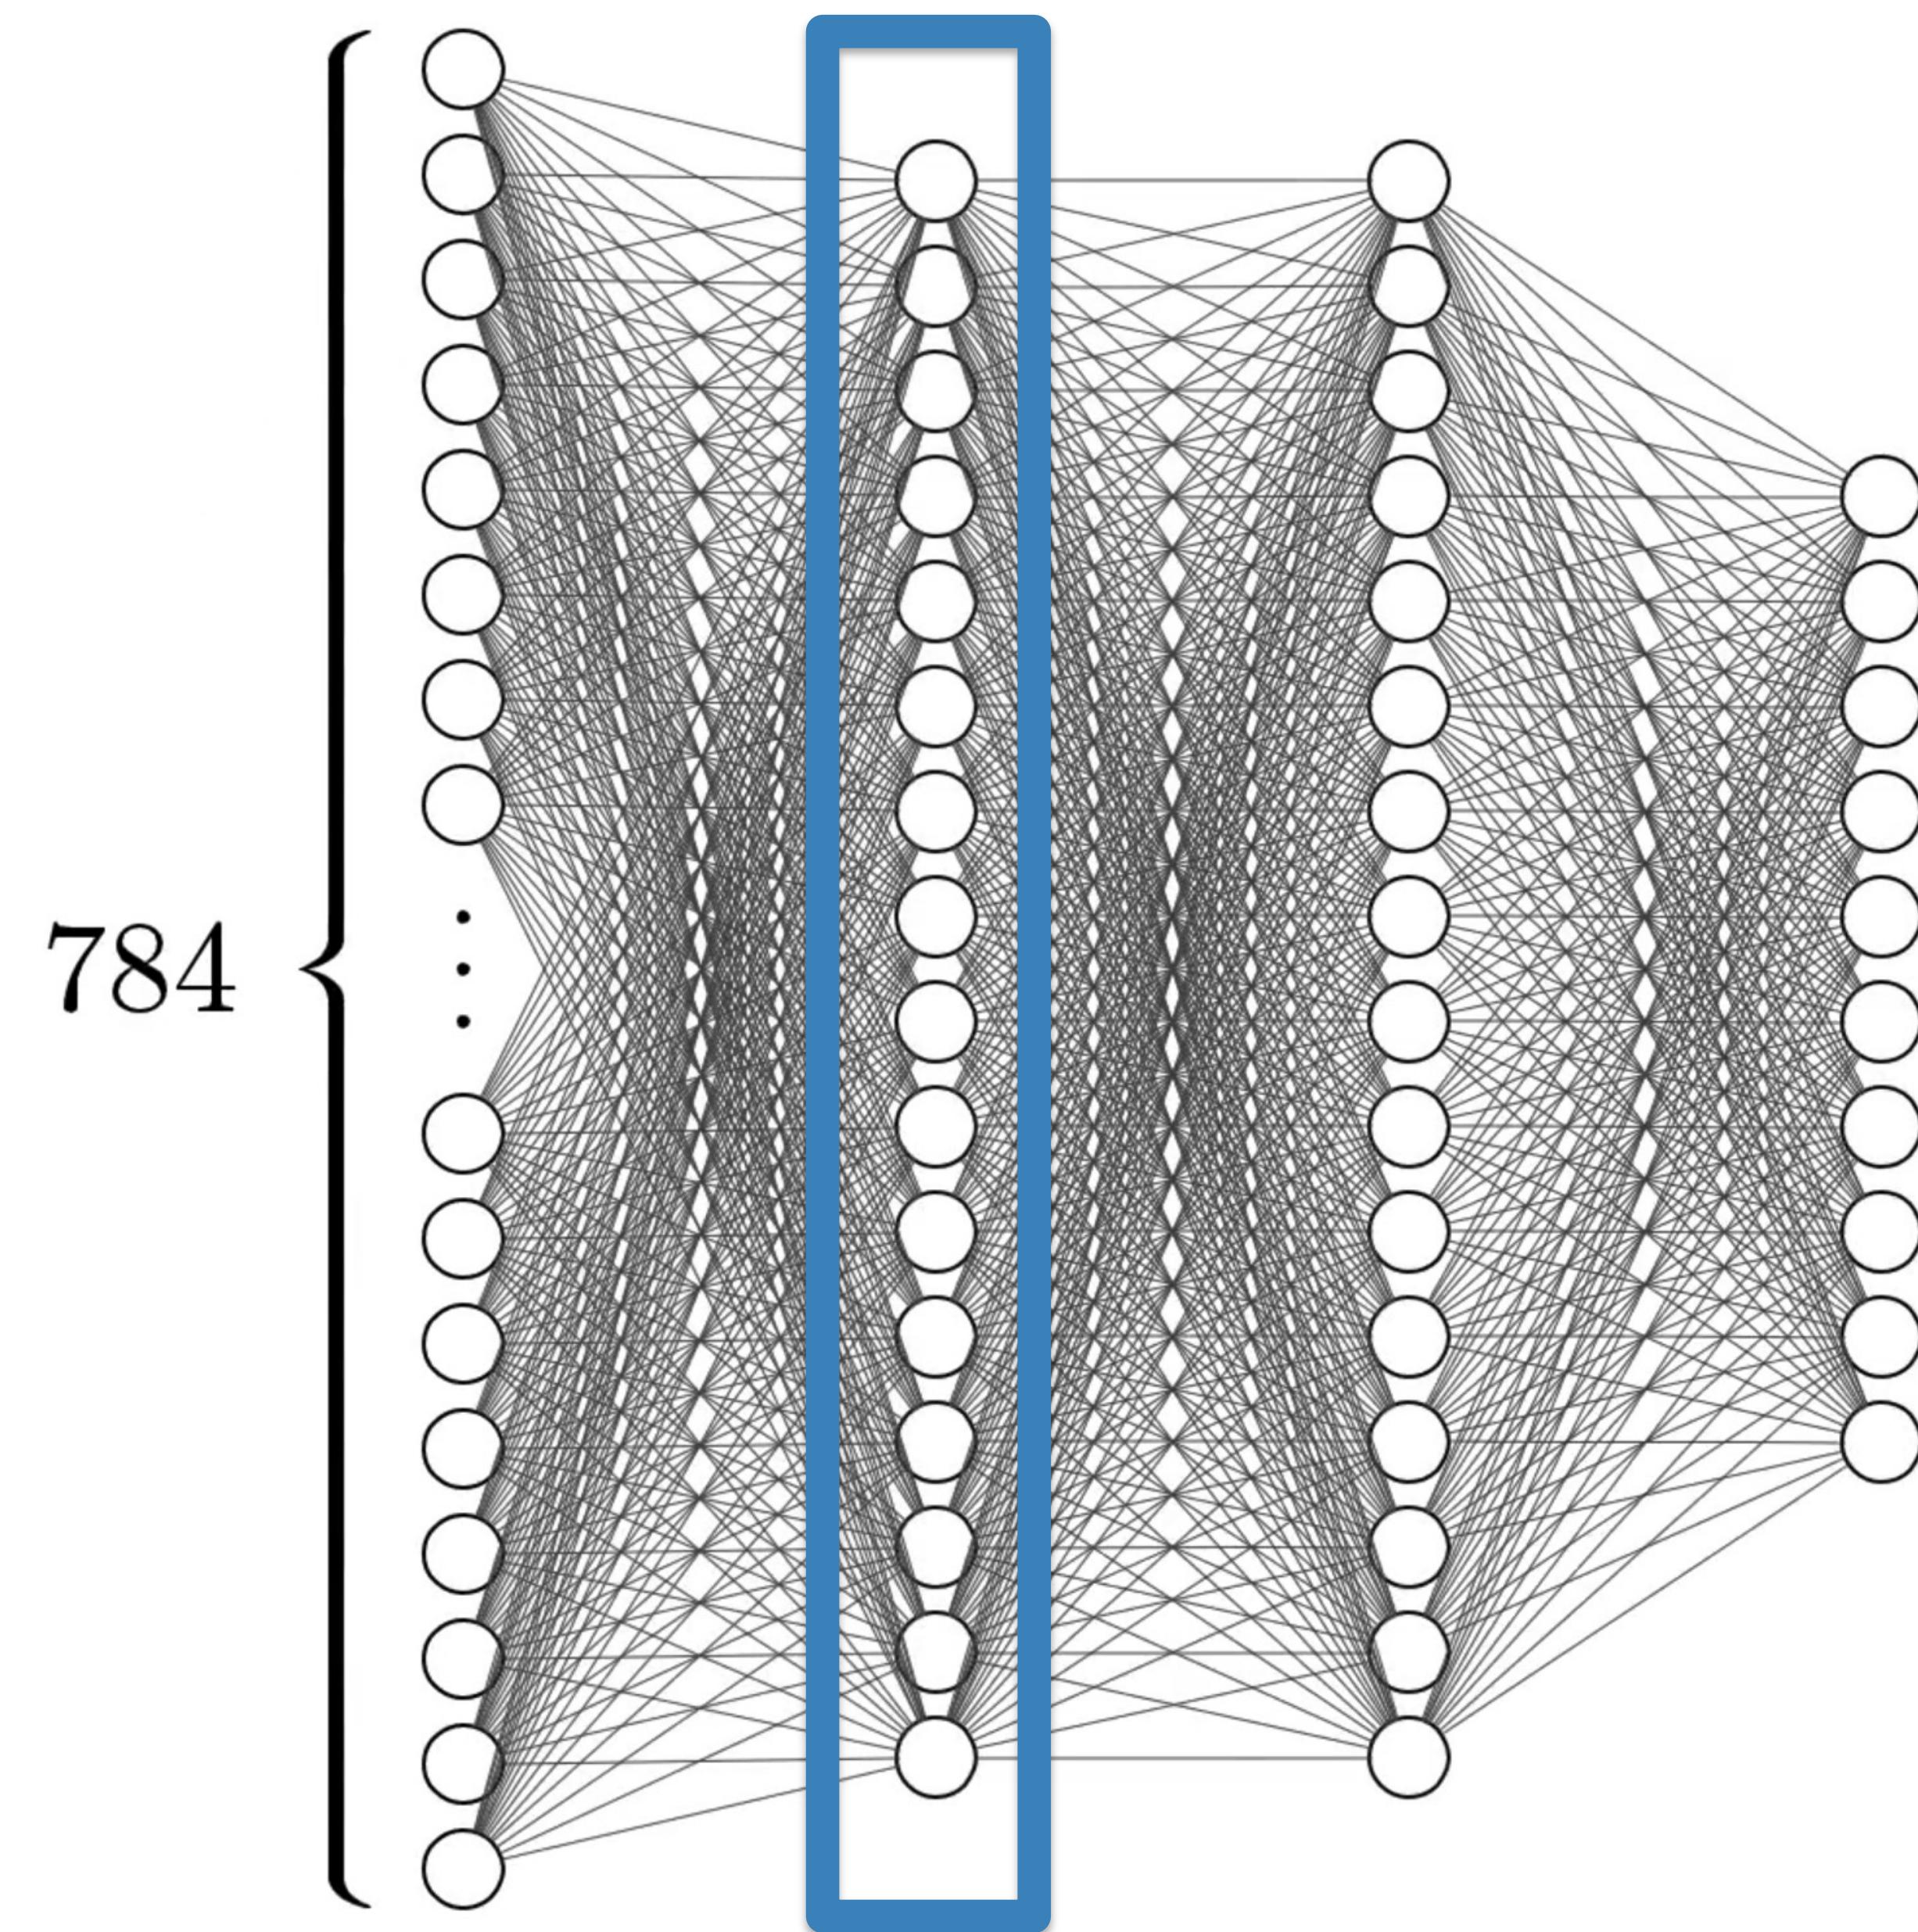


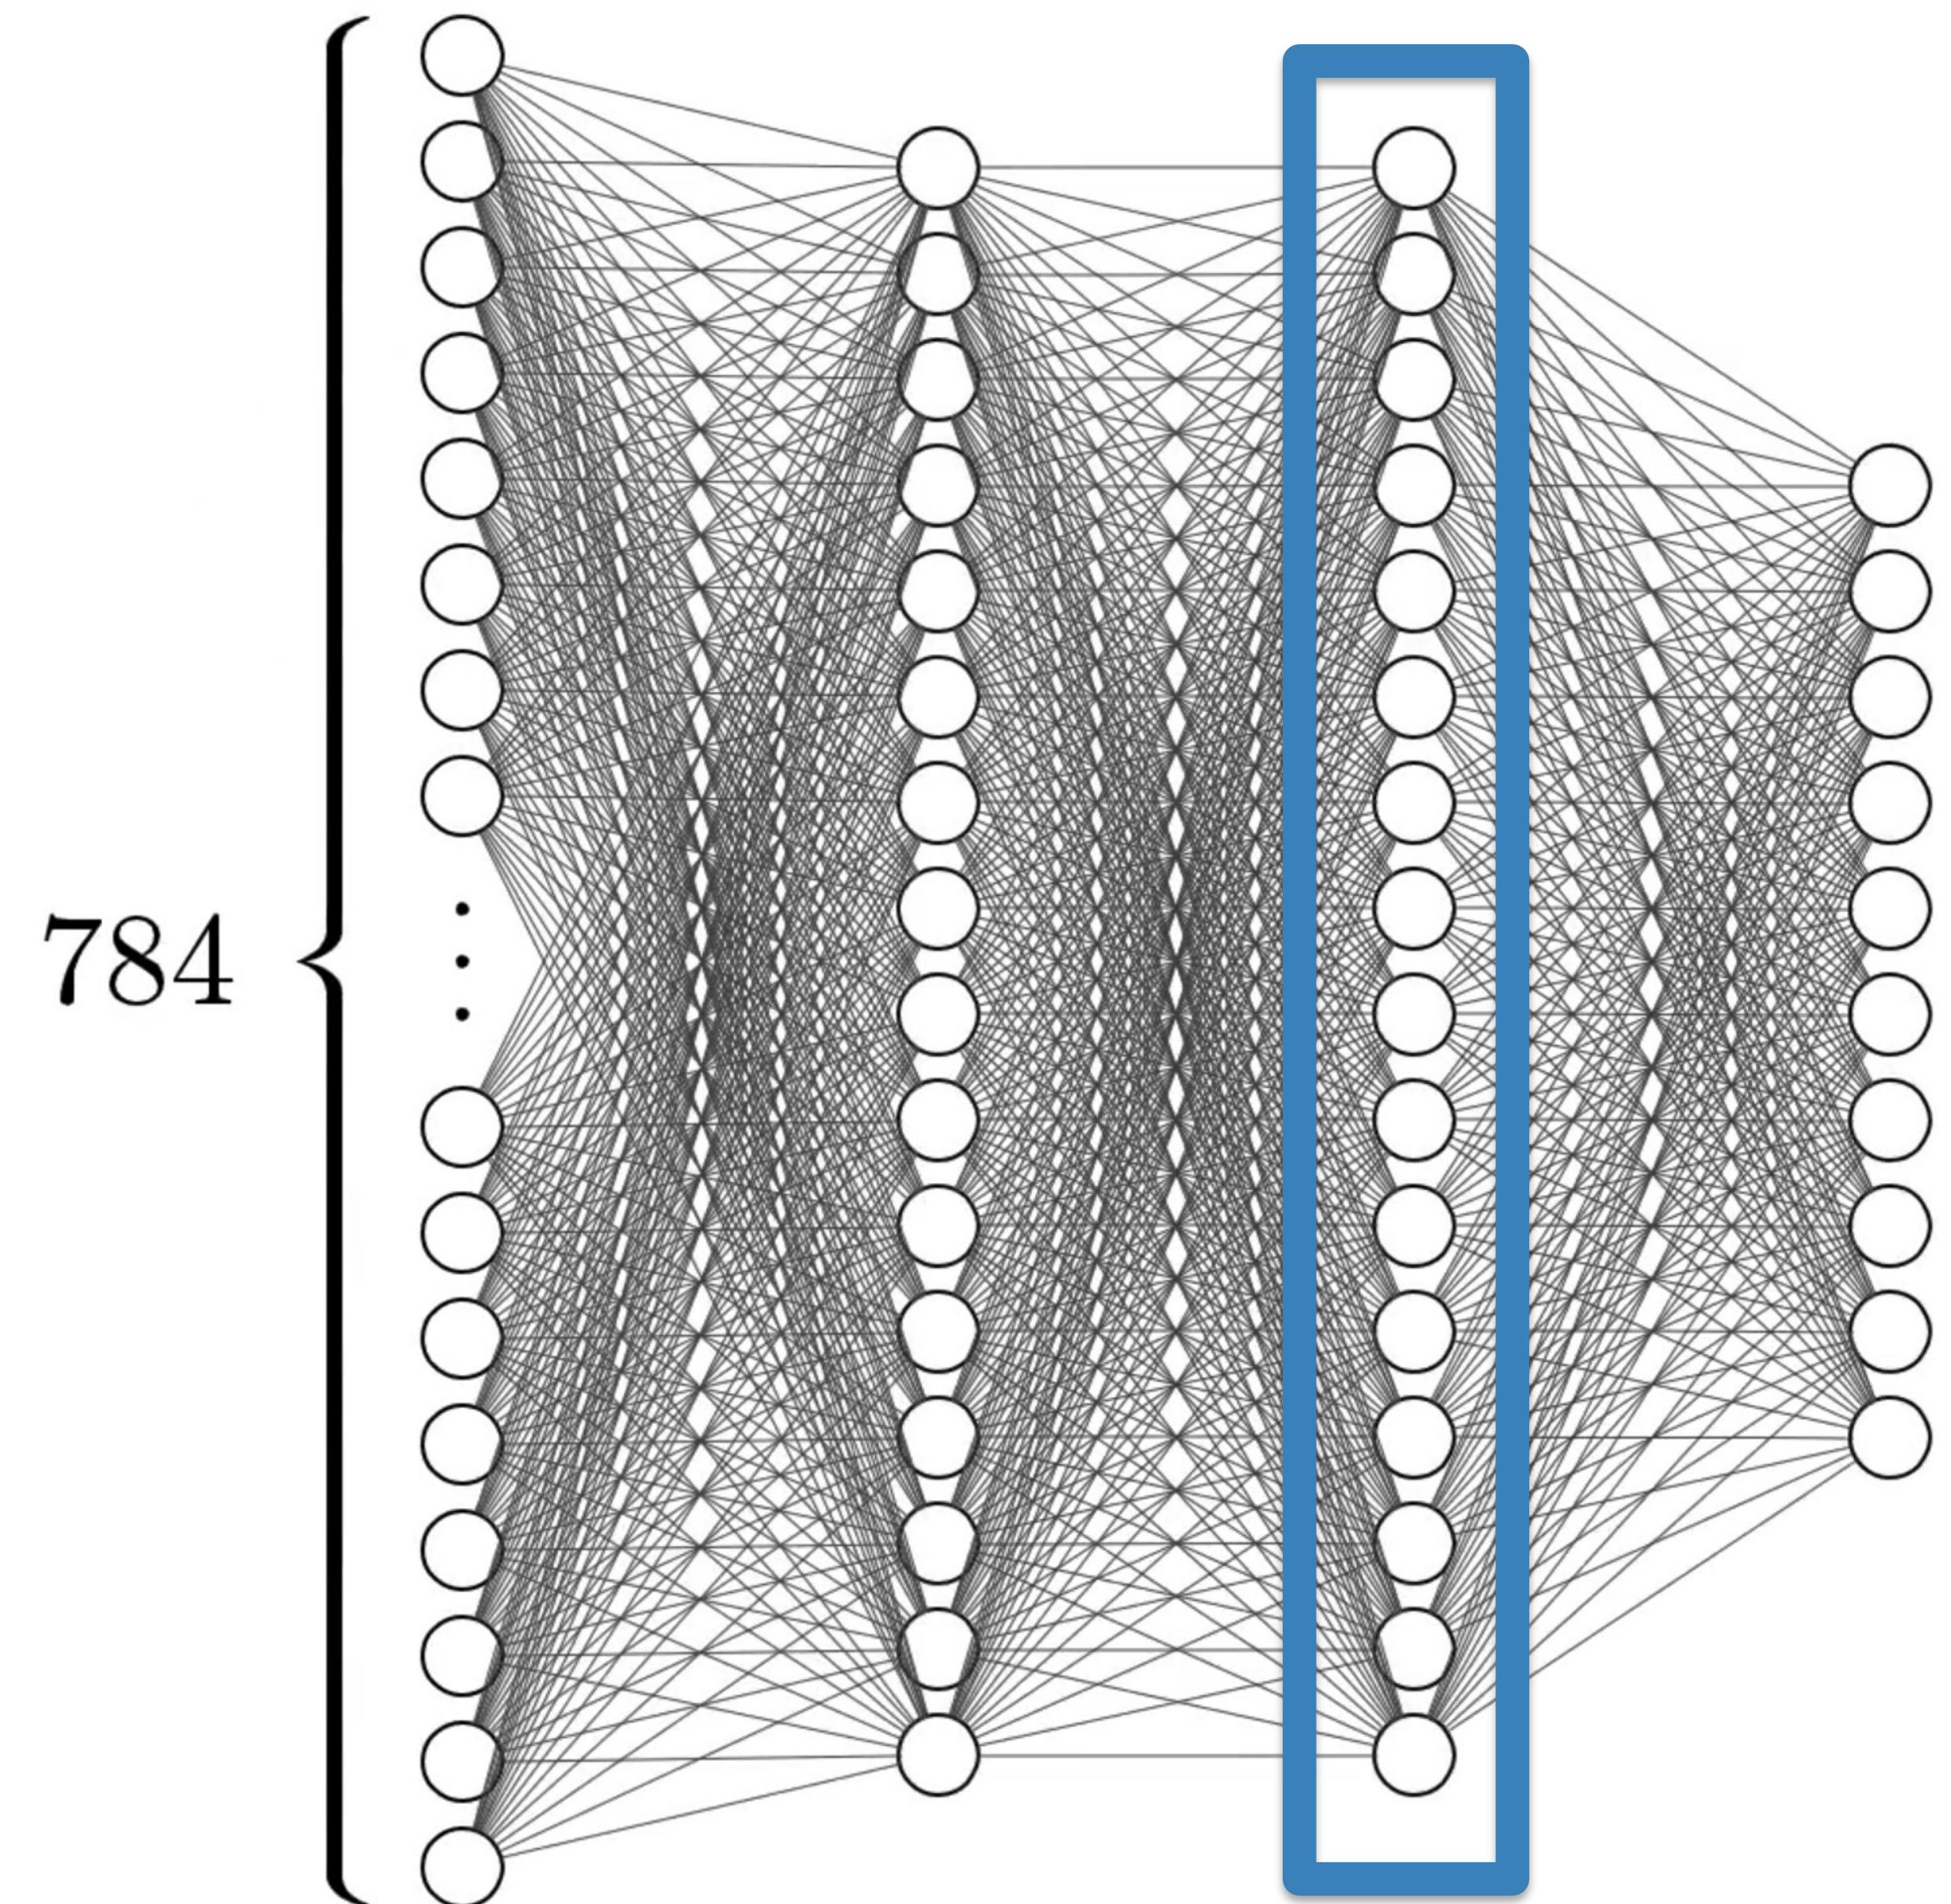
Input Layer

784

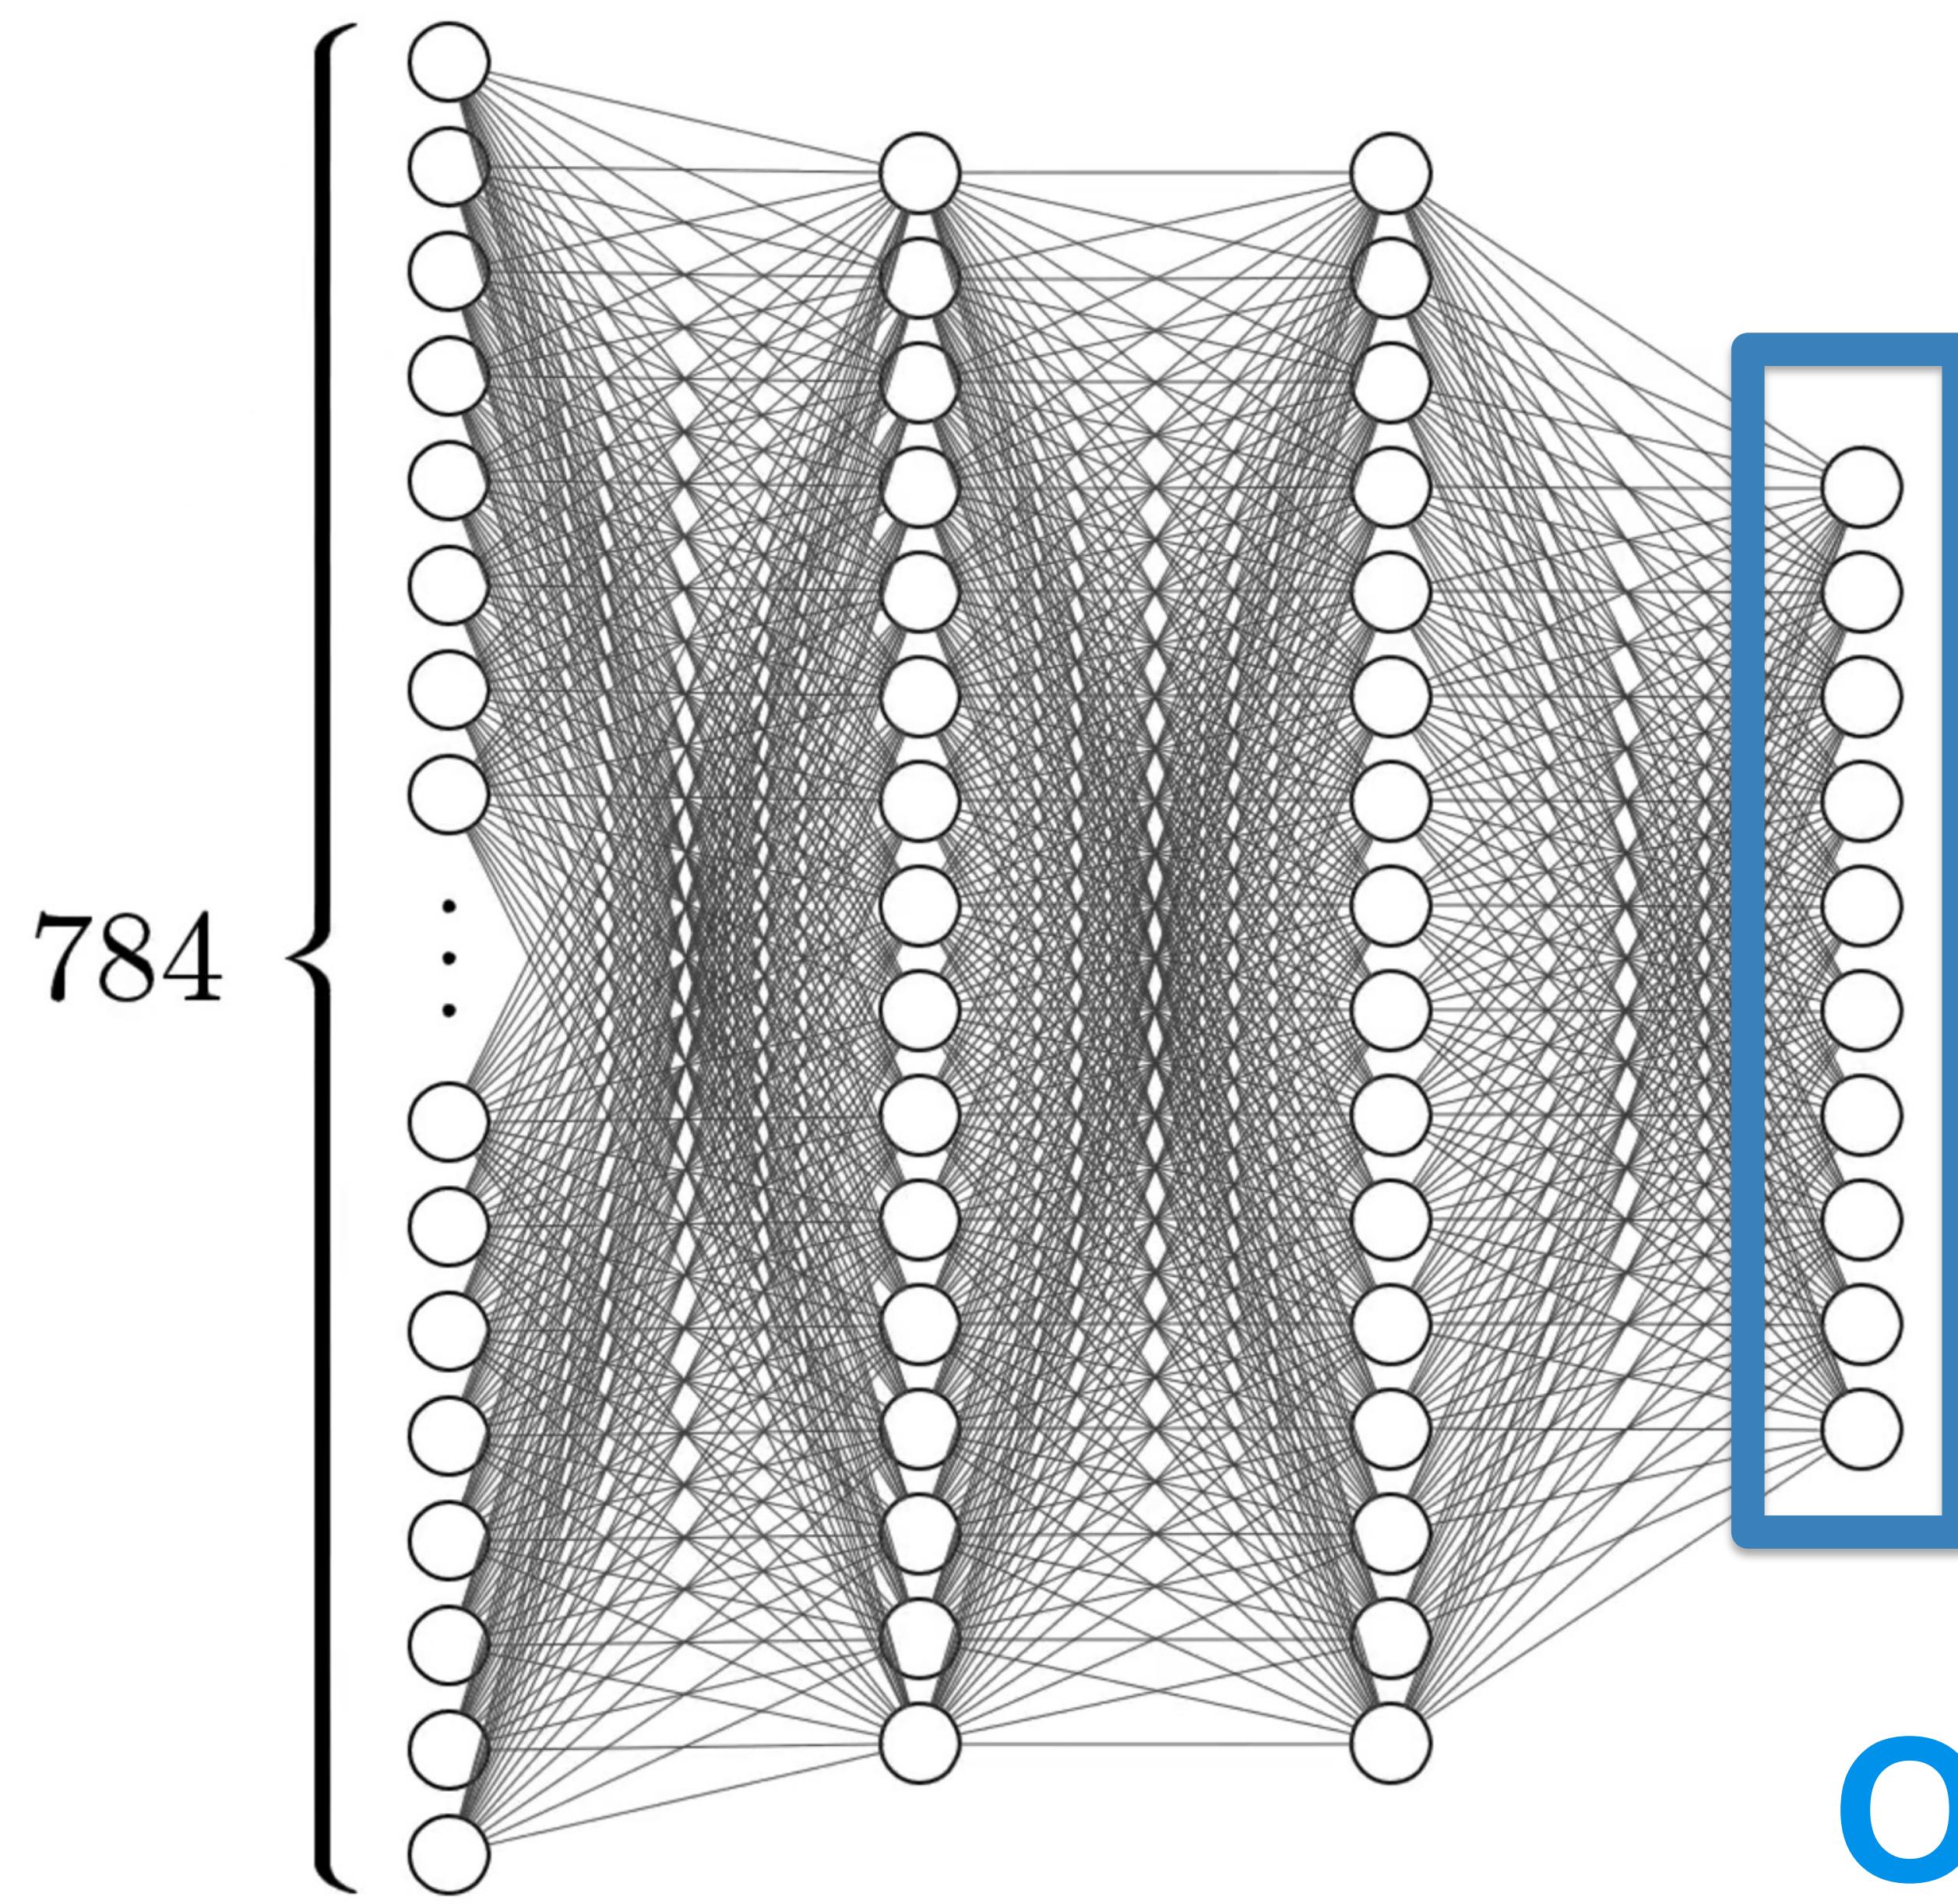


Hidden Layer 1

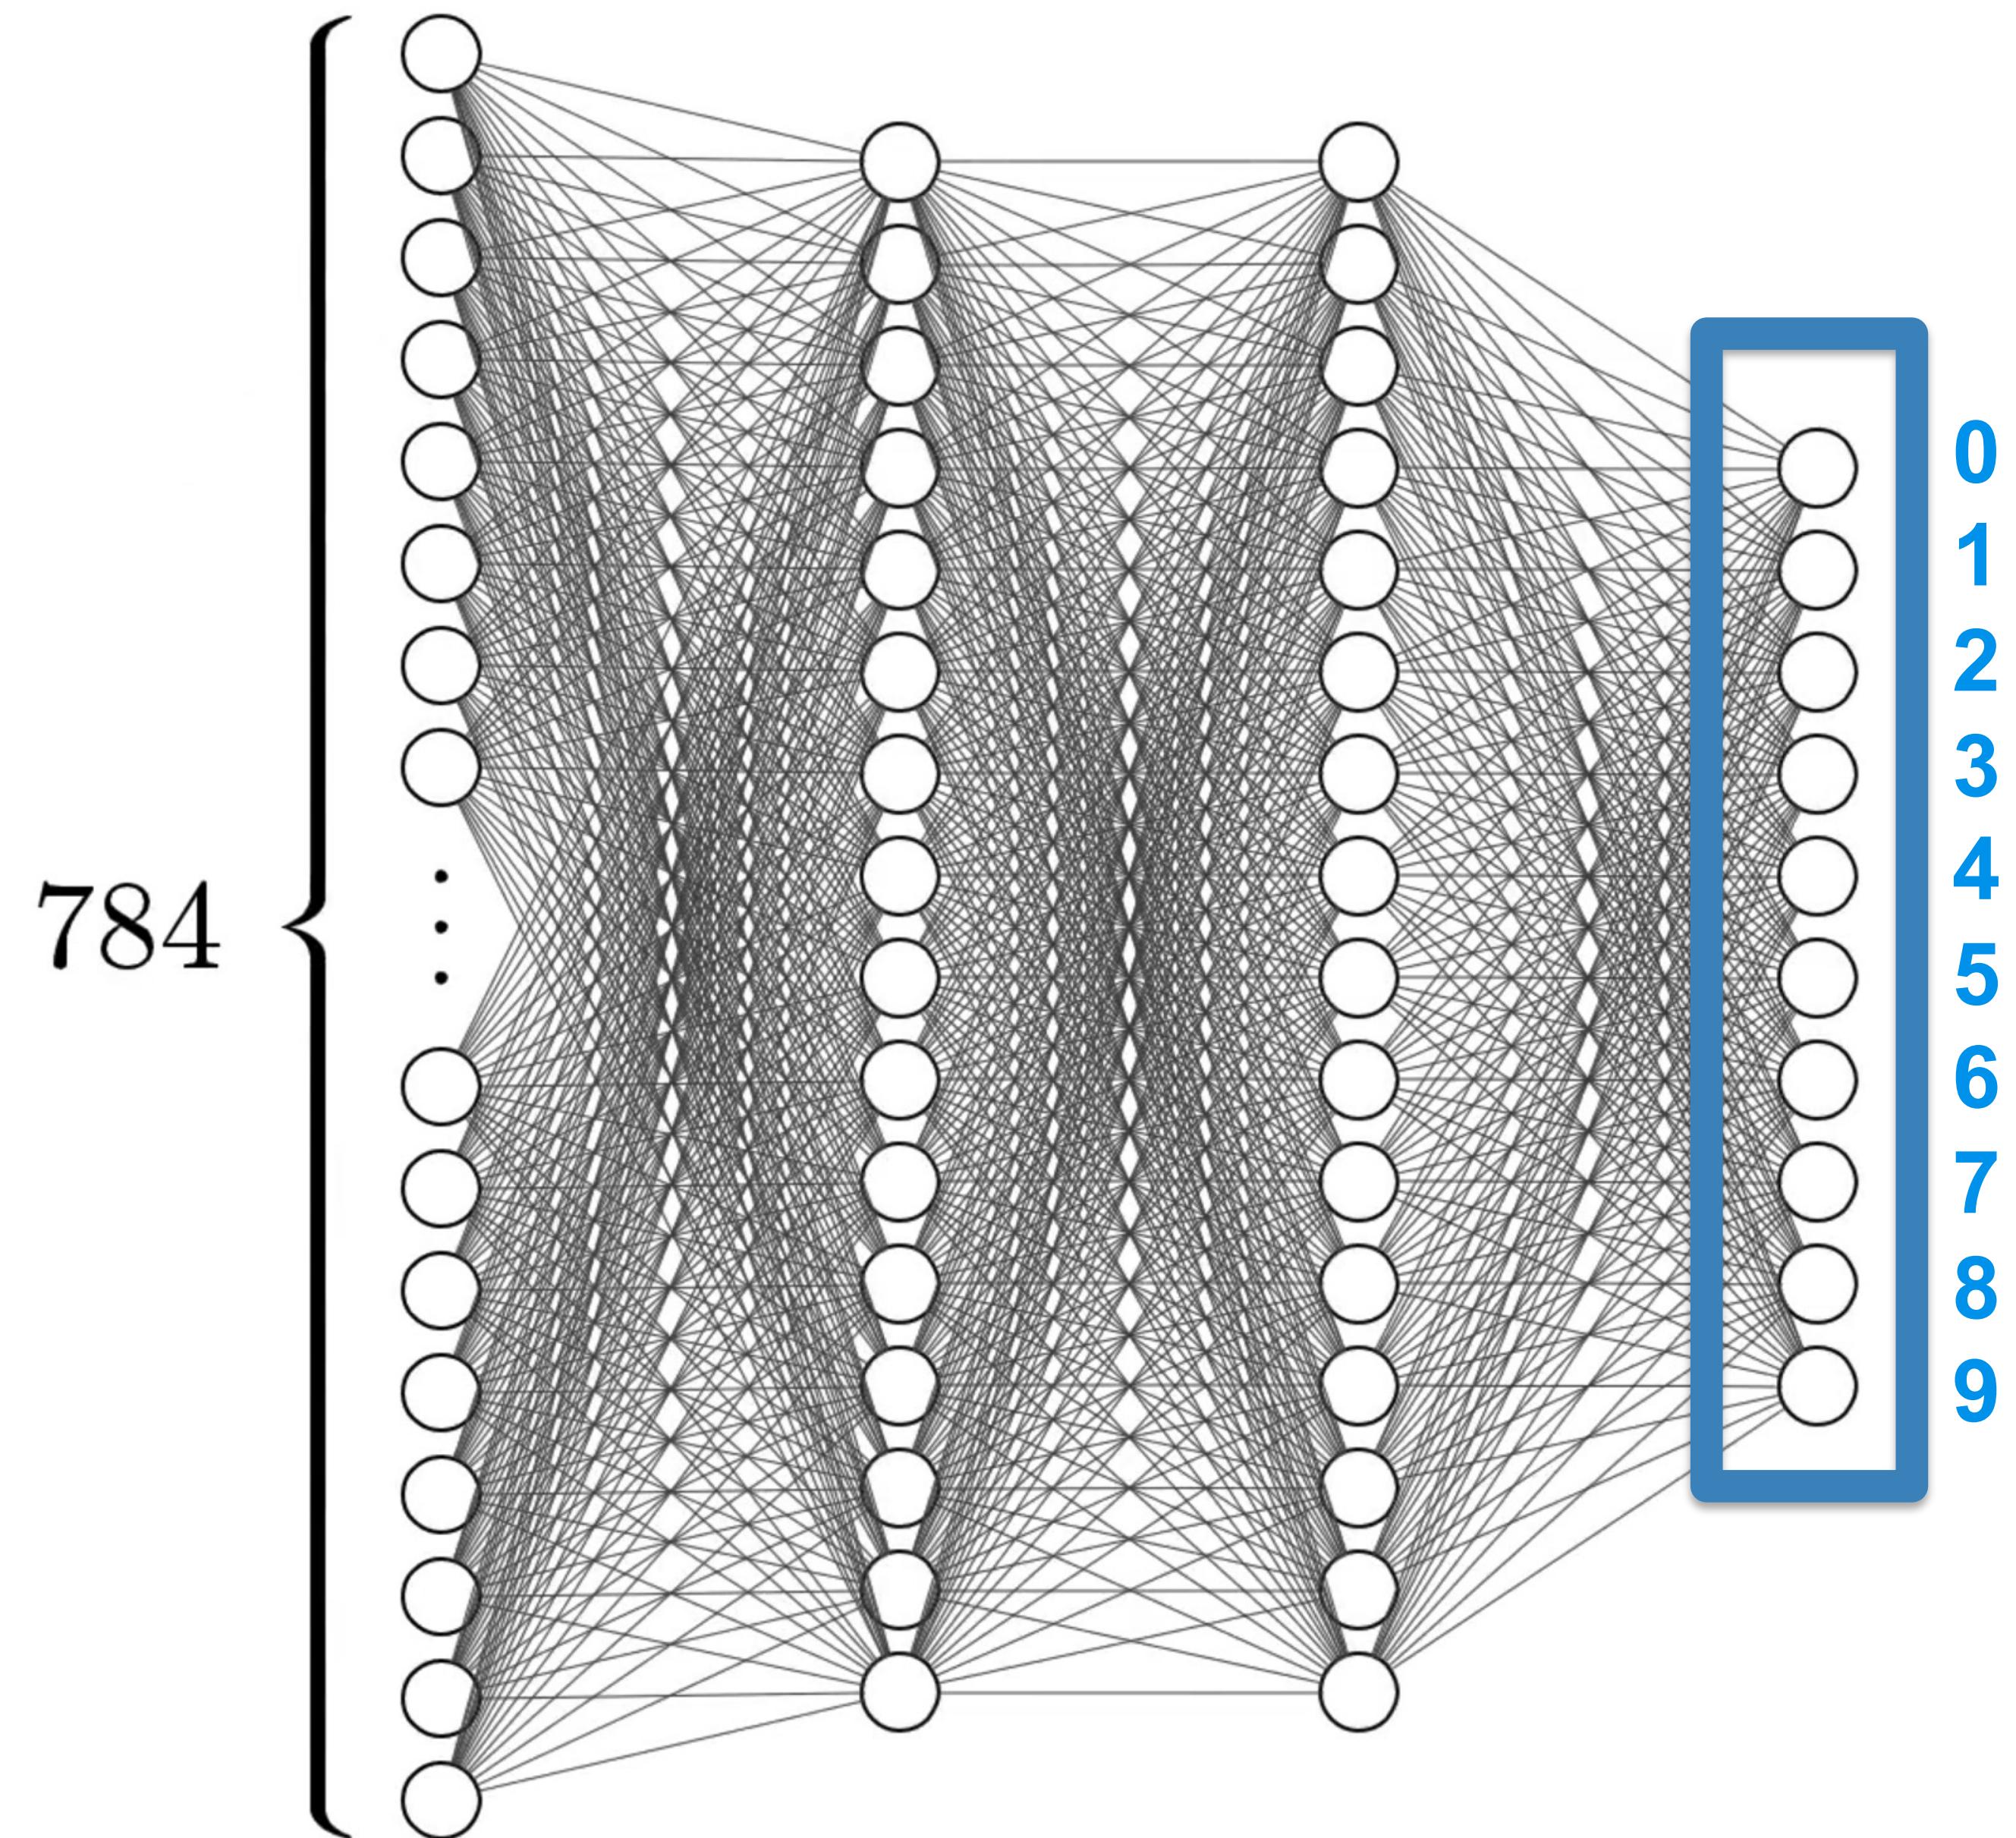




Hidden Layer 2

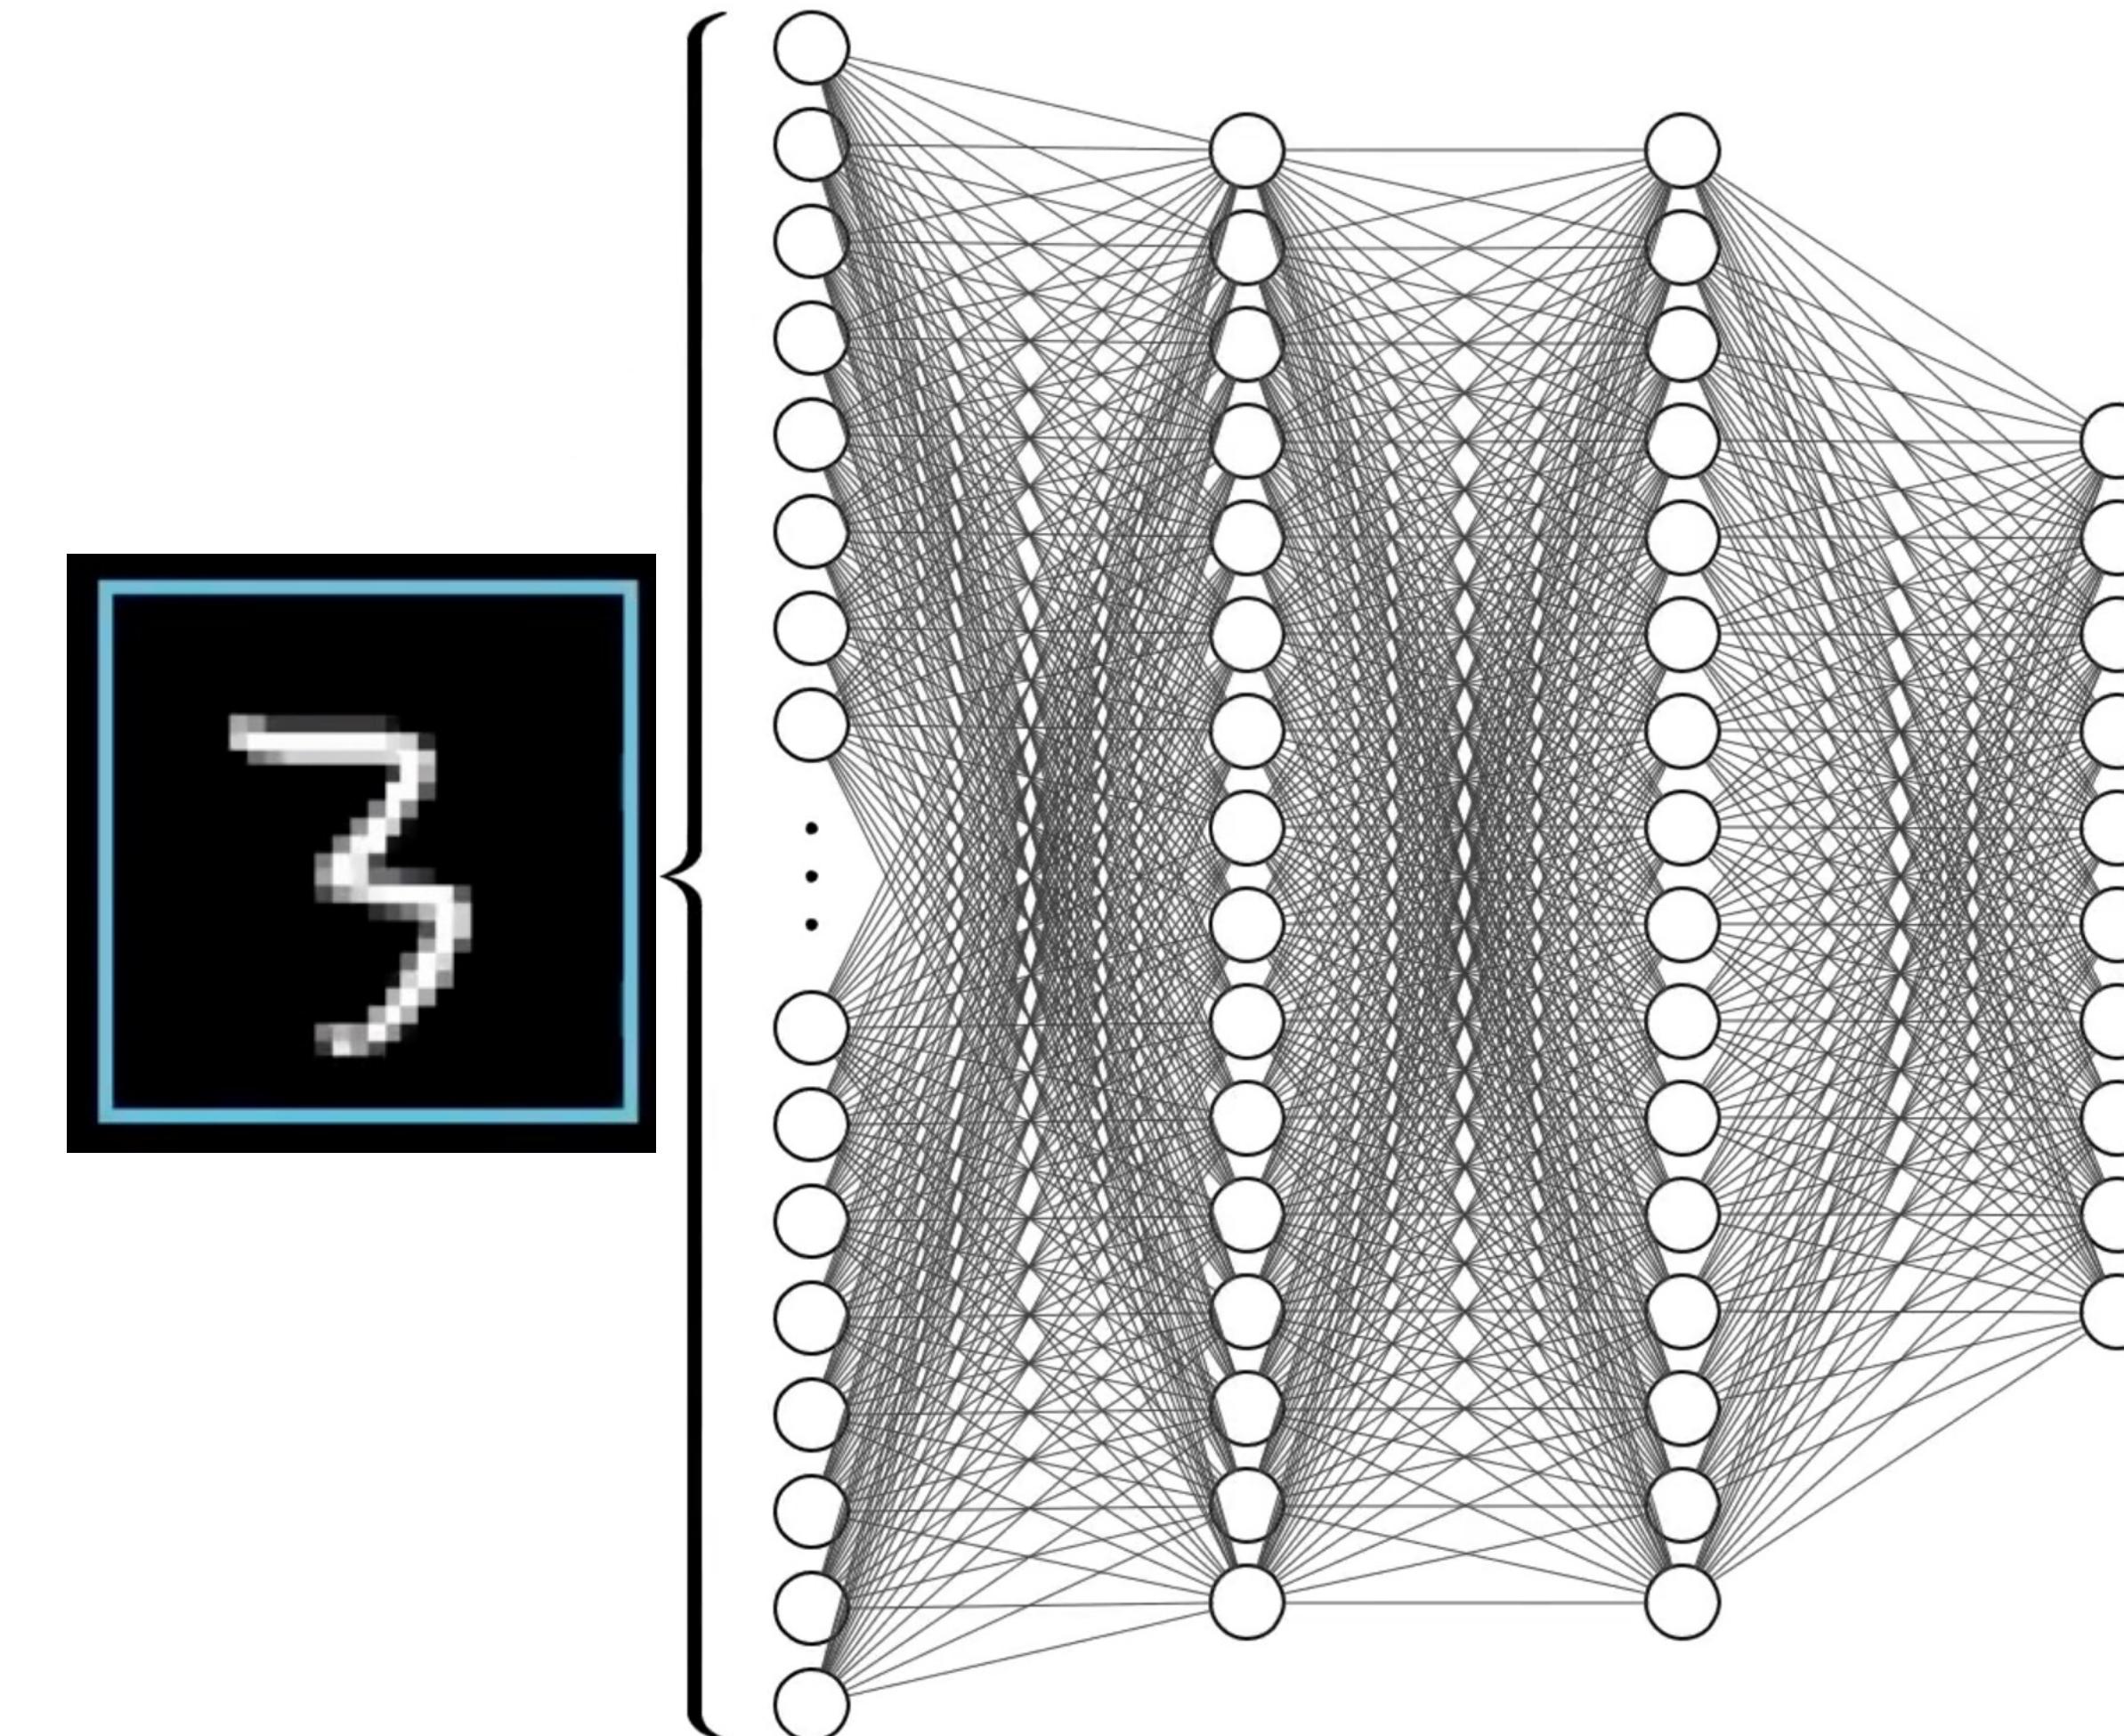


Output Layer



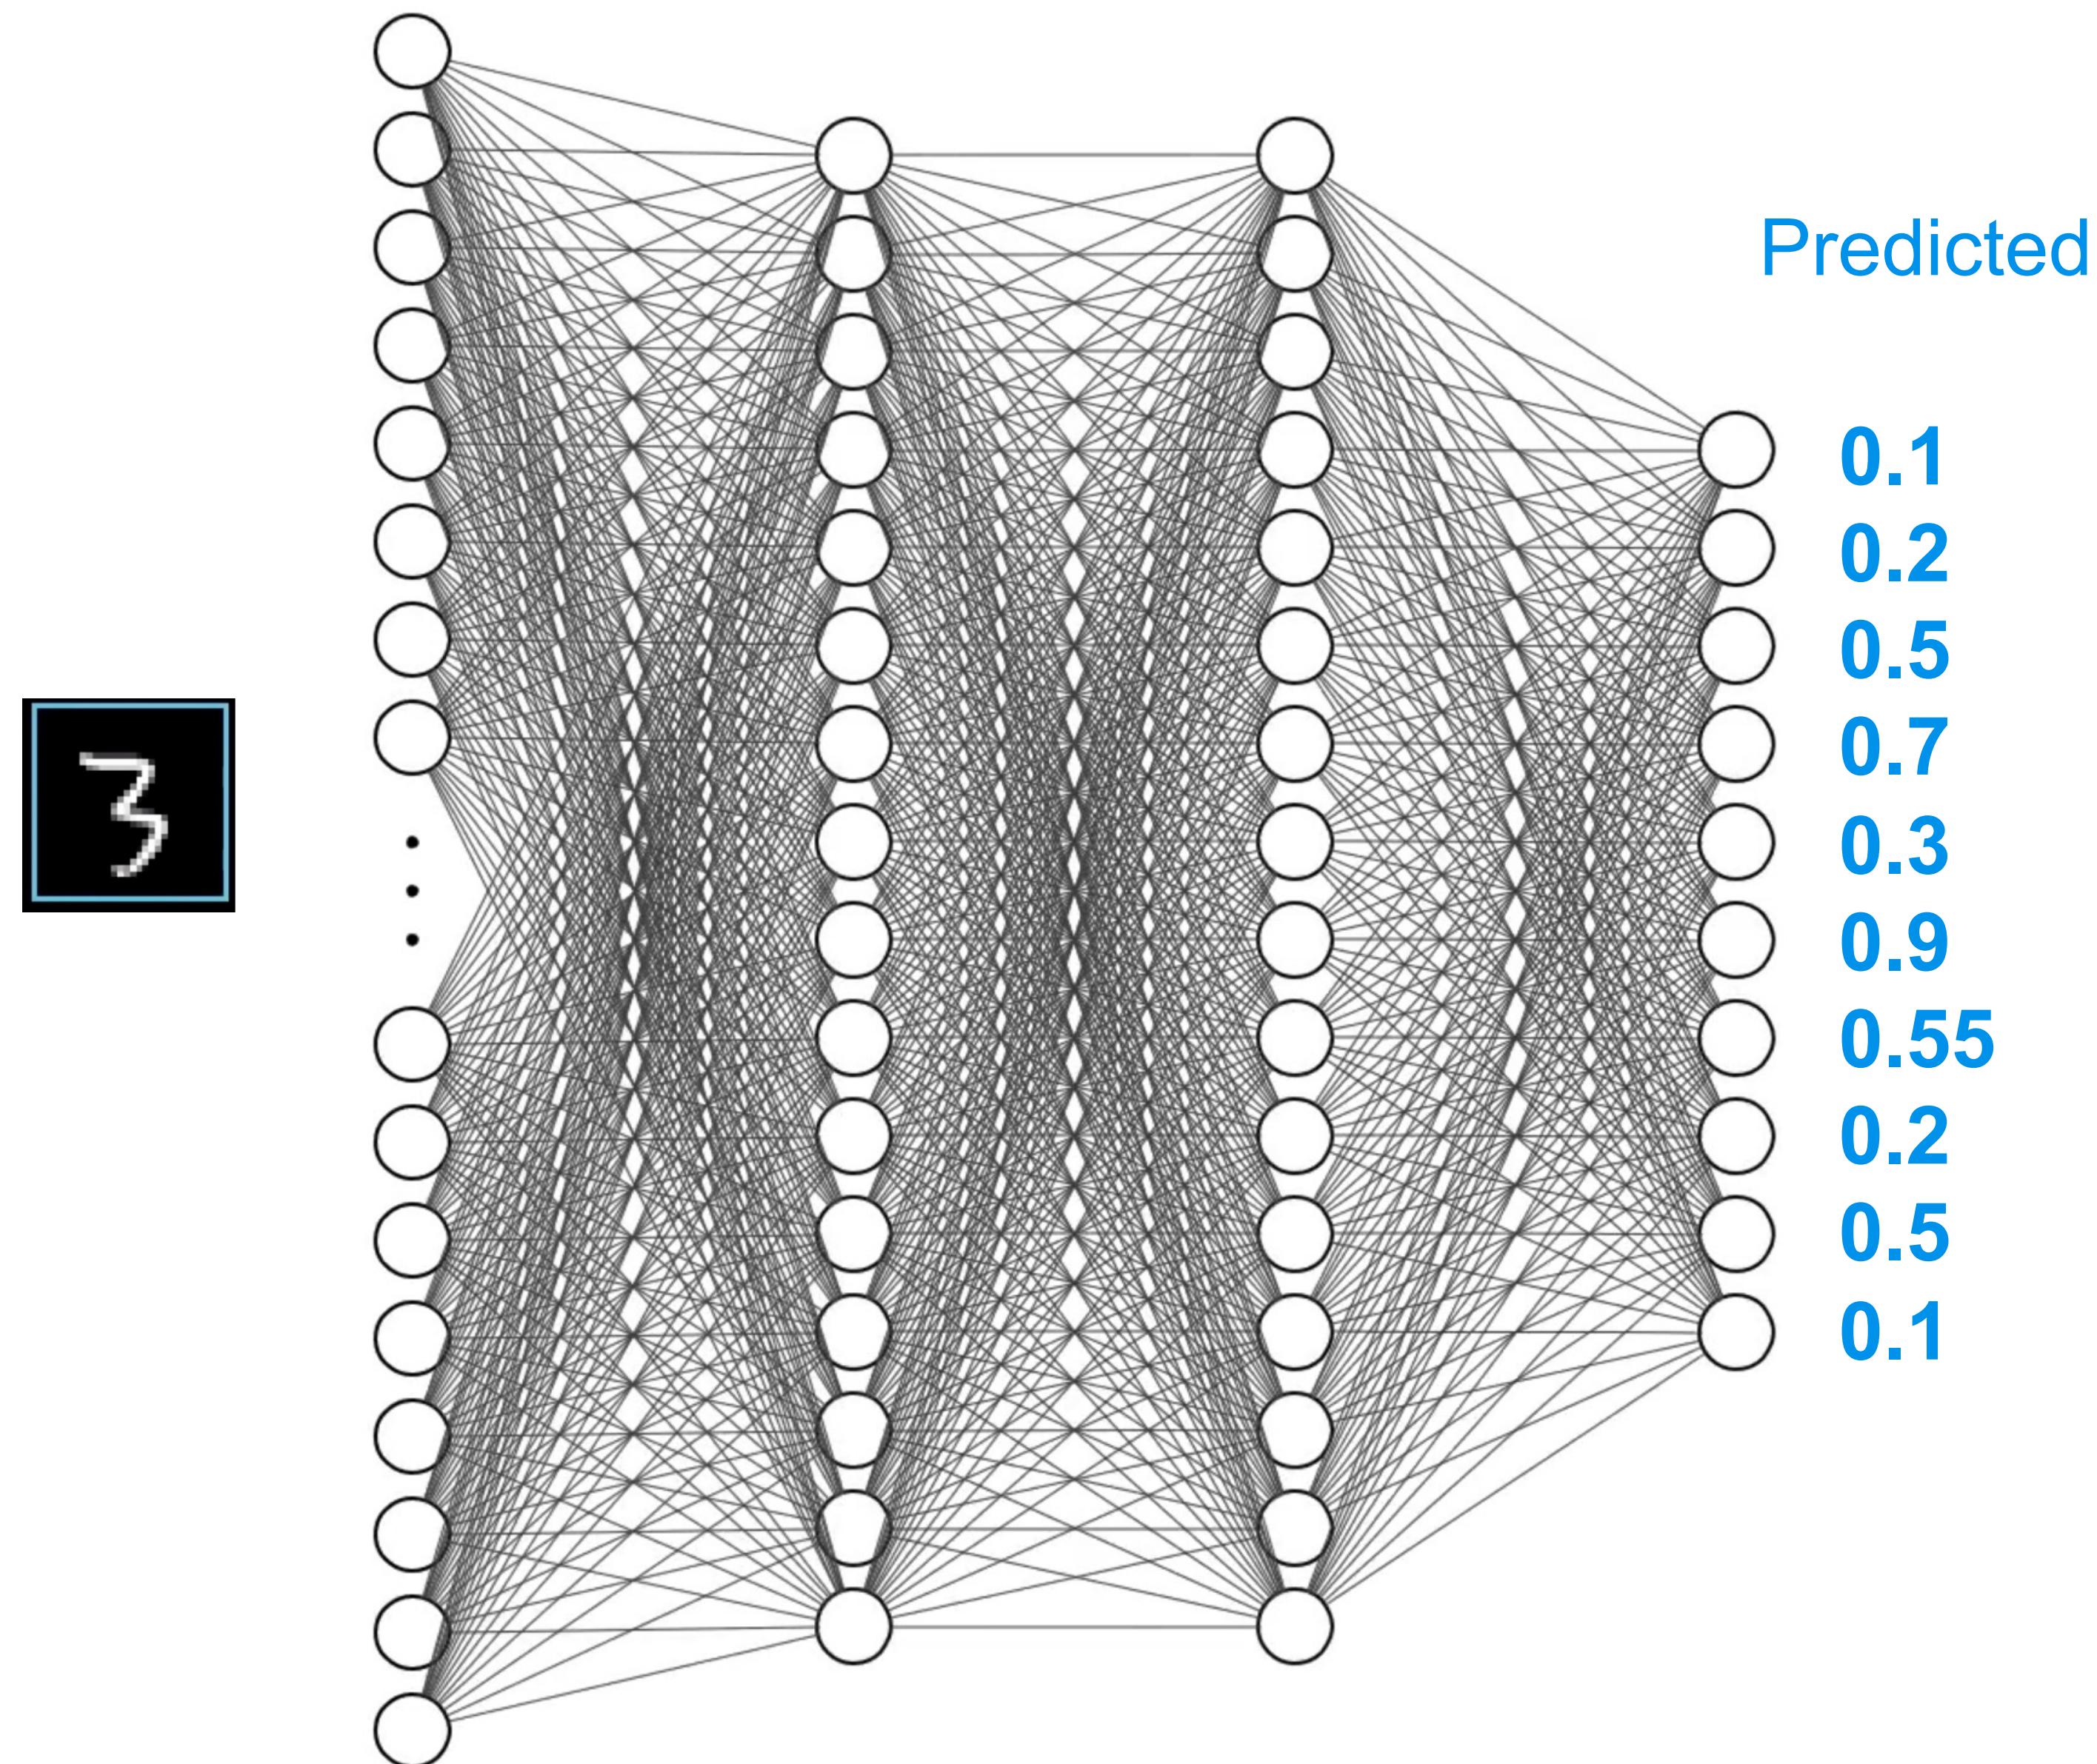
- Each node in the output layer corresponds to a label in the dataset

Training



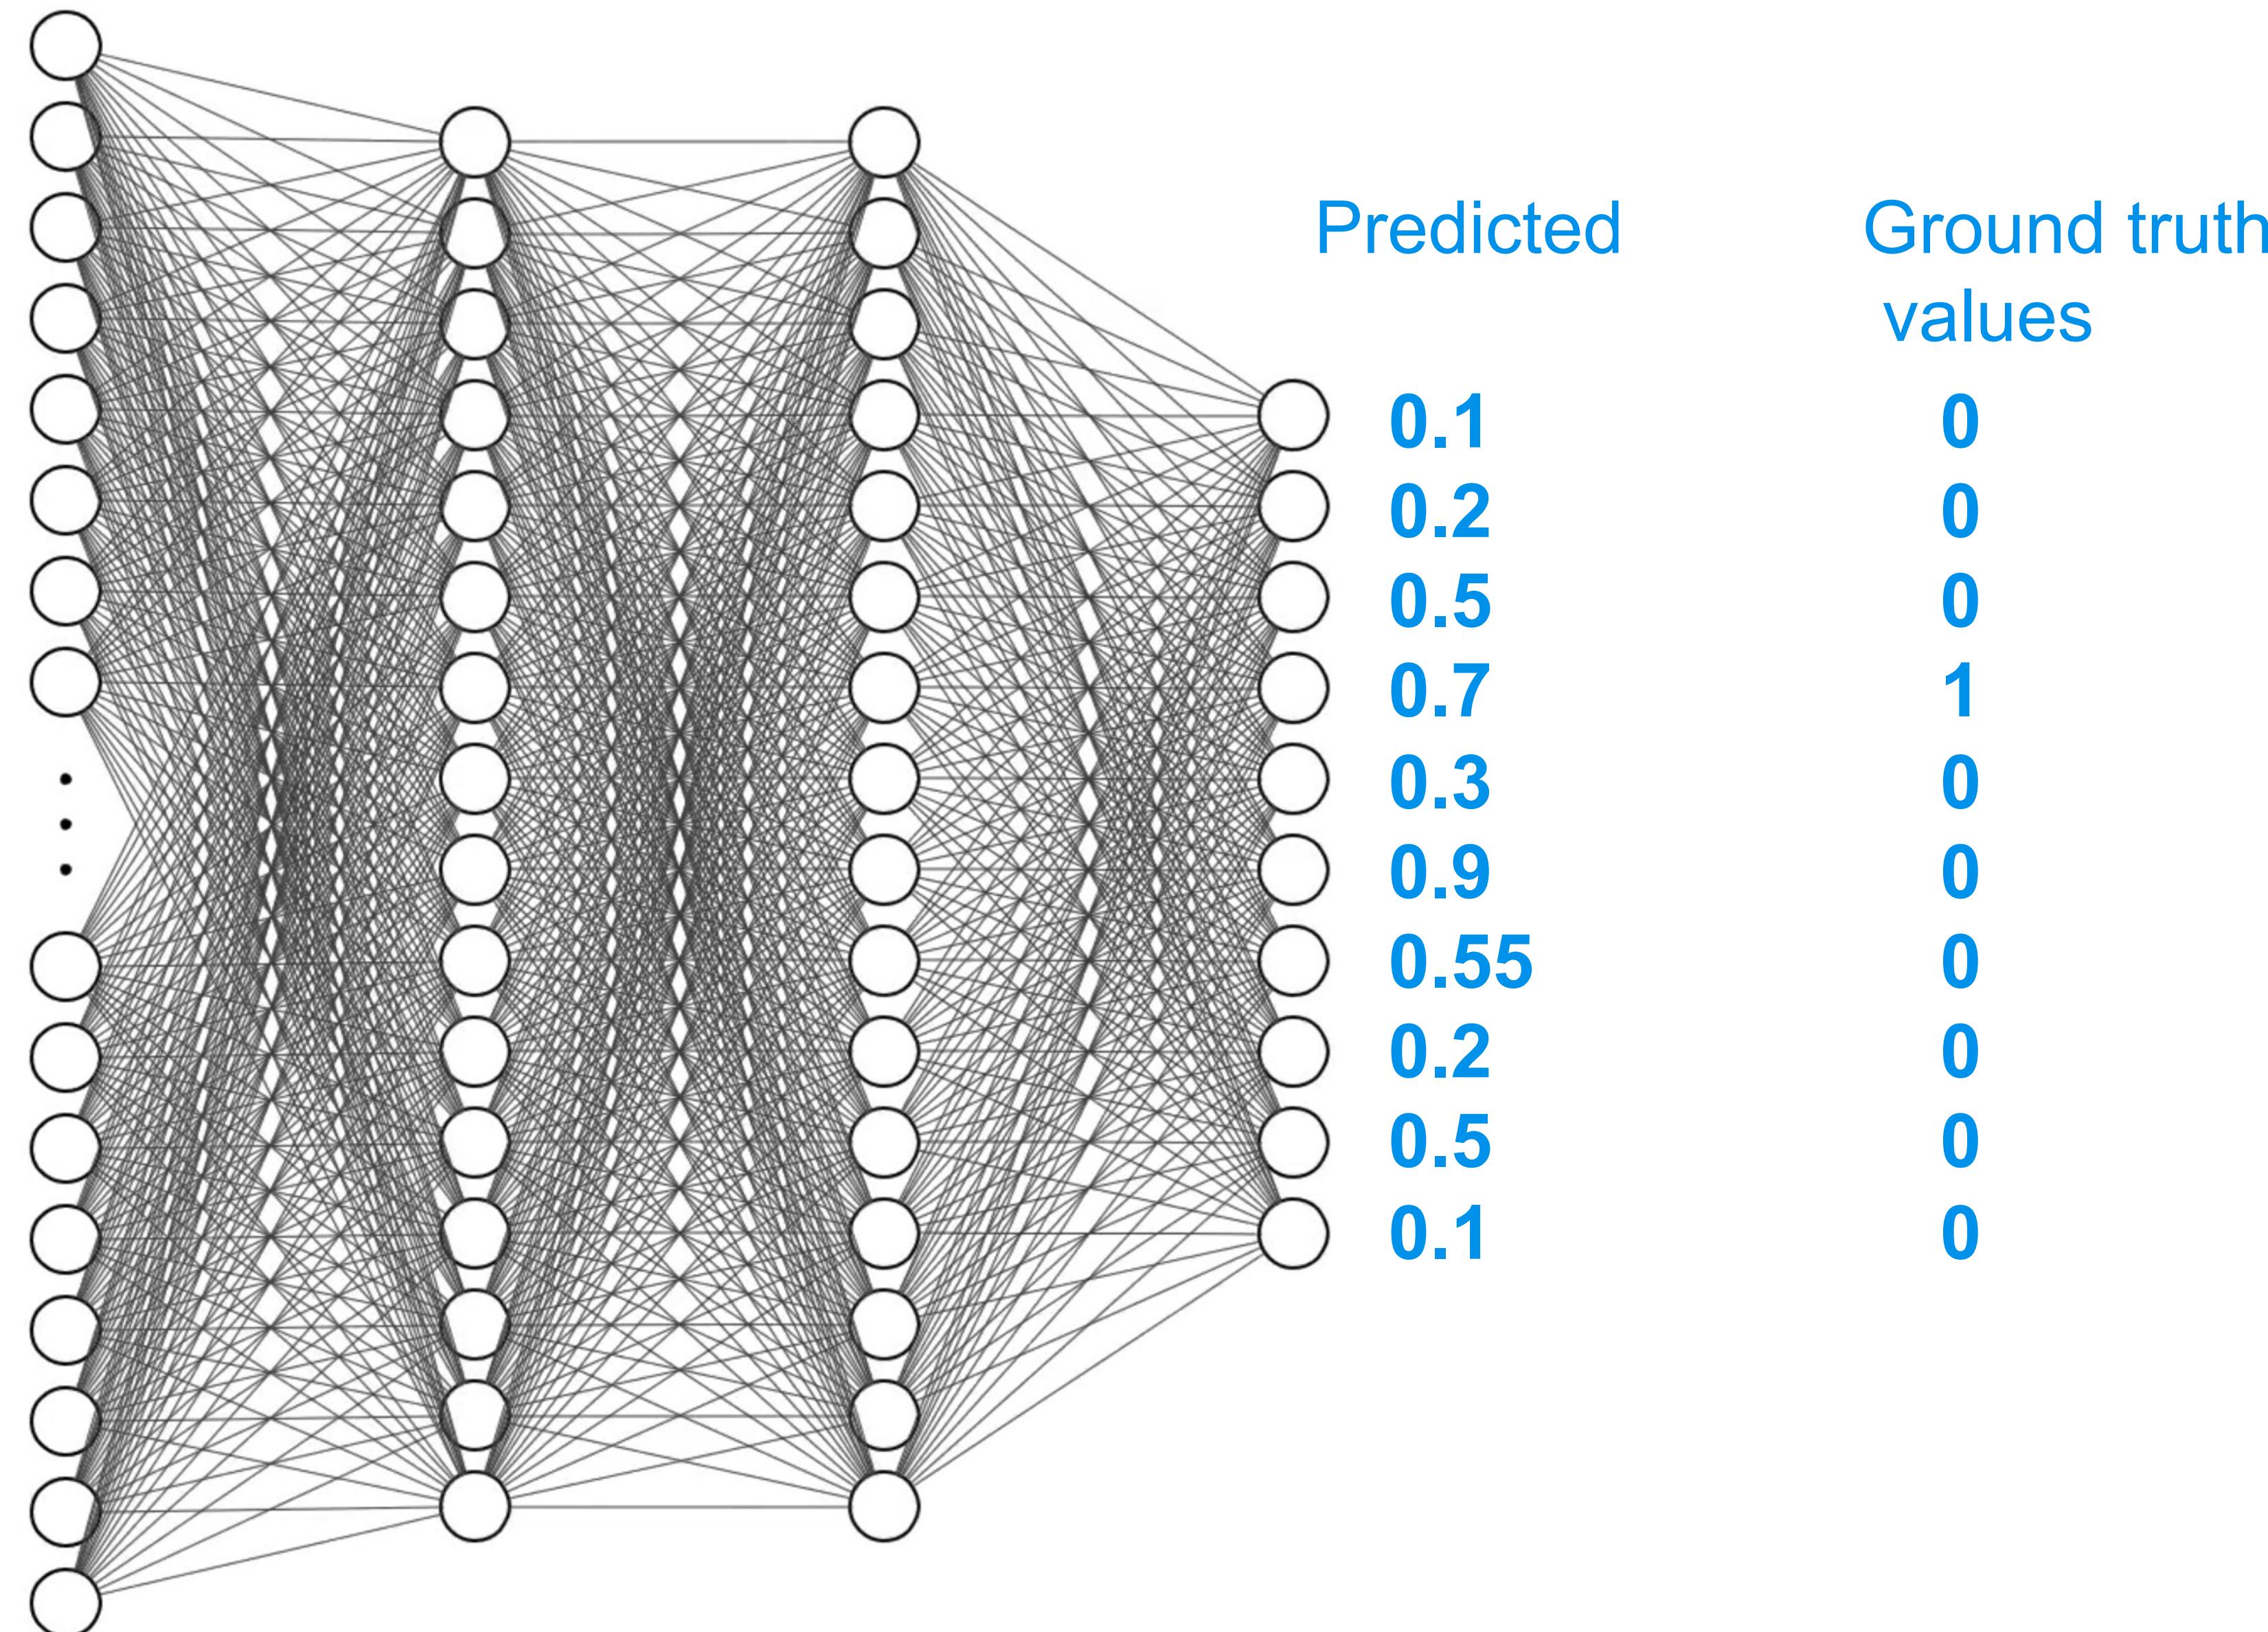
At the beginning of the training process, the weights are randomly initialised

Cost function



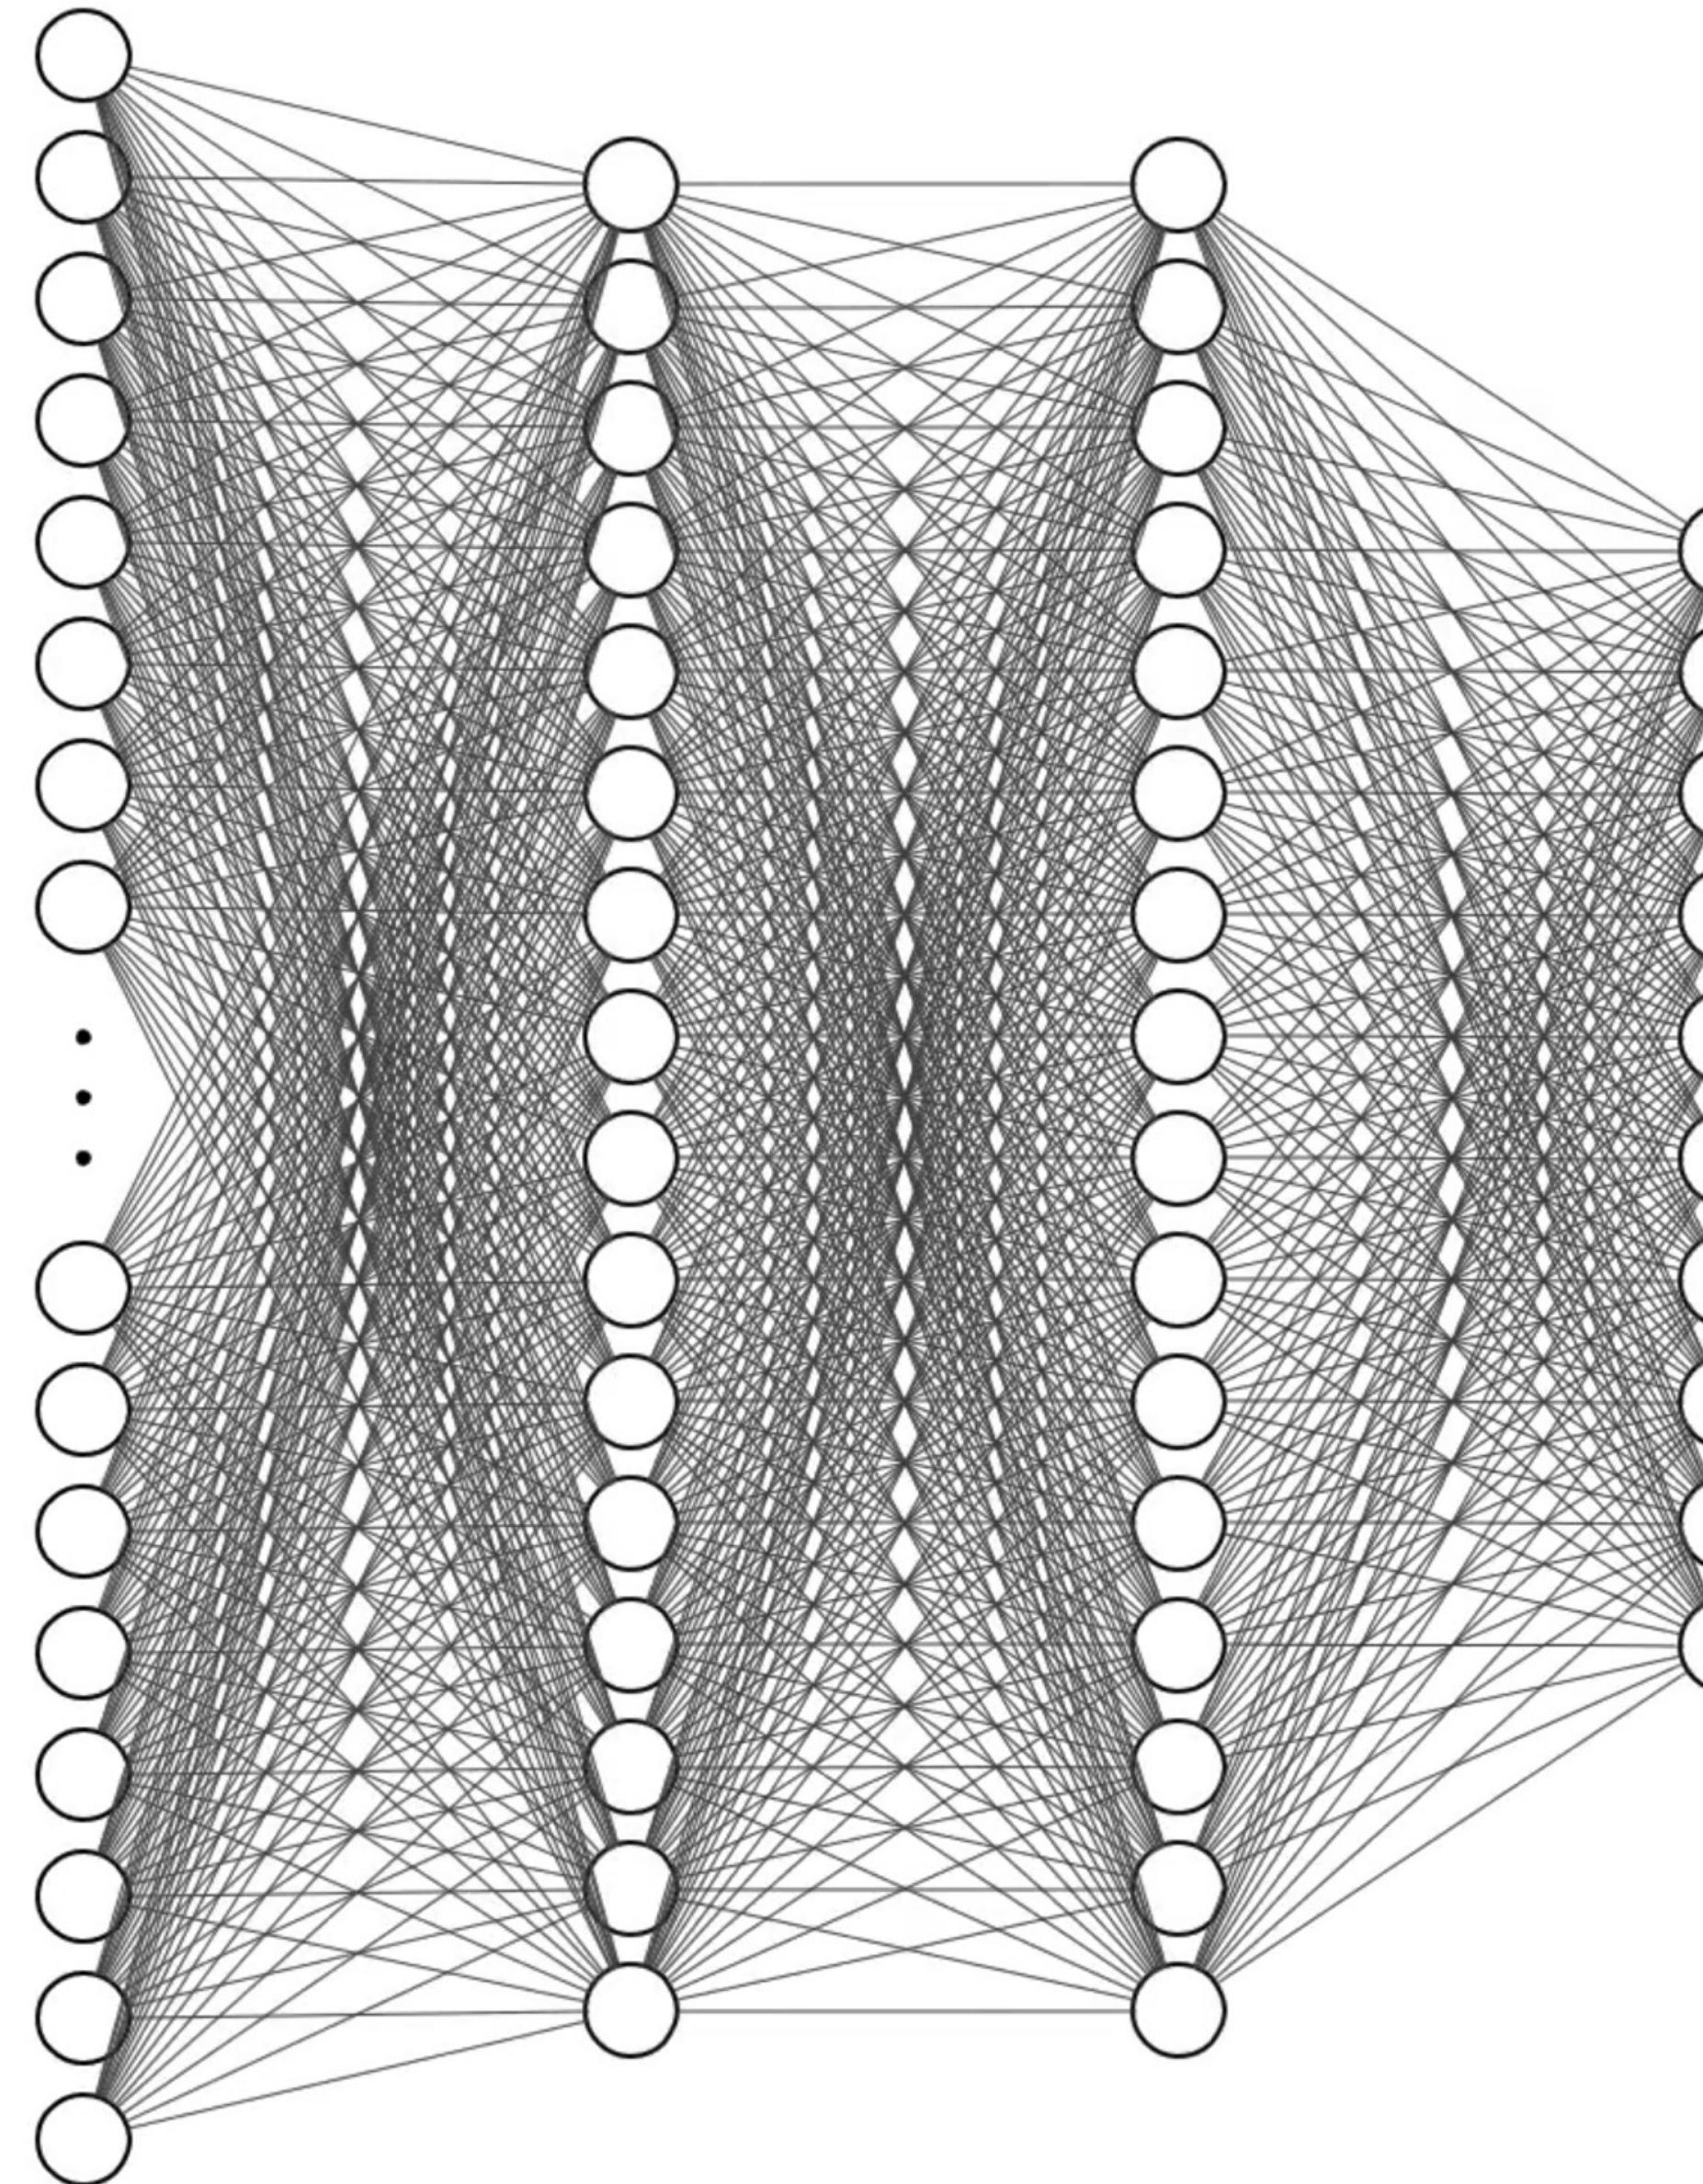
Cost function

3



Cost function

3



What is the cost difference?

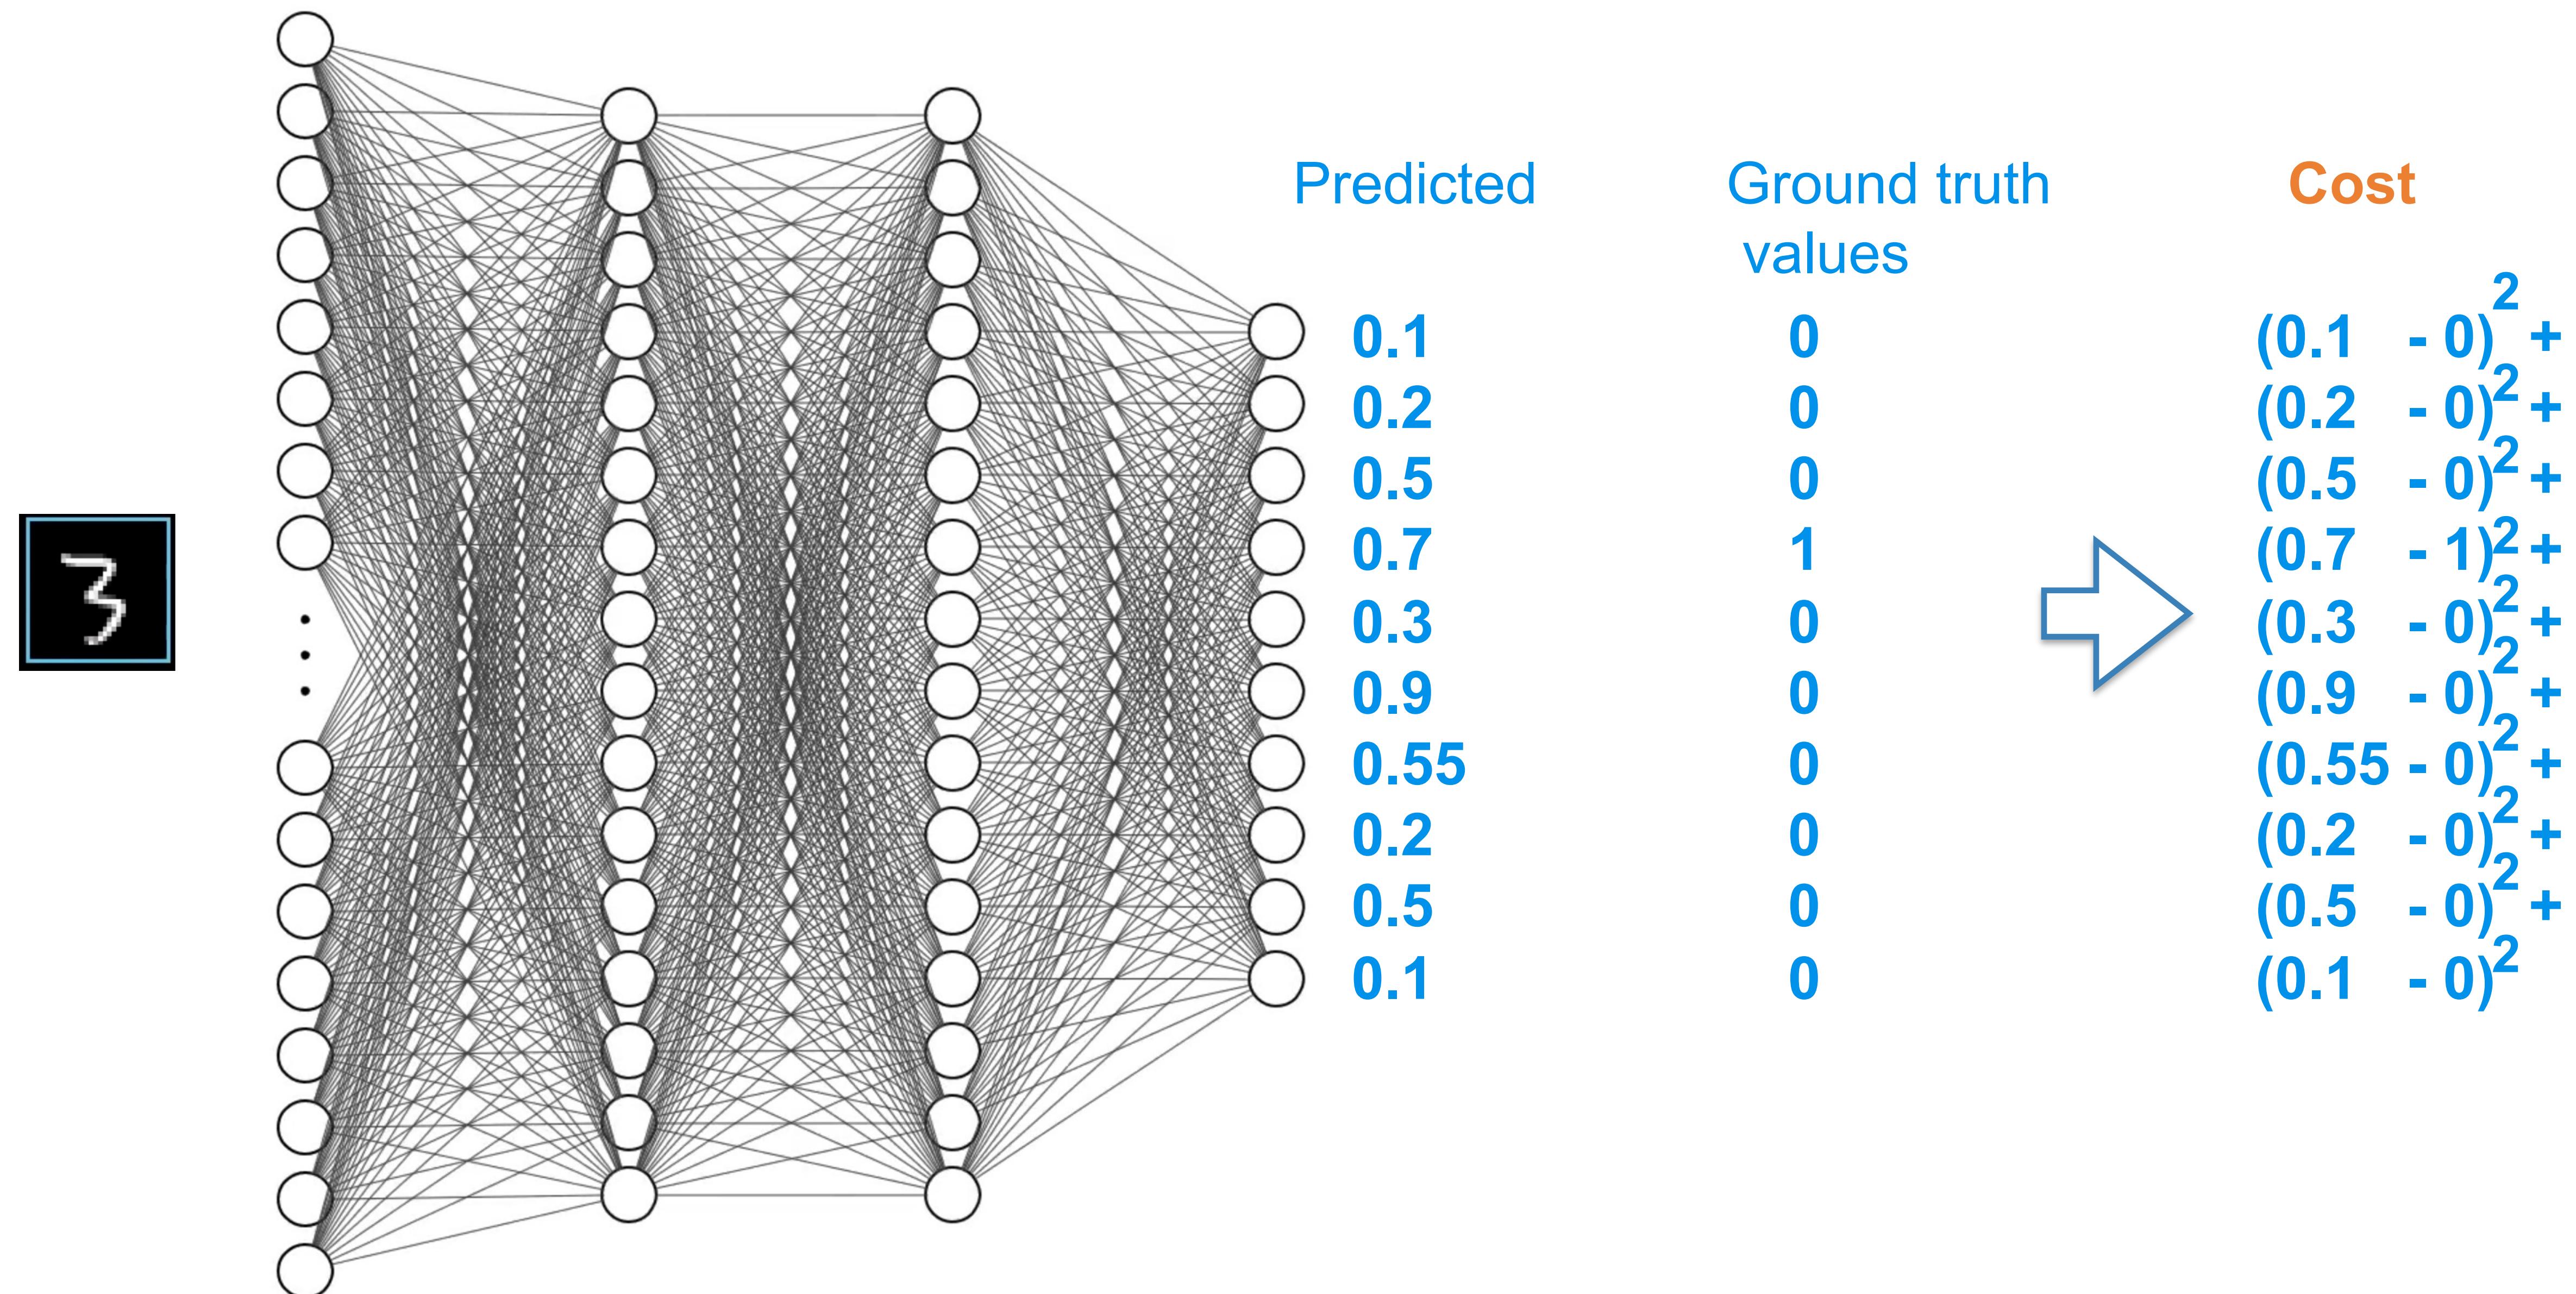
Predicted

0.1
0.2
0.5
0.7
0.3
0.9
0.55
0.2
0.5
0.1

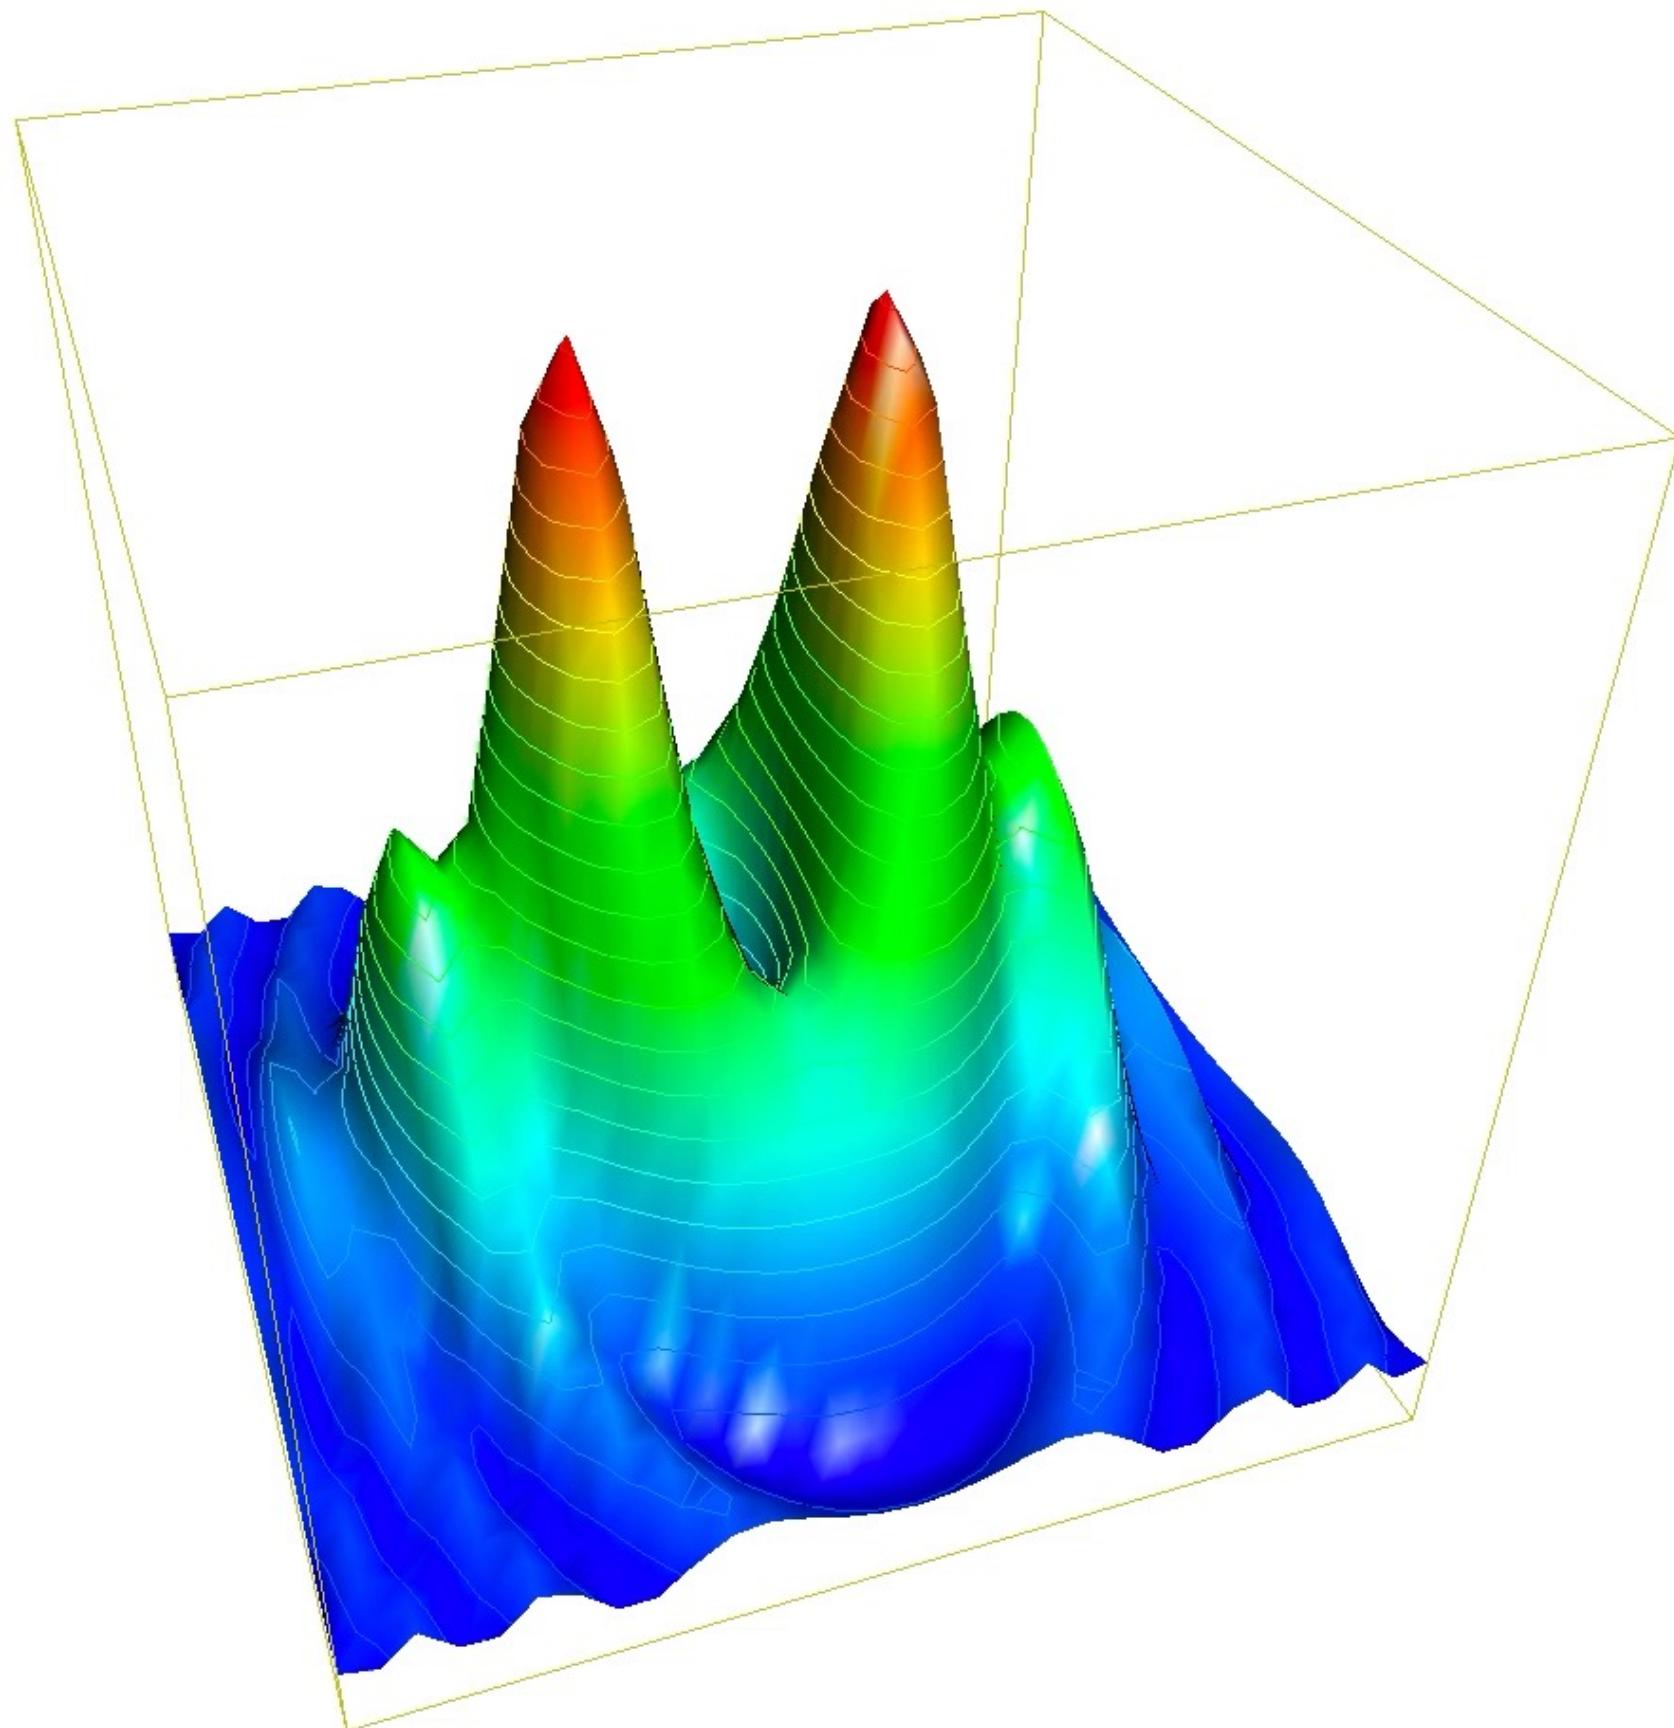
Ground truth
values

0
0
0
1
0
0
0
0
0
0

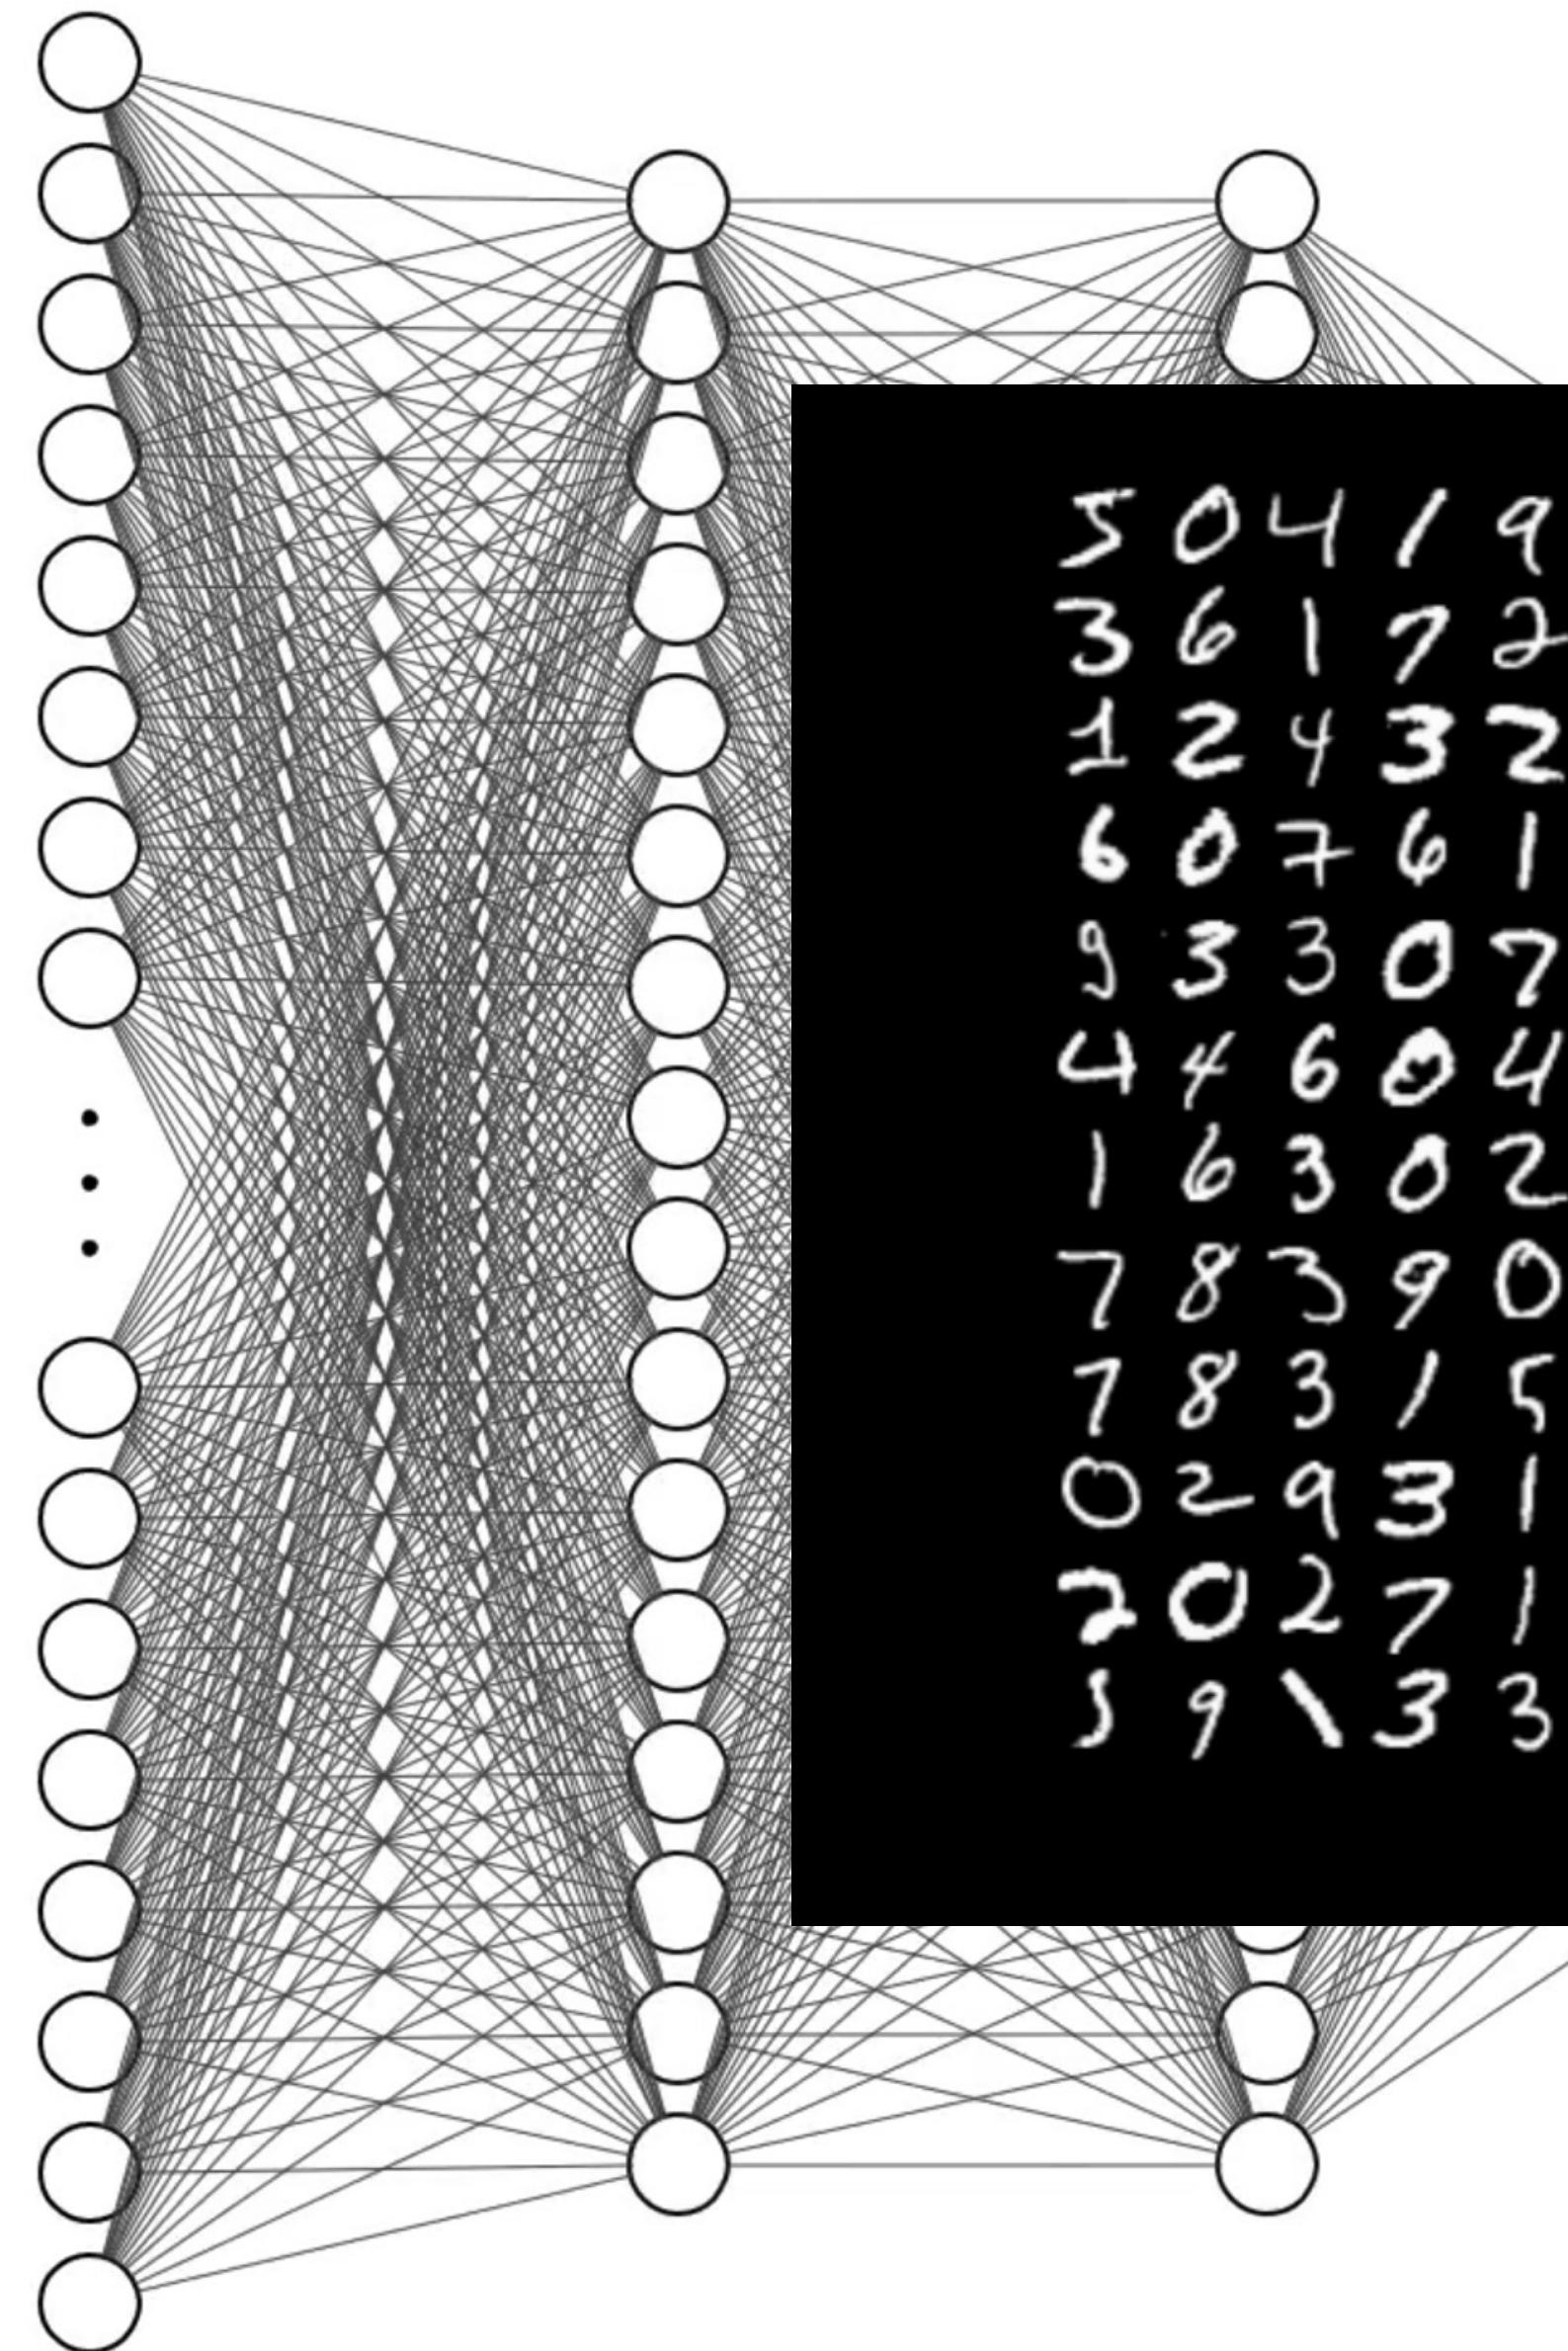
Cost function



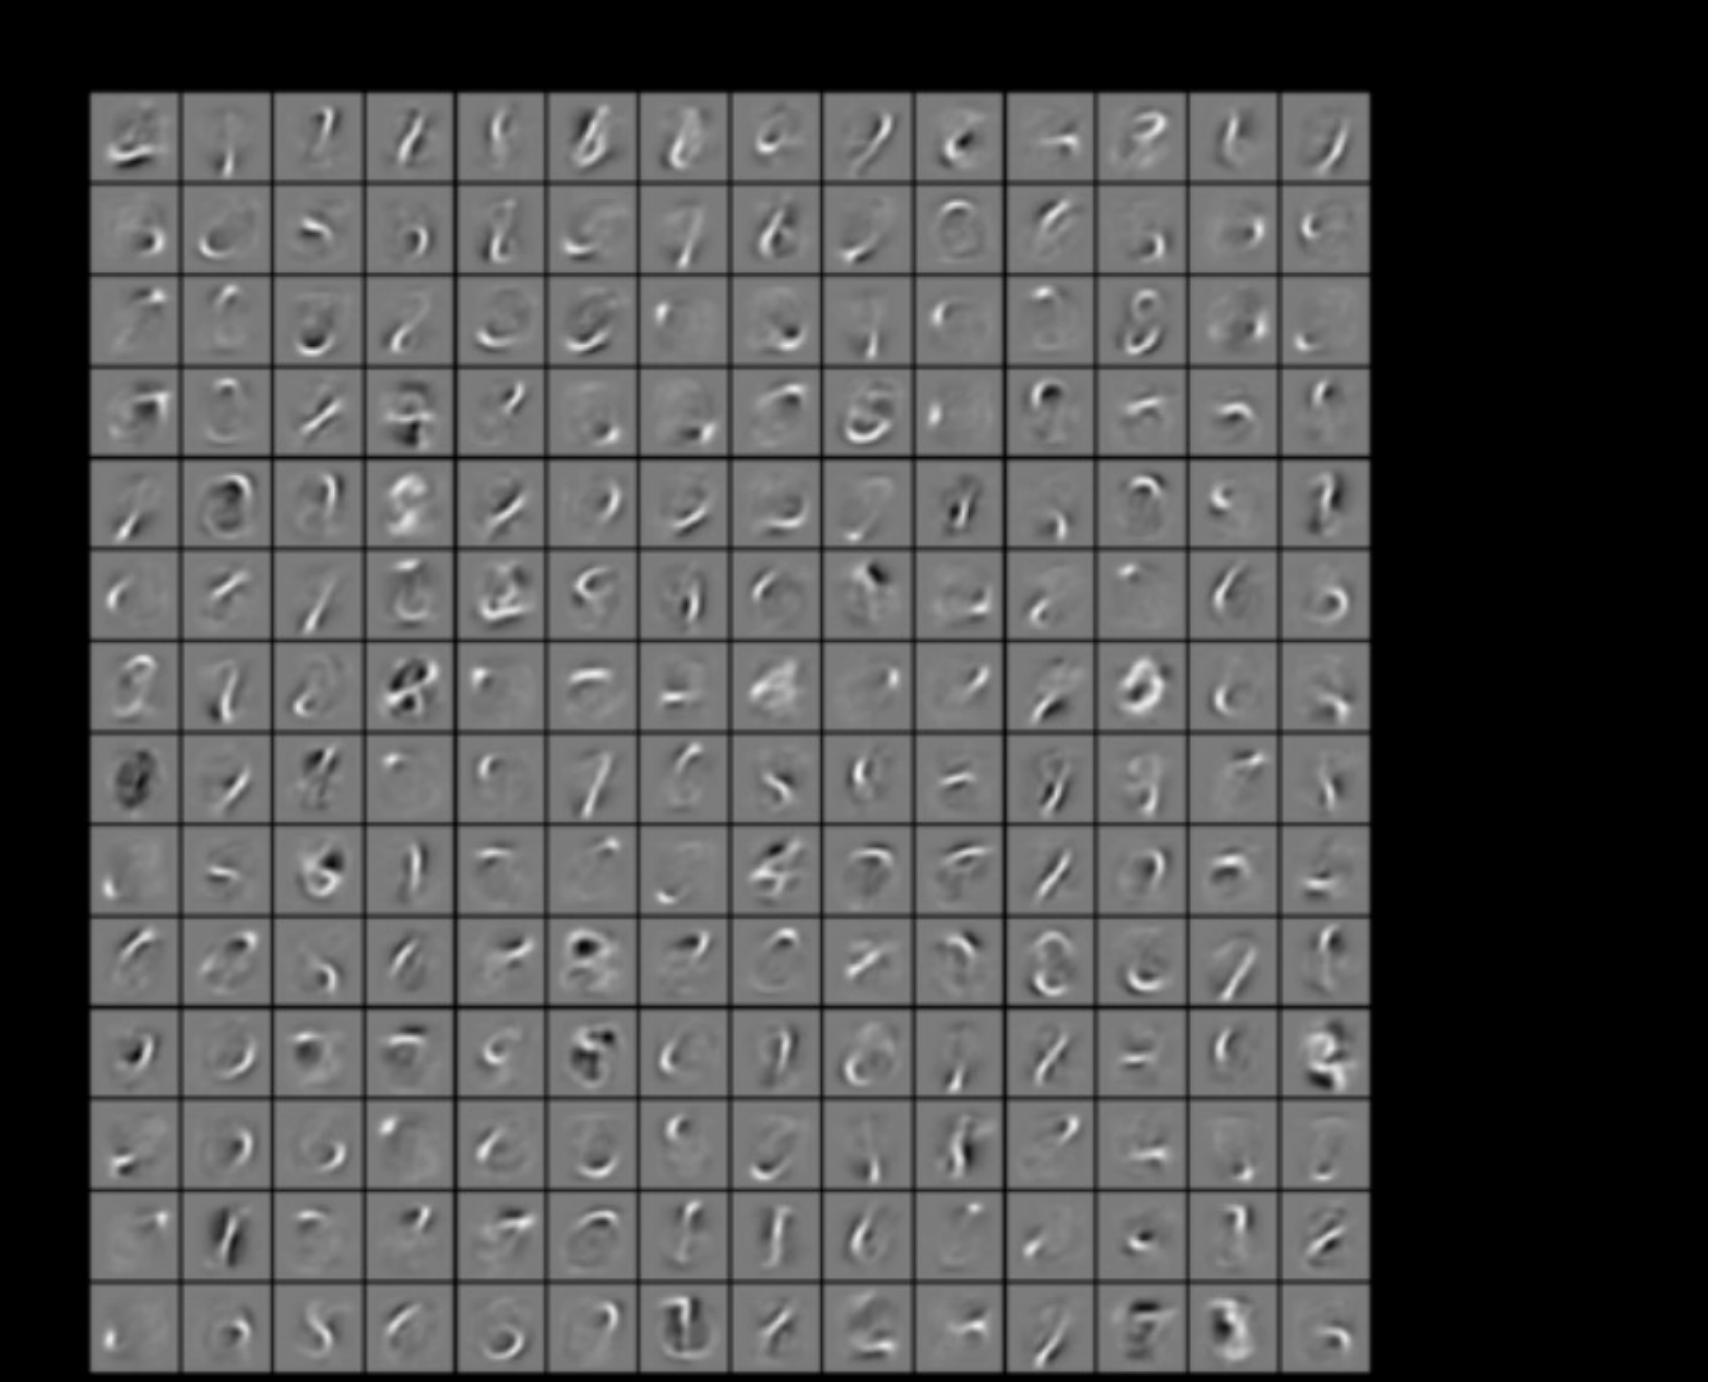
Gradient descent



3

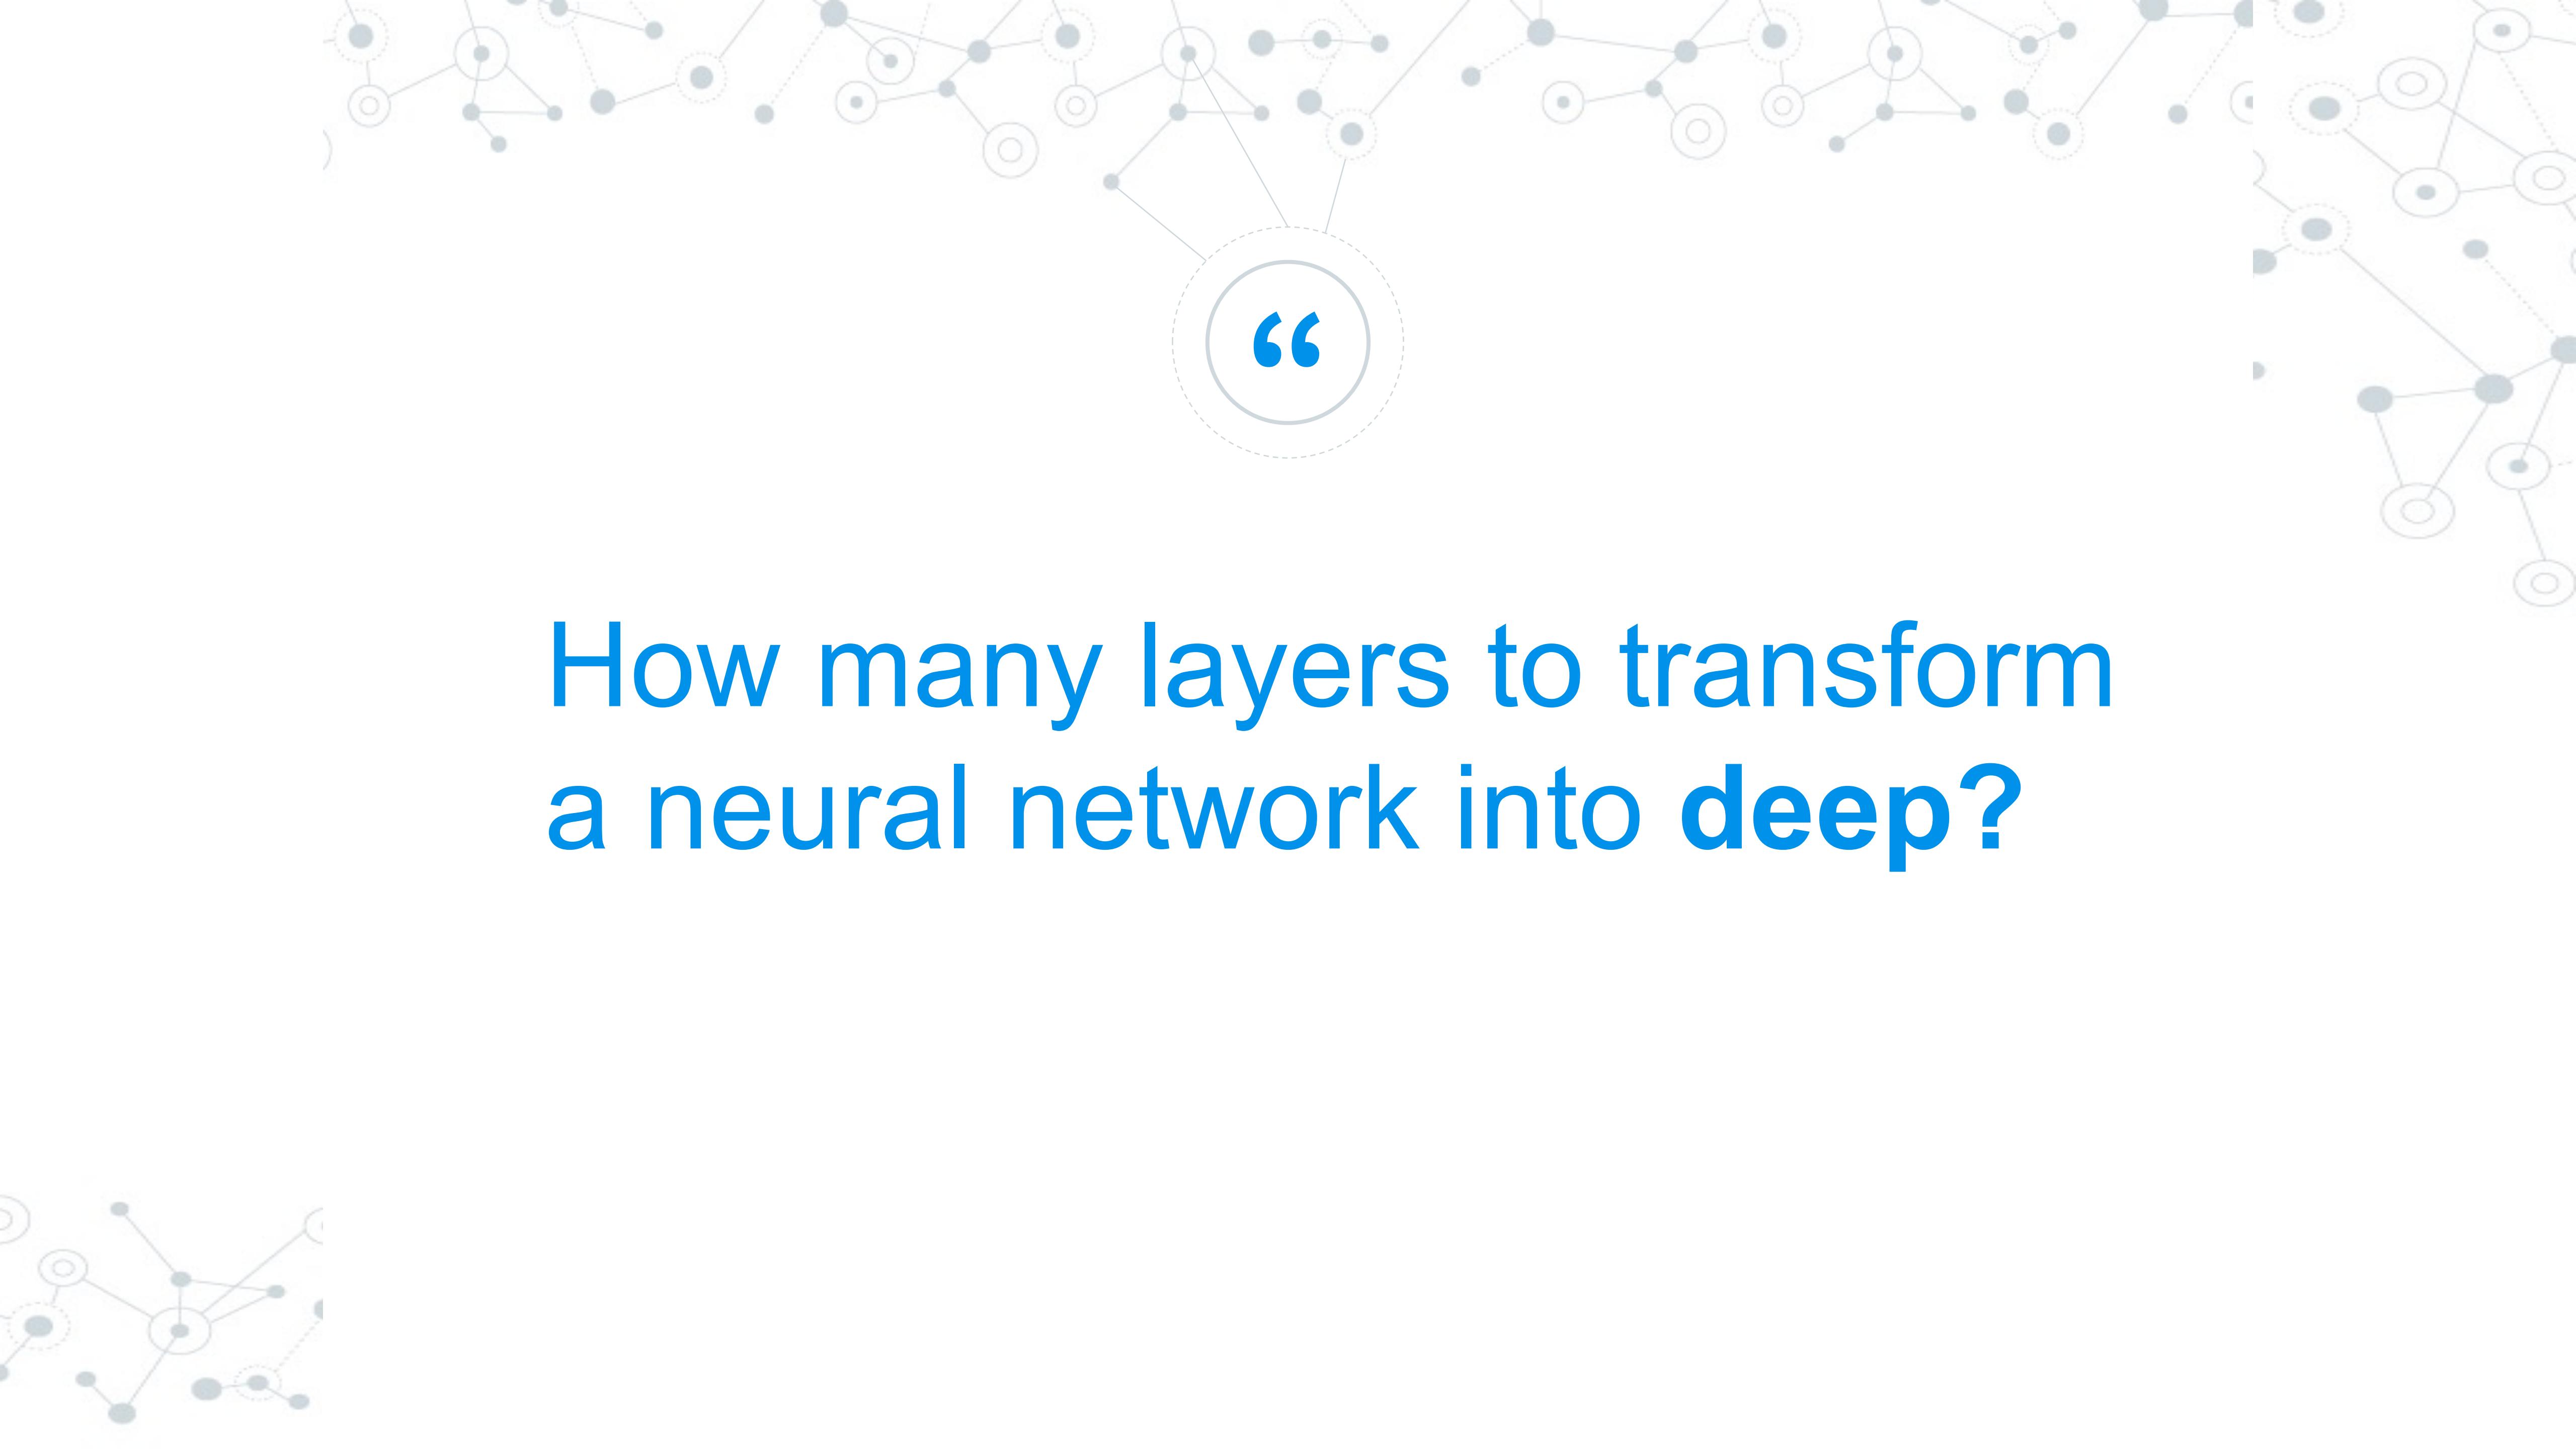


504192131435
361728694091
124327386905
607618793985
933074980941
446045610017
163021179026
783904674680
783157171163
029311049200
202718641634
39\338547742



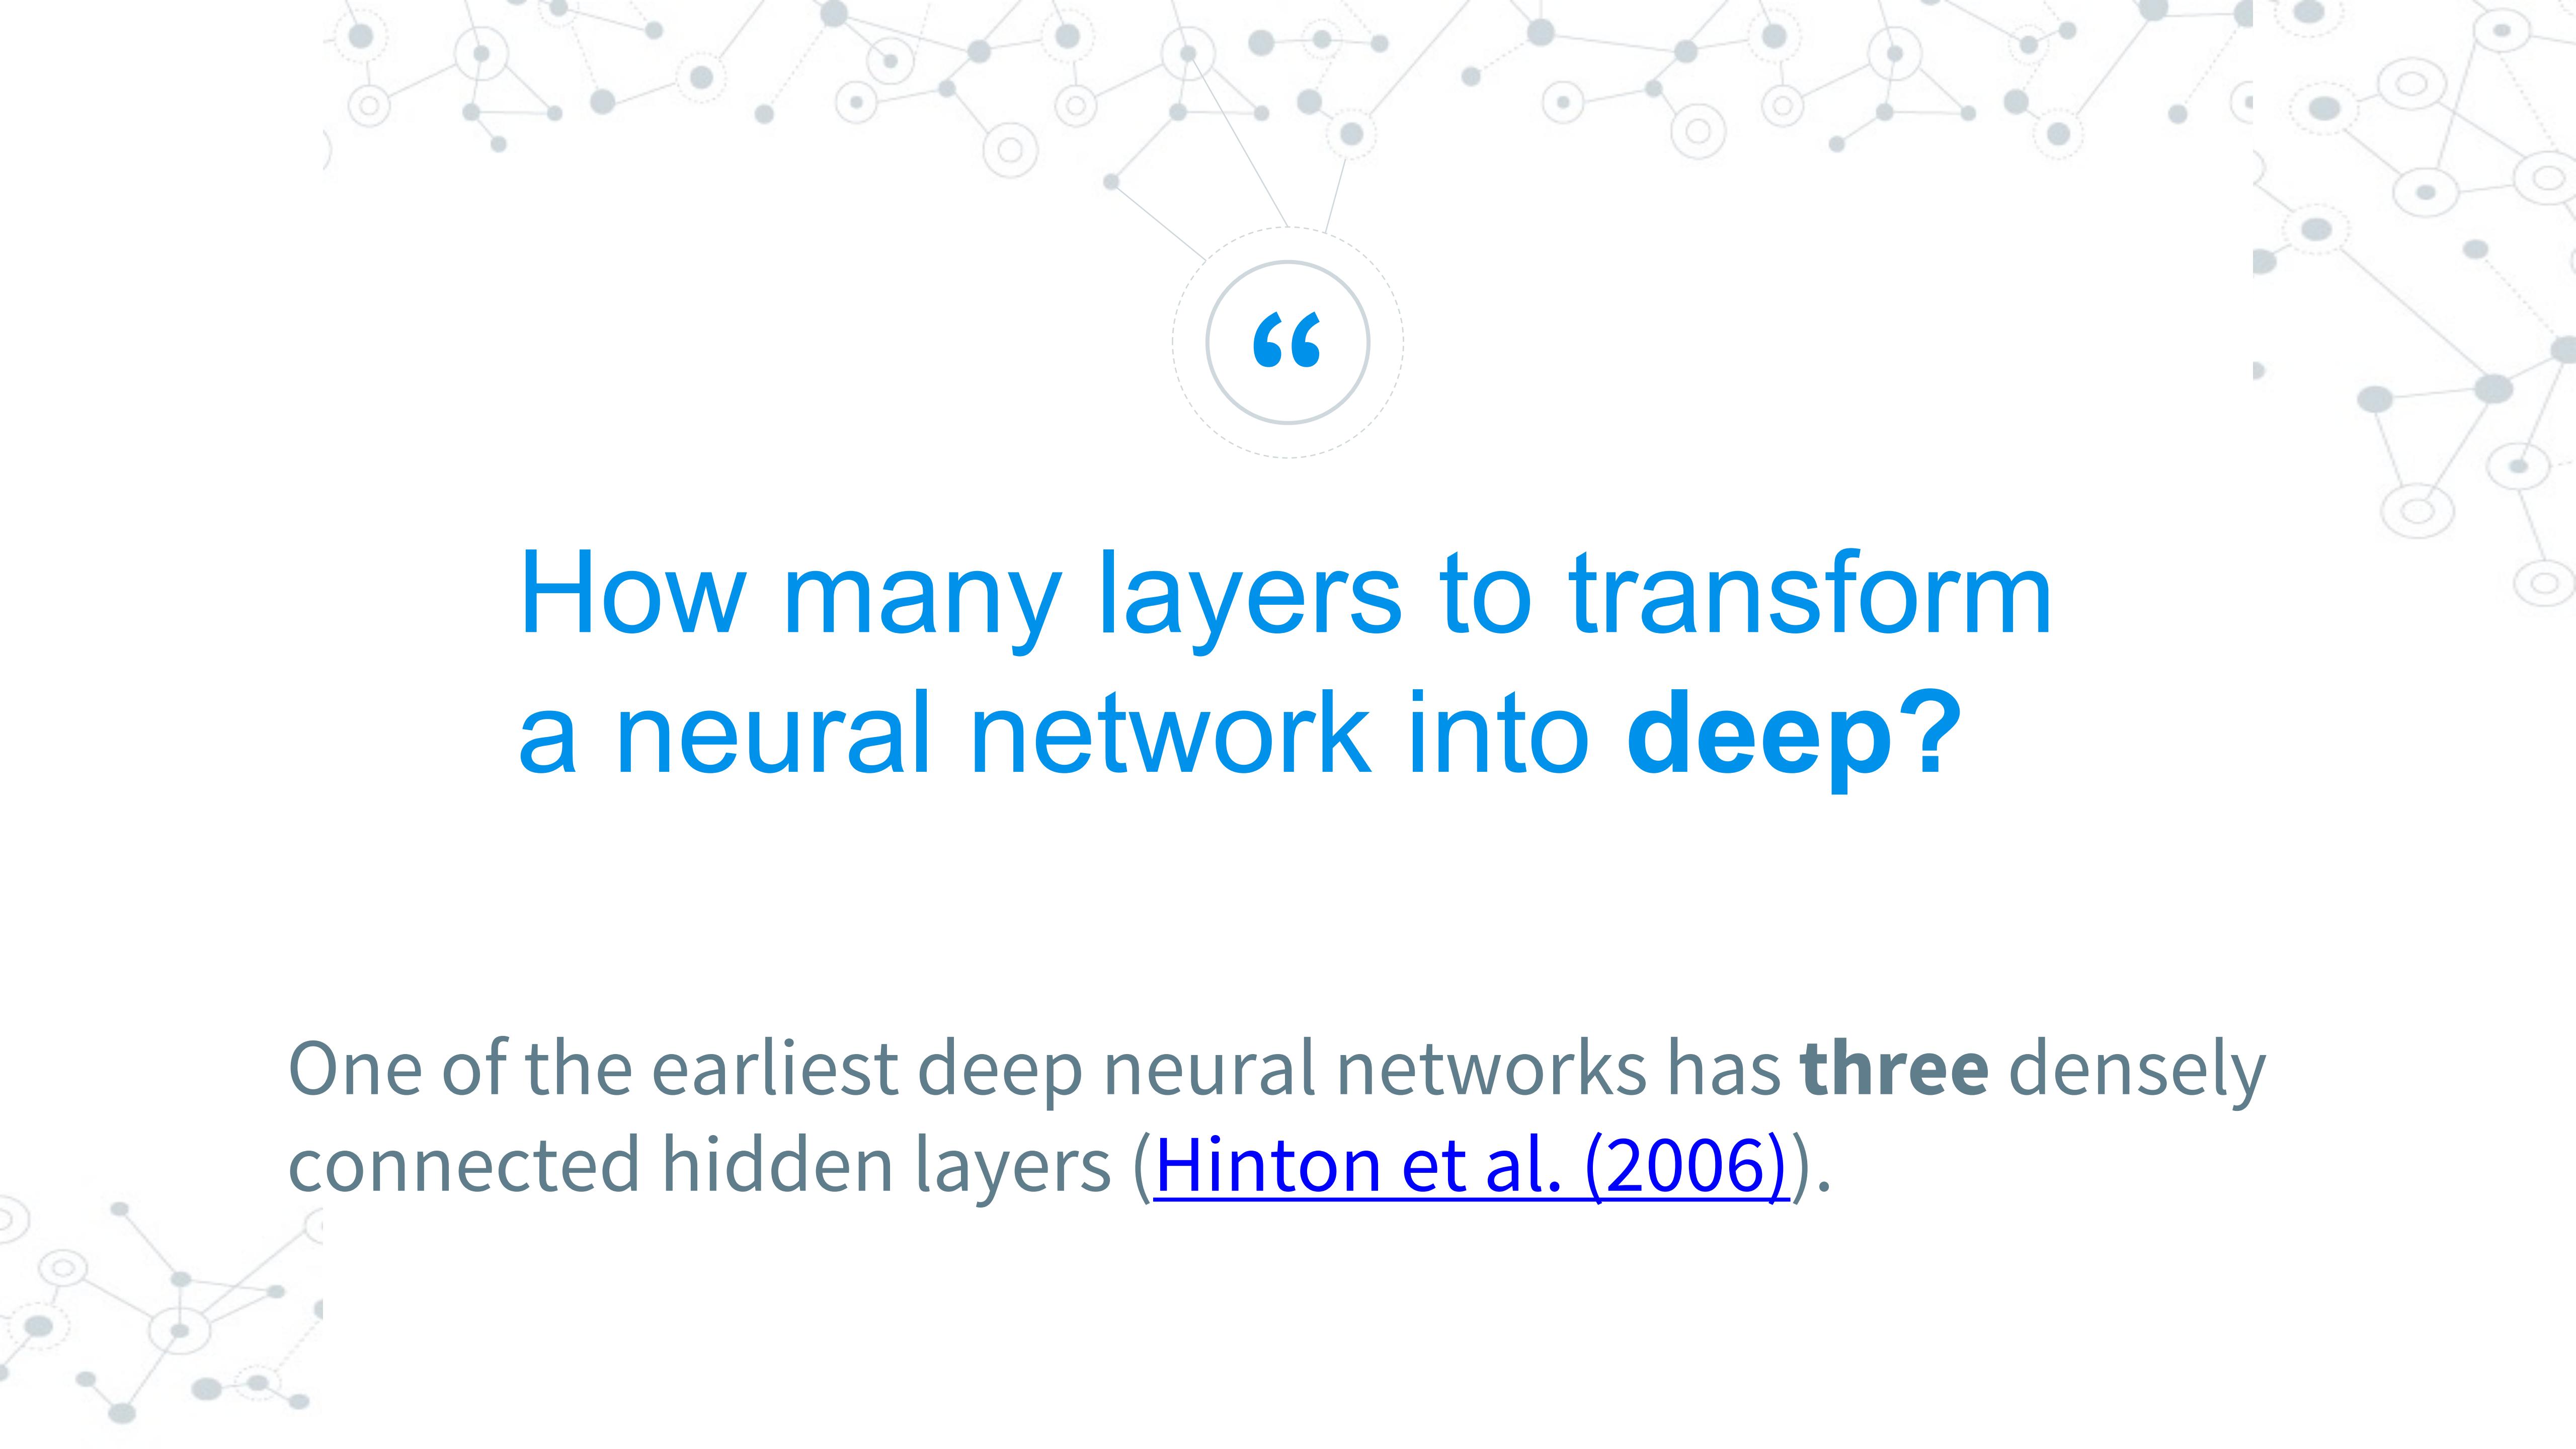


Deep learning



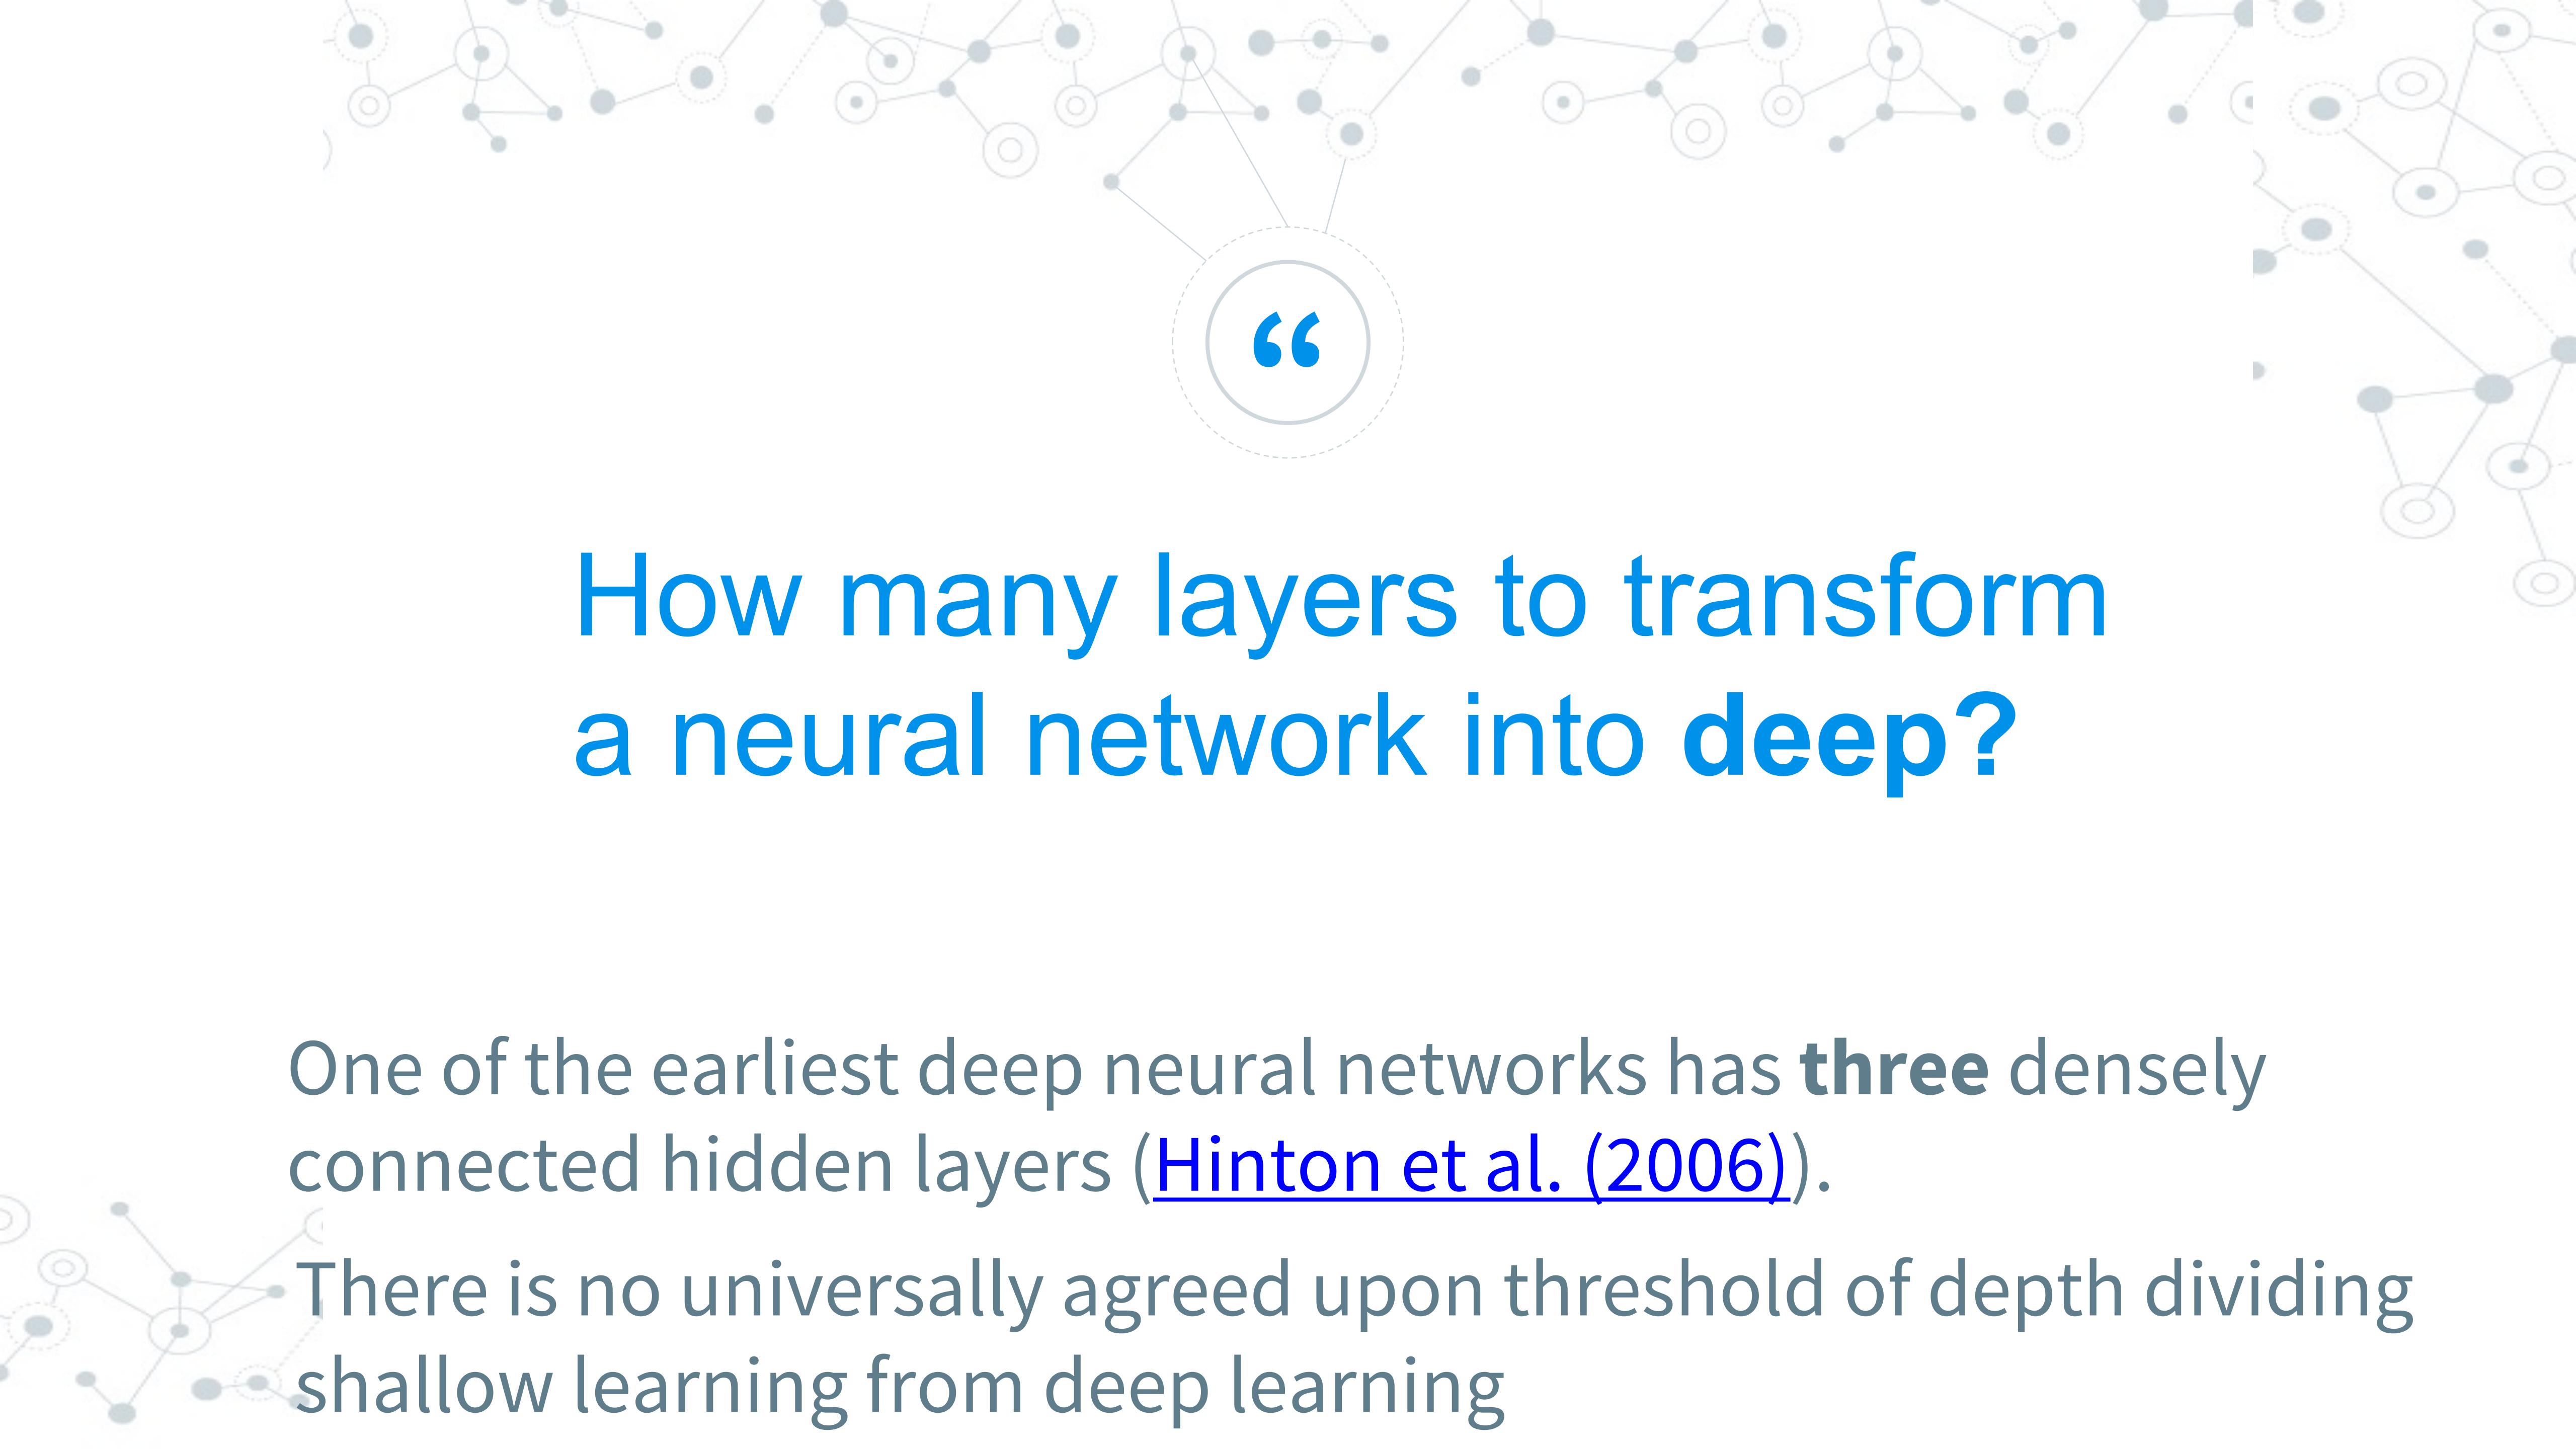
“

How many layers to transform a neural network into deep?



“ How many layers to transform a neural network into deep?

One of the earliest deep neural networks has **three** densely connected hidden layers ([Hinton et al. \(2006\)](#)).

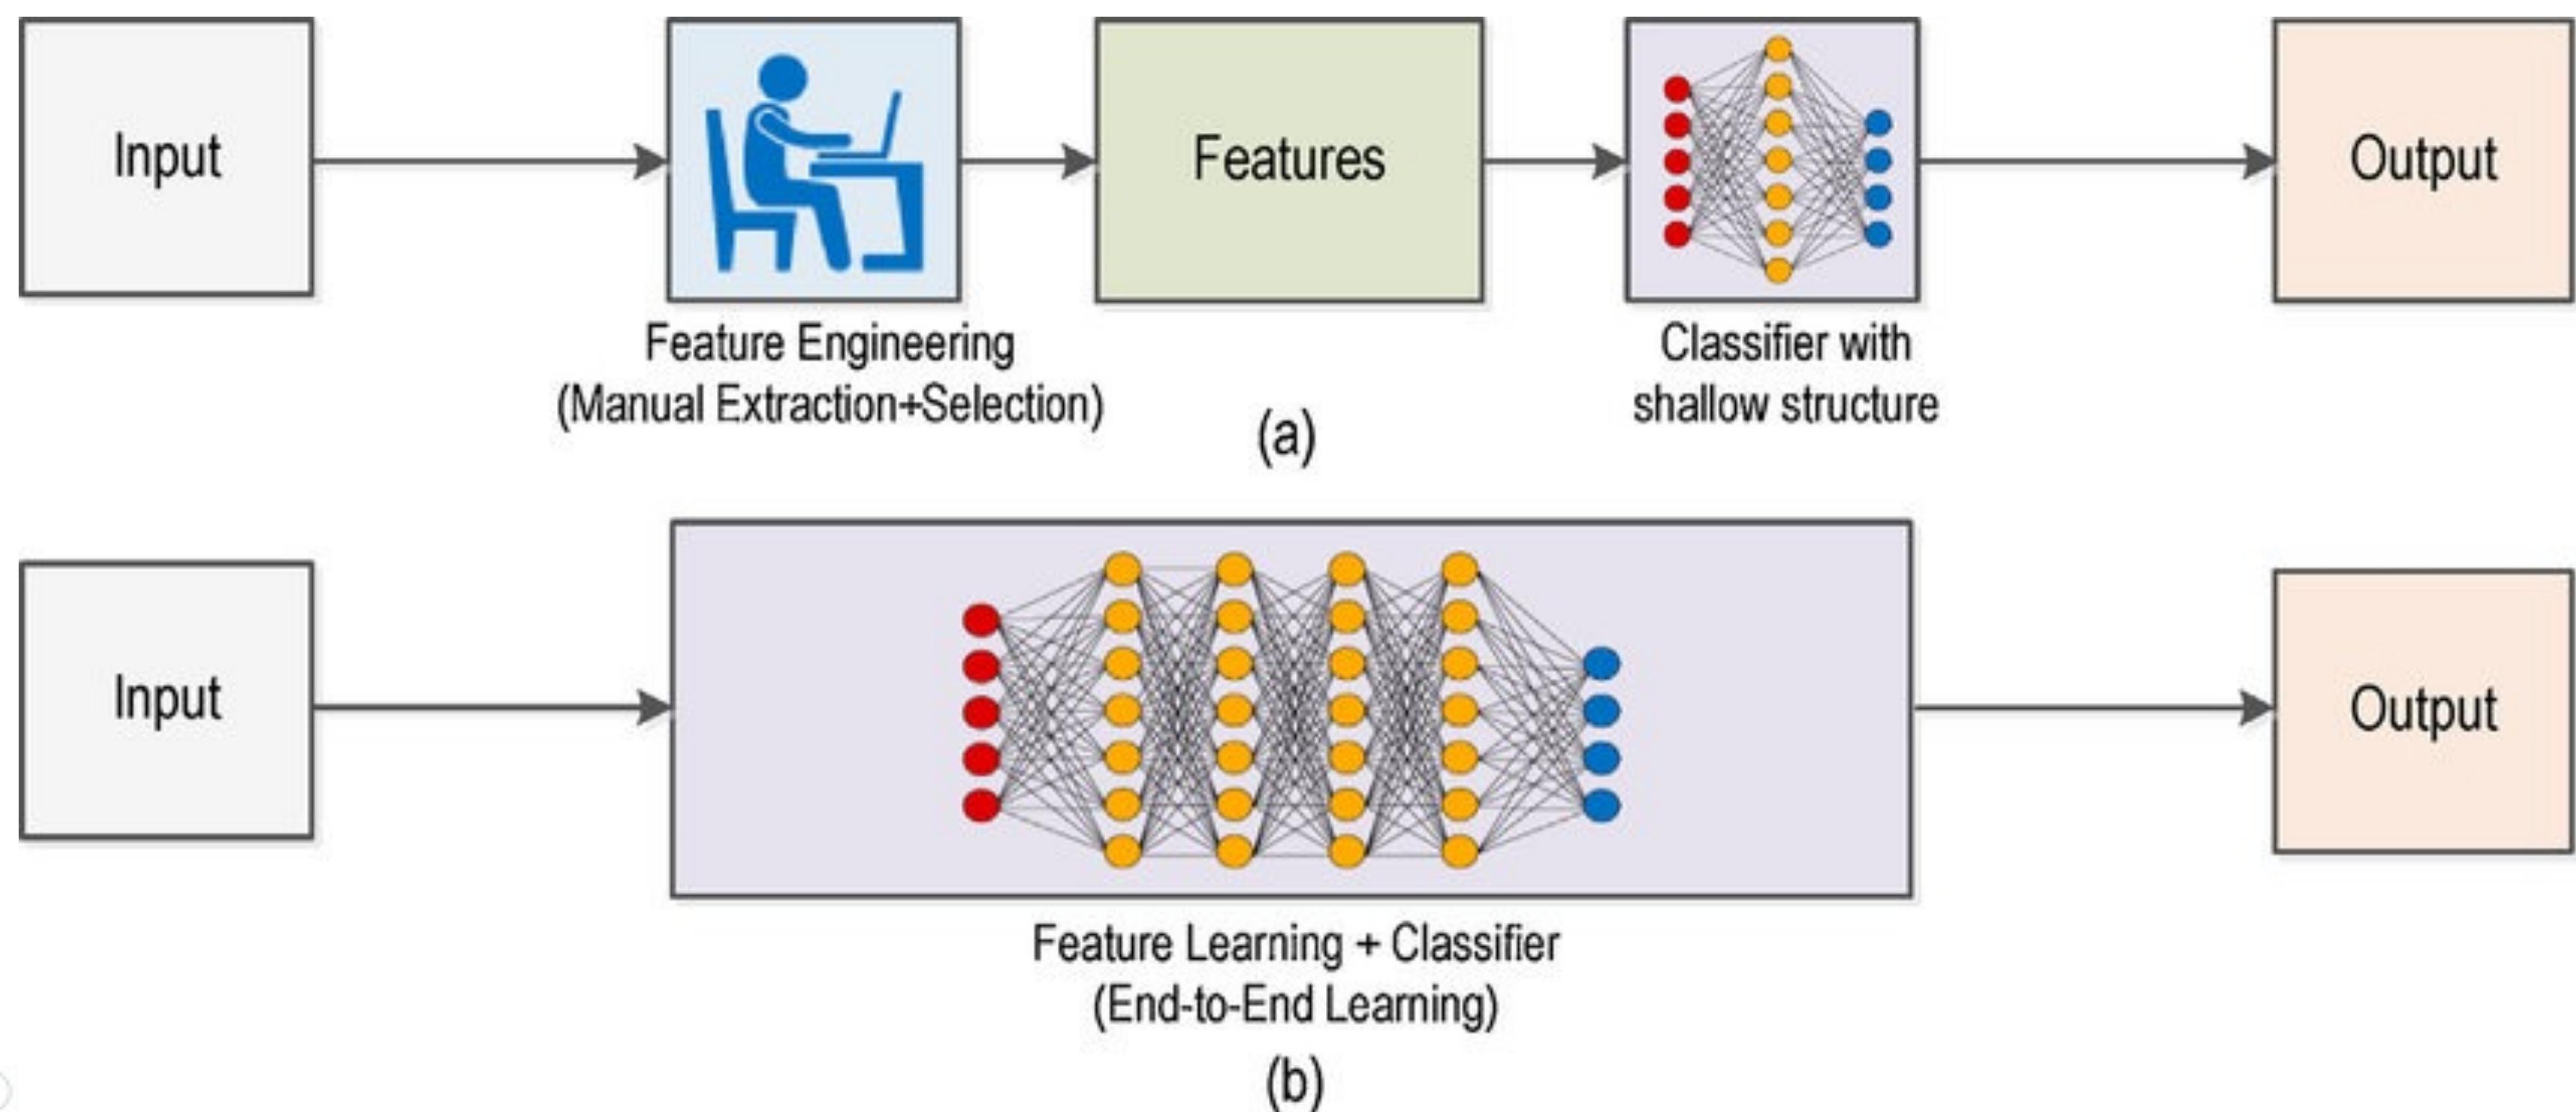


How many layers to transform a neural network into deep?

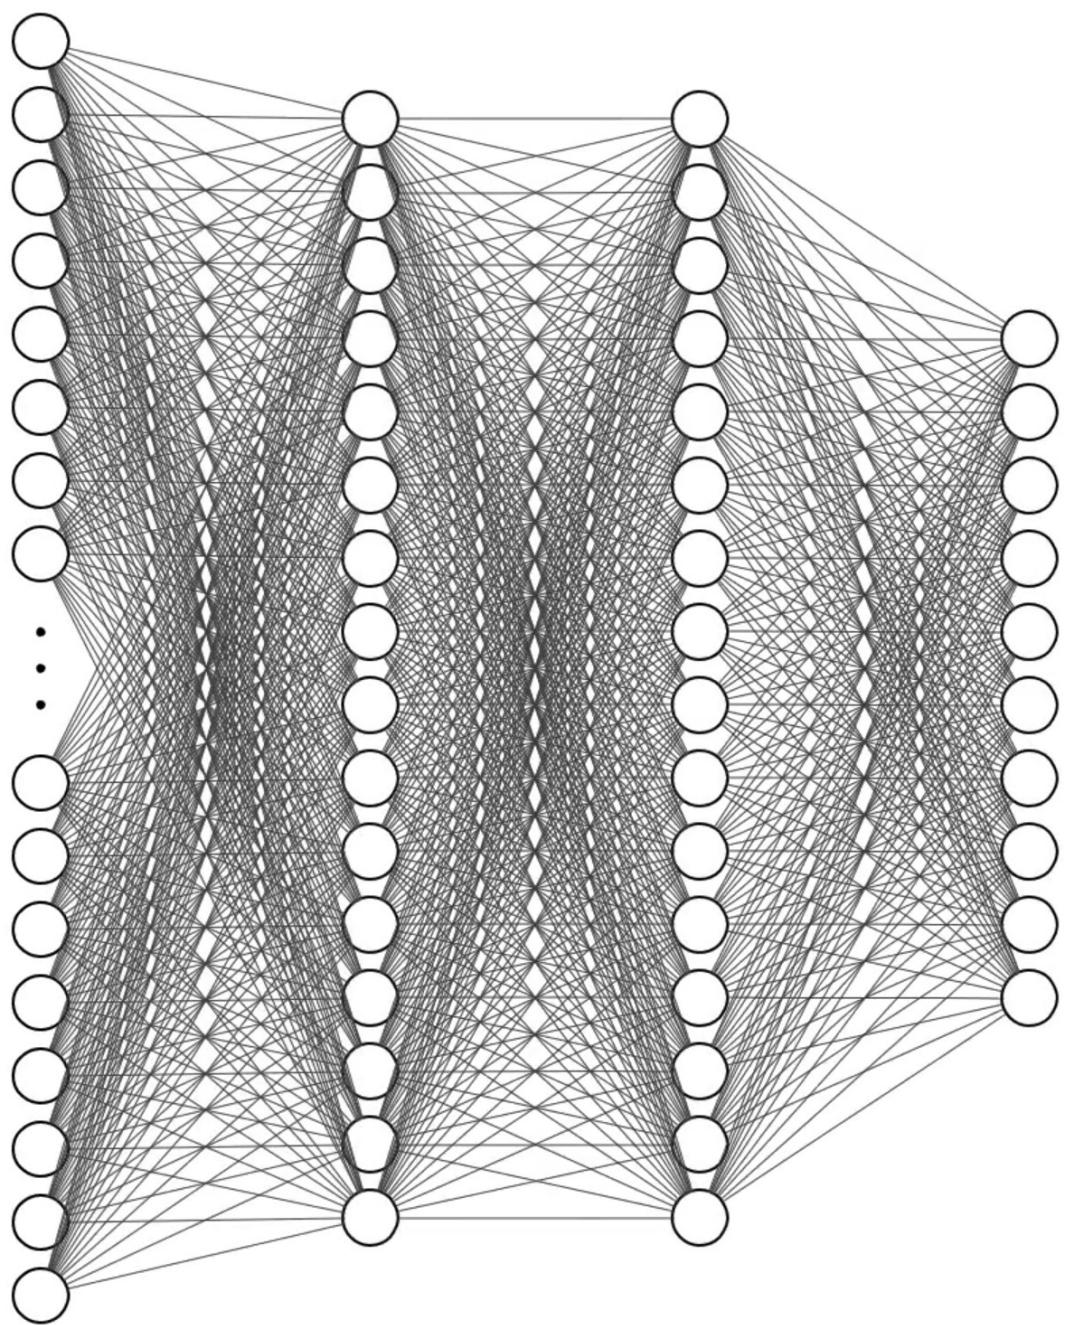
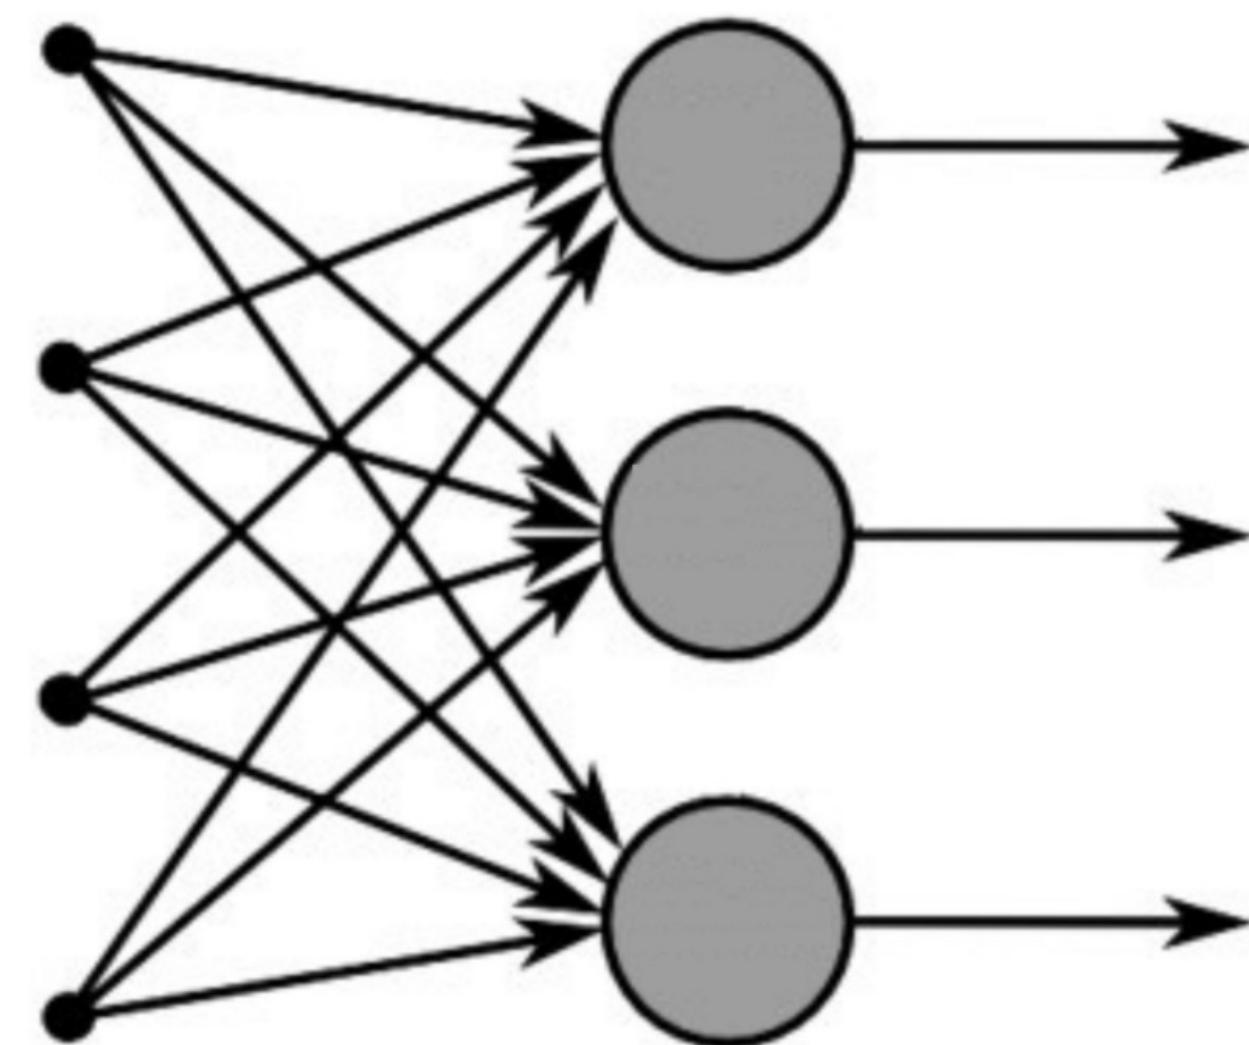
One of the earliest deep neural networks has **three** densely connected hidden layers ([Hinton et al. \(2006\)](#)).

There is no universally agreed upon threshold of depth dividing shallow learning from deep learning

Deep Learning vs Other Machine Learning Techniques

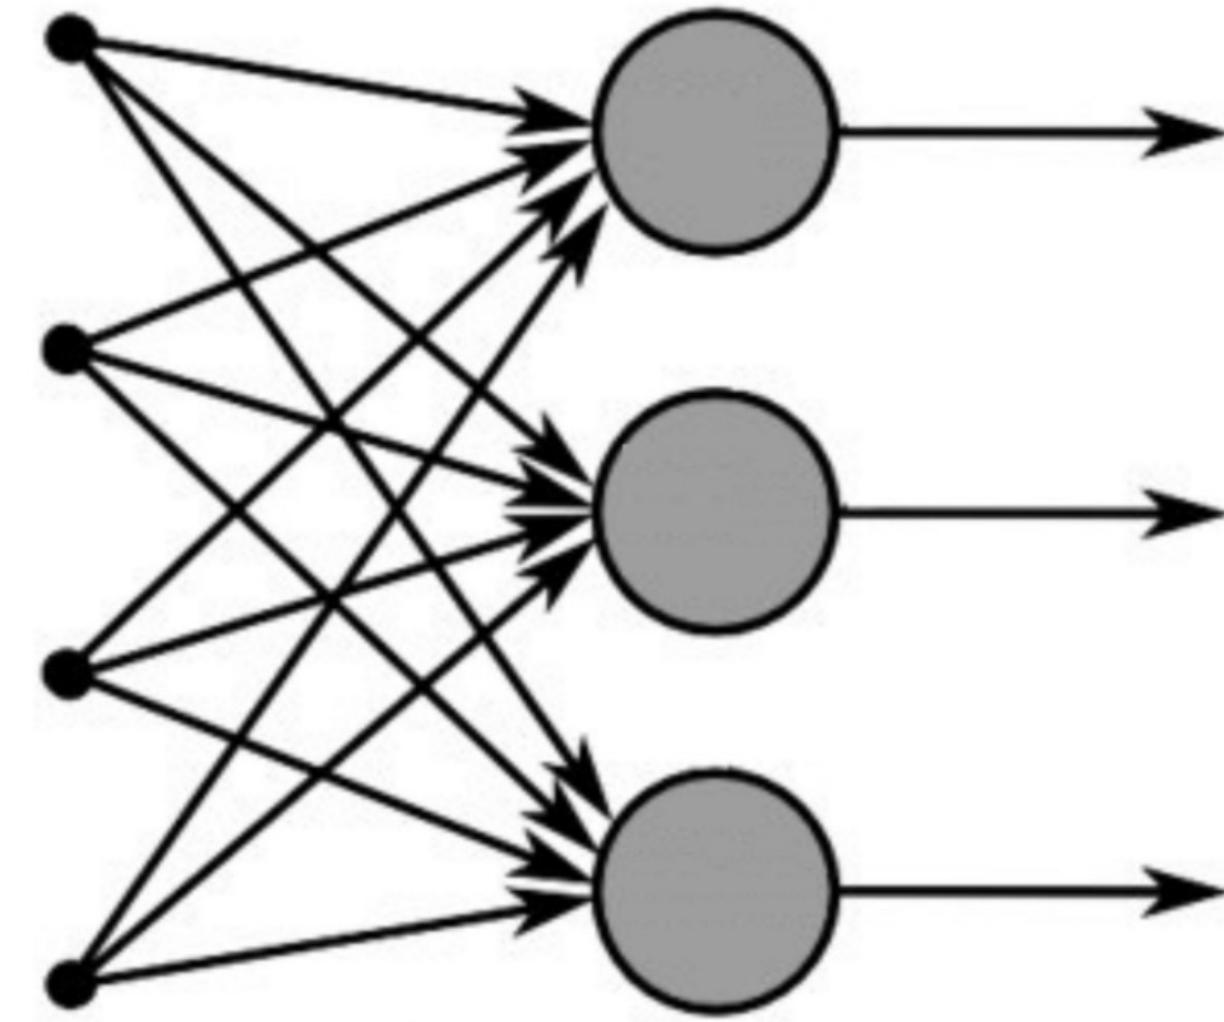


Examples of Neural Networks: Feed forward neural network



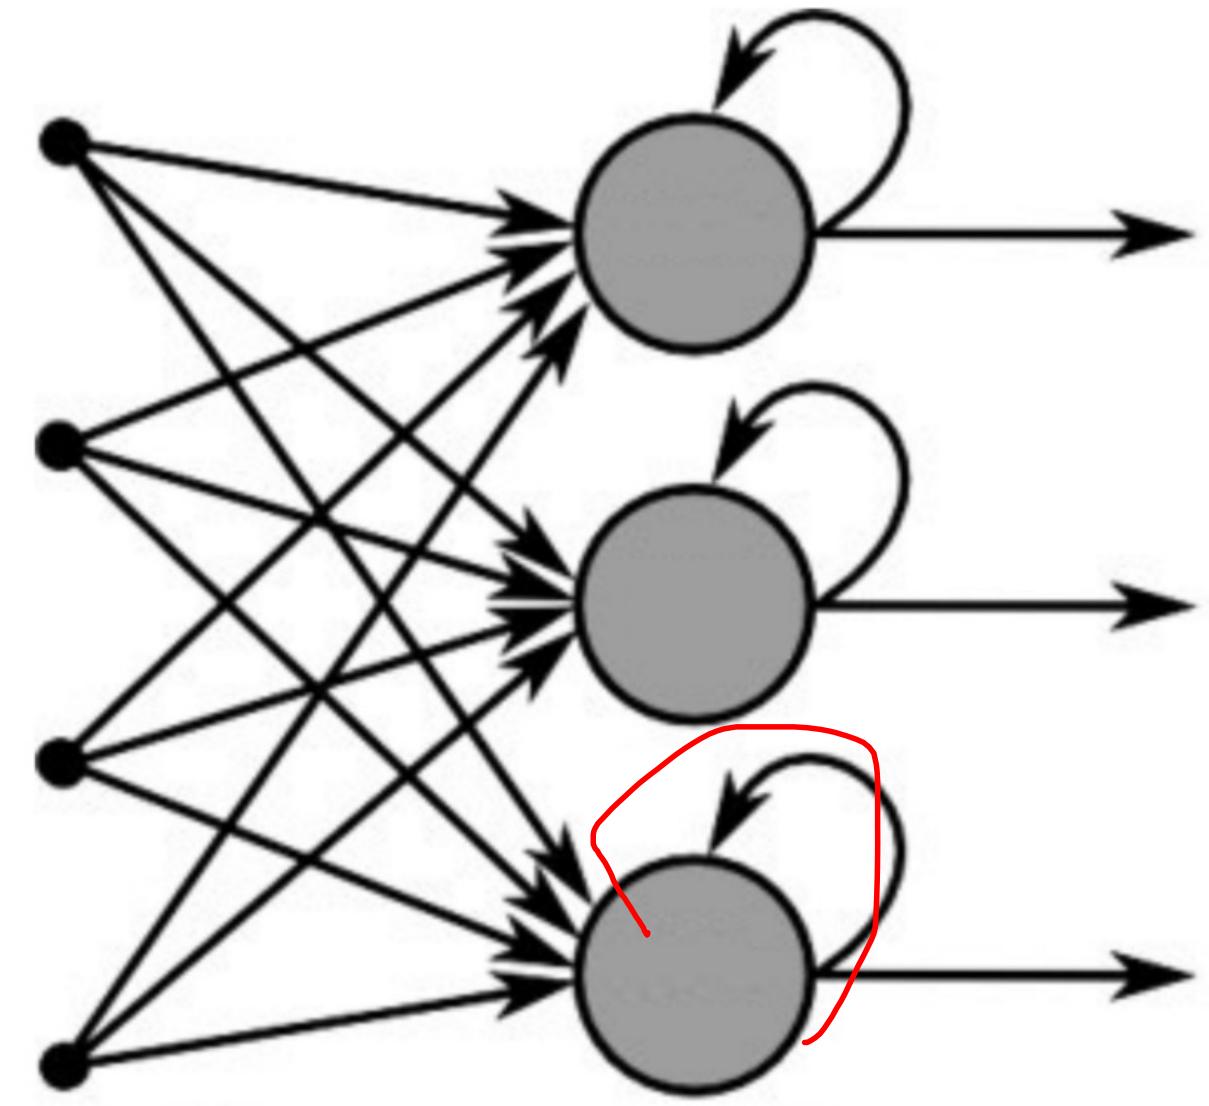
Examples of Neural Networks

Recurrent Neural Networks



Feed-Forward Neural Network

vs.

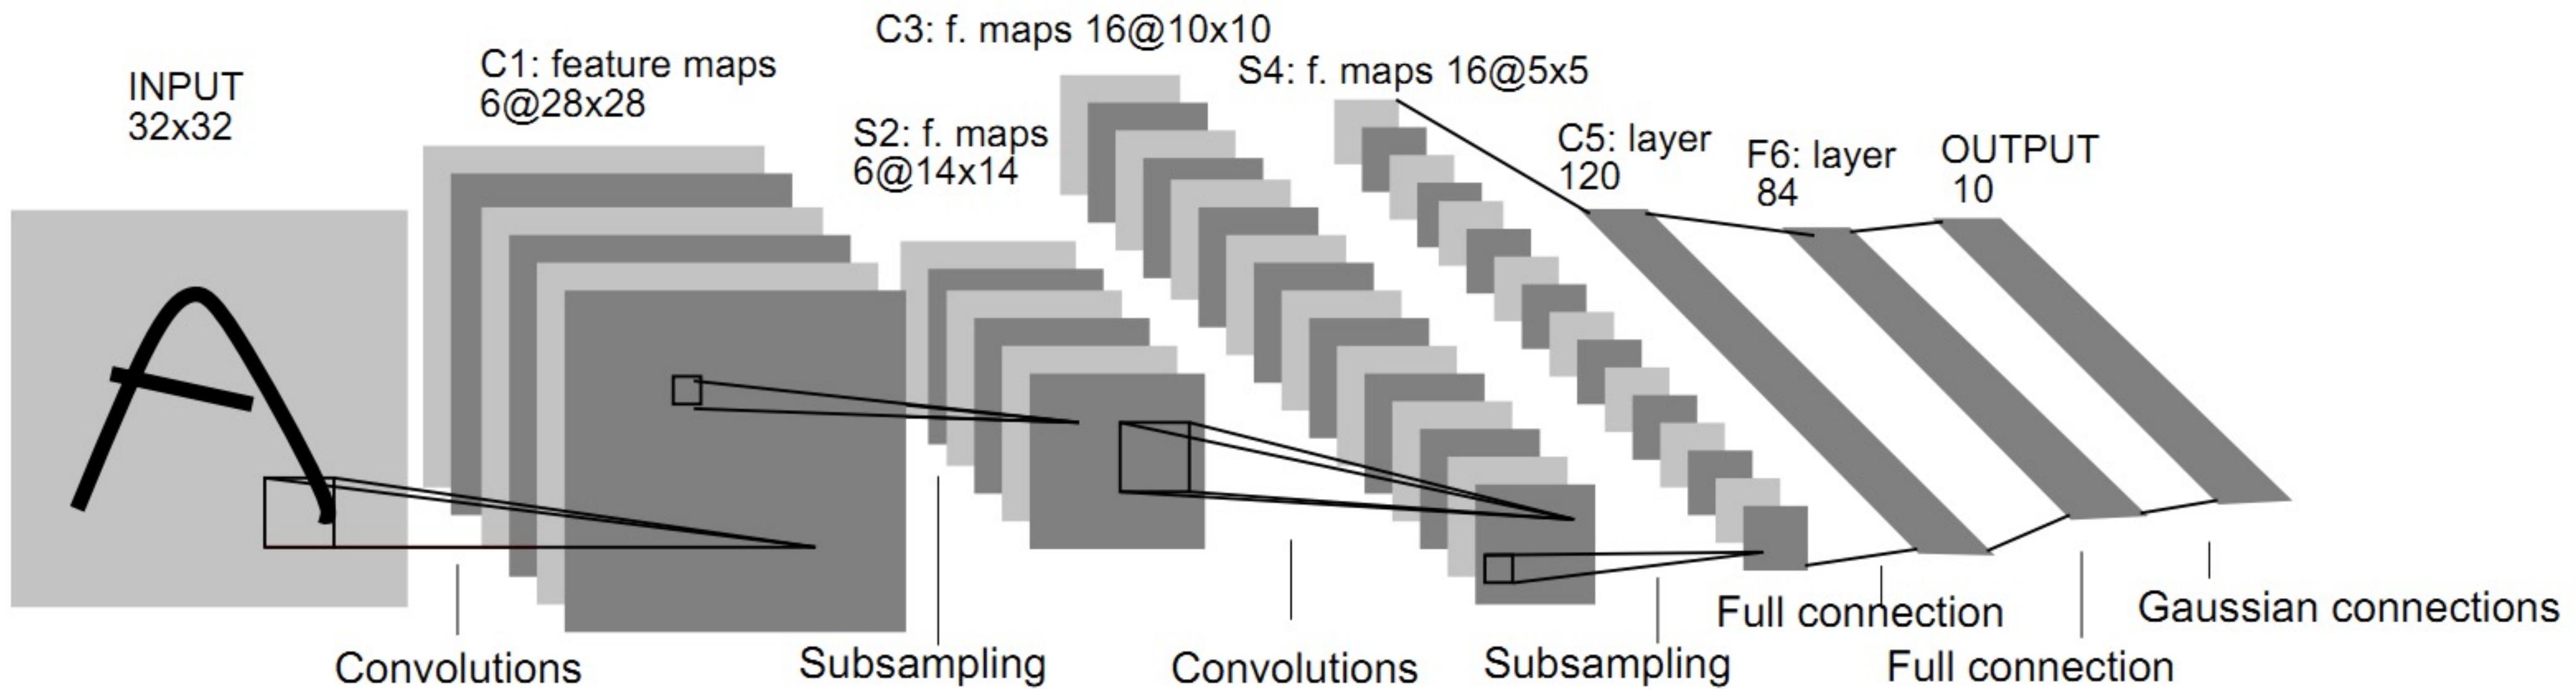


Recurrent Neural Network

Recurrent Neural Networks

-> good for speech recognition

Examples of Neural Networks: Convolutional Neural Networks



Convolutional Neural Networks -> good for image recognition

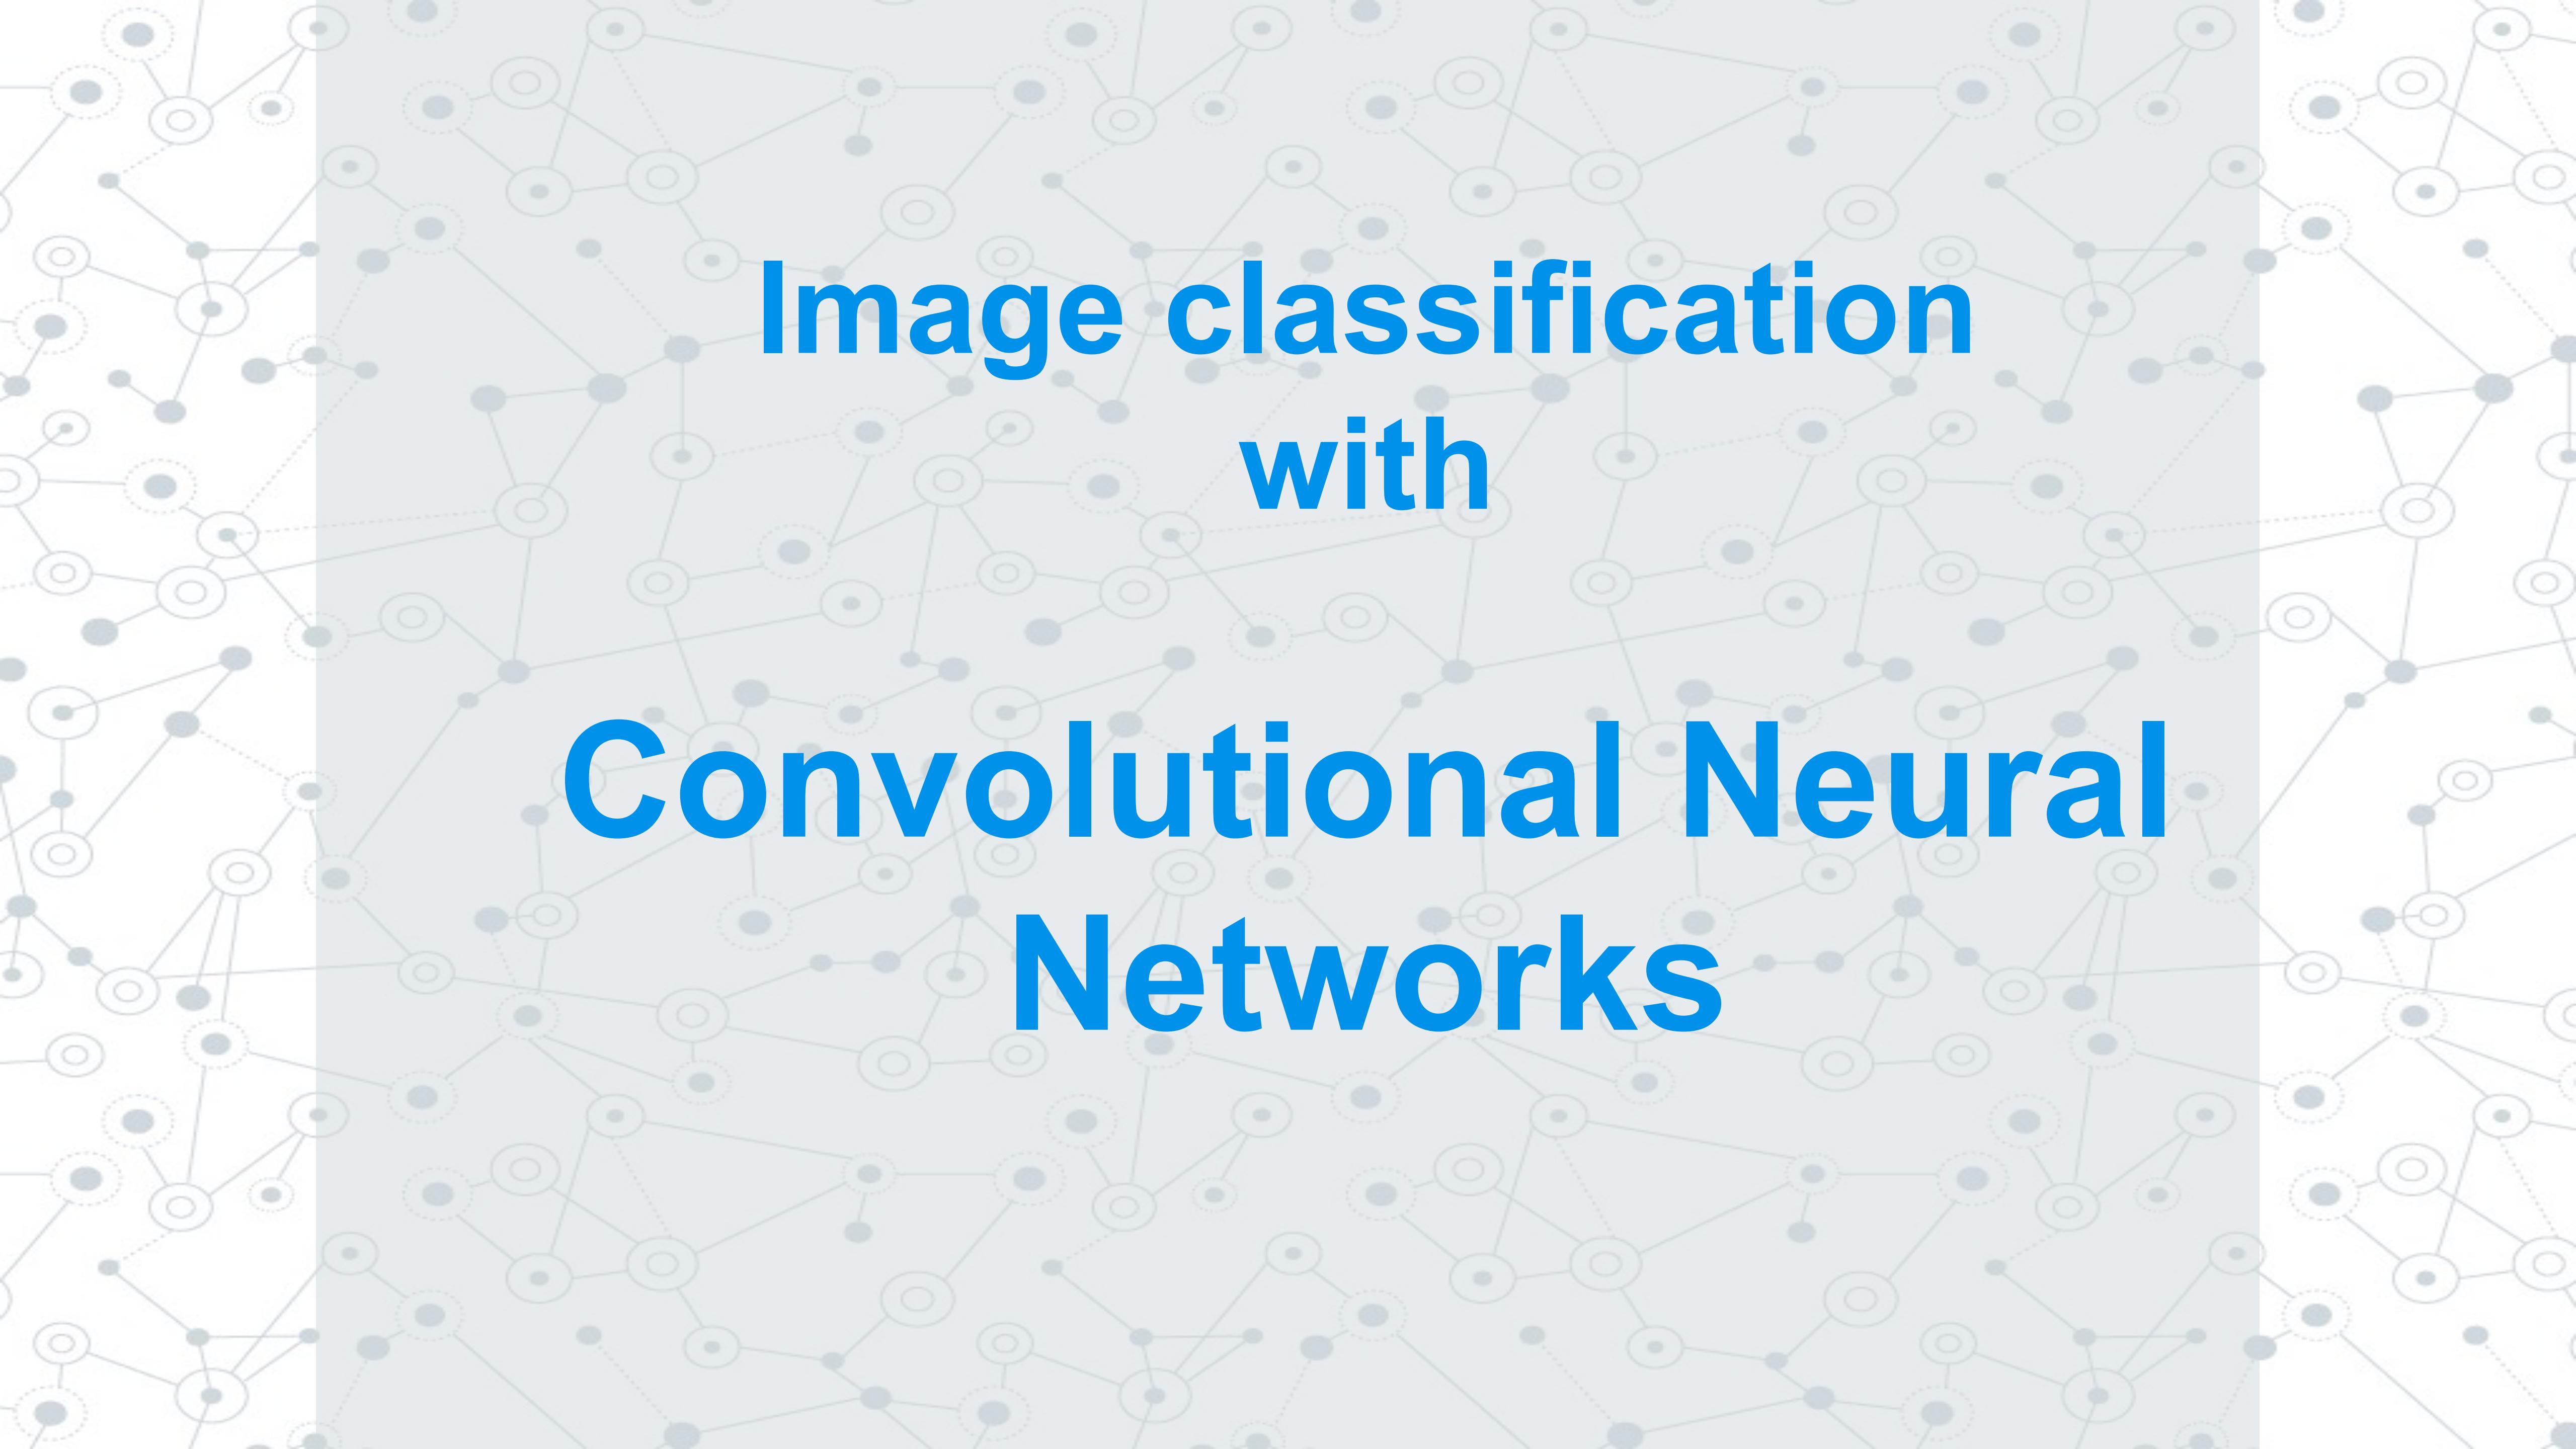
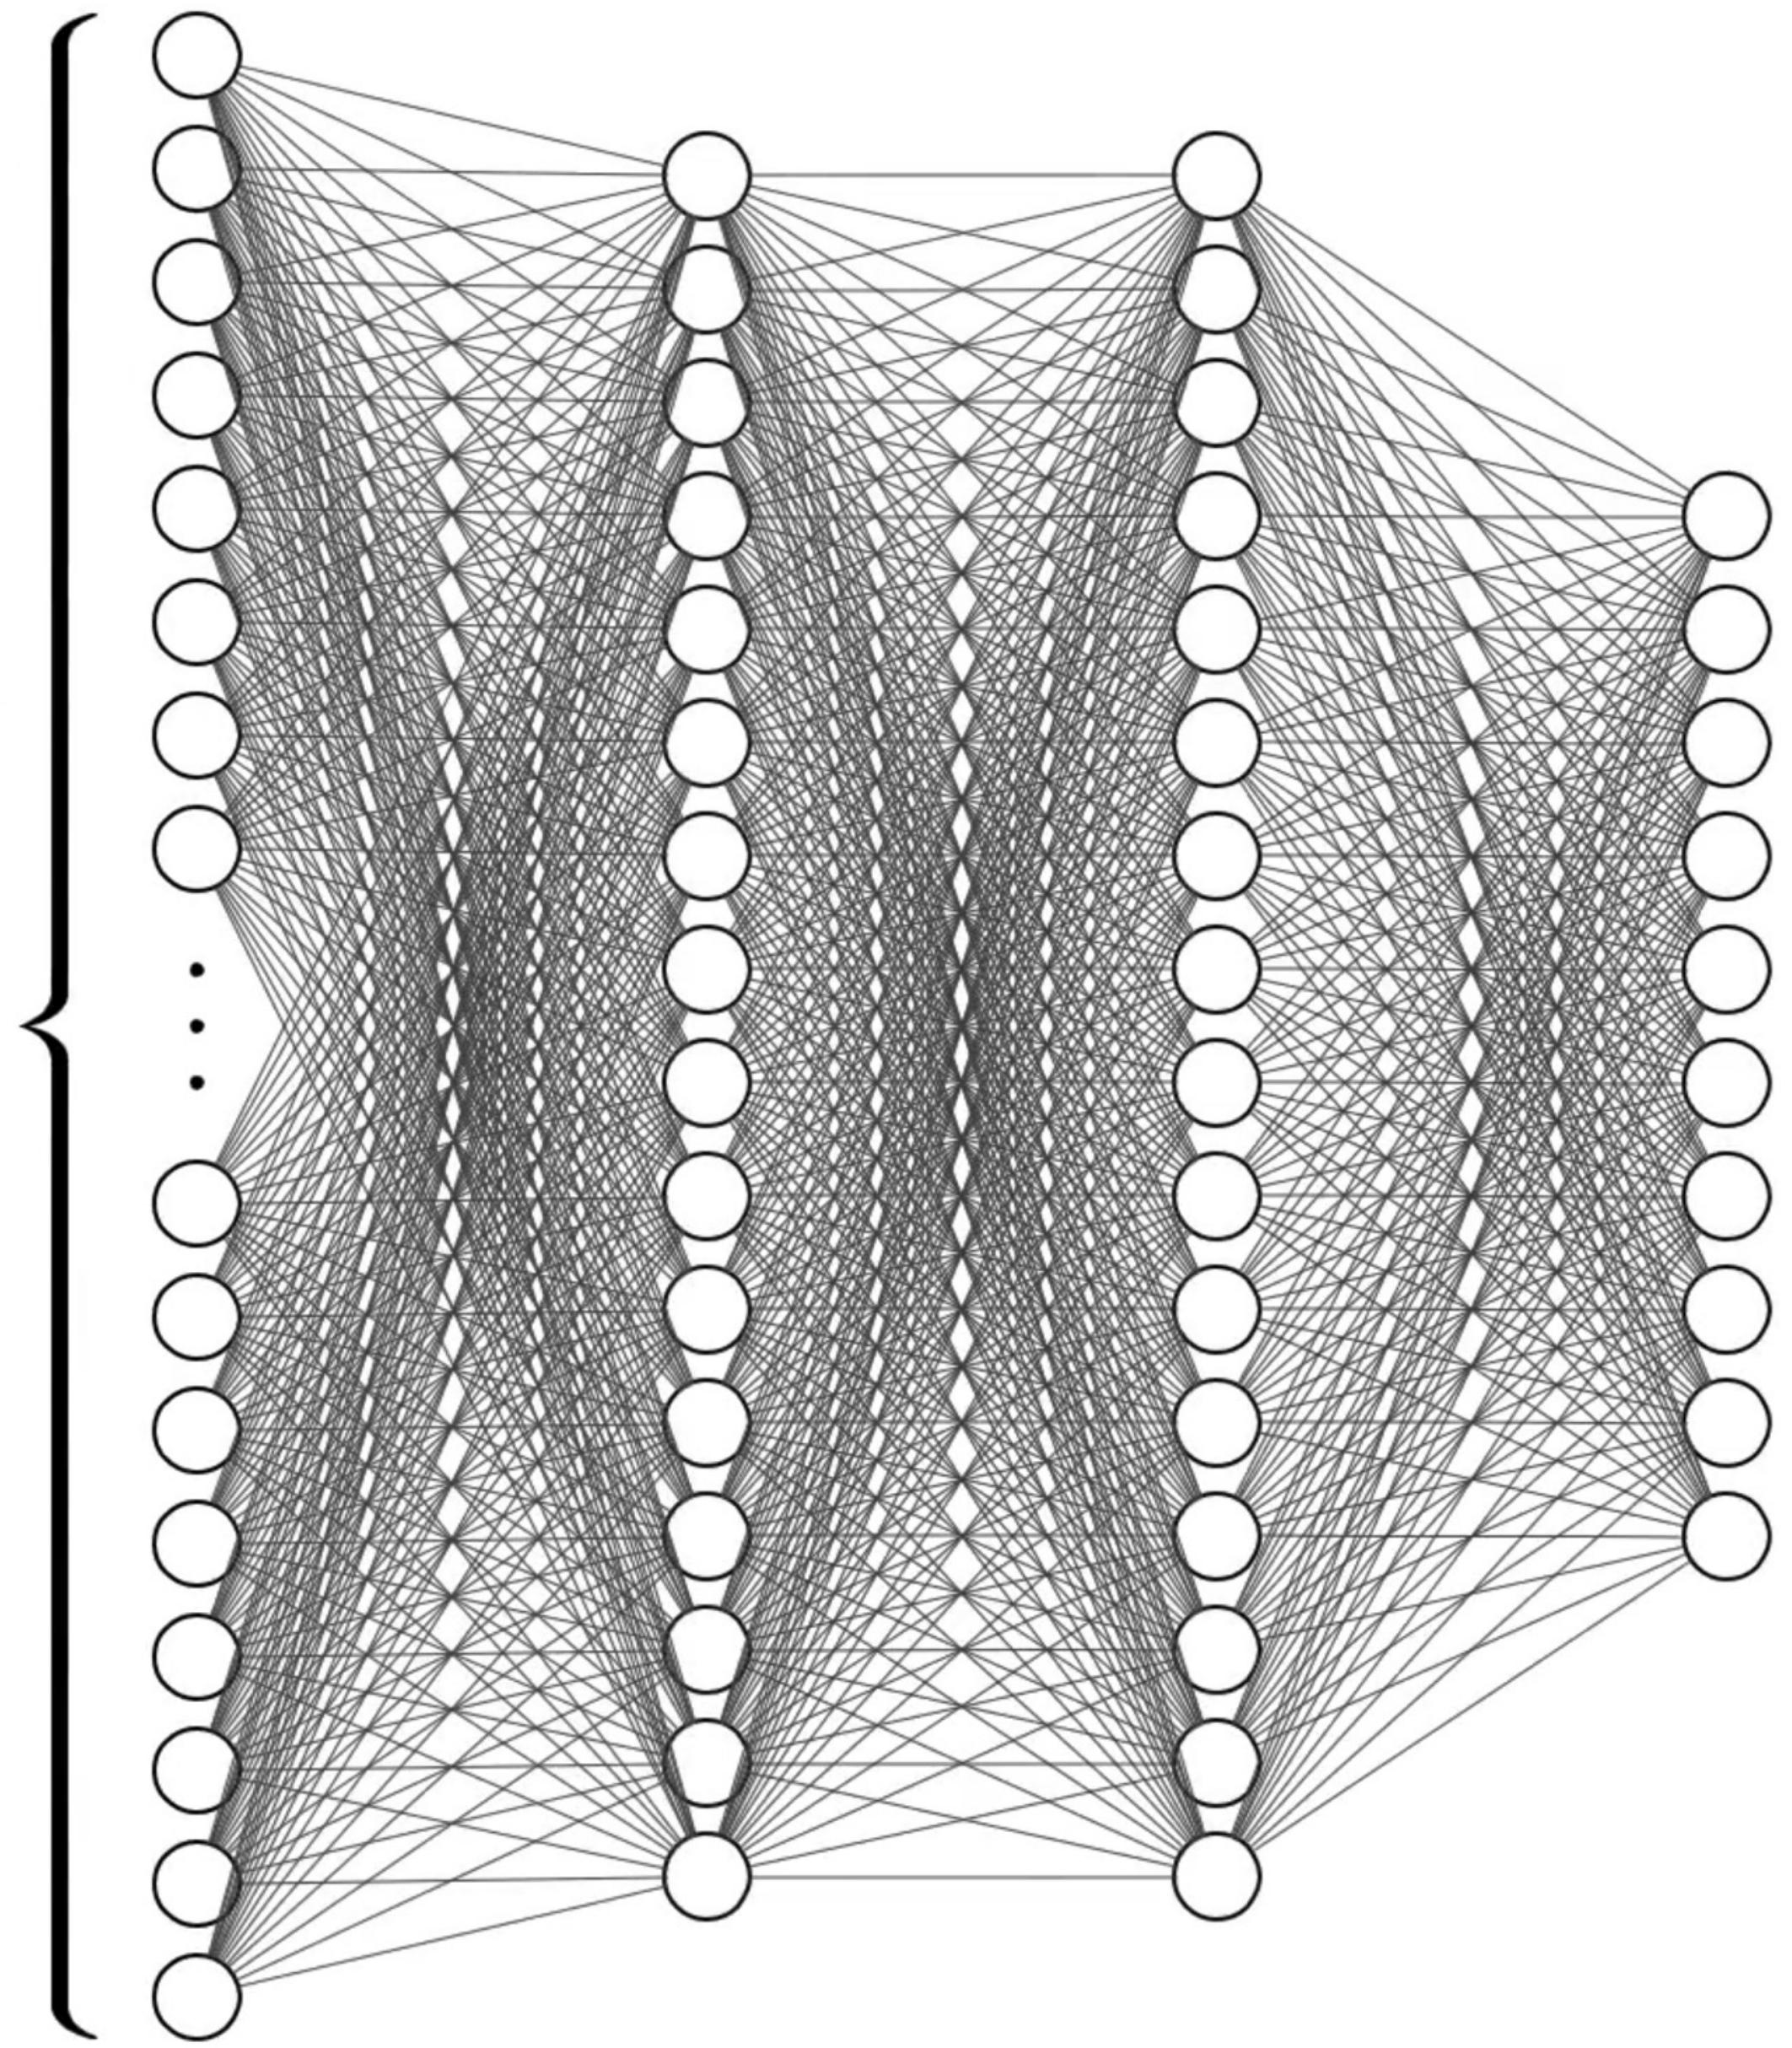


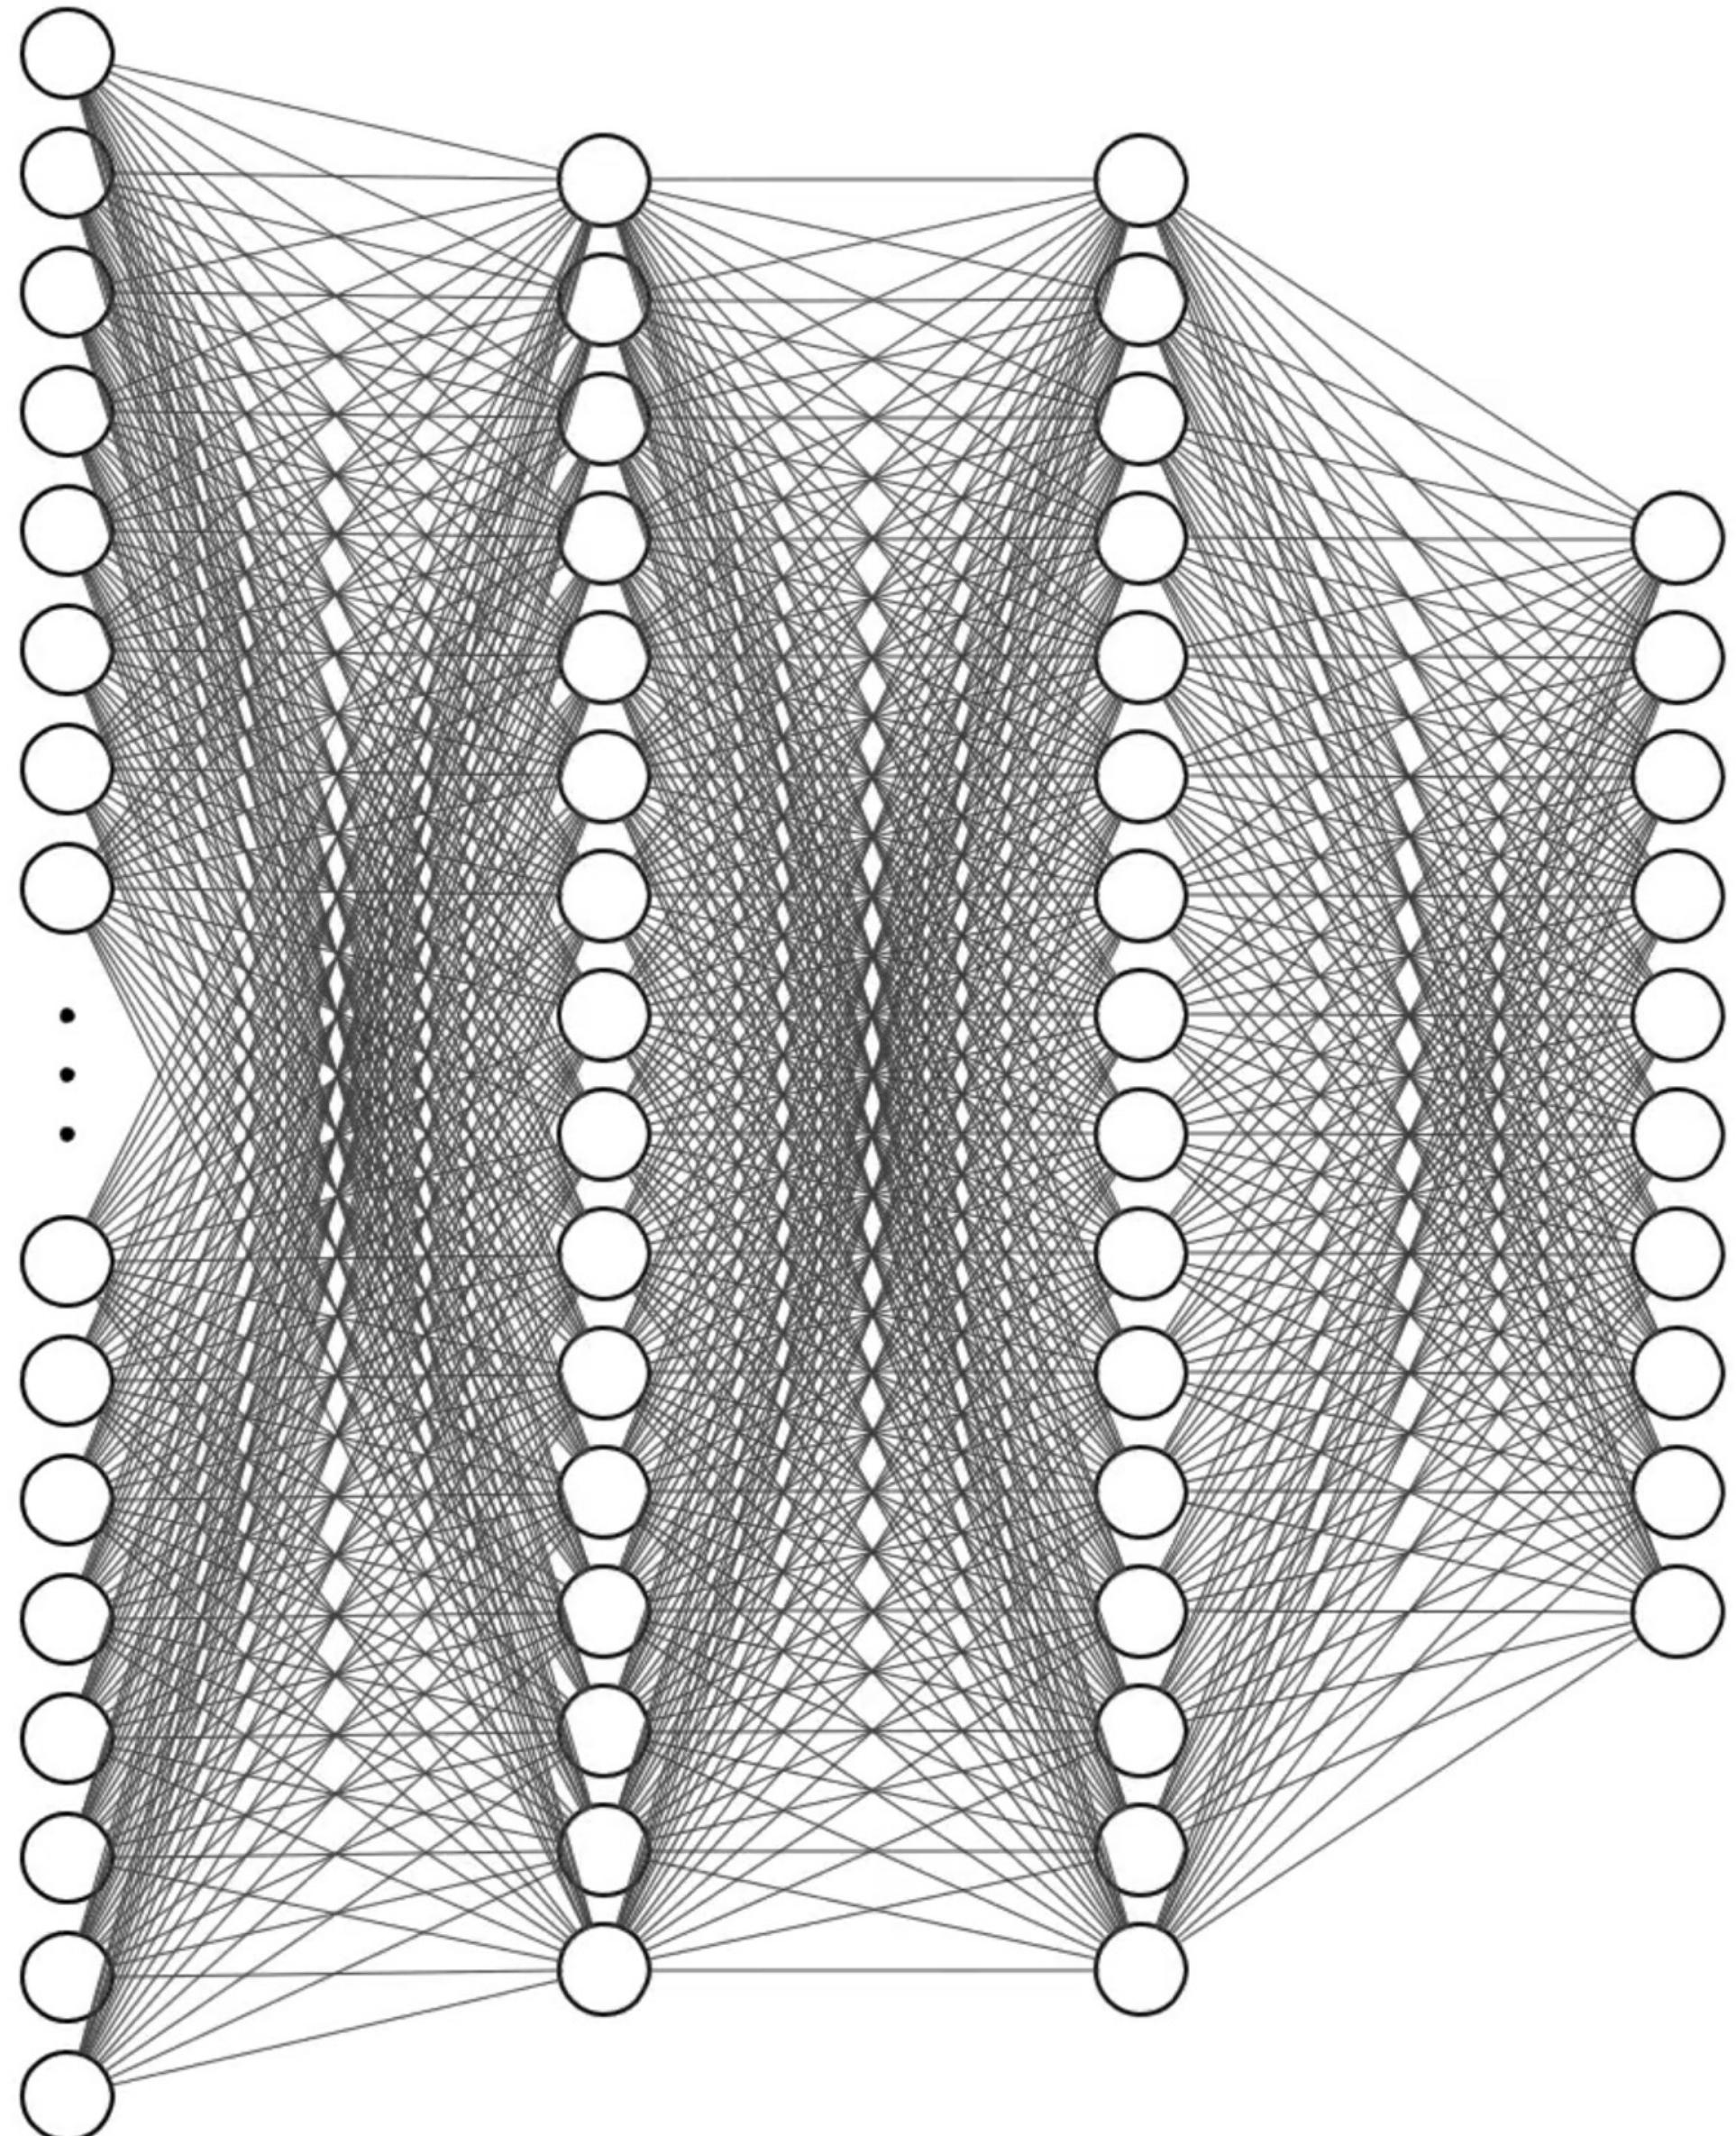
Image classification with Convolutional Neural Networks

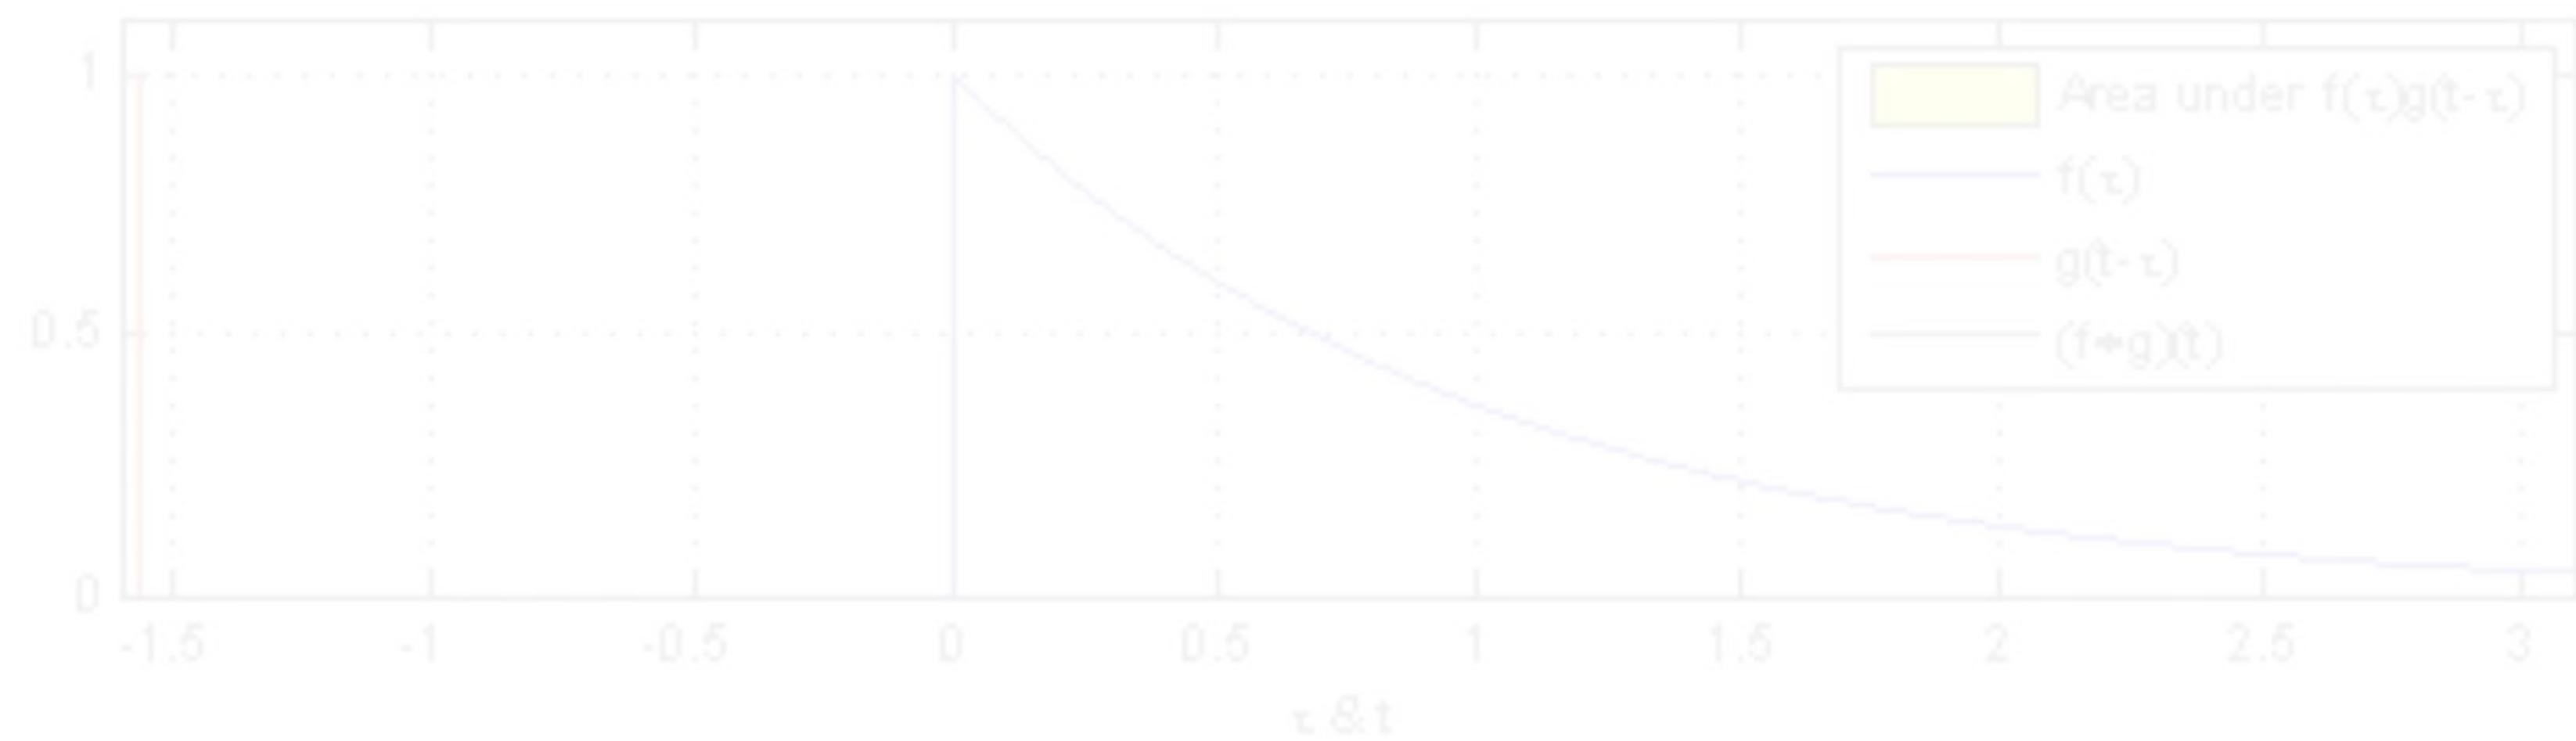
28 → 784



Why do we need convolutions?

- Even for small images, the number of a fully connected network is computationally expensive
- Here comes in play the convolution and pooling processes



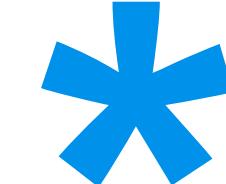


Brian Amberg derivative work: [Tinos \(talk\)](#)

What is a convolution?

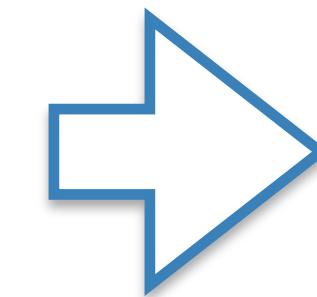
Original Image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0



Filter (or Kernel)

1	0	1
0	1	0
1	0	1



Feature Map
(Activation Map;
Convolved feature)

4	3	4
2	4	3
2	3	4

What is a convolution?

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

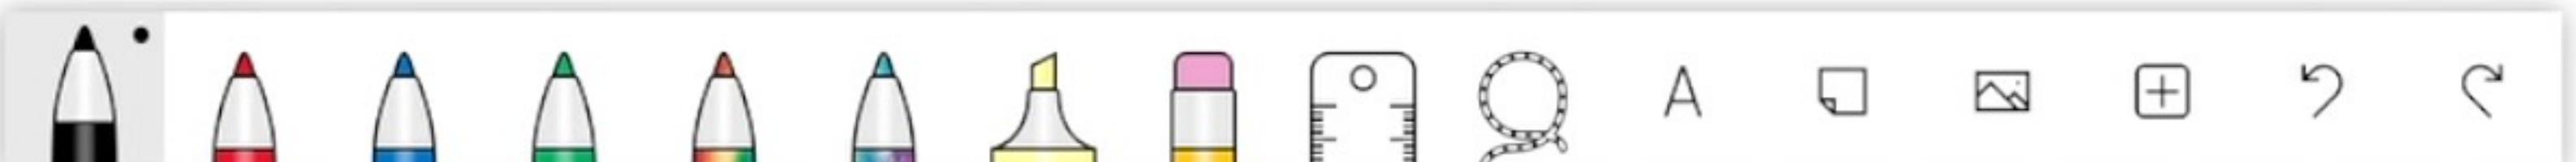
/

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

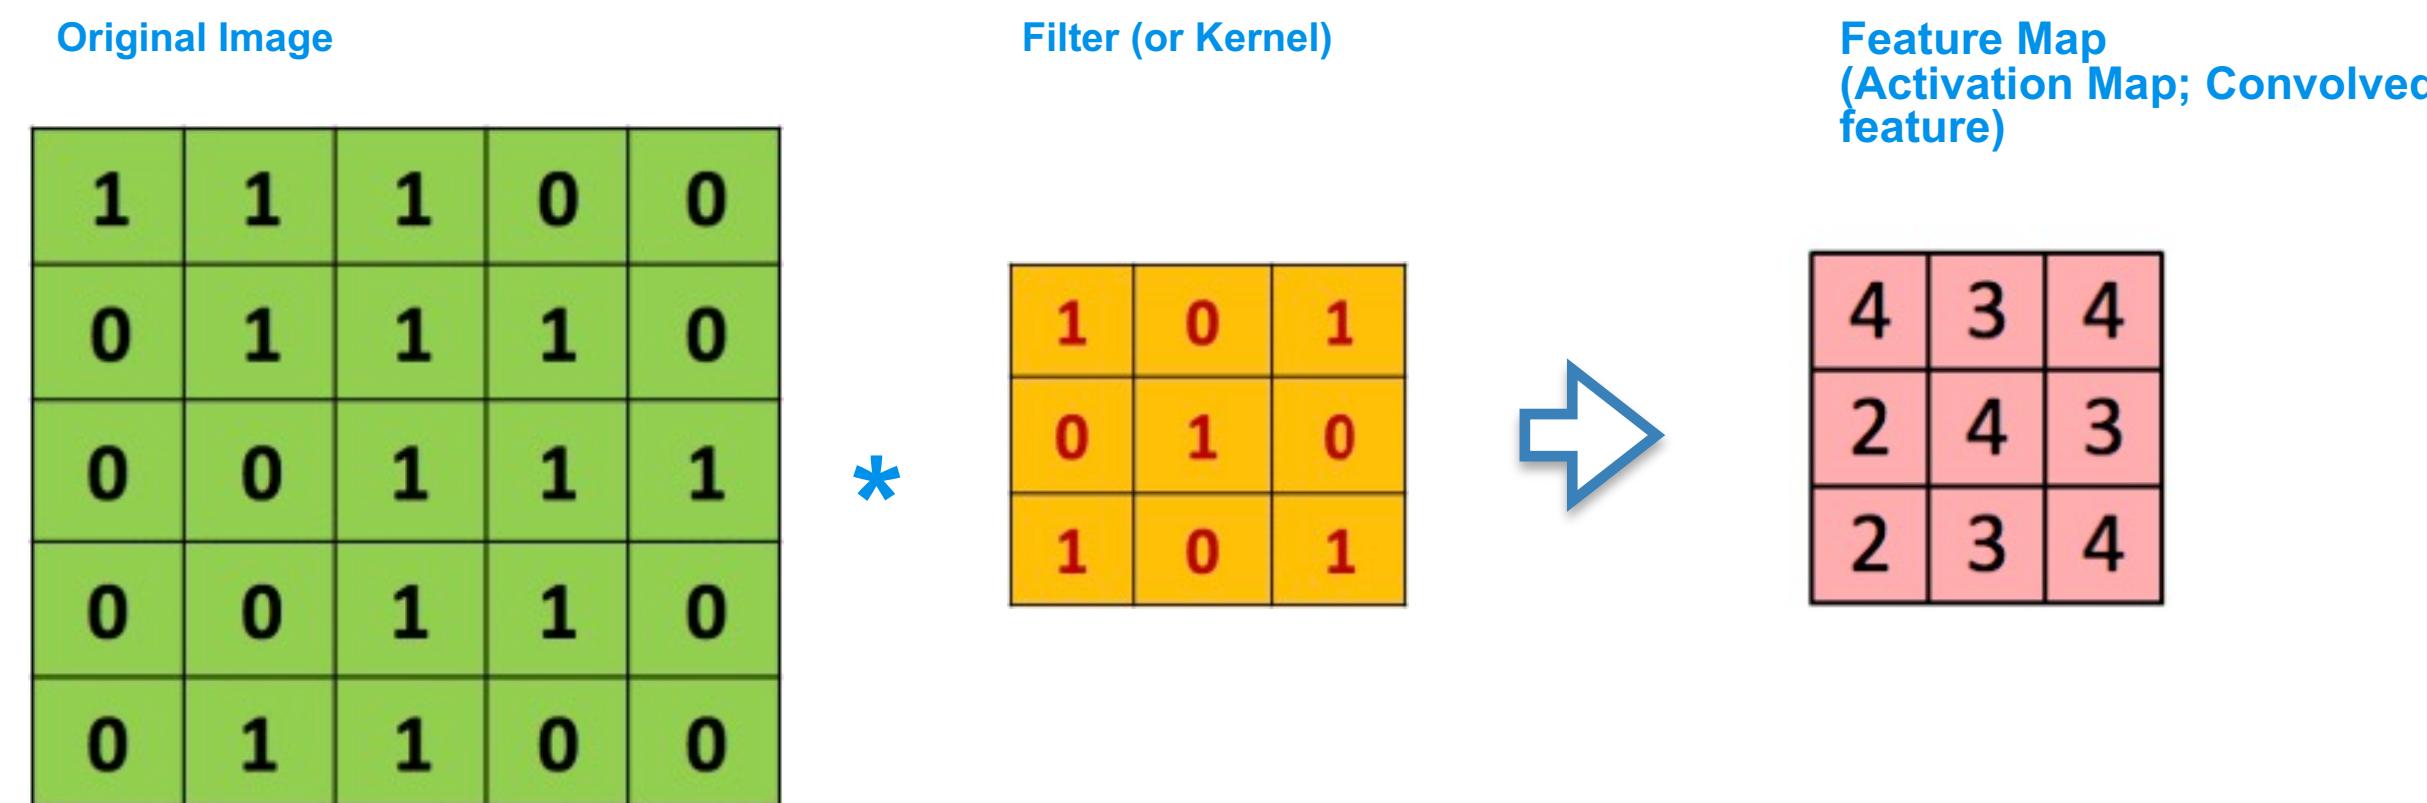
Image

4	3	4
2	4	3
2	3	4

Convolved
Feature



Convolution & training



- Think of the Filter (Kernel) as a pattern detector!
- When training the CNN we will learn the parameters of these filters
- For a filter of size 3×3 , we have to learn 10 parameters (9 weights + 1 bias)
- You have to define **the Filter size** and the **Step size**

w1	w2	w3
w4	w5	w6
w7	w8	w9

What is a convolution?

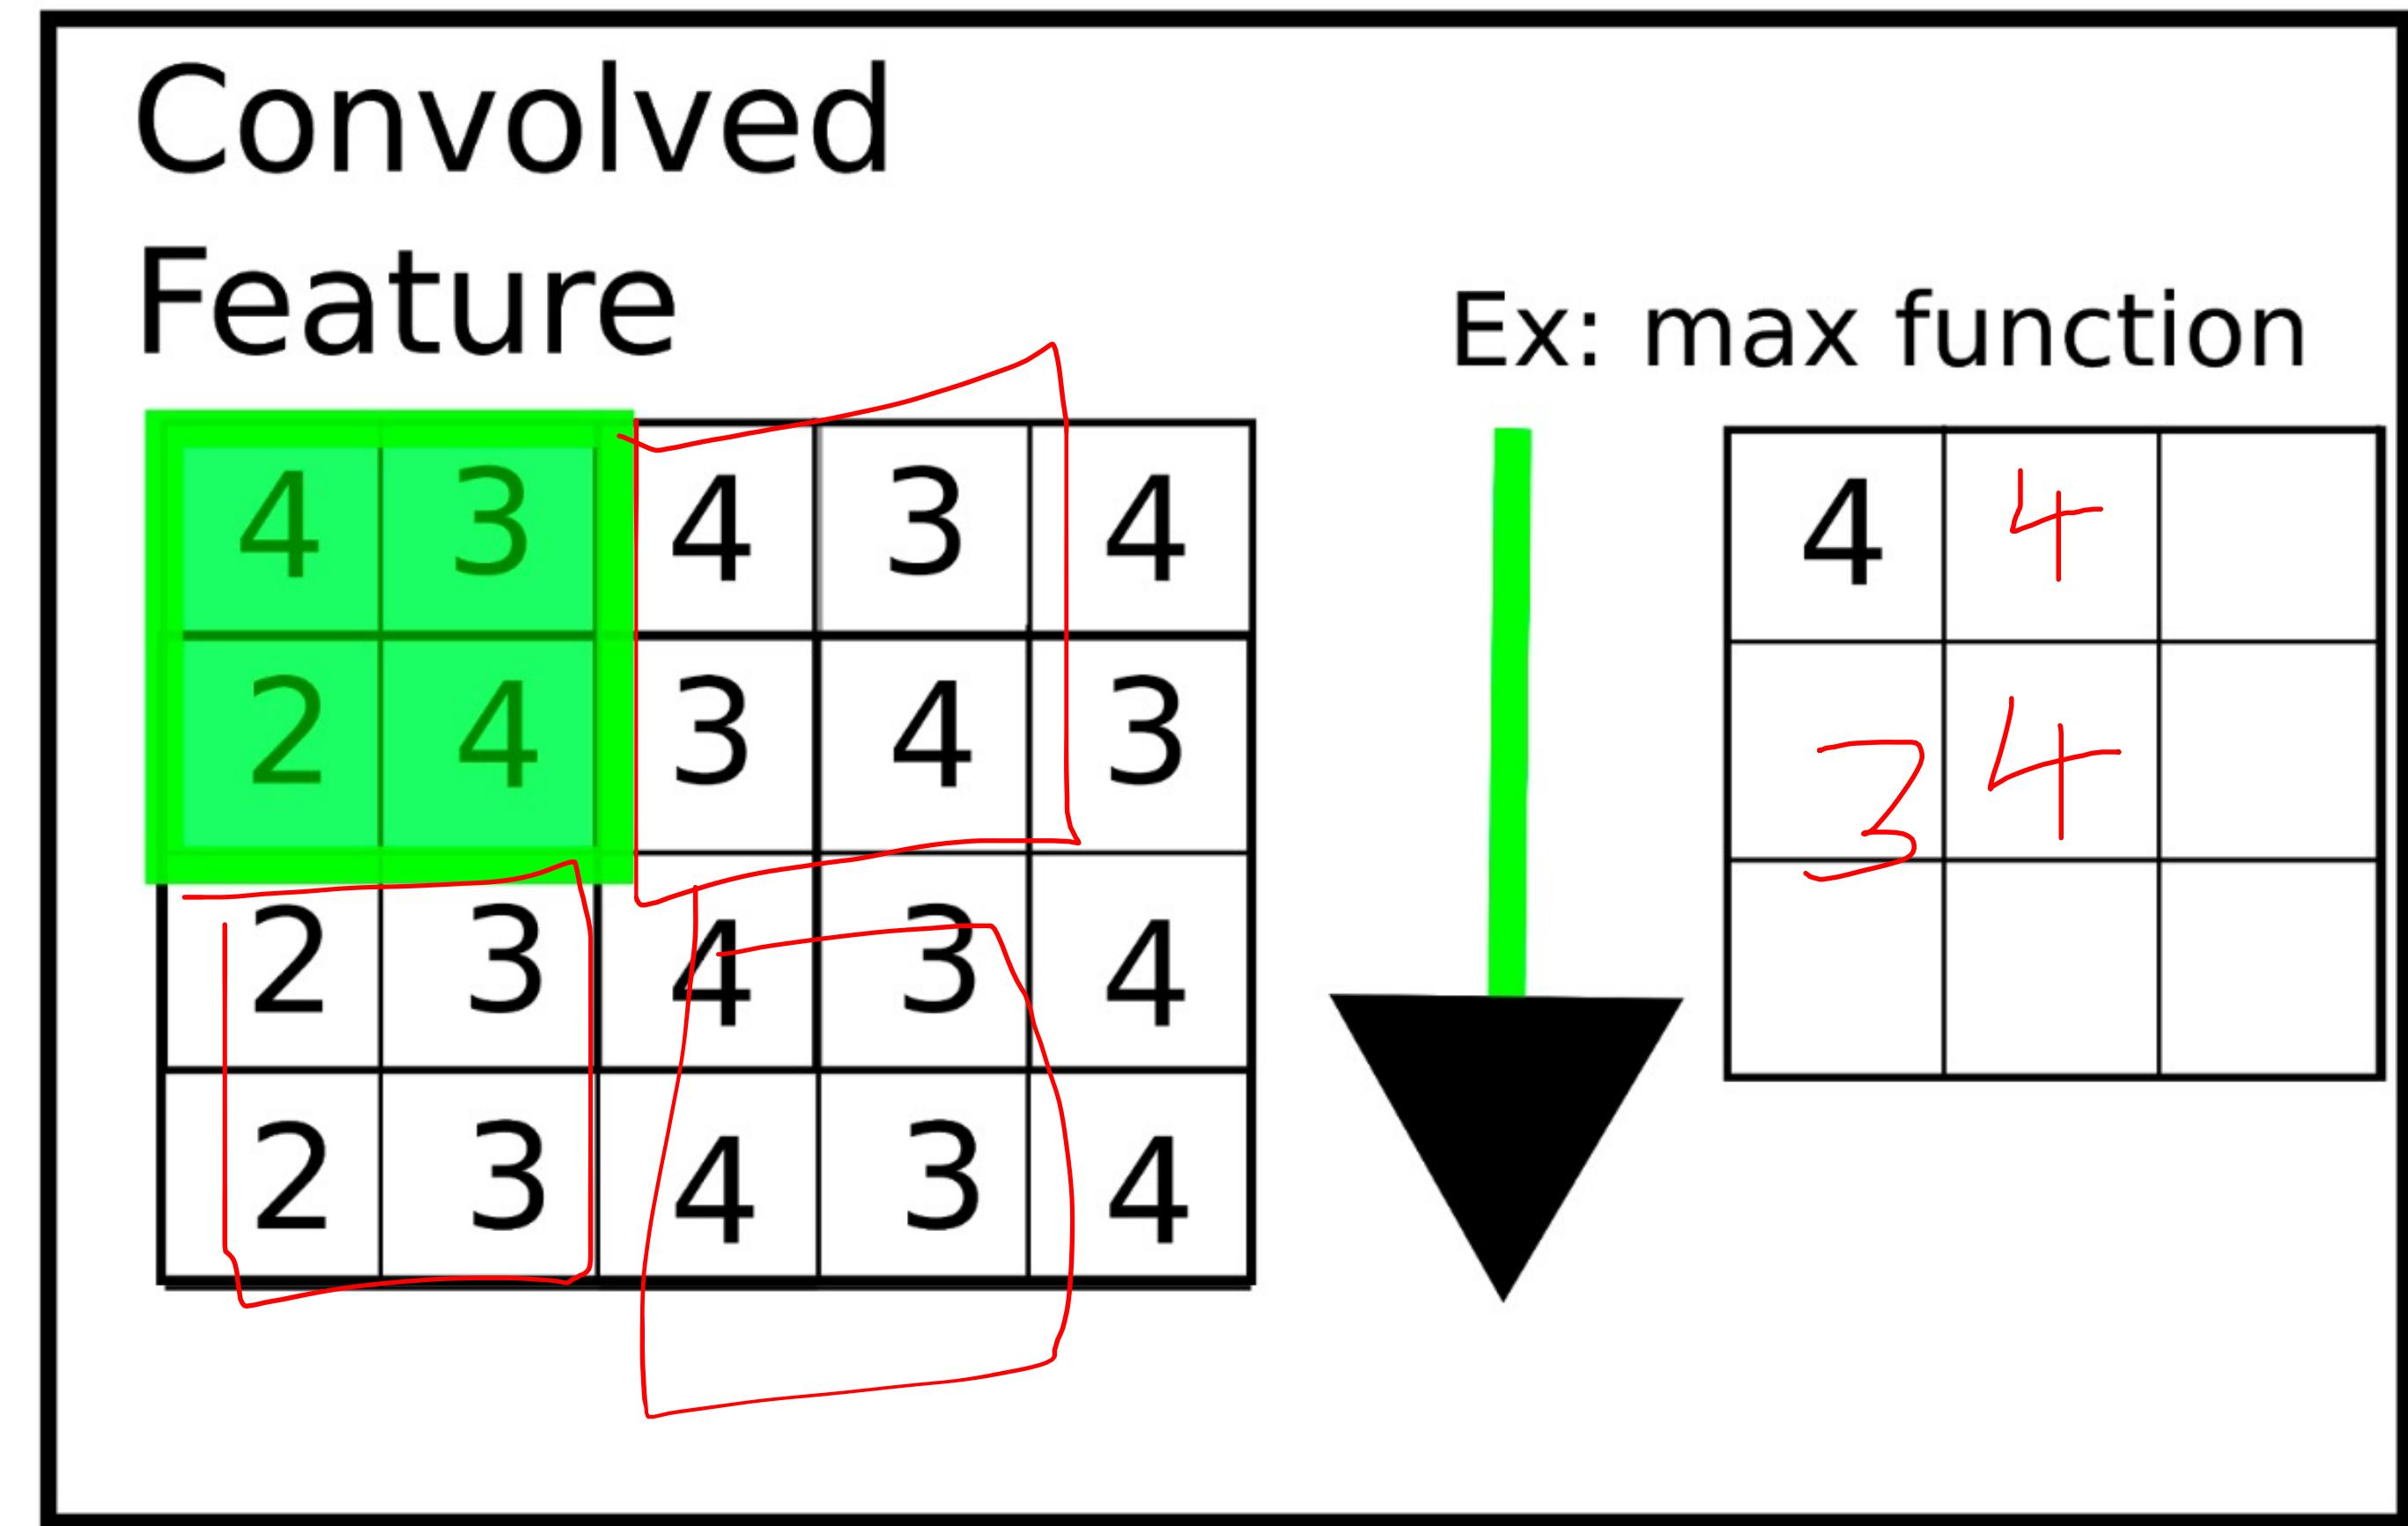


Pooling/subsampling?

- Pooling is a way of sub-sampling, meaning reducing the dimension of the input
- The role of pooling is to reduce the resolution of the feature map, while retaining the important features for classification
- The pooling will have a pooling size, step and type of operator
- Operators can be different, the most popular are max or average

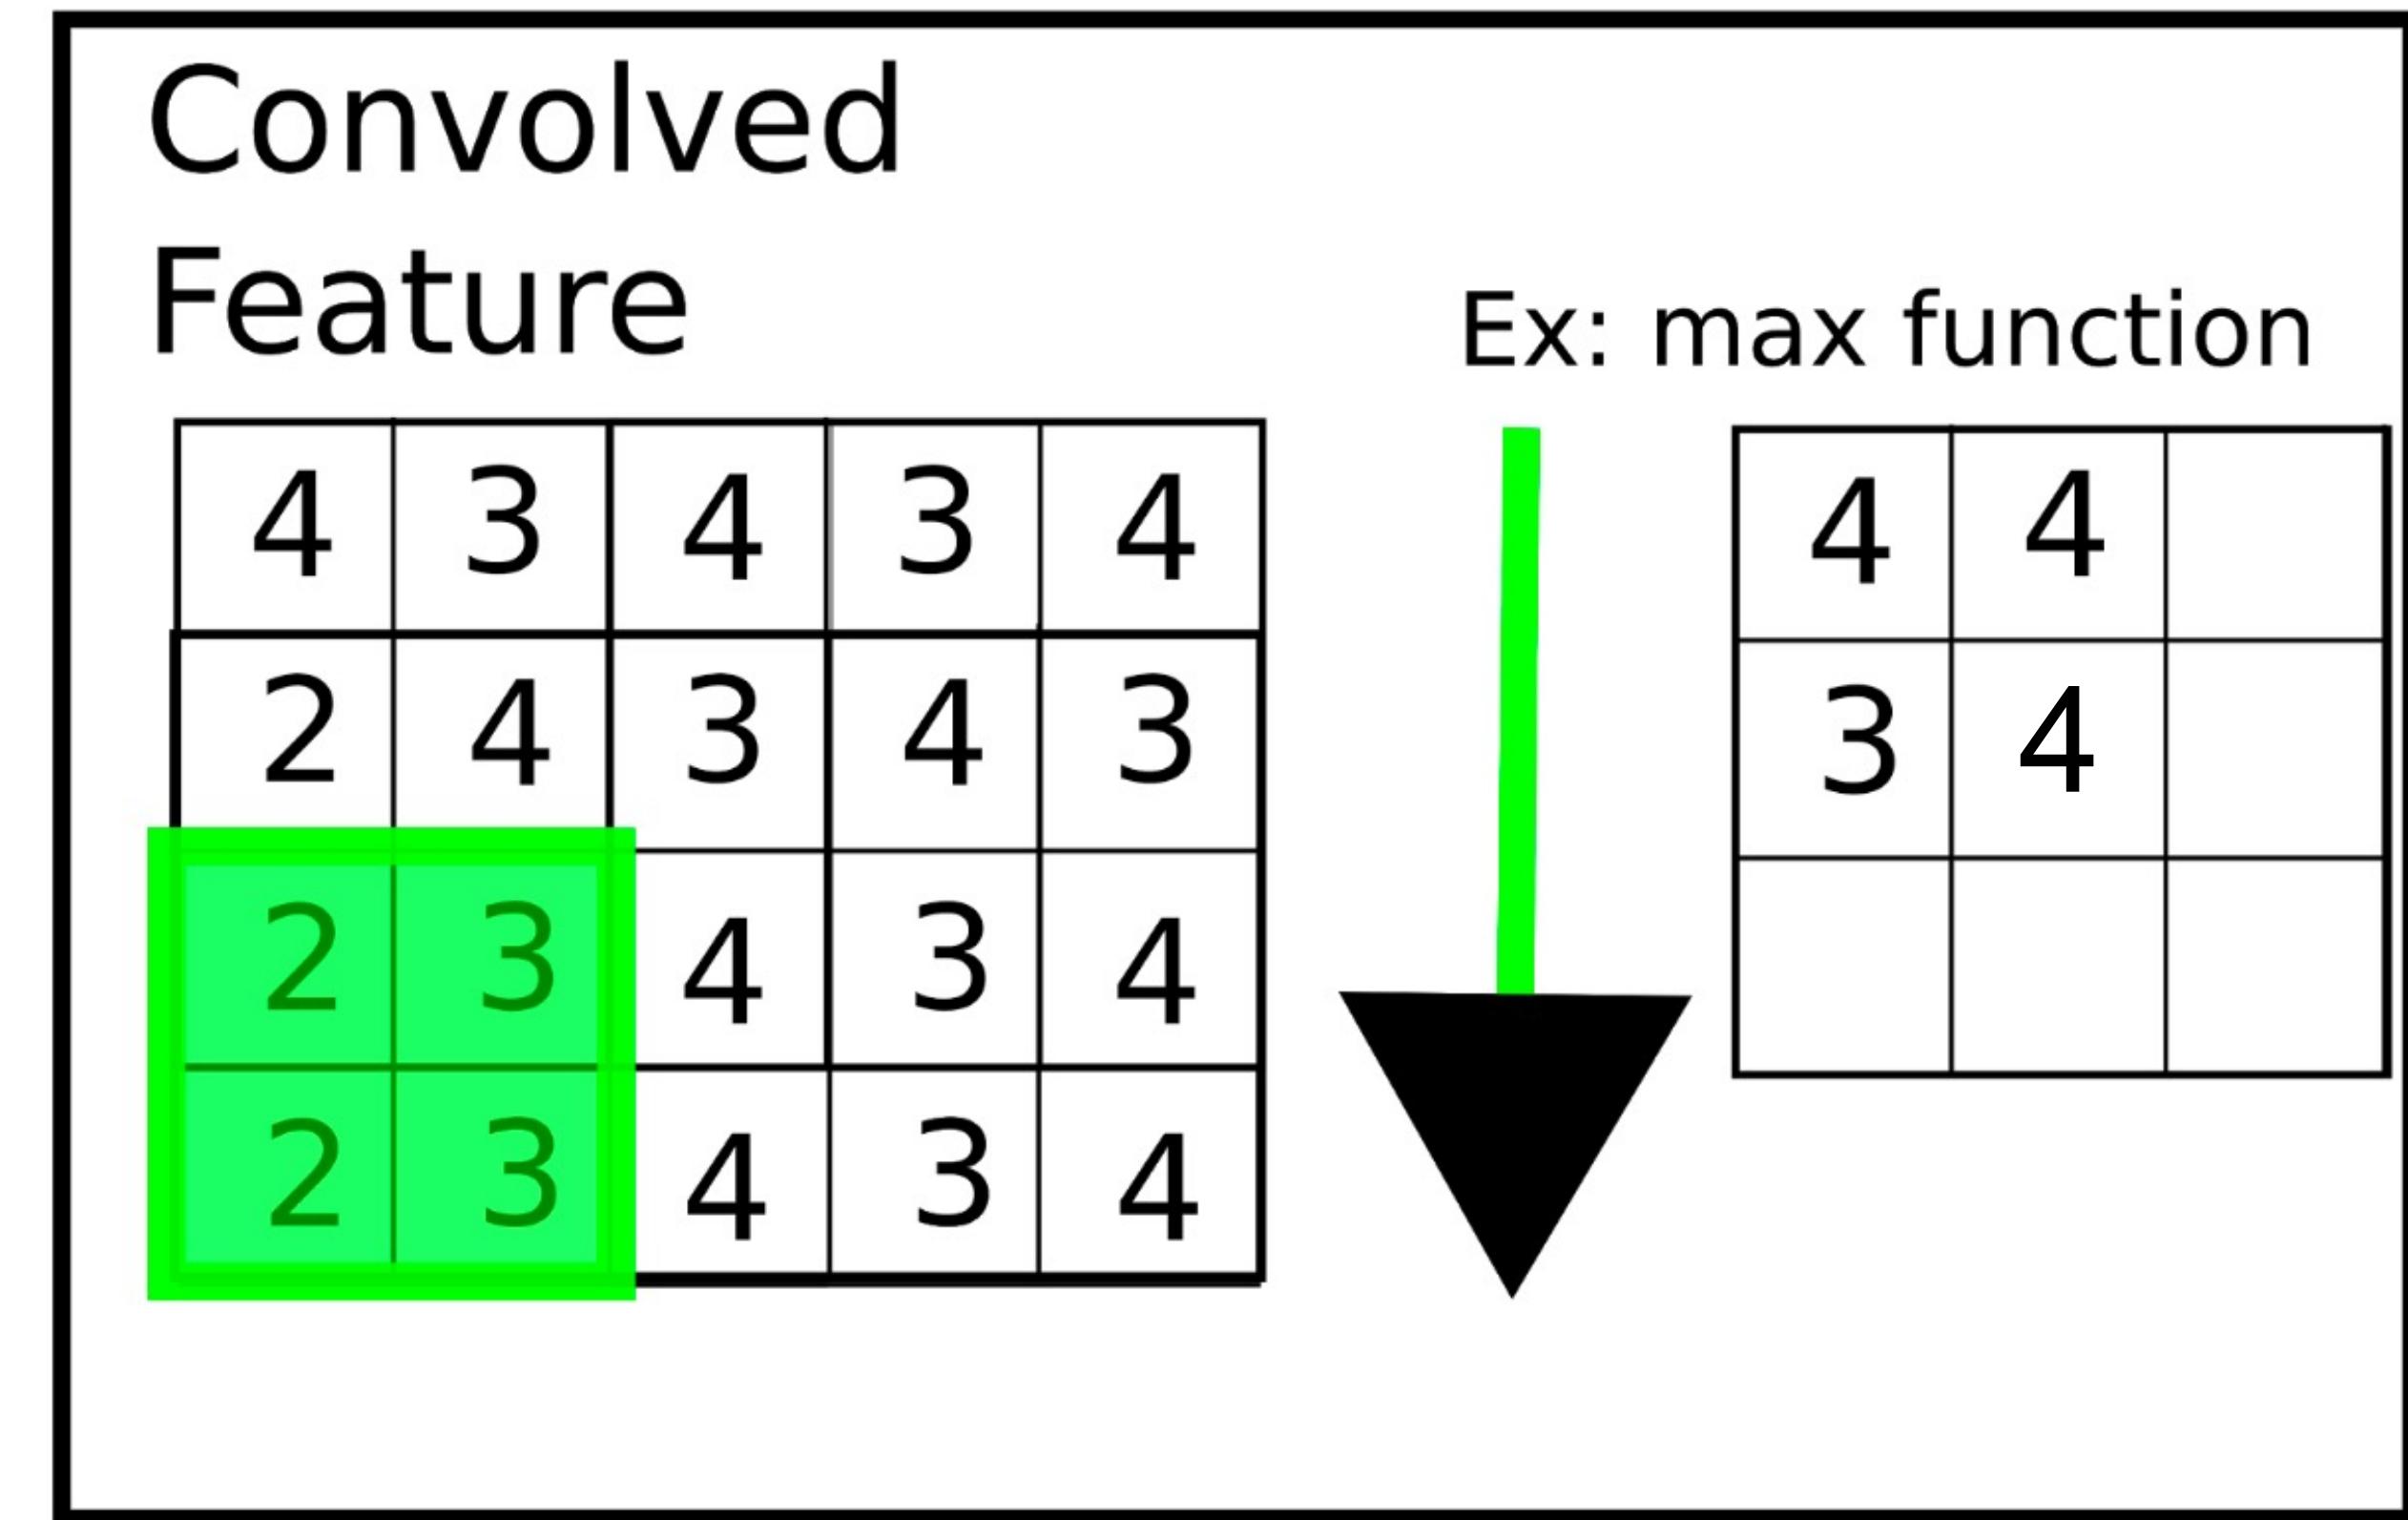
Pooling/subsampling : max pooling

- Filter size: 2x2
- Step: 2
- Operator: max



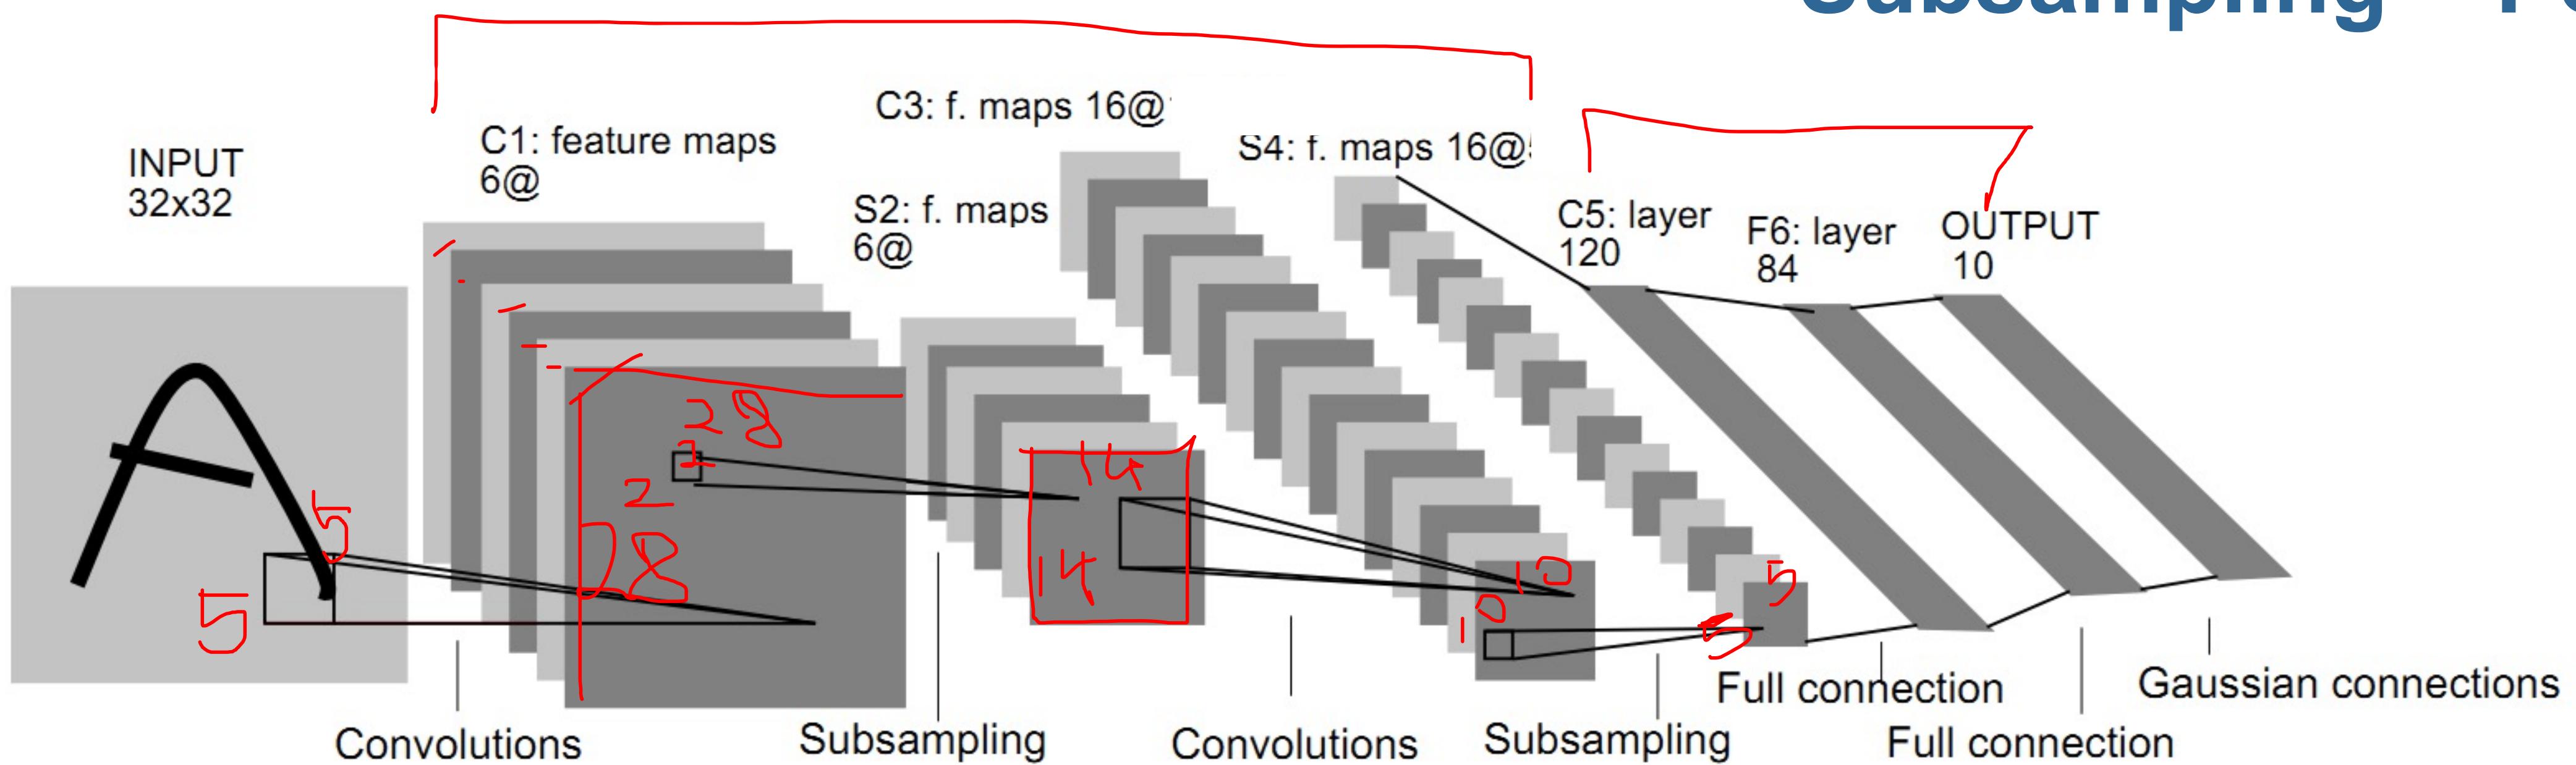
Pooling/subsampling : max pooling

- Filter size: 2x2
- Step: 2
- Operator: max



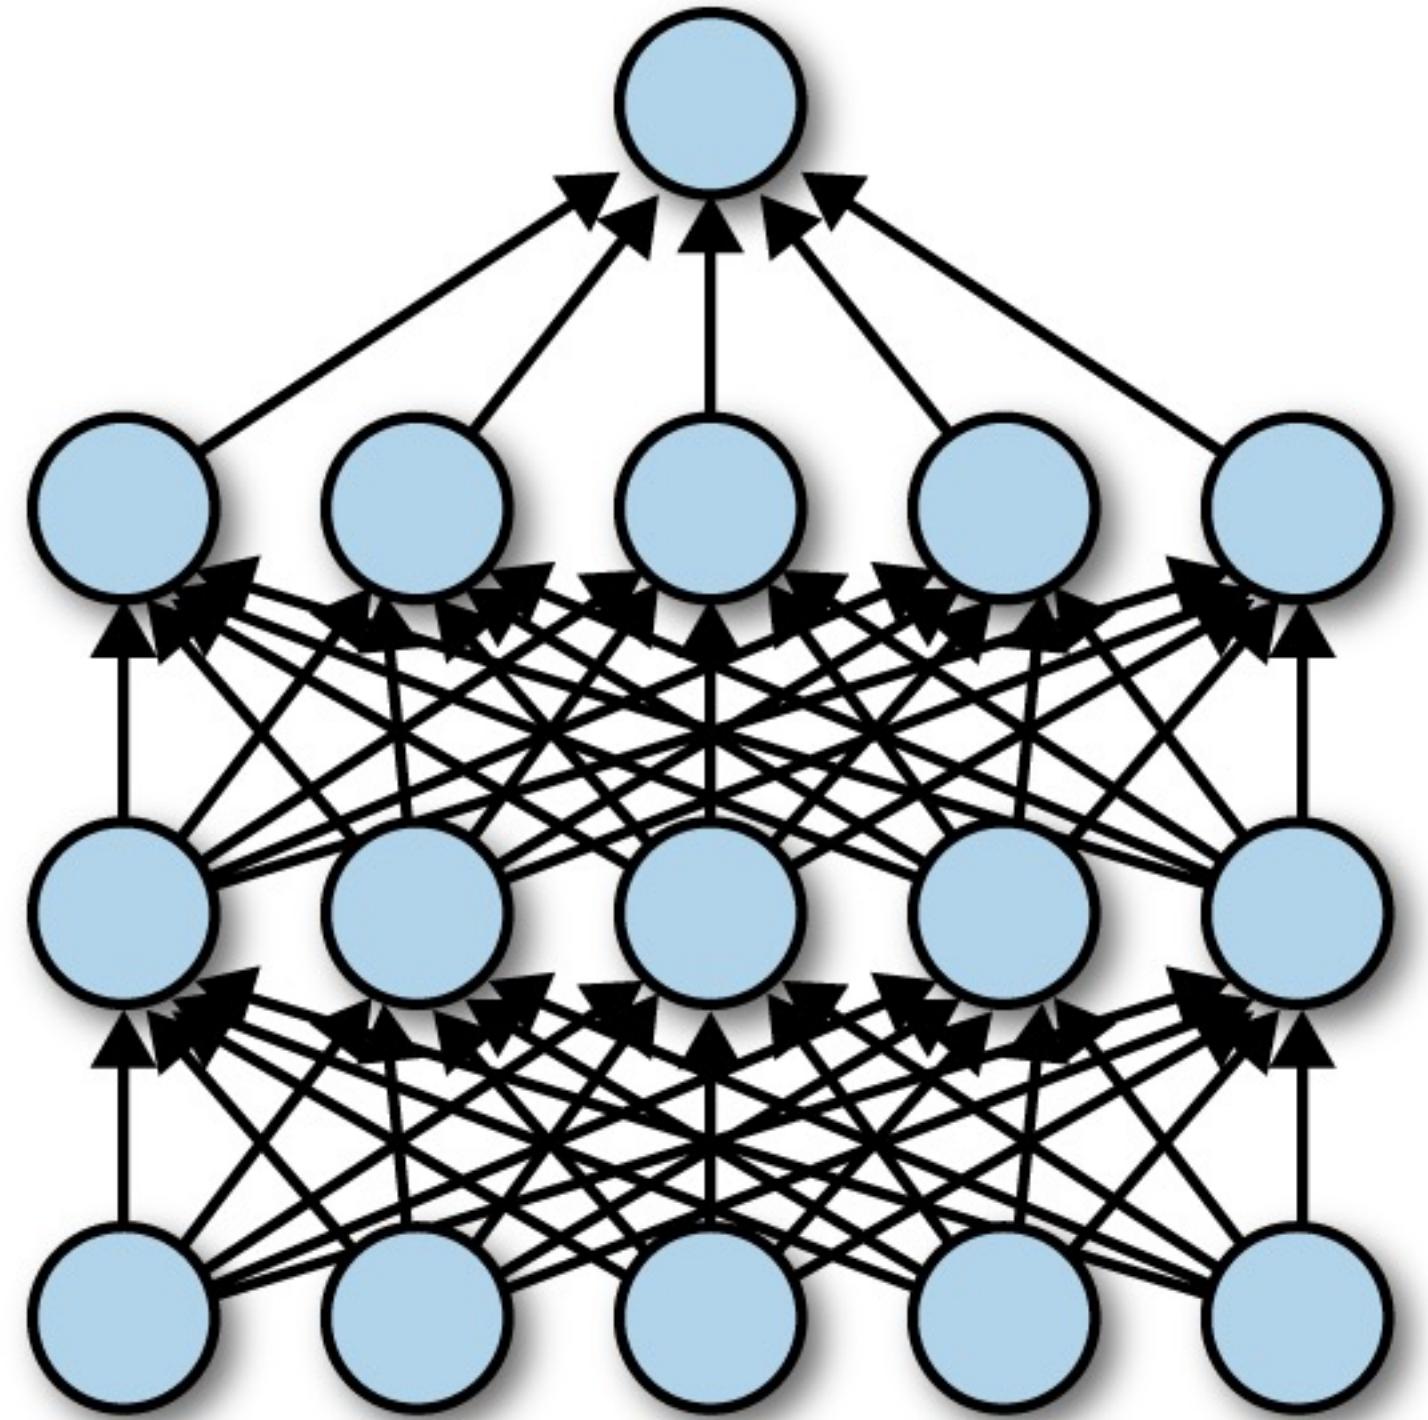
Examples of Neural Networks: Convolutional Neural Networks

Subsampling = Pooling

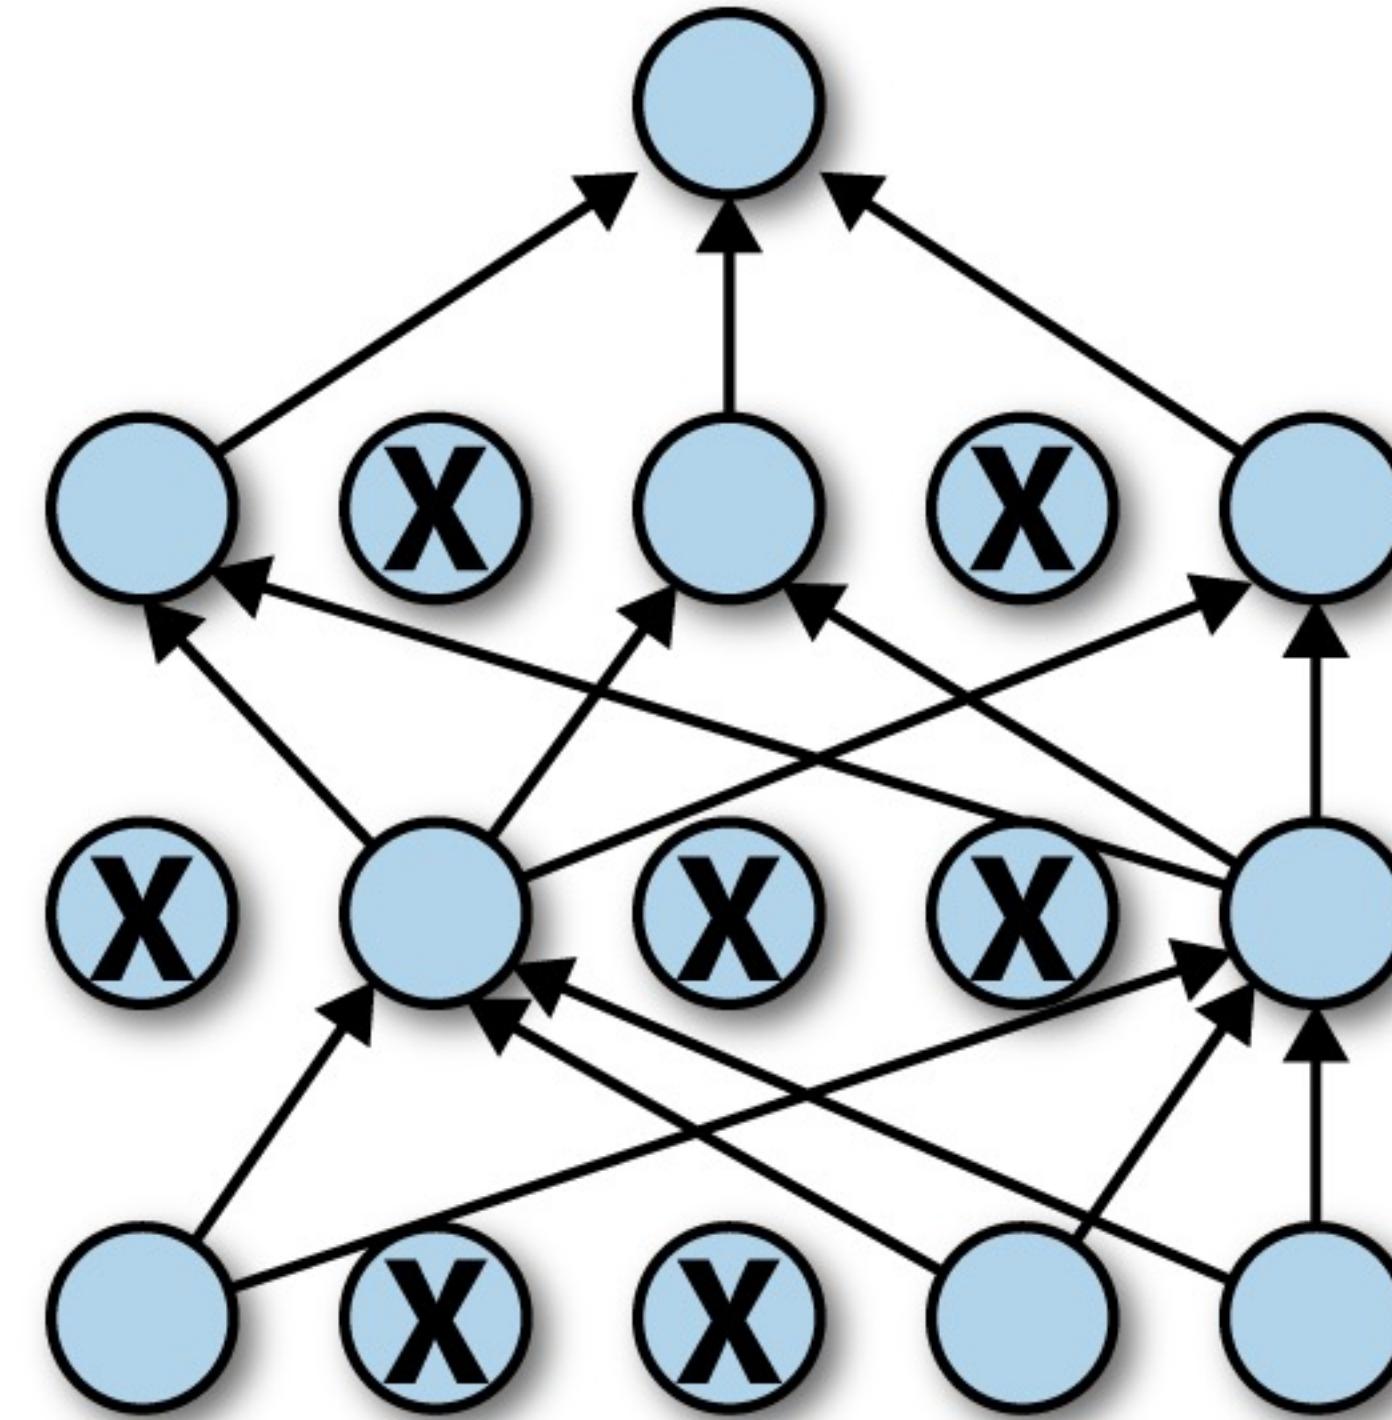


Convolutional Neural Networks -> good for image recognition

Dropout?

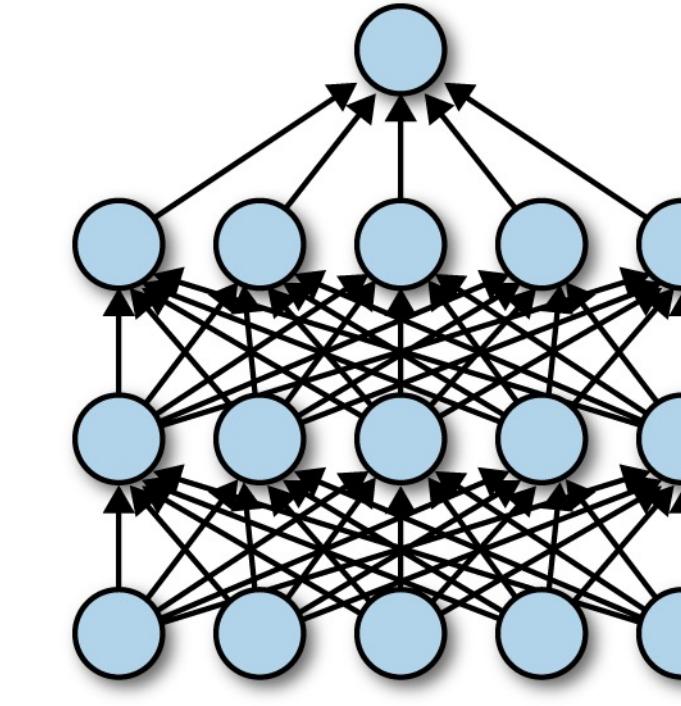


(a) Standard Neural Net

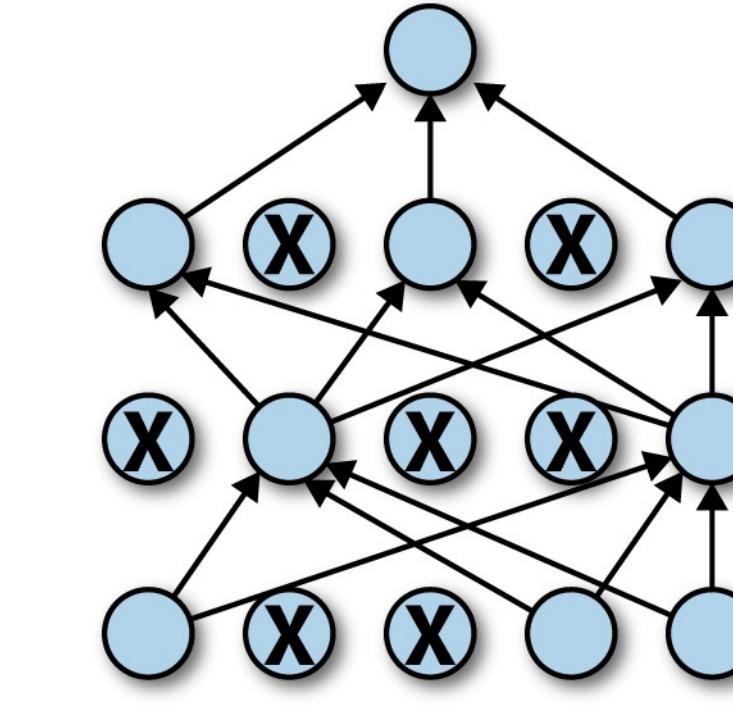


(b) After applying dropout

Dropout?



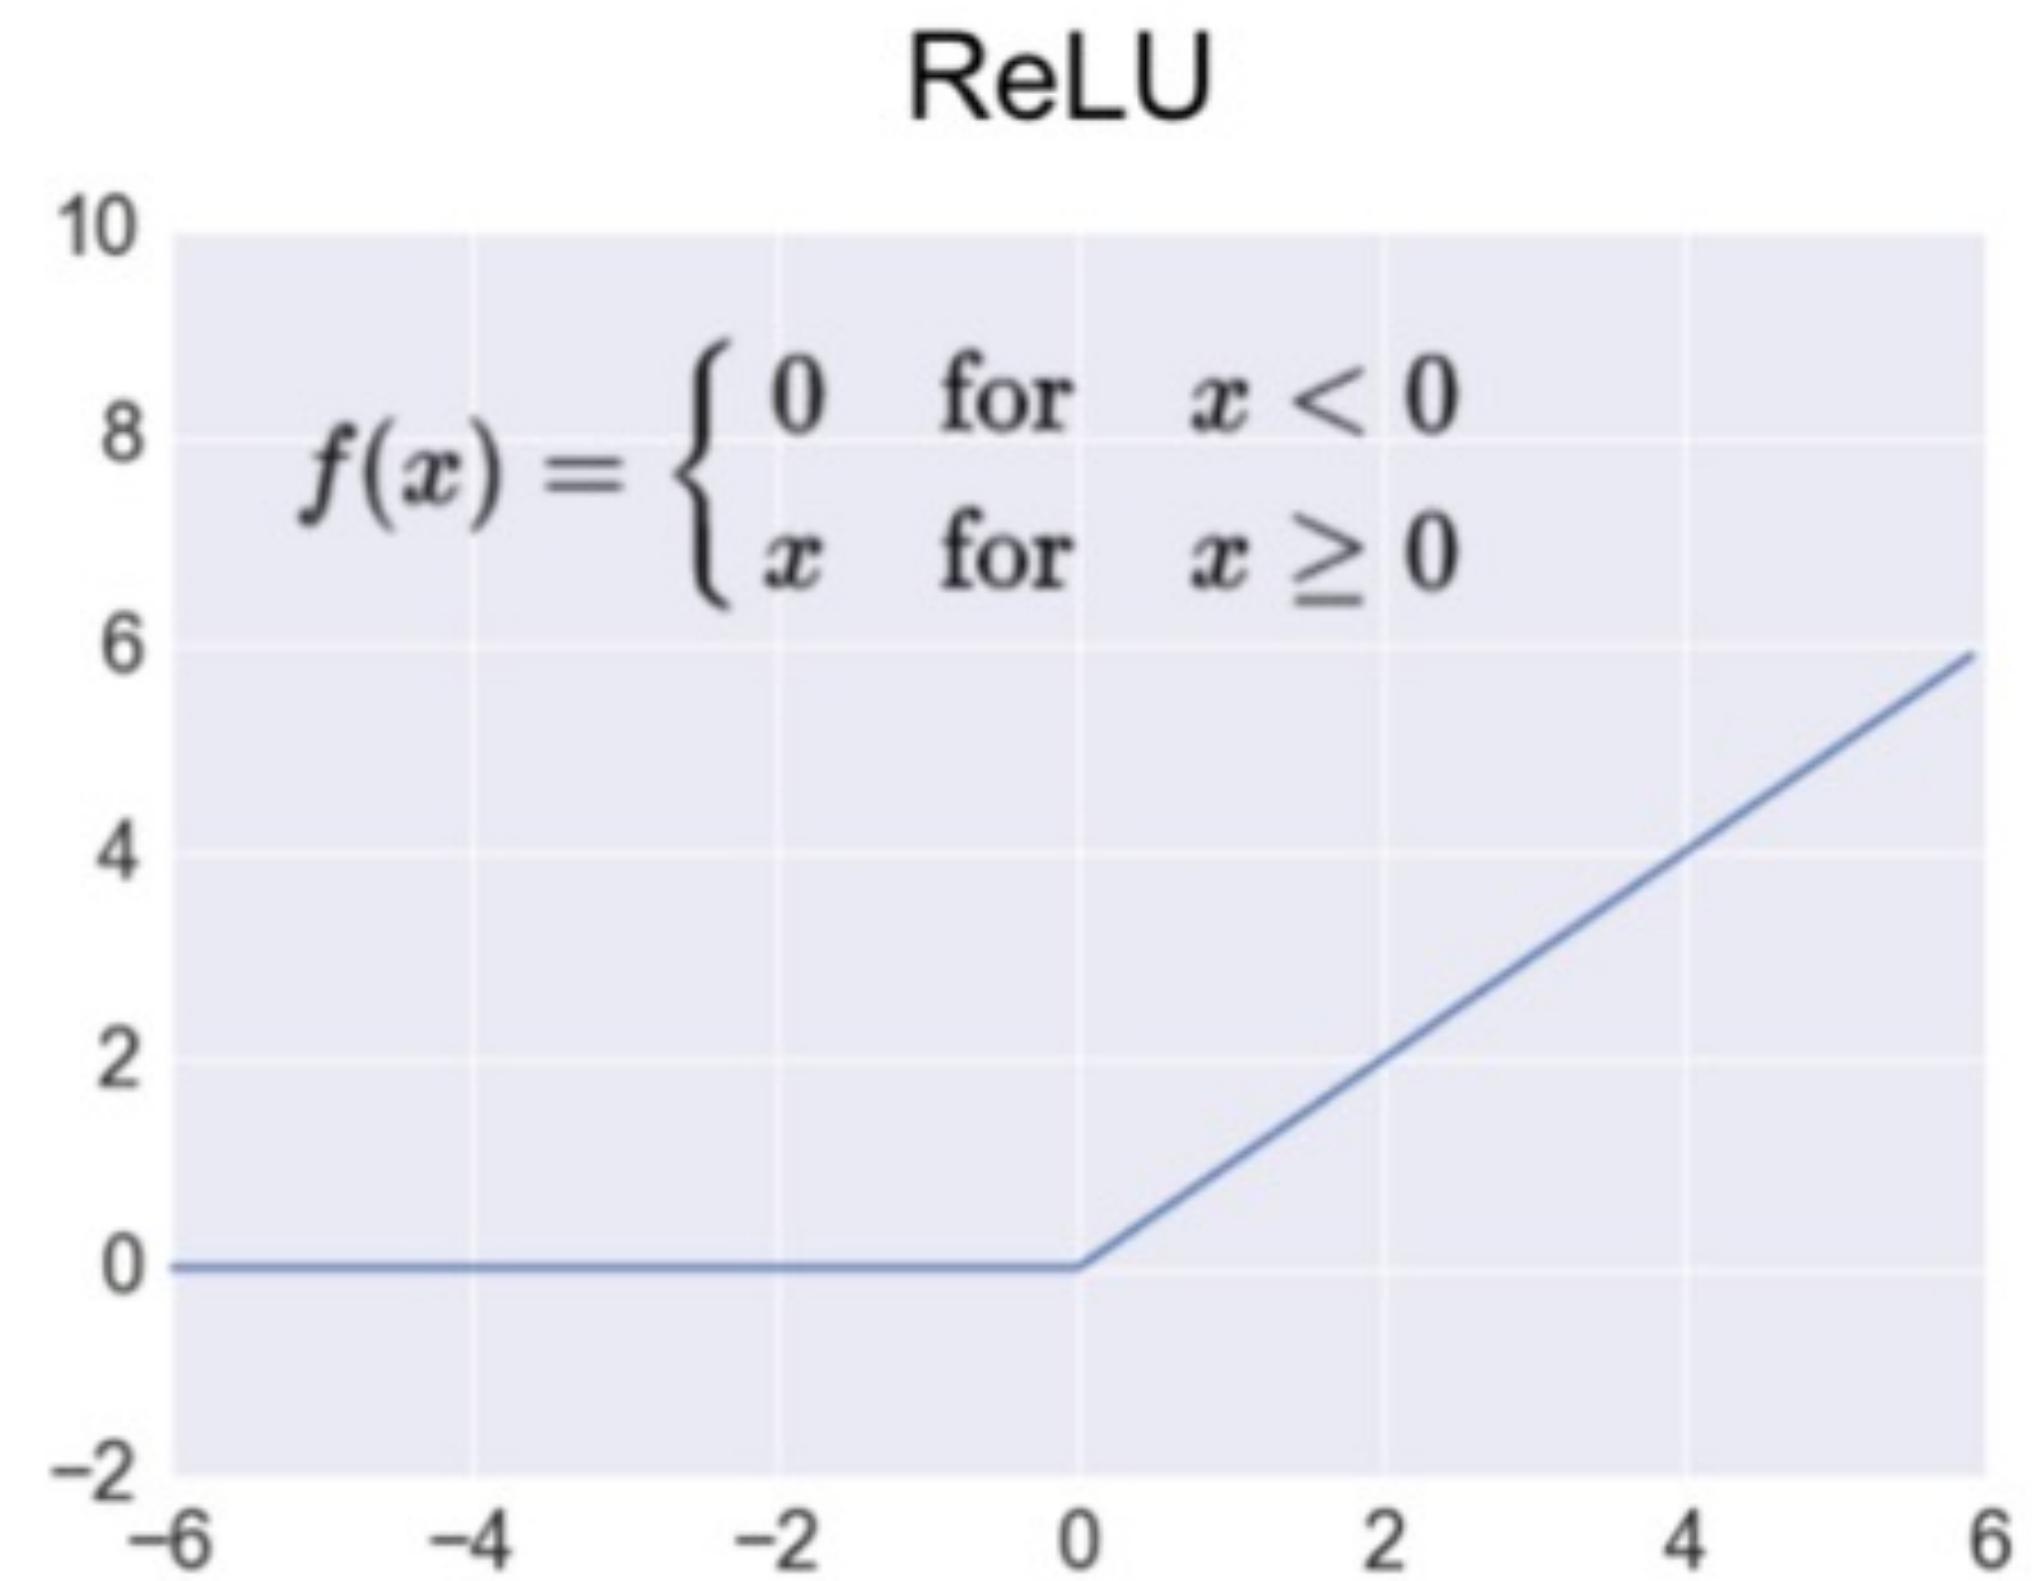
(a) Standard Neural Net



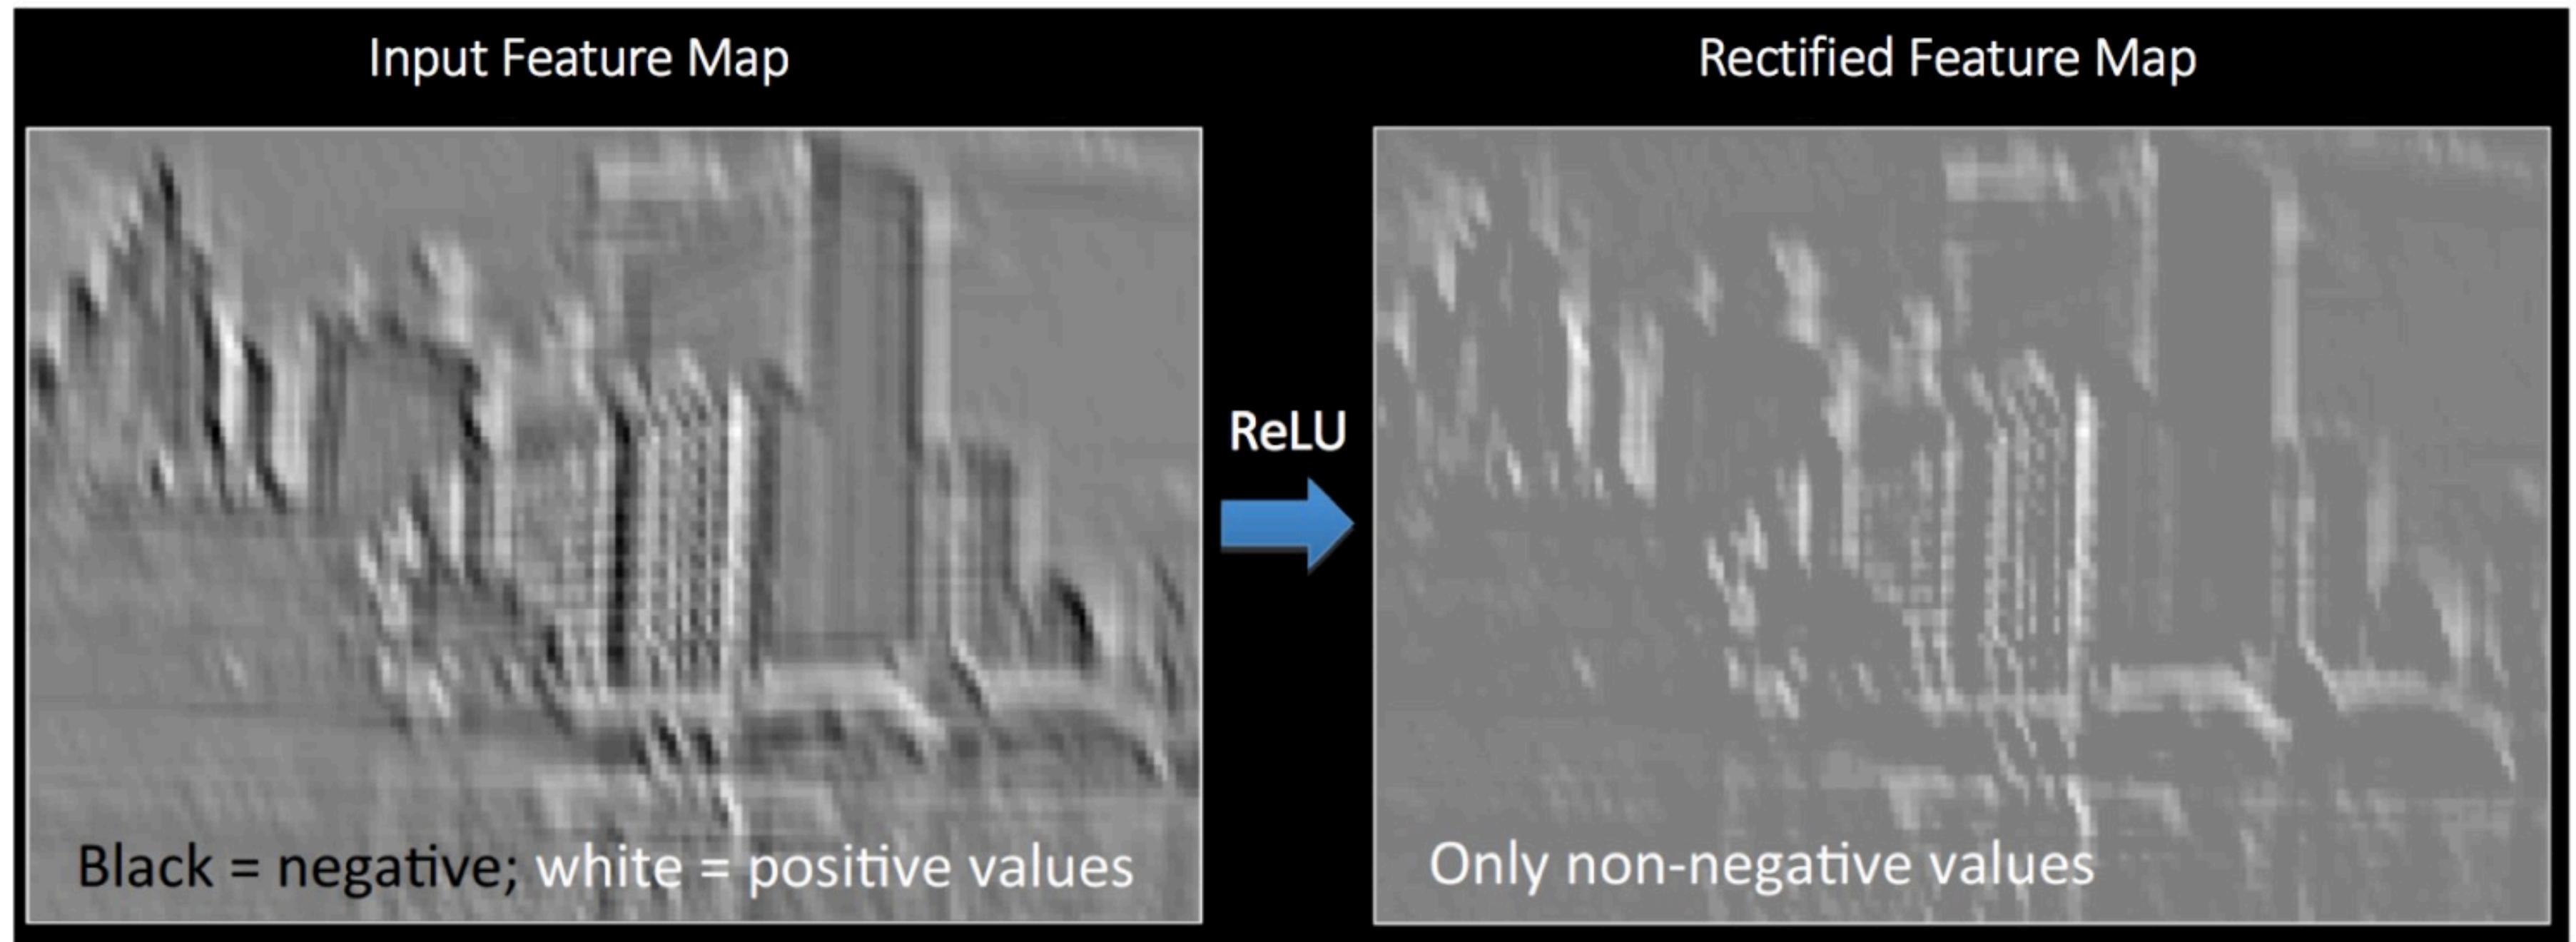
(b) After applying dropout

- Dropping a node means that we consider its activation 0
- It prevents the network memorising the training data (i.e. overfitting the data)
- It will help the network generalising better
- It is used only in training! When making predictions dropout is turned off.

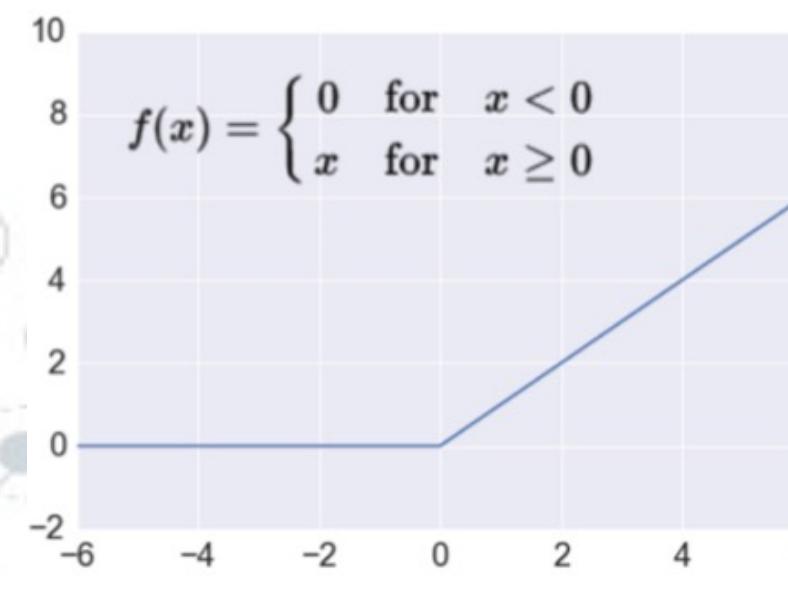
Activation function:RELU function



Activation function: RELU function



ReLU



A short history of deep learning

Ivaknenko:
“working
algorithm
algorithm for
supervised,
deep,
feedforward,
multilayer
perceptrons”

1965

Fukushima: deep
learning architecture
for computer vision

1980

The term of deep
learning is introduced

1989

LSTM
networks
proposed for
speech
recognition

1986

LeCunn:
Applied back-
propagation
algorithm, to a
deep neural
network with the
purpose of
recognising
handwritten ZIP
codes

2000

The term of “deep” is
Used in the context
of neural networks

1997

Ciresan: Proposed
CNN architecture
trained on GPU to
improve the results
of image
classification

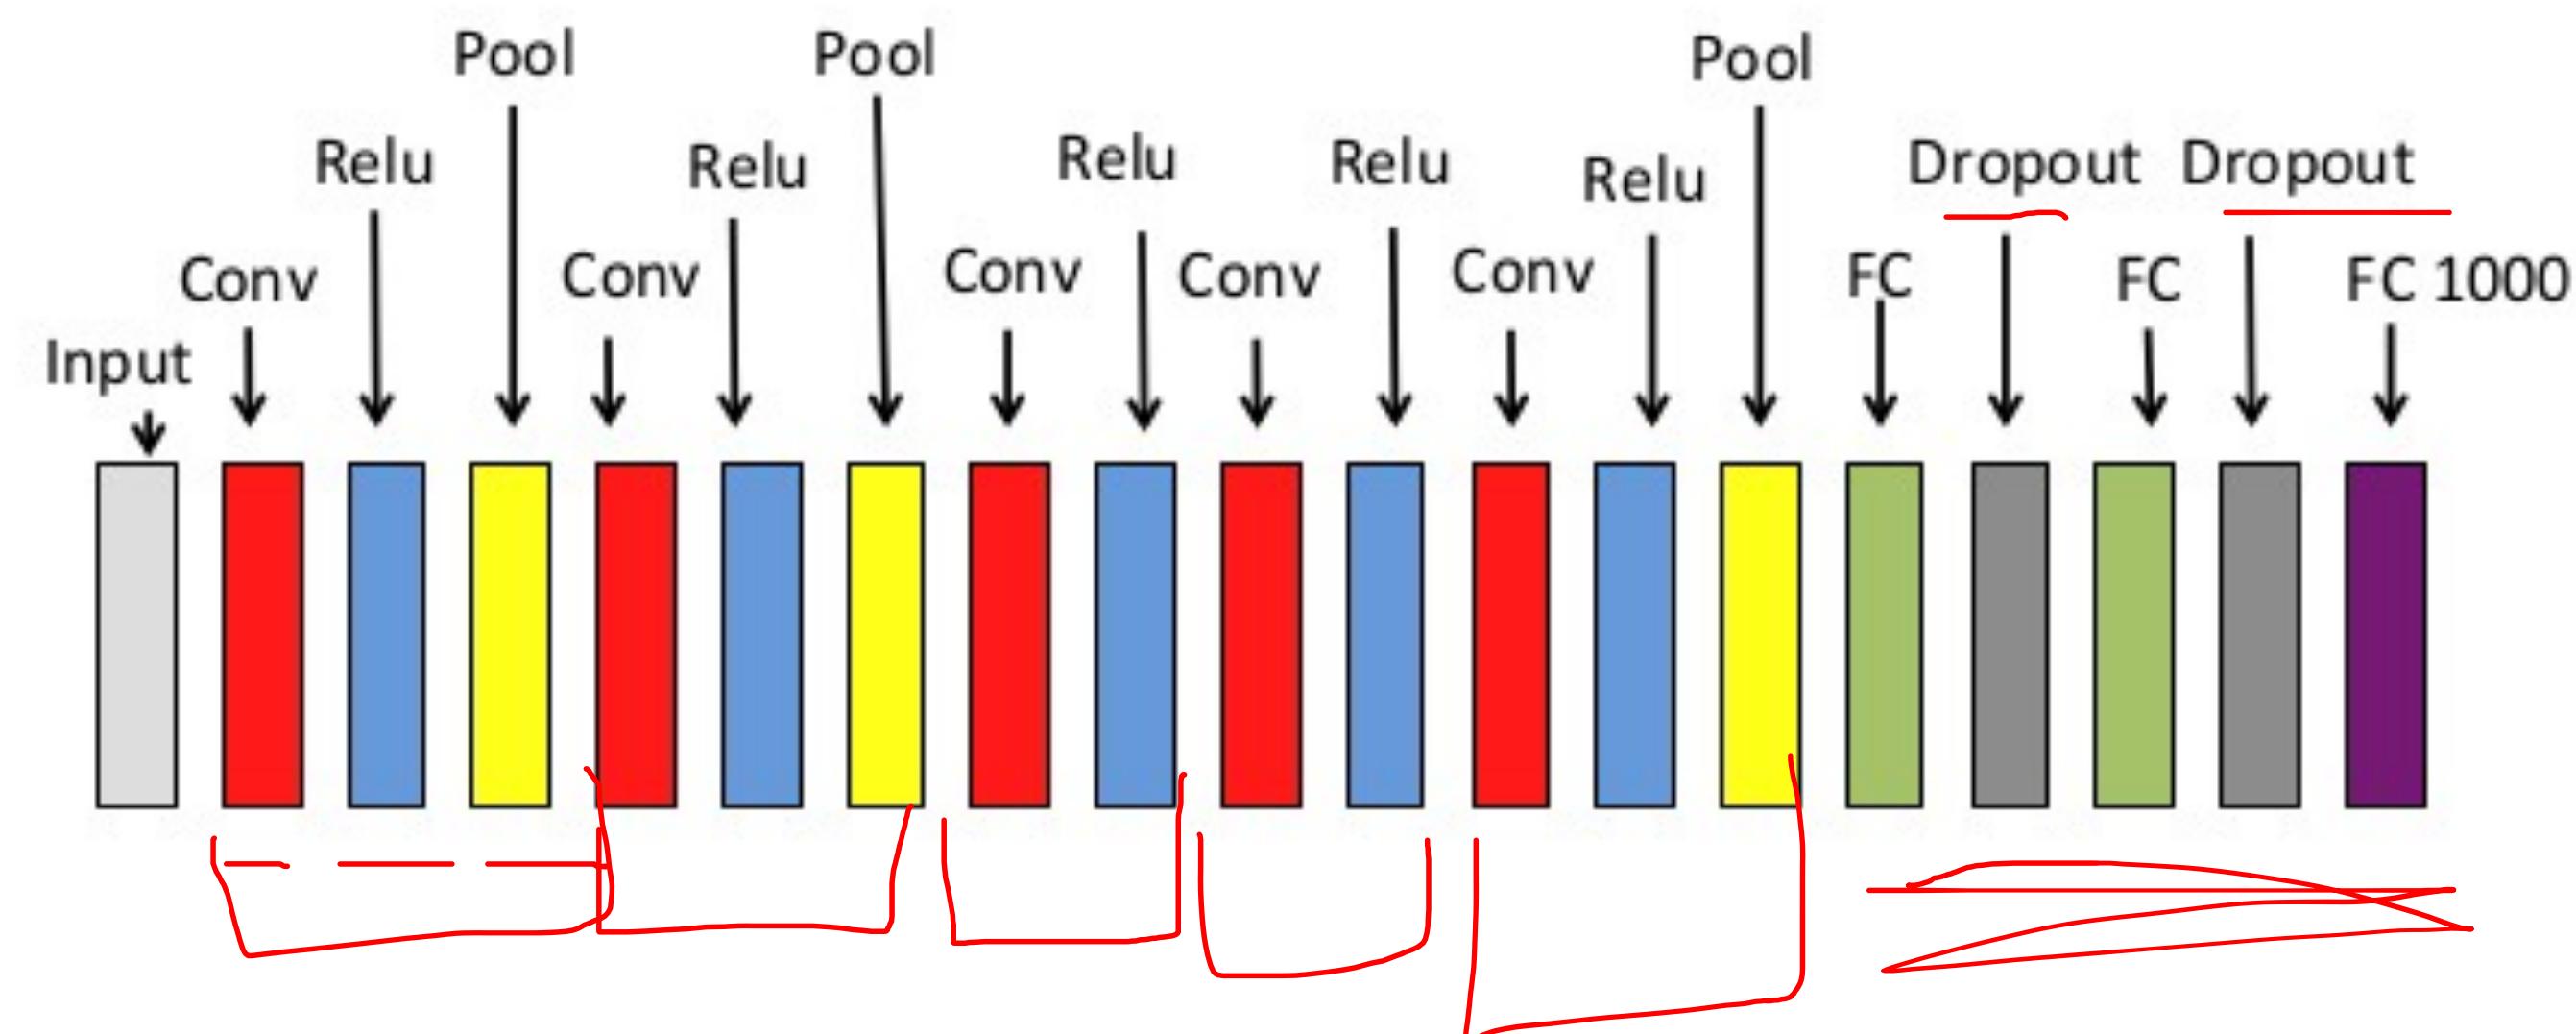
2011

AlexNet
architecture
won the
ImageNet
competition
by a large
margin

2012

The network that showed the power of convolutions: AlexNet 2012

Alexnet Architecture - 2012

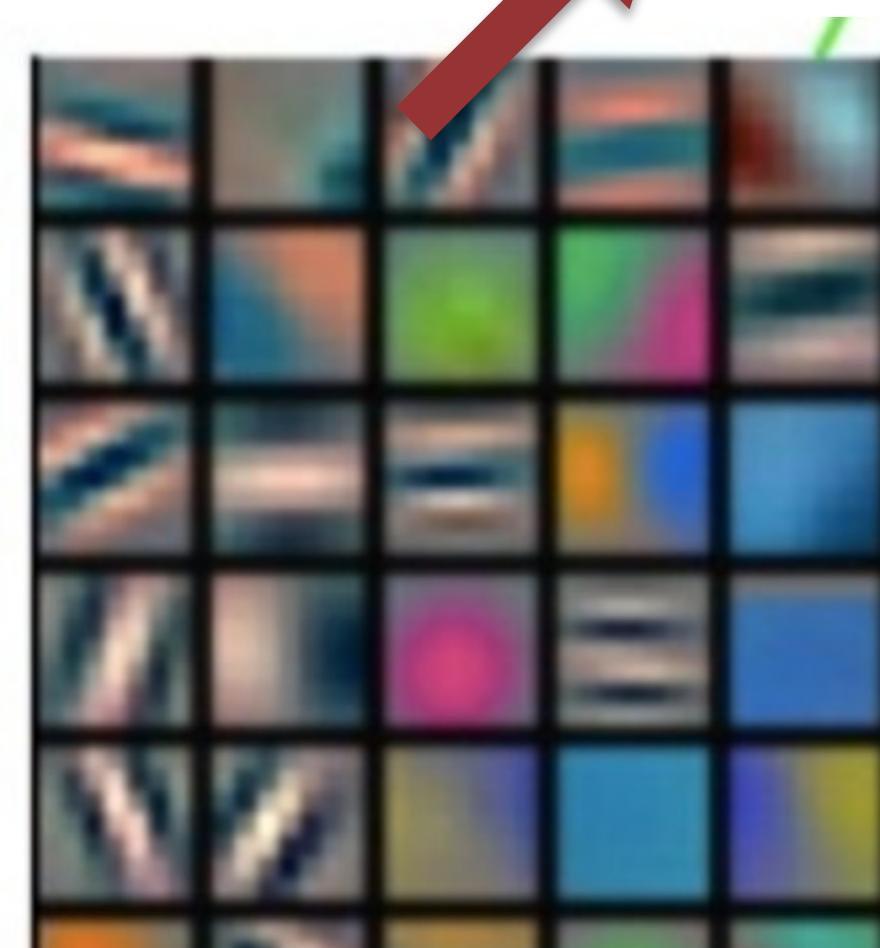
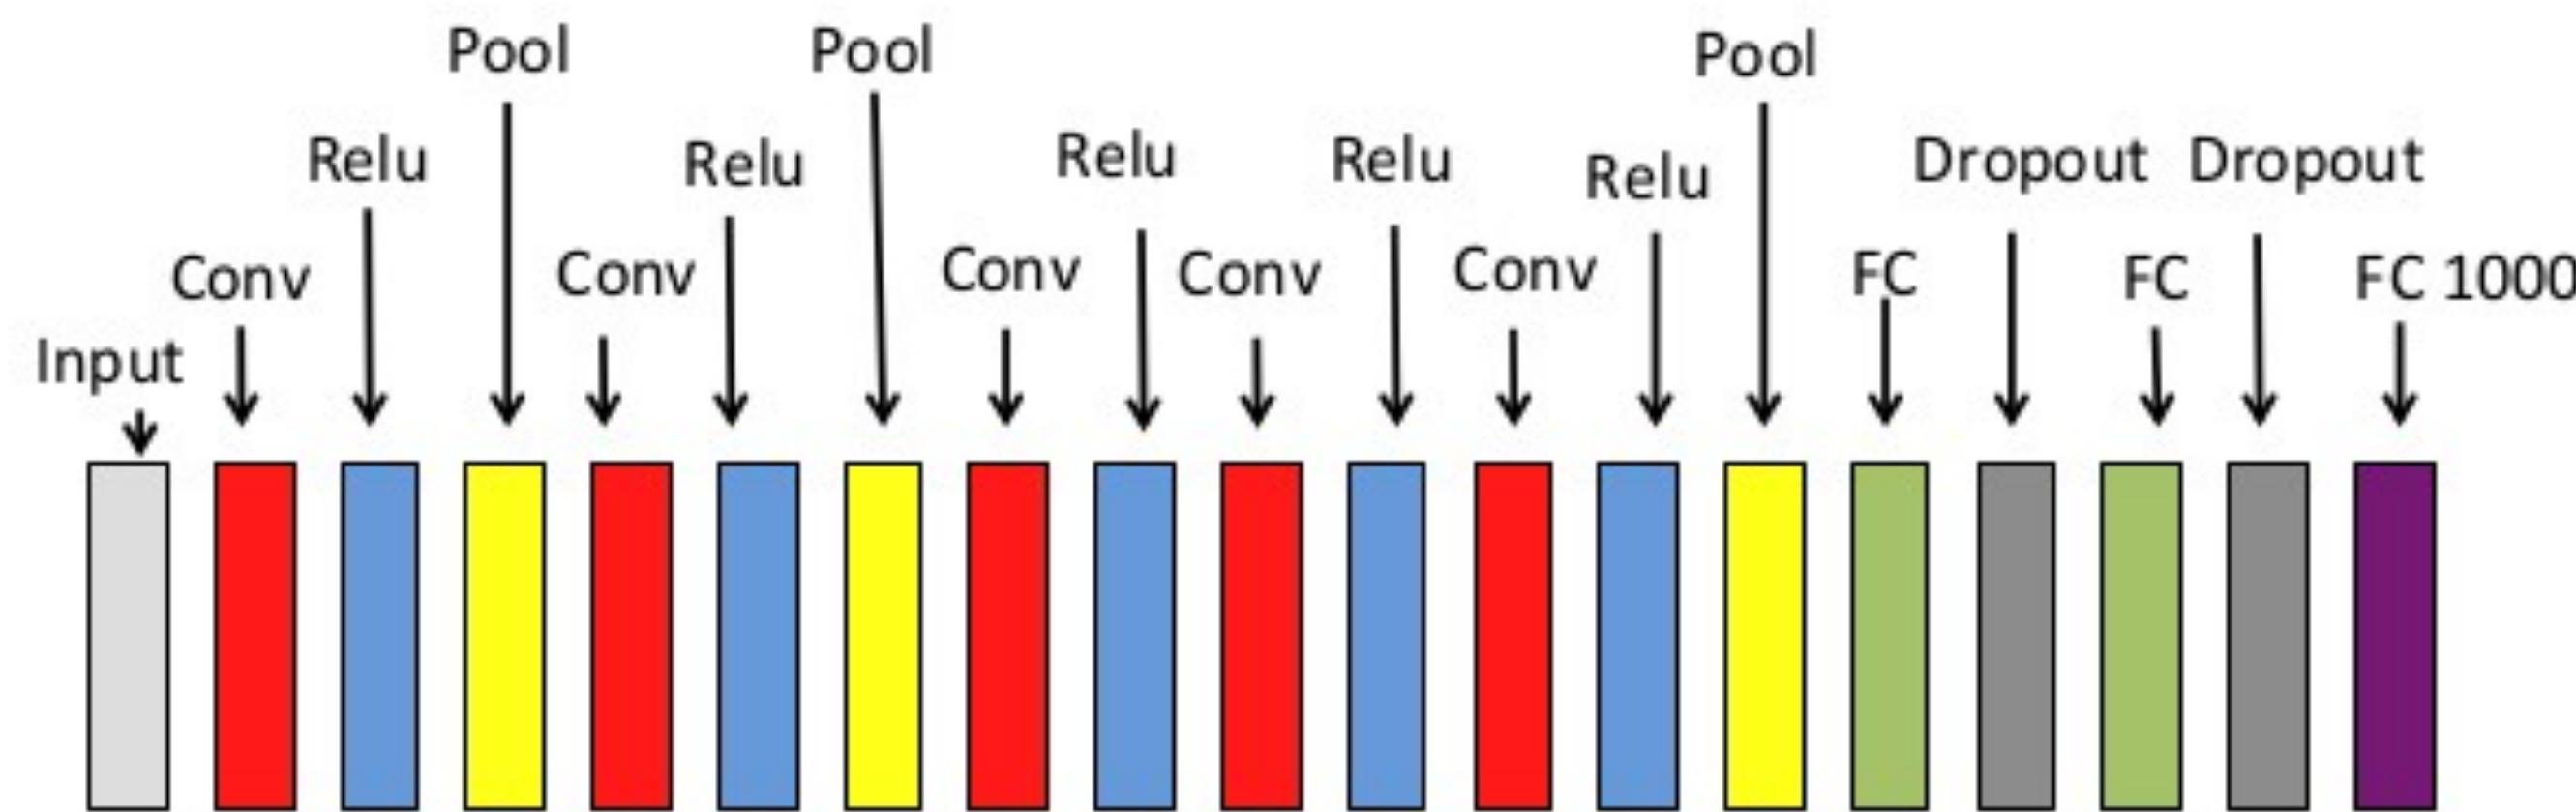


ImageNet Classification with Deep Convolutional Neural Networks Alex Krizhevsky, Ilya Sutskever and Geoffrey E. Hinton
Advances in Neural Information Processing Systems 25 eds. F. Pereira, C.J.C. Burges, L. Bottou and K.Q. Weinberger pp.
1097-1105, 2012

<http://adilmoujahid.com/posts/2016/06/introduction-deep-learning-python-cafe/>

AlexNet 2012

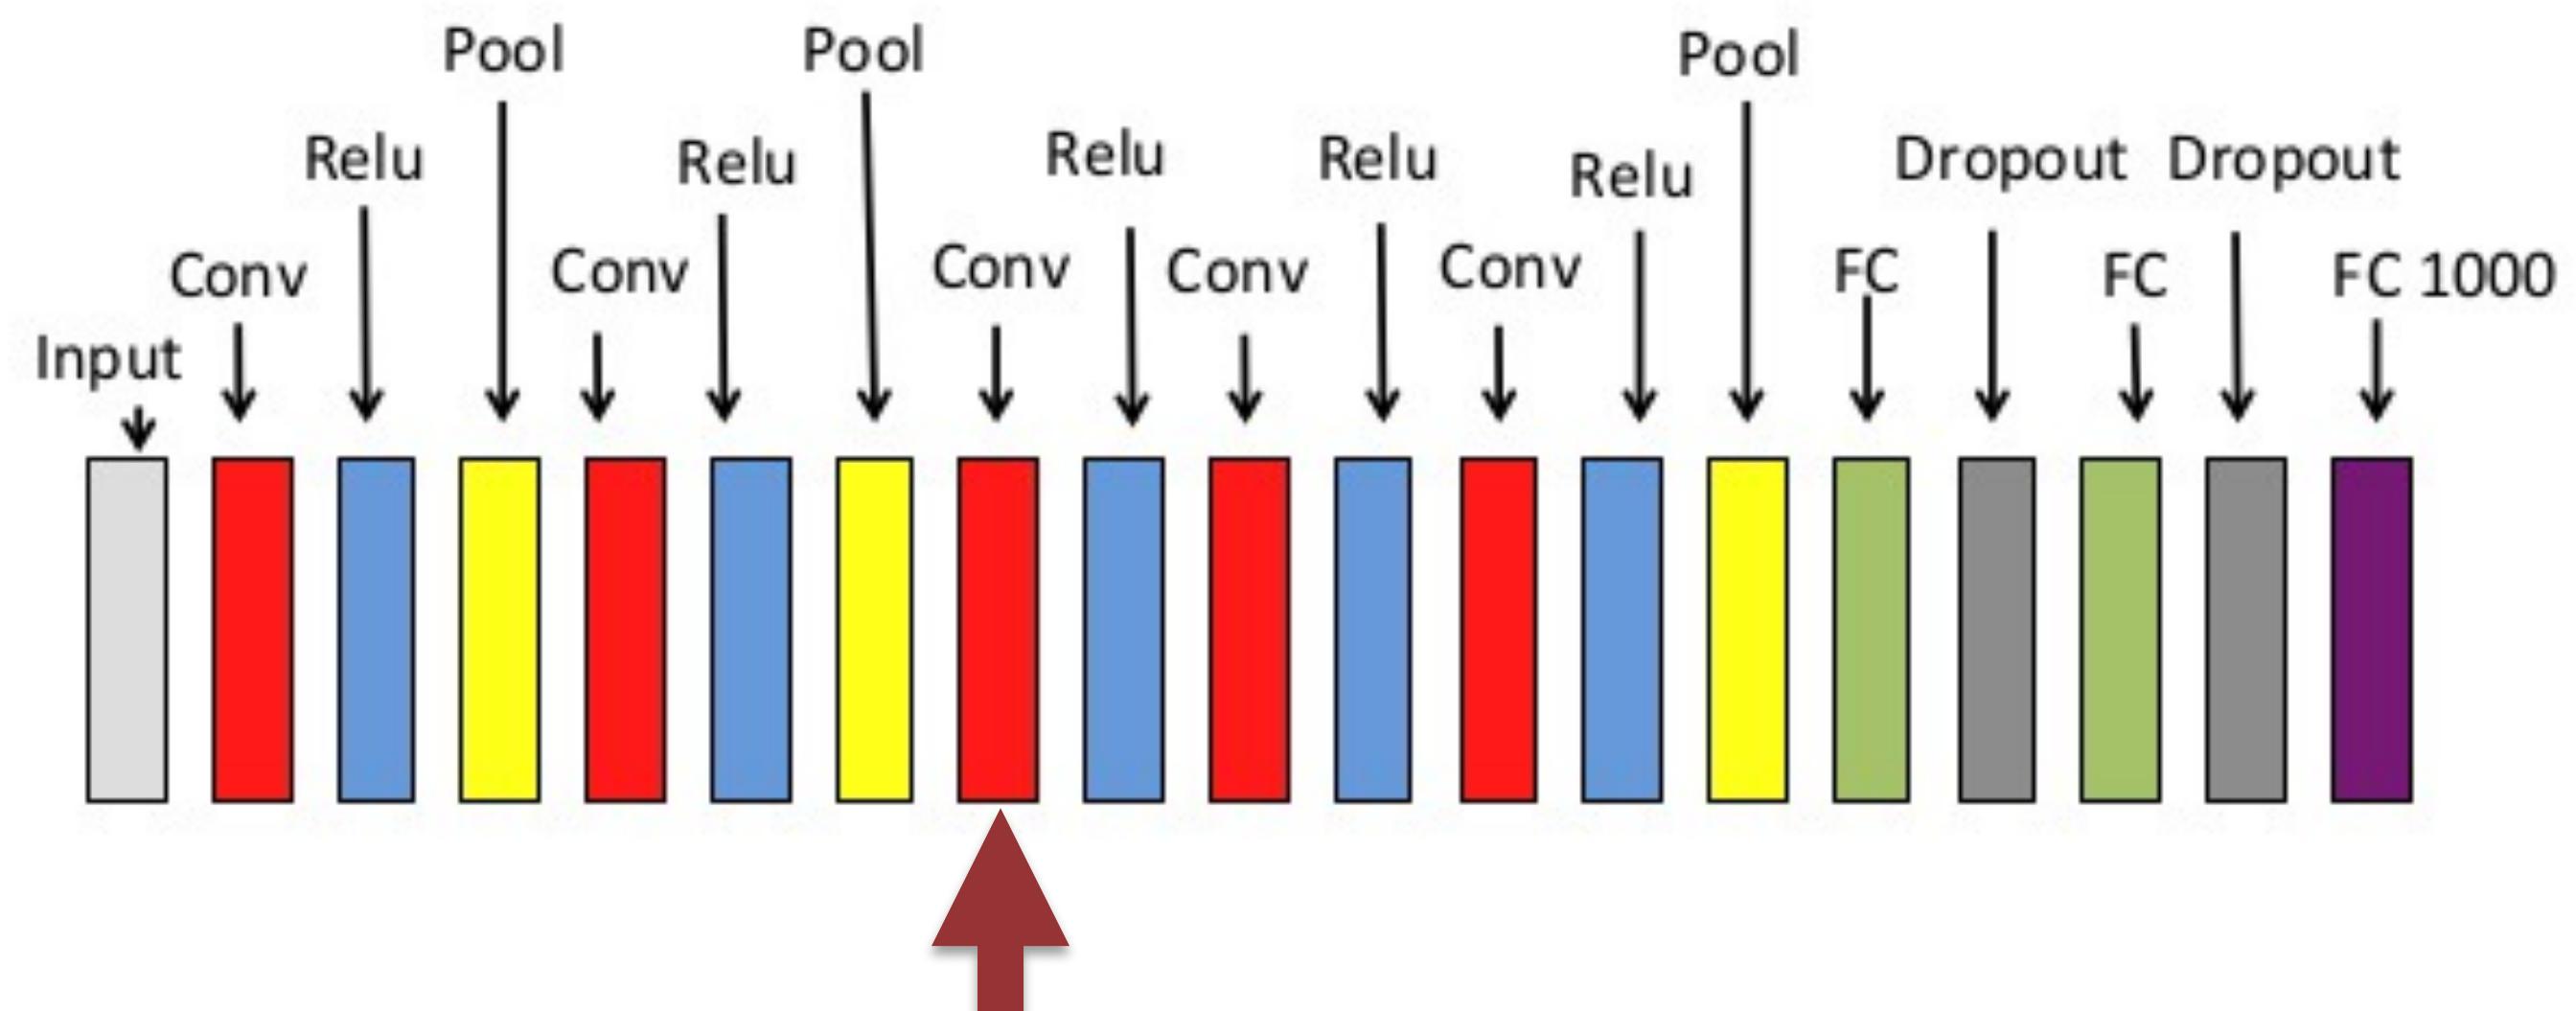
Alexnet Architecture - 2012



Computational Neural Networks Alex Krizhevsky, Ilya Sutskever and Geoffrey E. Hinton
in: Advances in Neural Information Processing Systems 25 eds. F. Pereira, C.J.C. Burges, L. Bottou and K.Q. Weinberger pp.
1097-1105, 2012

AlexNet 2012

Alexnet Architecture - 2012

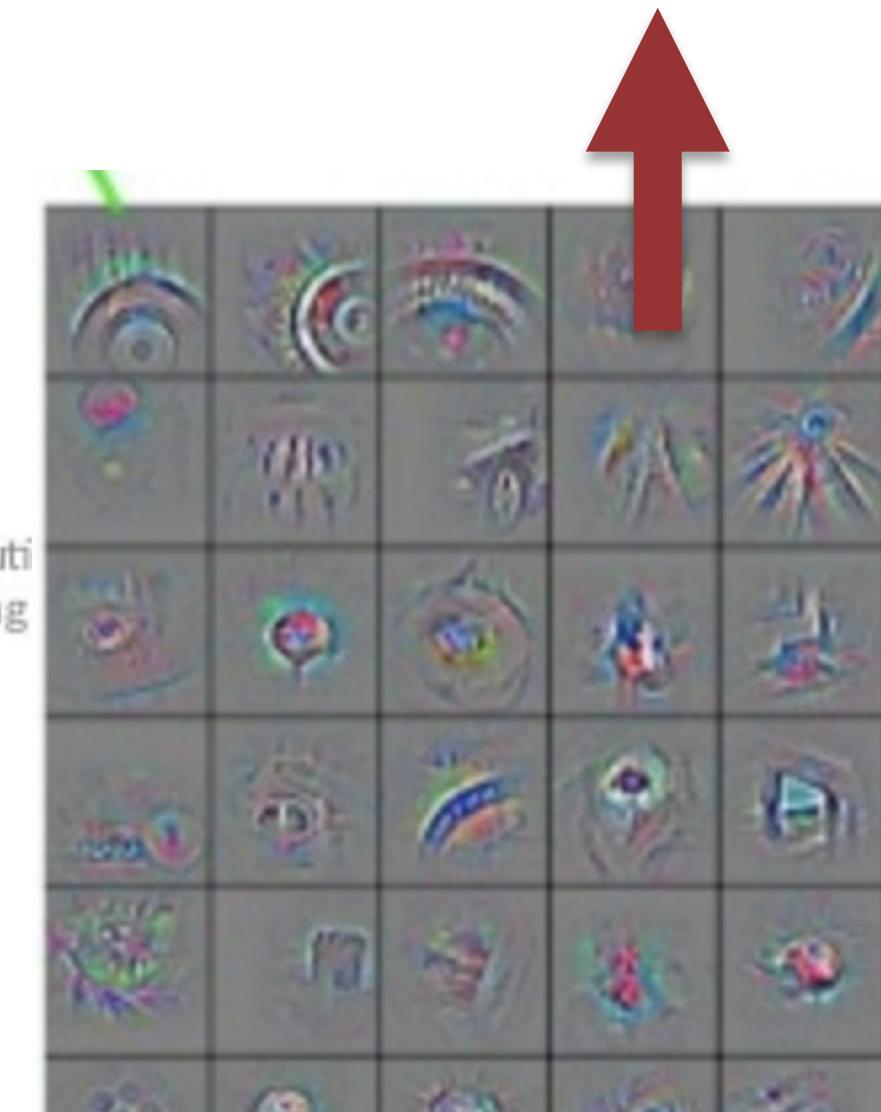
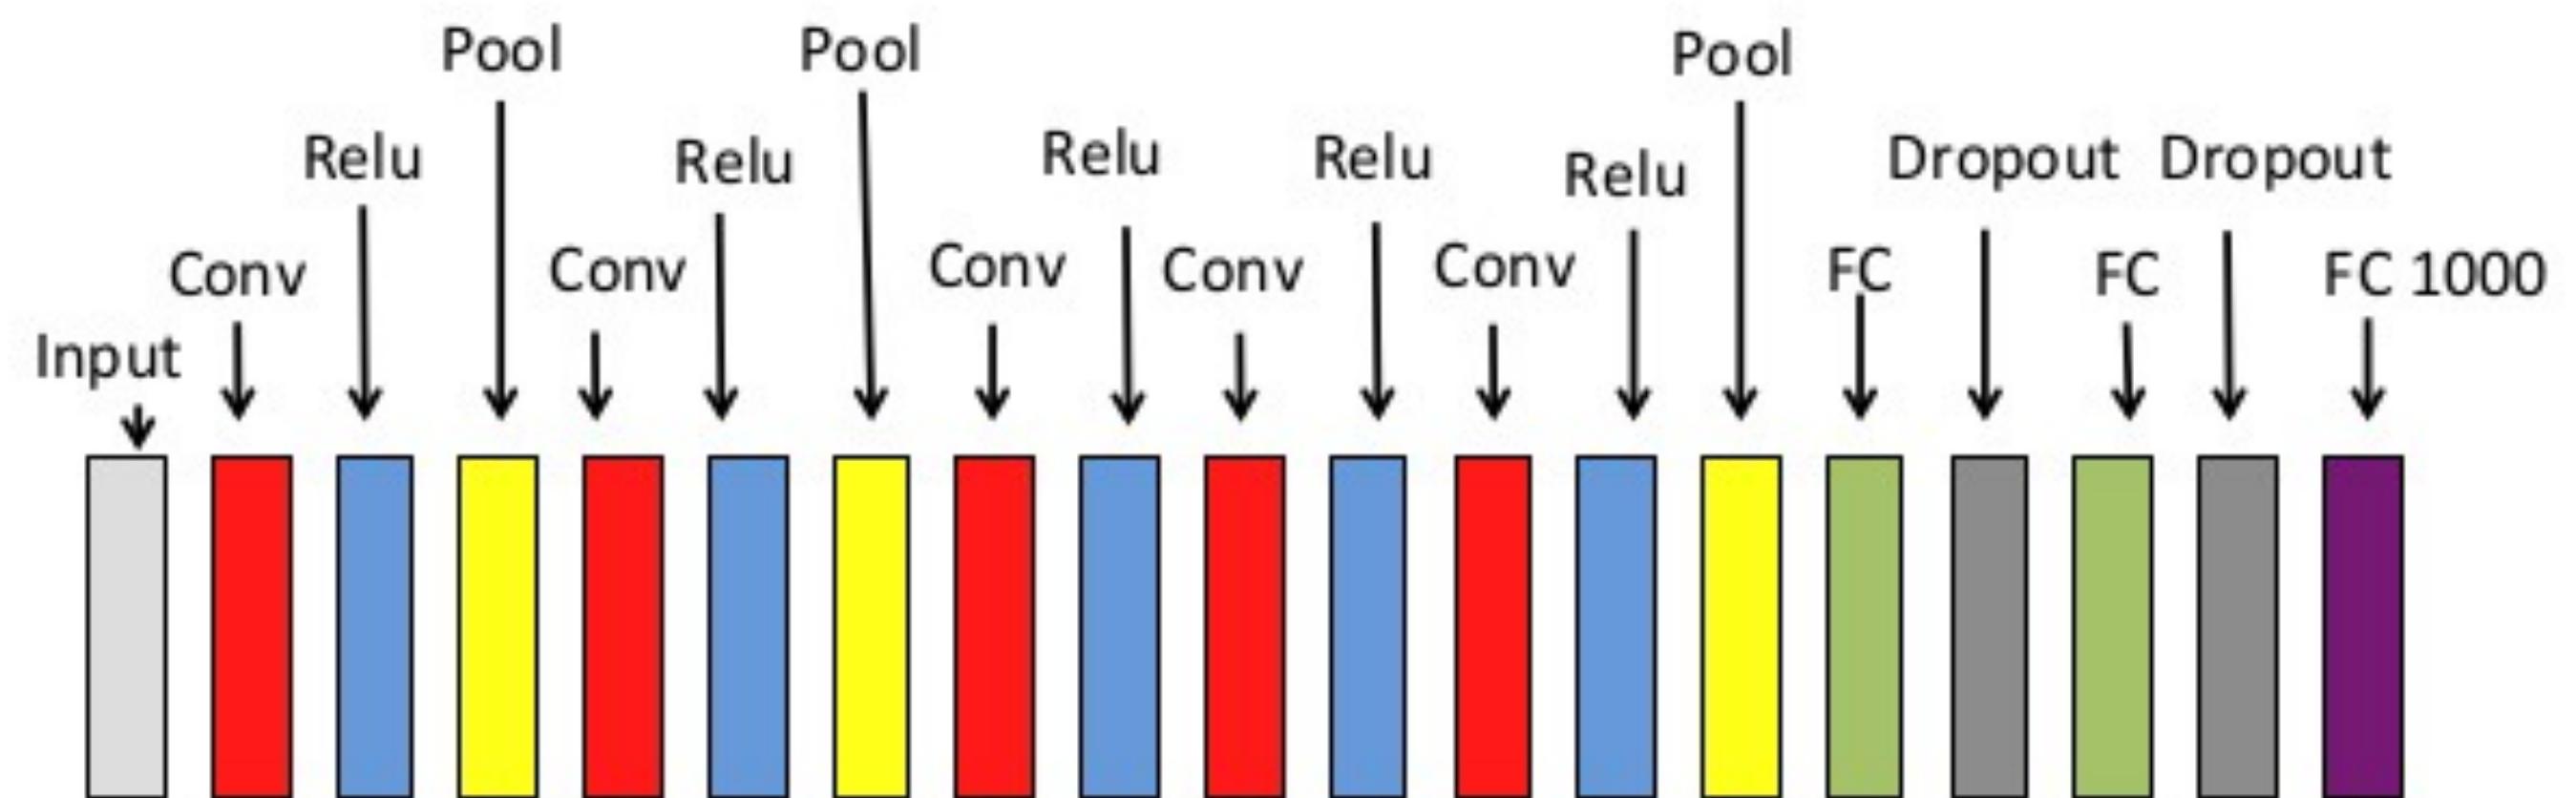


ImageNet Classification with Deep C
Advances in Neural Information Pr

ya Sutskever and Geoffrey E. Hinton
L. Bottou and K.Q. Weinberger pp.

AlexNet 2012

Alexnet Architecture - 2012

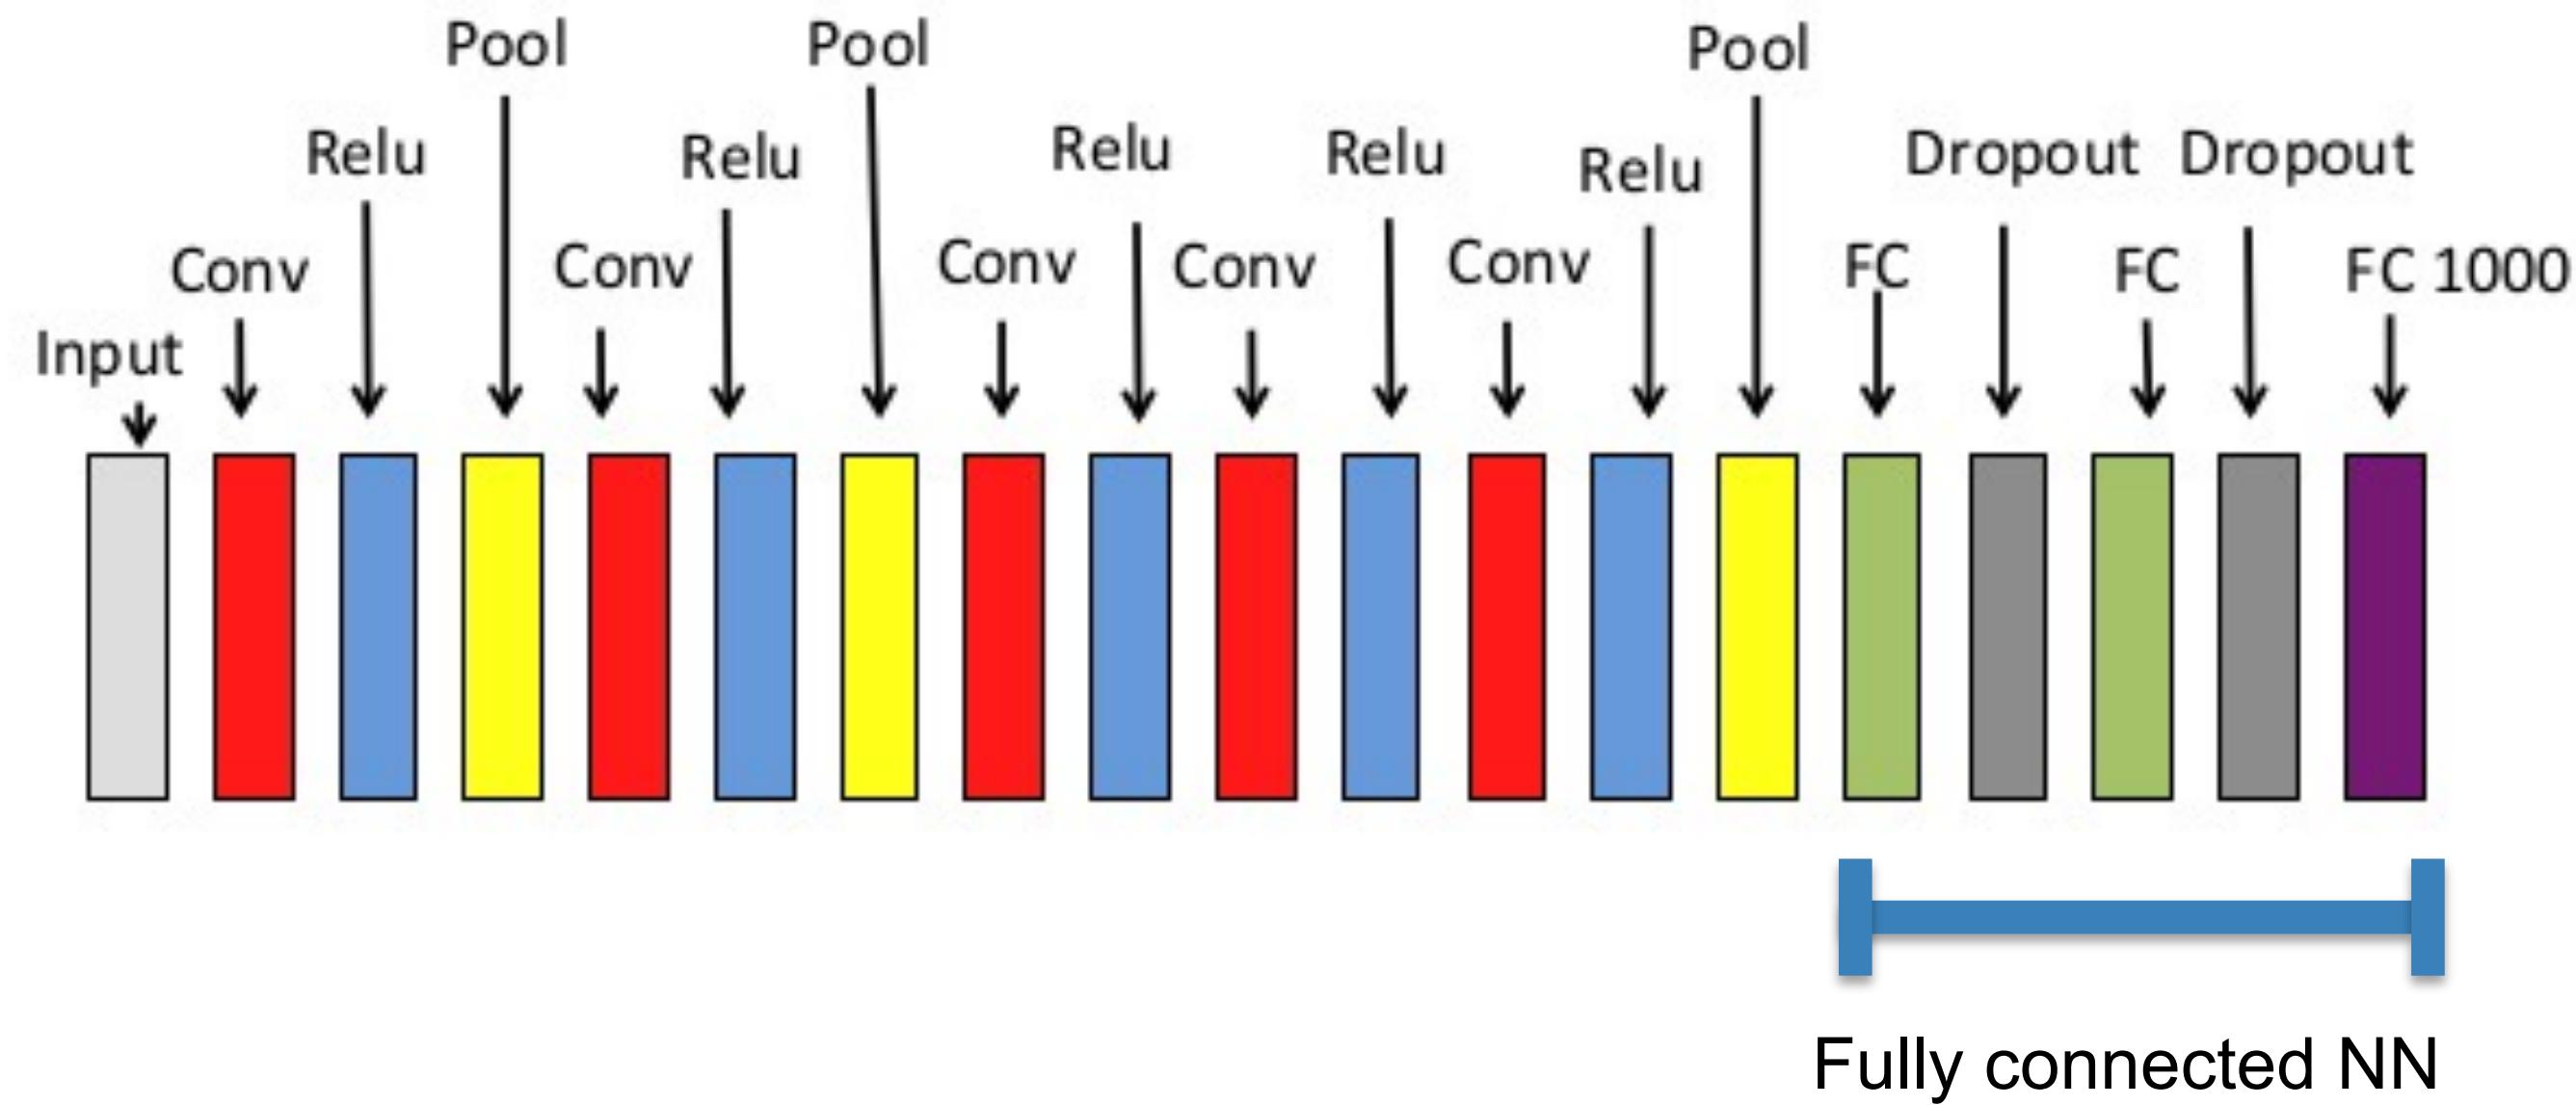


ImageNet Classification with Deep Convolutional Networks
Advances in Neural Information Processing Systems 25

Ruslan R. Salakhutdinov and Geoffrey E. Hinton
Yann LeCun, Yoshua Bengio, and K.Q. Weinberger pp.

AlexNet 2012

Alexnet Architecture - 2012



ImageNet Classification with Deep Convolutional Neural Networks Alex Krizhevsky, Ilya Sutskever and Geoffrey E. Hinton
Advances in Neural Information Processing Systems 25 eds. F. Pereira, C.J.C. Burges, L. Bottou and K.Q. Weinberger pp.
1097-1105, 2012

How many layers can we have?

Deep Residual Learning for Image Recognition

Kaiming He

Xiangyu Zhang

Shaoqing Ren

Jian Sun

Microsoft Research

{kahe, v-xiangz, v-shren, jiansun}@microsoft.com

v1 [CS.CV] 10 Dec 2015

Abstract

Deeper neural networks are more difficult to train. We present a residual learning framework to ease the training of networks that are substantially deeper than those used previously. We explicitly reformulate the layers as learning residual functions with reference to the layer inputs, instead of learning unreferenced functions. We provide comprehensive empirical evidence showing that these residual networks are easier to optimize, and can gain accuracy from considerably increased depth. On the ImageNet dataset we evaluate residual nets with a depth of up to 152 layers—8× deeper than VGG nets [41] but still having lower complexity. An ensemble of these residual nets achieves 3.57% error on the ImageNet test set. This result won the 1st place on the ILSVRC 2015 classification task. We also present analysis

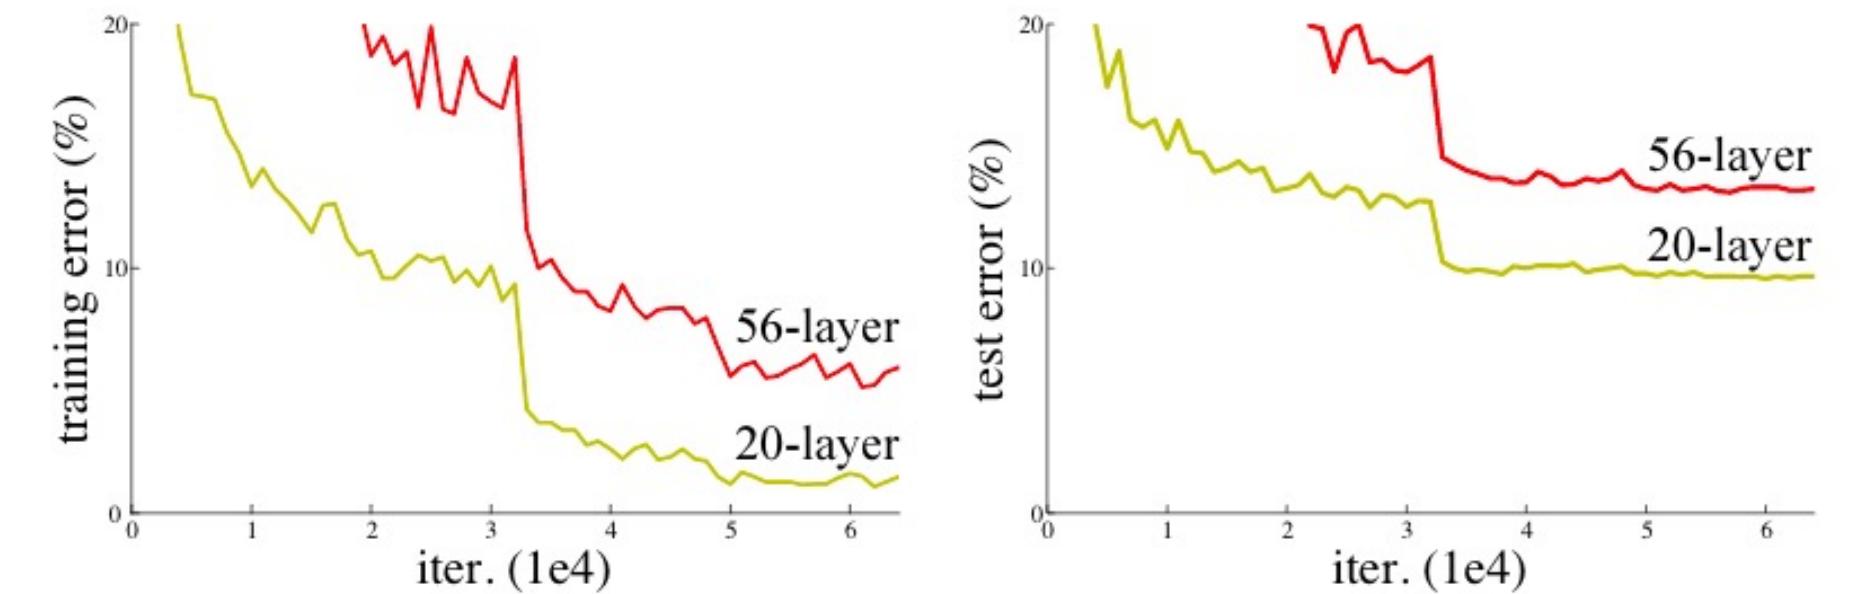


Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.

greatly benefited from very deep models.

Driven by the significance of depth, a question arises: *Is learning better networks as easy as stacking more layers?* An obstacle to answering this question was the notorious

How many layers can we have?

Deep Residual Learning for Image Recognition

Kaiming He

Xiangyu Zhang

Shaoqing Ren

Jian Sun

Microsoft Research

{kahe, v-xiangz, v-shren, jiansun}@microsoft.com

stead of learning unreferenced functions. We provide comprehensive empirical evidence showing that these residual networks are easier to optimize, and can gain accuracy from considerably increased depth. On the ImageNet dataset we evaluate residual nets with a depth of up to 152 layers—8× deeper than VGG nets [41] but still having lower complexity. An ensemble of these residual nets achieves 3.57% error on the ImageNet test set. This result won the 1st place on the ILSVRC 2015 classification task. We also present analysis on CIFAR-10 with 100 and 1000 layers.

The depth of representations is of central importance for many visual recognition tasks. Solely due to our extremely deep representations, we obtain a 28% relative im-

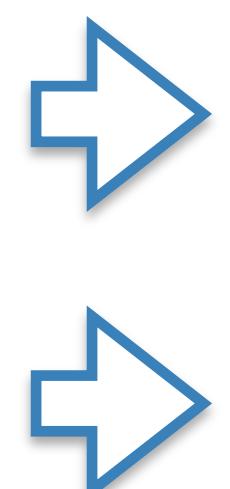
Software libraries

Existing frameworks and libraries

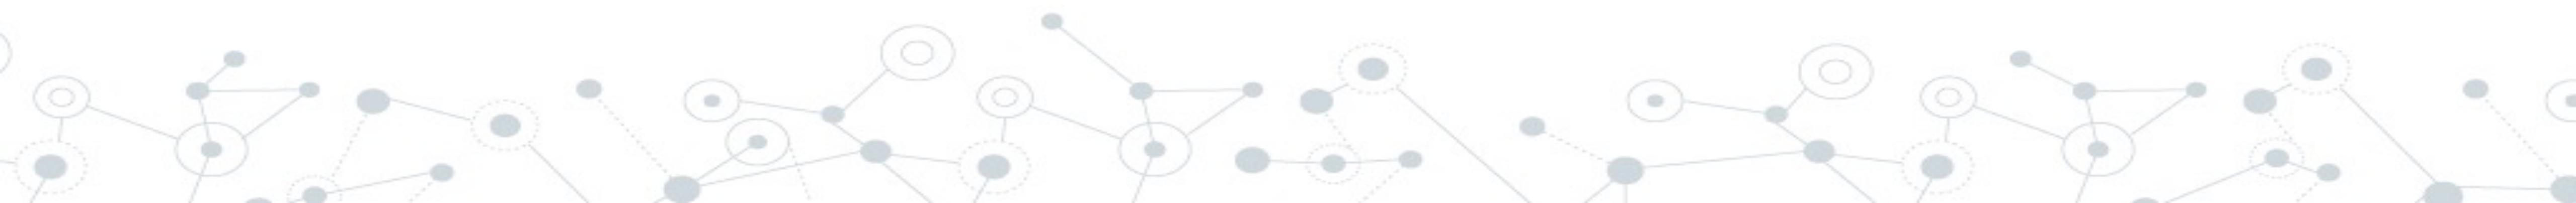
Library	Language(s)	Licence	Author	URL
TensorFlow	Python/C++/Java	Apache 2.0	Google et al.	tensorflow.org
Keras	Python	MIT	François Chollet et al.	keras.io
Torch	Lua/C++	Modified BSD	Idiap Research Institute et al.	torch.ch
PyTorch	Python/C++	Modified BSD	Facebook et al.	pytorch.org
Microsoft Cognitive Toolkit (previously known as CNTK)	Python/C#/C++/Brain Script	MIT	Microsoft	docs.microsoft.com/en-gb/cognitive-toolkit/
Caffe	C++/Python/Matlab	BSD 2-Clause	Berkeley AI Research (BAIR)	caffe.berkeleyvision.org
Apache MXNet	Python/C++/Julia/Closure/R/Scala/Perl	Apache 2.0	Apache Foundation	mxnet.incubator.apache.org
DeepLearning4j	Java	Apache 2.0	Skydnn	deeplearning4j.org
Theano	Python	BSD 3-Clause	MILA	deeplearning.net/software/theano
TFLearn	Python	MIT	Aymeric Damien	tflearn.org
Cafe2	Python/C++	Apache 2.0	Facebook	caffe2.ai
Deep Learning Toolbox	Matlab/C++	Proprietary	MathWorks	uk.mathworks.com/products/deep-learning.html
SKLearn	Python	BSD	scikit-learn.org/stable/about.html	scikit-learn.org



Existing frameworks and libraries

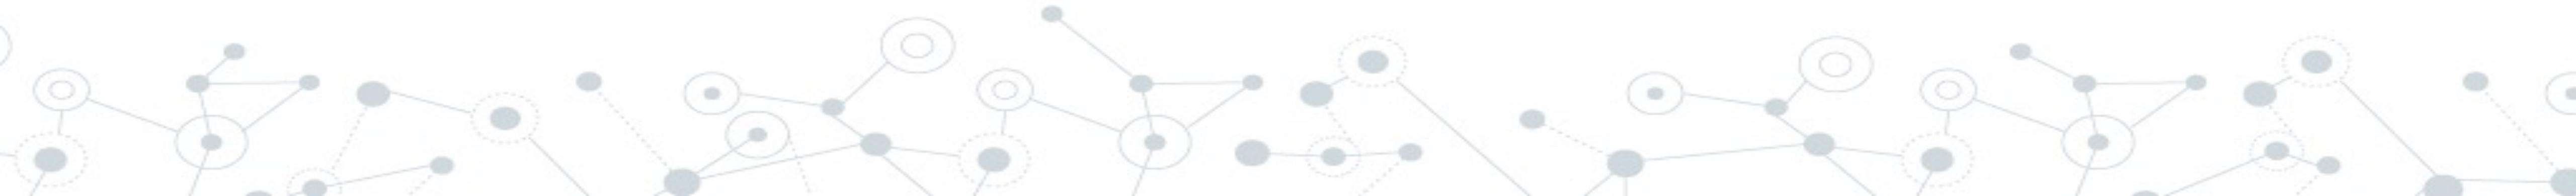


Library	Language(s)	Licence	Author	URL
TensorFlow	Python/C++/Java	Apache 2.0	Google et al.	tensorflow.org
Keras	Python	MIT	François Chollet et al.	keras.io
Torch	Lua/C++	Modified BSD	Idiap Research Institute et al.	torch.ch
PyTorch	Python/C++	Modified BSD	Facebook et al.	pytorch.org
Microsoft Cognitive Toolkit (previously known as CNTK)	Python/C#/C++/Brain Script	MIT	Microsoft	docs.microsoft.com/en-gb/cognitive-toolkit/
Caffe	C++/Python/Matlab	BSD 2-Clause	Berkeley AI Research (BAIR)	caffe.berkeleyvision.org
Apache MXNet	Python/C++/Julia/Closure/R/Scala/Perl	Apache 2.0	Apache Foundation	mxnet.incubator.apache.org
DeepLearning4j	Java	Apache 2.0	Skydnn	deeplearning4j.org
Theano	Python	BSD 3-Clause	MILA	deeplearning.net/software/theano
TFLearn	Python	MIT	Aymeric Damien	tflearn.org
Cafe2	Python/C++	Apache 2.0	Facebook	caffe2.ai
Deep Learning Toolbox	Matlab/C++	Proprietary	MathWorks	uk.mathworks.com/products/deep-learning.html
SKLearn	Python	BSD	scikit-learn.org/stable/about.html	scikit-learn.org



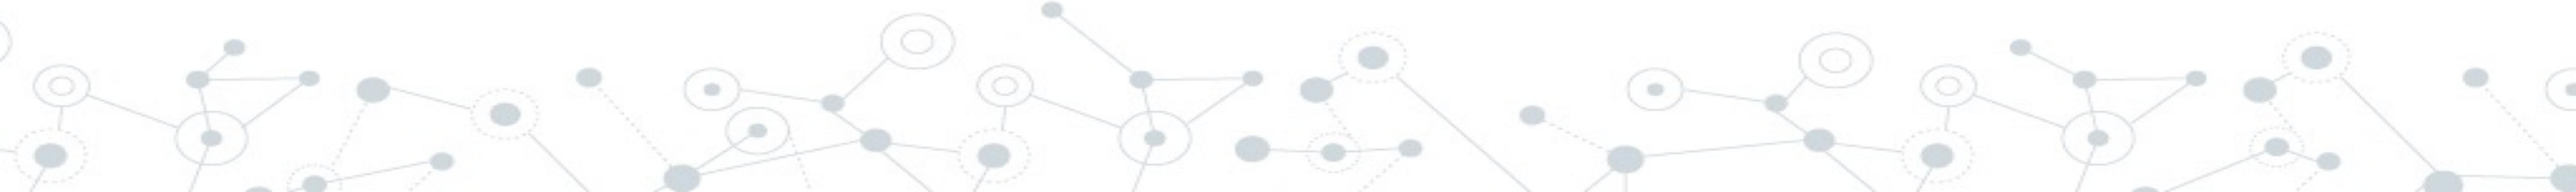
TensorFlow

- DeepLearning Framework developed by Google
- Widely available and used by companies in production
- Many models (pre-trained or code) available on Github
- Models trained on GPU can be easily queried on CPU
- Tools for distributed training, online parameters tracking, pre-processing, etc.



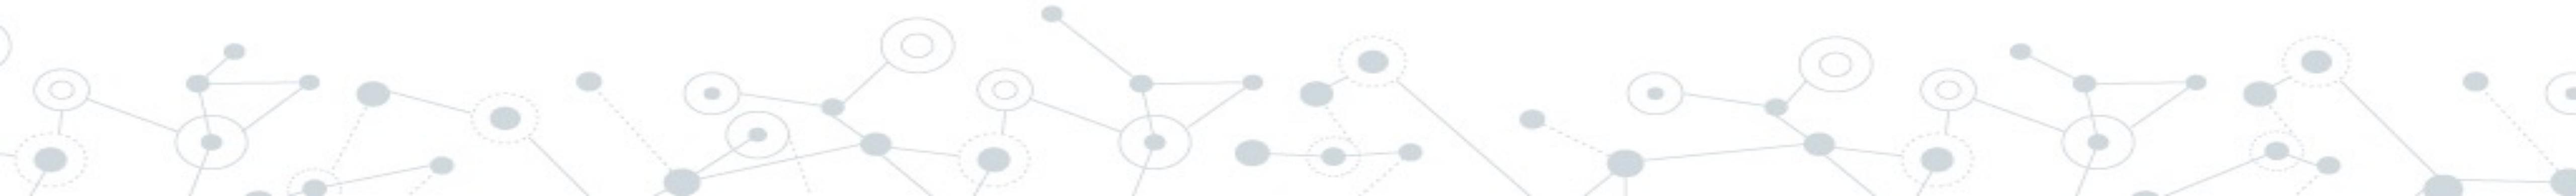
Deep Learning Advantages

- It works better than other methods for image recognition, speech recognition etc.
- Reduces the need of feature engineering
- The network architectures can be reused for a variety of problems



Deep Learning Disadvantages

- Computational expensive to train
- Requires a LARGE amount of data (with few examples, deep learning is very unlikely to outperform other machine learning methods)
- Sometimes it is a black box, difficult to debug and understand what is actually going on



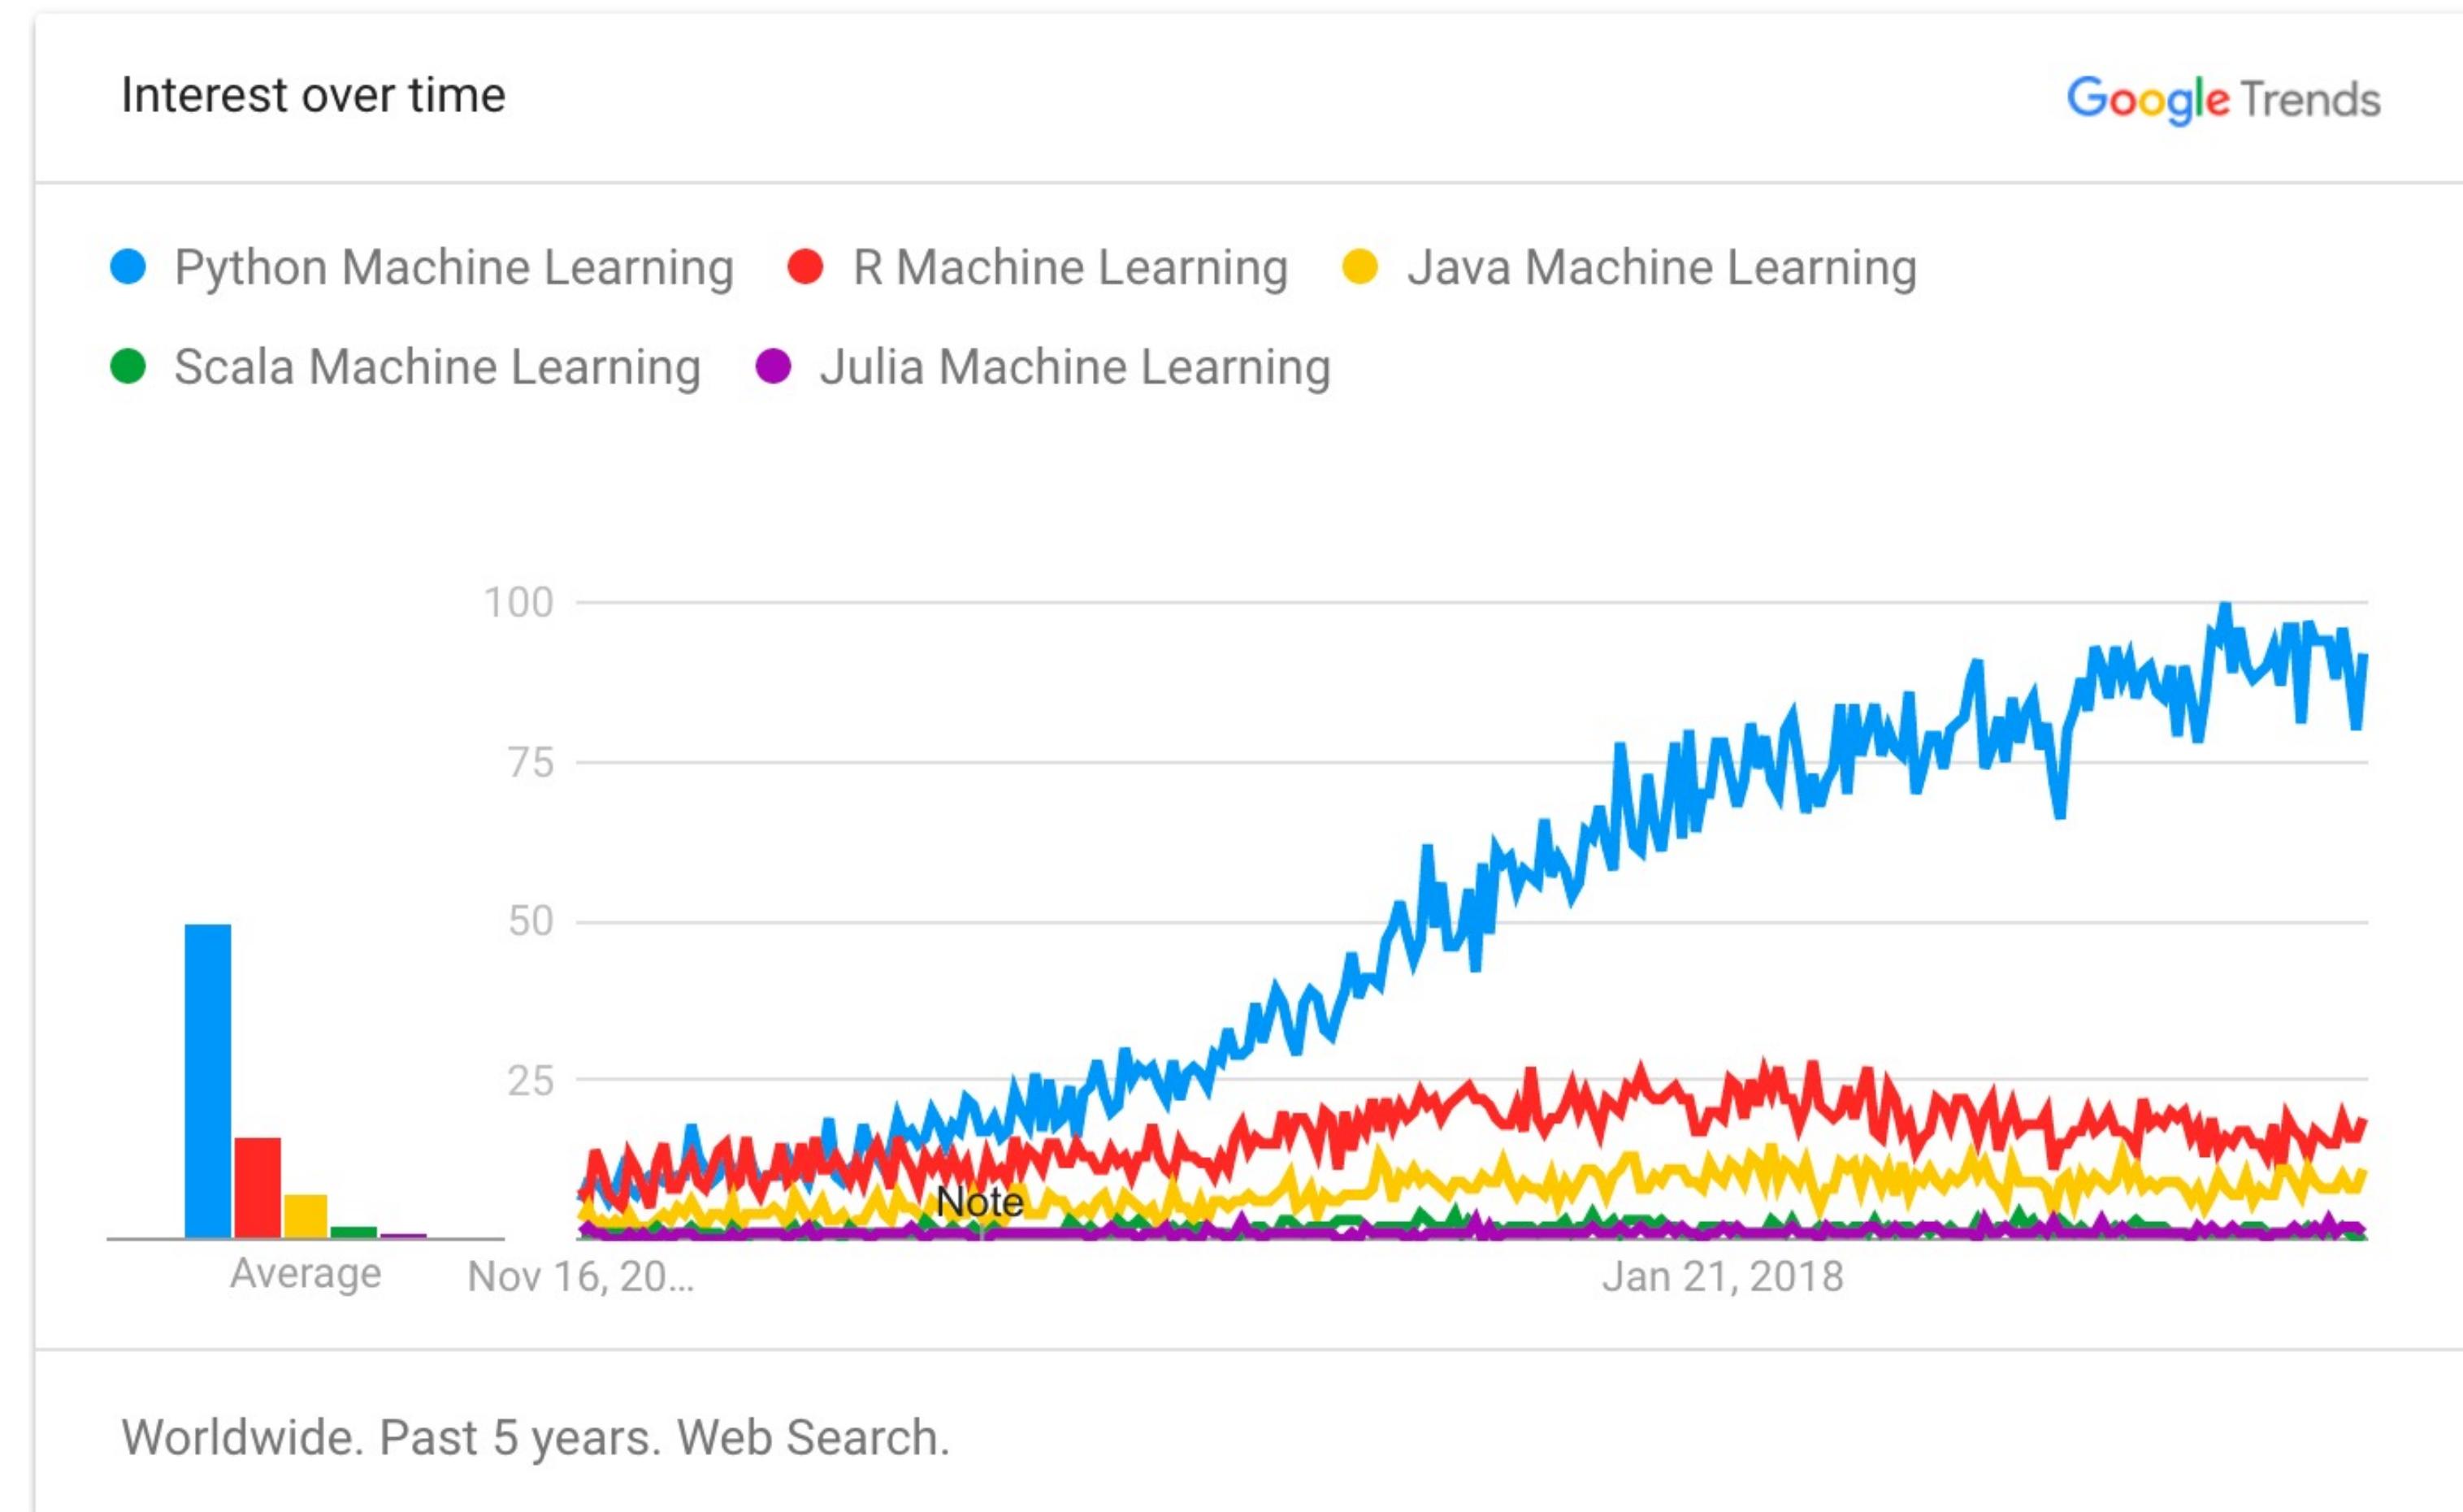
If your DeepLearning project fails

- You might have too few examples
- You didn't trained it long enough
- You trained it too long (overfitting)
- Testing data distribution is different from the training data
- It will not fail because your neural network achieved sentience!
- It will not fail because daemons and zombies*!

*https://en.wikipedia.org/wiki/Zombie_process

Python: a short overview

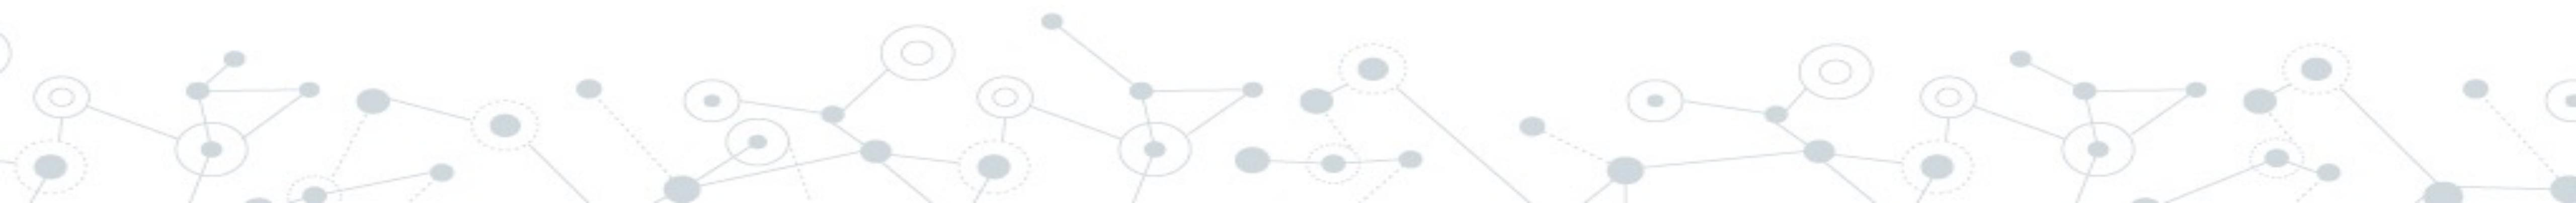
Python for machine learning?



<https://www.geeksforgeeks.org/why-is-python-the-best-suited-programming-language-for-machine-learning/>

Python vs Java

- Python and Java are both object-oriented languages
- Python is “duck typed” i.e. dynamically typed, while java is a statically typed
 - Ducked typed: ”If it looks like a duck and quacks like a duck, it's a duck”
 - you don't need a type in order to invoke an existing method on an object
- In Python, it is easy to write and read, but can be difficult to analyse



Python syntax: indentation

```
if 5 > 2:  
    print("Five is greater than two!")
```

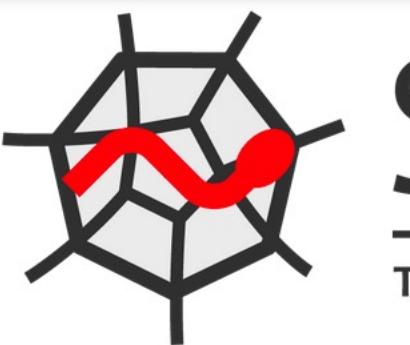
```
if 5 > 2:  
print("Five is greater than two!")
```

=> Syntax error

Python syntax: data types

Example	Data Type
x = "Hello World"	str
x = 20	int
x = 20.5	float
x = 1j	complex
x = ["apple", "banana", "cherry"]	list
x = ("apple", "banana", "cherry")	tuple
x = range(6)	range
x = {"name" : "John", "age" : 36}	dict
x = {"apple", "banana", "cherry"}	set
x = frozenset({"apple", "banana", "cherry"})	frozenset
x = True	bool

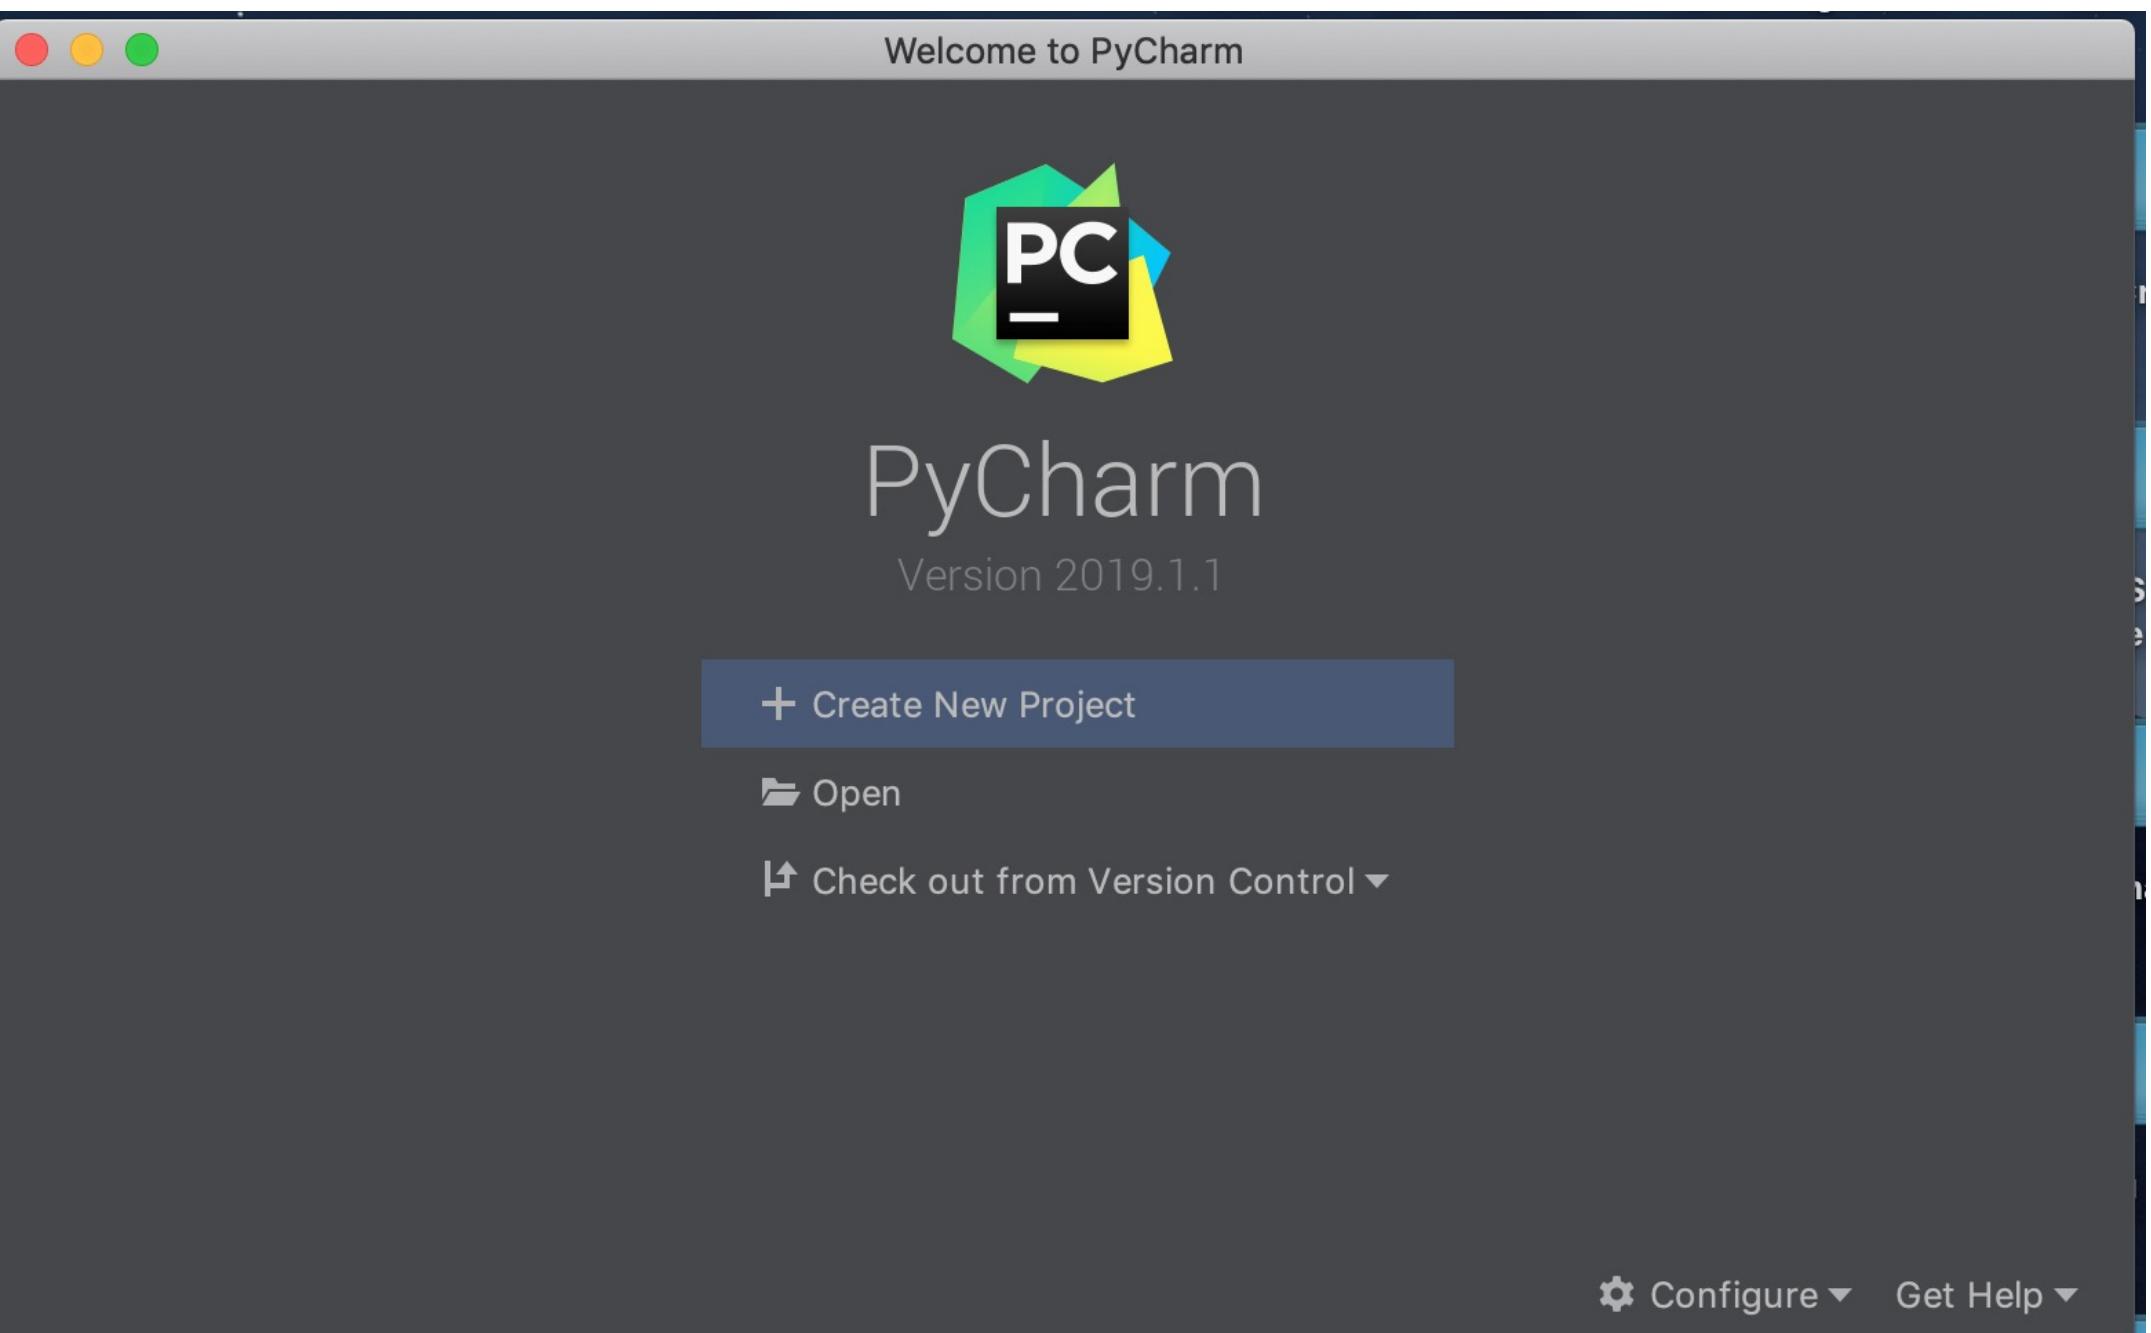
Python IDEs



SPYDER

The Scientific Python Development Environment

The screenshot shows the Spyder 3.6 Python IDE interface. The top menu bar includes File, Edit, Search, Source, Run, Debug, Consoles, Projects, Tools, View, and Help. The left sidebar contains the Project explorer, showing a tree view of the current workspace with files like temp.py, interpolation.py, _init_.py, umd_helper.py, umd_main.py, and README.md. The main area features an Editor tab with the content of interpolation.py, which includes importing pylab, numpy, and scipy.interpolate modules, generating data for analysis, creating a spiral in 3-space, adding noise, performing calculations using splprep and splev, and plotting results. A sidebar on the right displays the Outline and Variable explorer. The Variable explorer lists variables such as bars (a BarContainer object), df (a DataFrame), filename (a string), list_test (a list), nrows (an integer), r (a float64), radii (a float64 array), region (a tuple), rgb (a float64 array), series (a Series object), and test_none (NoneType). Below the variable list is a toolbar with Help, File explorer, Find in files, Breakpoints, Static code analysis, Profiler, and Online help. The bottom right corner shows two plots: a 3D surface plot of a shaded surface calculated from 'shade' and a polar plot showing data points at various angles and radii. The bottom status bar indicates permissions (RW), end-of-lines (LF), encoding (UTF-8), line (26), column (4), memory (48%), and CPU (16%).



Python Tutorials

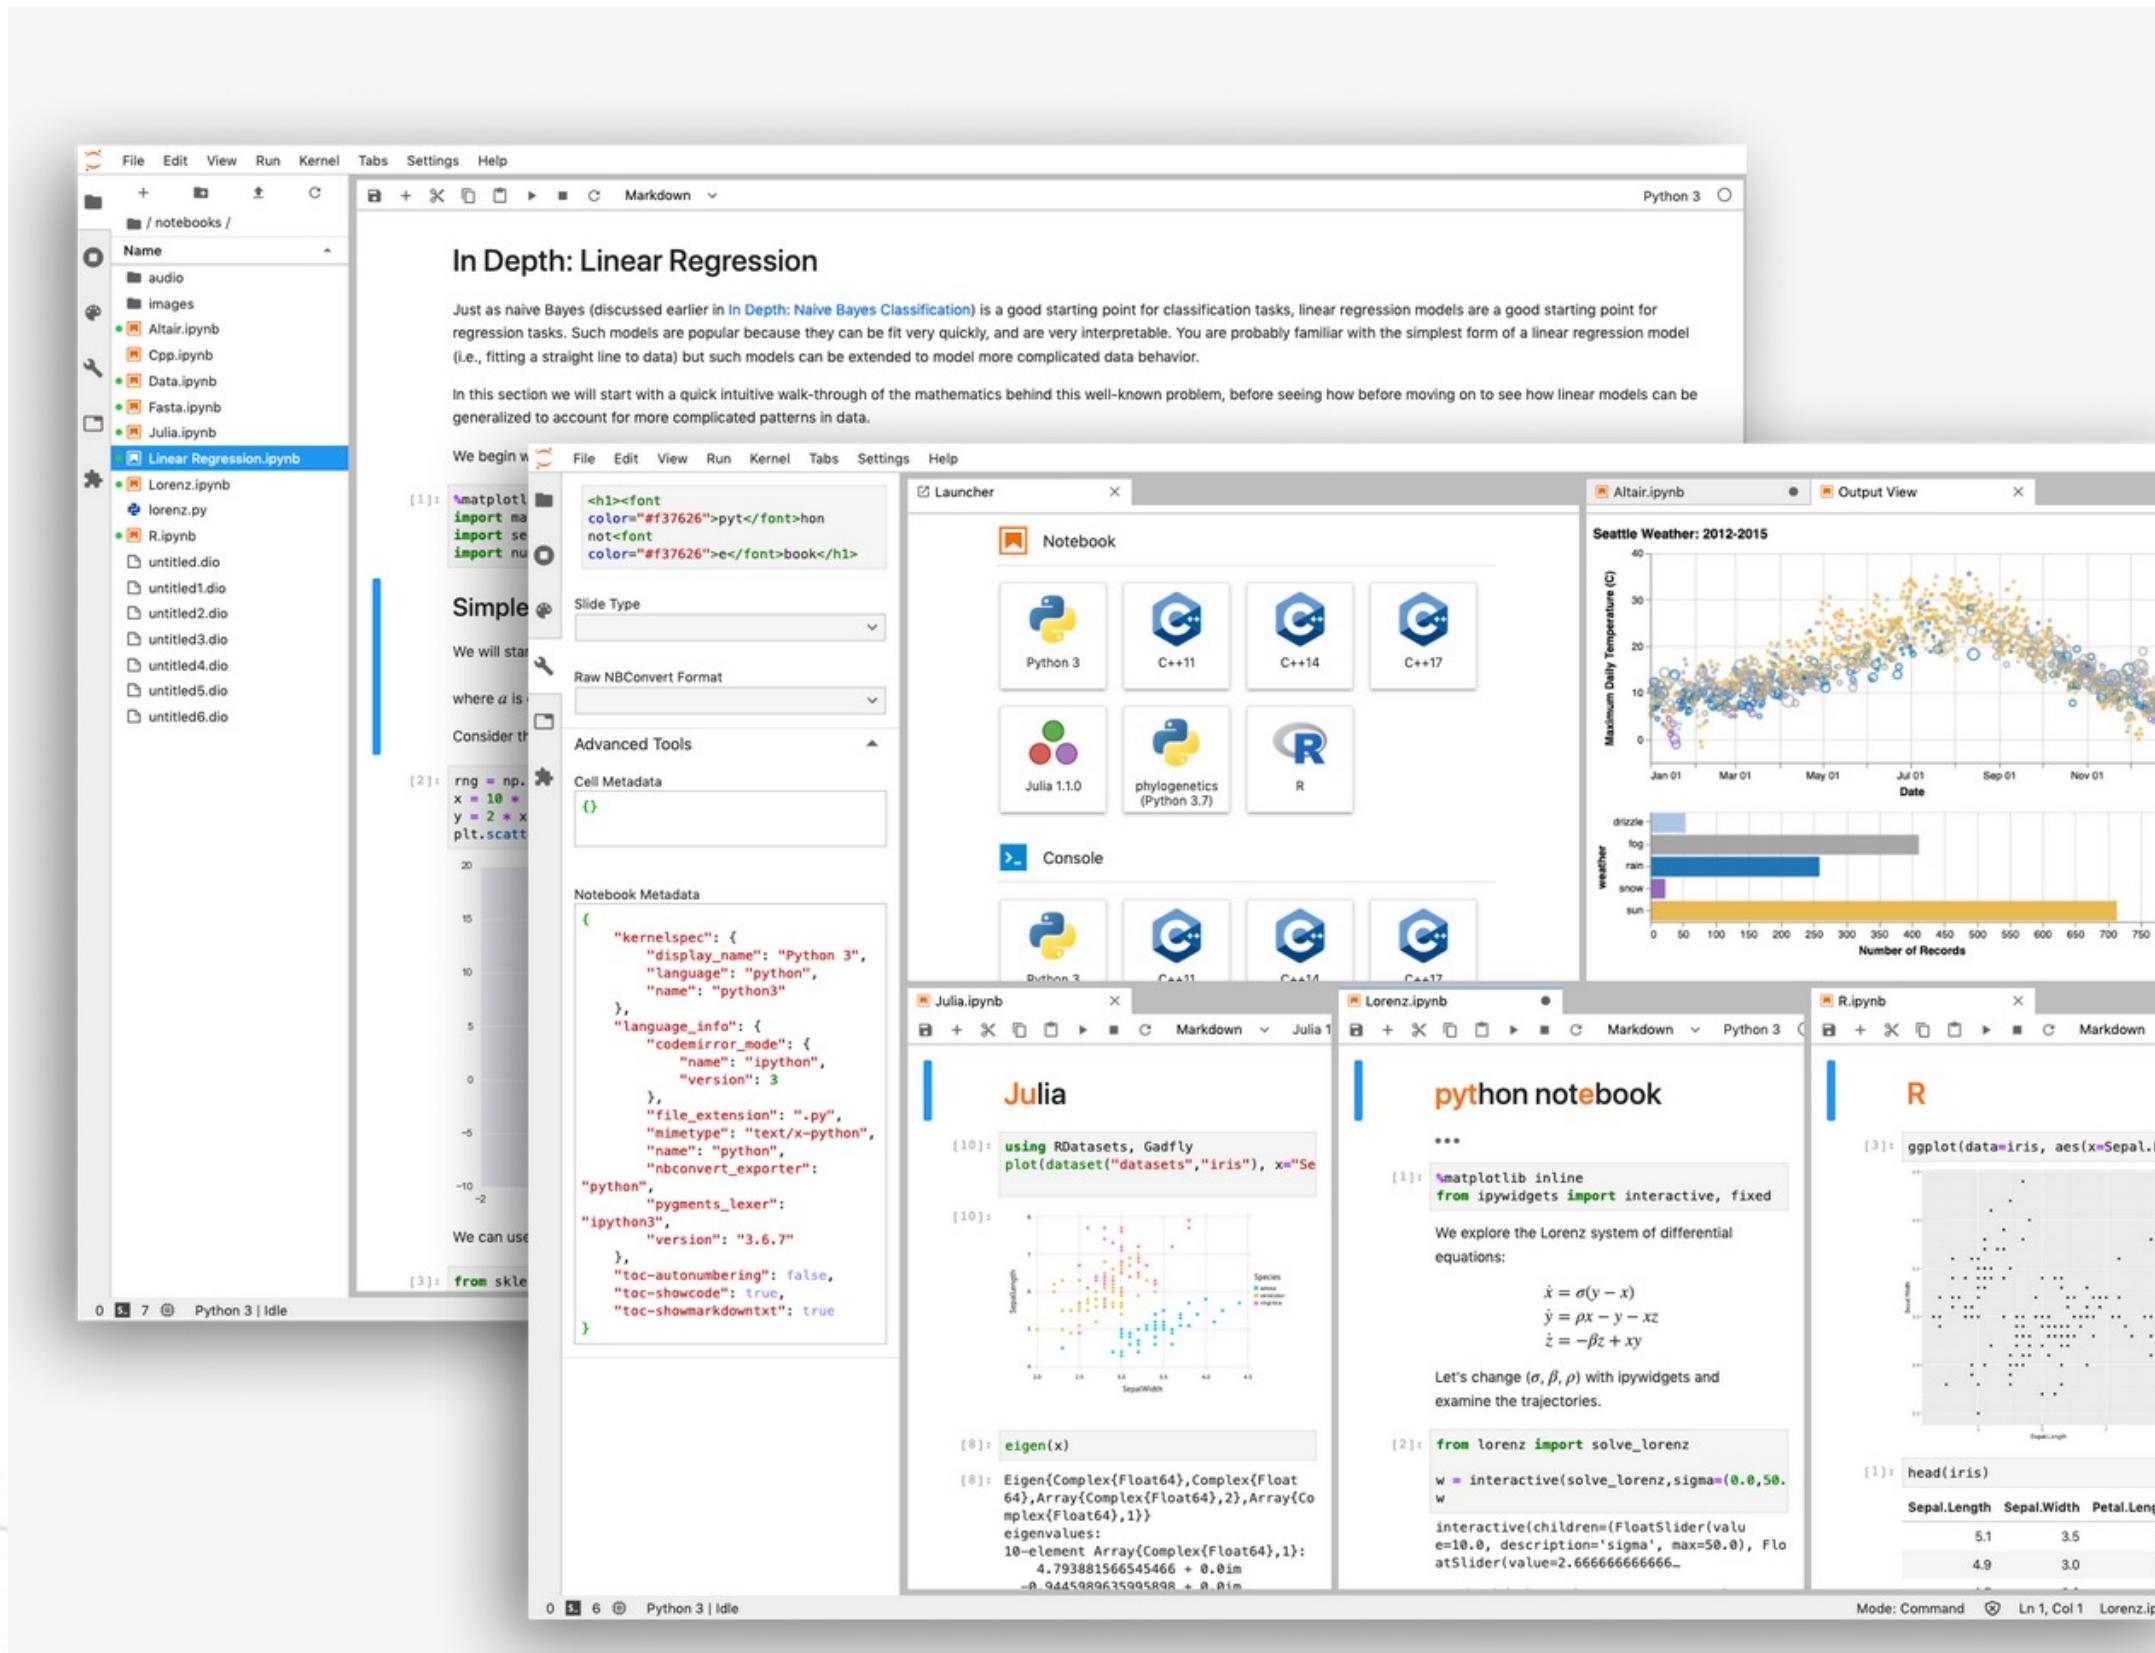
<https://www.w3schools.com/python/default.asp>

<https://docs.python.org/3/tutorial/>

<https://www.datacamp.com/courses/intro-to-python-for-data-science>

Python IDEs

“Jupyter Notebook is a web-based interactive computational environment



JupyterLab 1.0: Jupyter's Next-Generation Notebook Interface

JupyterLab is a web-based interactive development environment for Jupyter notebooks, code, and data. JupyterLab is flexible: configure and arrange the user interface to support a wide range of workflows in data science, scientific computing, and machine learning. JupyterLab is extensible and modular: write plugins that add new components and integrate with existing ones.

Try it in your browser

Install JupyterLab

<https://jupyter.org/>

TensorFlow: a short overview

TensorFlow V2.0 was released on the September 30, 2019

The screenshot shows the official TensorFlow website. At the top, there's a navigation bar with links for 'Install', 'Learn', 'API', 'Resources', 'Community', 'Why TensorFlow', a search bar, a language selector, 'GitHub', and 'Sign in'. The main headline on the left reads 'An end-to-end open source machine learning platform'. Below it, a call-to-action button says 'Get started with TensorFlow'. To the right, there's a large graphic illustrating TensorFlow's ecosystem with various icons like a laptop, smartphone, car, airplane, butterfly, and a flask, all interconnected by a network of lines.

An end-to-end open source machine learning platform

TensorFlow For JavaScript For Mobile & IoT For Production

The core open source library to help you develop and train ML models. Get started quickly by running Colab notebooks directly in your browser.

Get started with TensorFlow

TensorFlow V2.0

Tensor

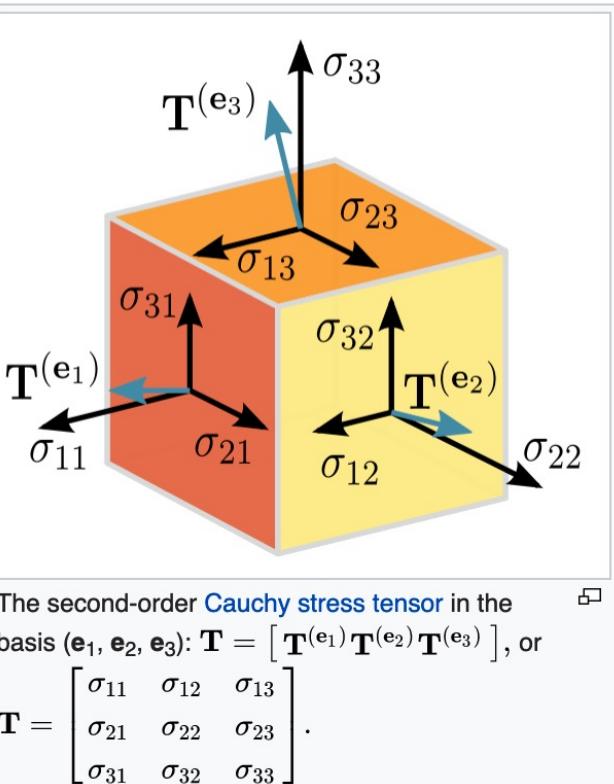
From Wikipedia, the free encyclopedia

This article is about tensors on a single [vector space](#). For vector fields, see [vector field](#). For tensor fields, see [Tensor field](#). For other uses, see [Tensor \(disambiguation\)](#).

In [mathematics](#), a **tensor** is an algebraic object that describes a [linear mapping](#) from one set of algebraic objects to another. Objects that tensors may map between include, but are not limited to, [vectors](#) and [scalars](#), and, recursively, even other tensors (for example, a matrix is a map between vectors, and is thus a tensor. Therefore a linear map between matrices is also a tensor). Tensors are inherently related to [vector spaces](#) and their [dual spaces](#), and can take several different forms – for example: a [scalar](#), a [vector](#), a [dual vector](#) at a point, or a [multi-linear](#) map between vector spaces. [Euclidean vectors](#) and scalars (which are often used in elementary physics and engineering applications where [general relativity](#) is irrelevant) are the simplest tensors.^[1] While tensors are defined [independent](#) of any [basis](#), the literature on physics often refers to them by their components in a basis related to a particular coordinate system.

An elementary example of a mapping describable as a tensor is the [dot product](#), which maps two vectors to a scalar. A more complex example is the [Cauchy stress tensor](#) \mathbf{T} , which takes a directional unit vector \mathbf{v} as input and maps it to the stress vector $\mathbf{T}^{(\mathbf{v})}$, which is the force (per unit area) exerted by material on the negative side of the plane orthogonal to \mathbf{v} against the material on the positive side of the plane, thus expressing a relationship between these two vectors, shown in the figure (right). The [cross product](#), where two vectors are mapped to a third one, is strictly speaking not a tensor because it changes its sign under those transformations that change the orientation of the coordinate system. The [totally anti-symmetric symbol](#) ϵ_{ijk} nevertheless allows a convenient handling of the cross product in equally oriented three dimensional coordinate systems.

Assuming a [basis](#) of a real vector space, e.g., a coordinate frame in the ambient space, a tensor can be represented as an organized [multidimensional array](#) of numerical values with respect to this specific basis. Changing the basis transforms the values in the array in a characteristic way that allows to *define* tensors as objects adhering to this transformational behavior. For example, there are invariants of



<https://en.wikipedia.org/wiki/Tensor>

Dataflow programming

From Wikipedia, the free encyclopedia

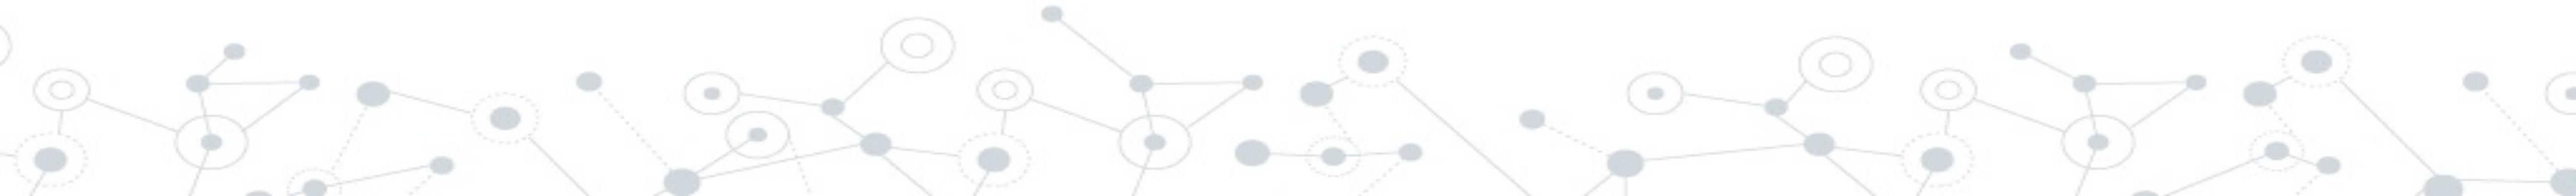
In [computer programming](#), **dataflow programming** is a [programming paradigm](#) that models a program as a [directed graph](#) of the data flowing between operations, thus implementing [dataflow](#) principles and architecture. Dataflow programming languages share some features of [functional languages](#), and were generally developed in order to bring some functional concepts to a language more suitable for numeric processing. Some authors use the term *datastream* instead of *dataflow* to avoid confusion with dataflow computing or [dataflow architecture](#), based on an indeterministic machine paradigm. Dataflow programming was pioneered by [Jack Dennis](#) and his graduate students at MIT in the 1960s.

Contents [hide]
1 Properties of dataflow programming languages
1.1 State
1.2 Representation
2 History
3 Languages
4 Application programming interfaces
5 See also
6 References
7 External links

https://en.wikipedia.org/wiki/Dataflow_programming

TensorFlow V2.0

- Can run on multiple GPUs & CPUs
- Works on Linux, MacOS, Windows, iOS, Android
- It has a stable Python & C implementation, but could be integrated with JavaScript, C++, Java, C#, Haskell, R (third-party packages)



TensorFlow V2.0

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

In TensorFlow V2.0 Keras is the official API !!!

TensorFlow comes with some simple datasets

In the original dataset a pixel value is a grayscale colour from 0-255. We need to normalise this in the interval [0,1]

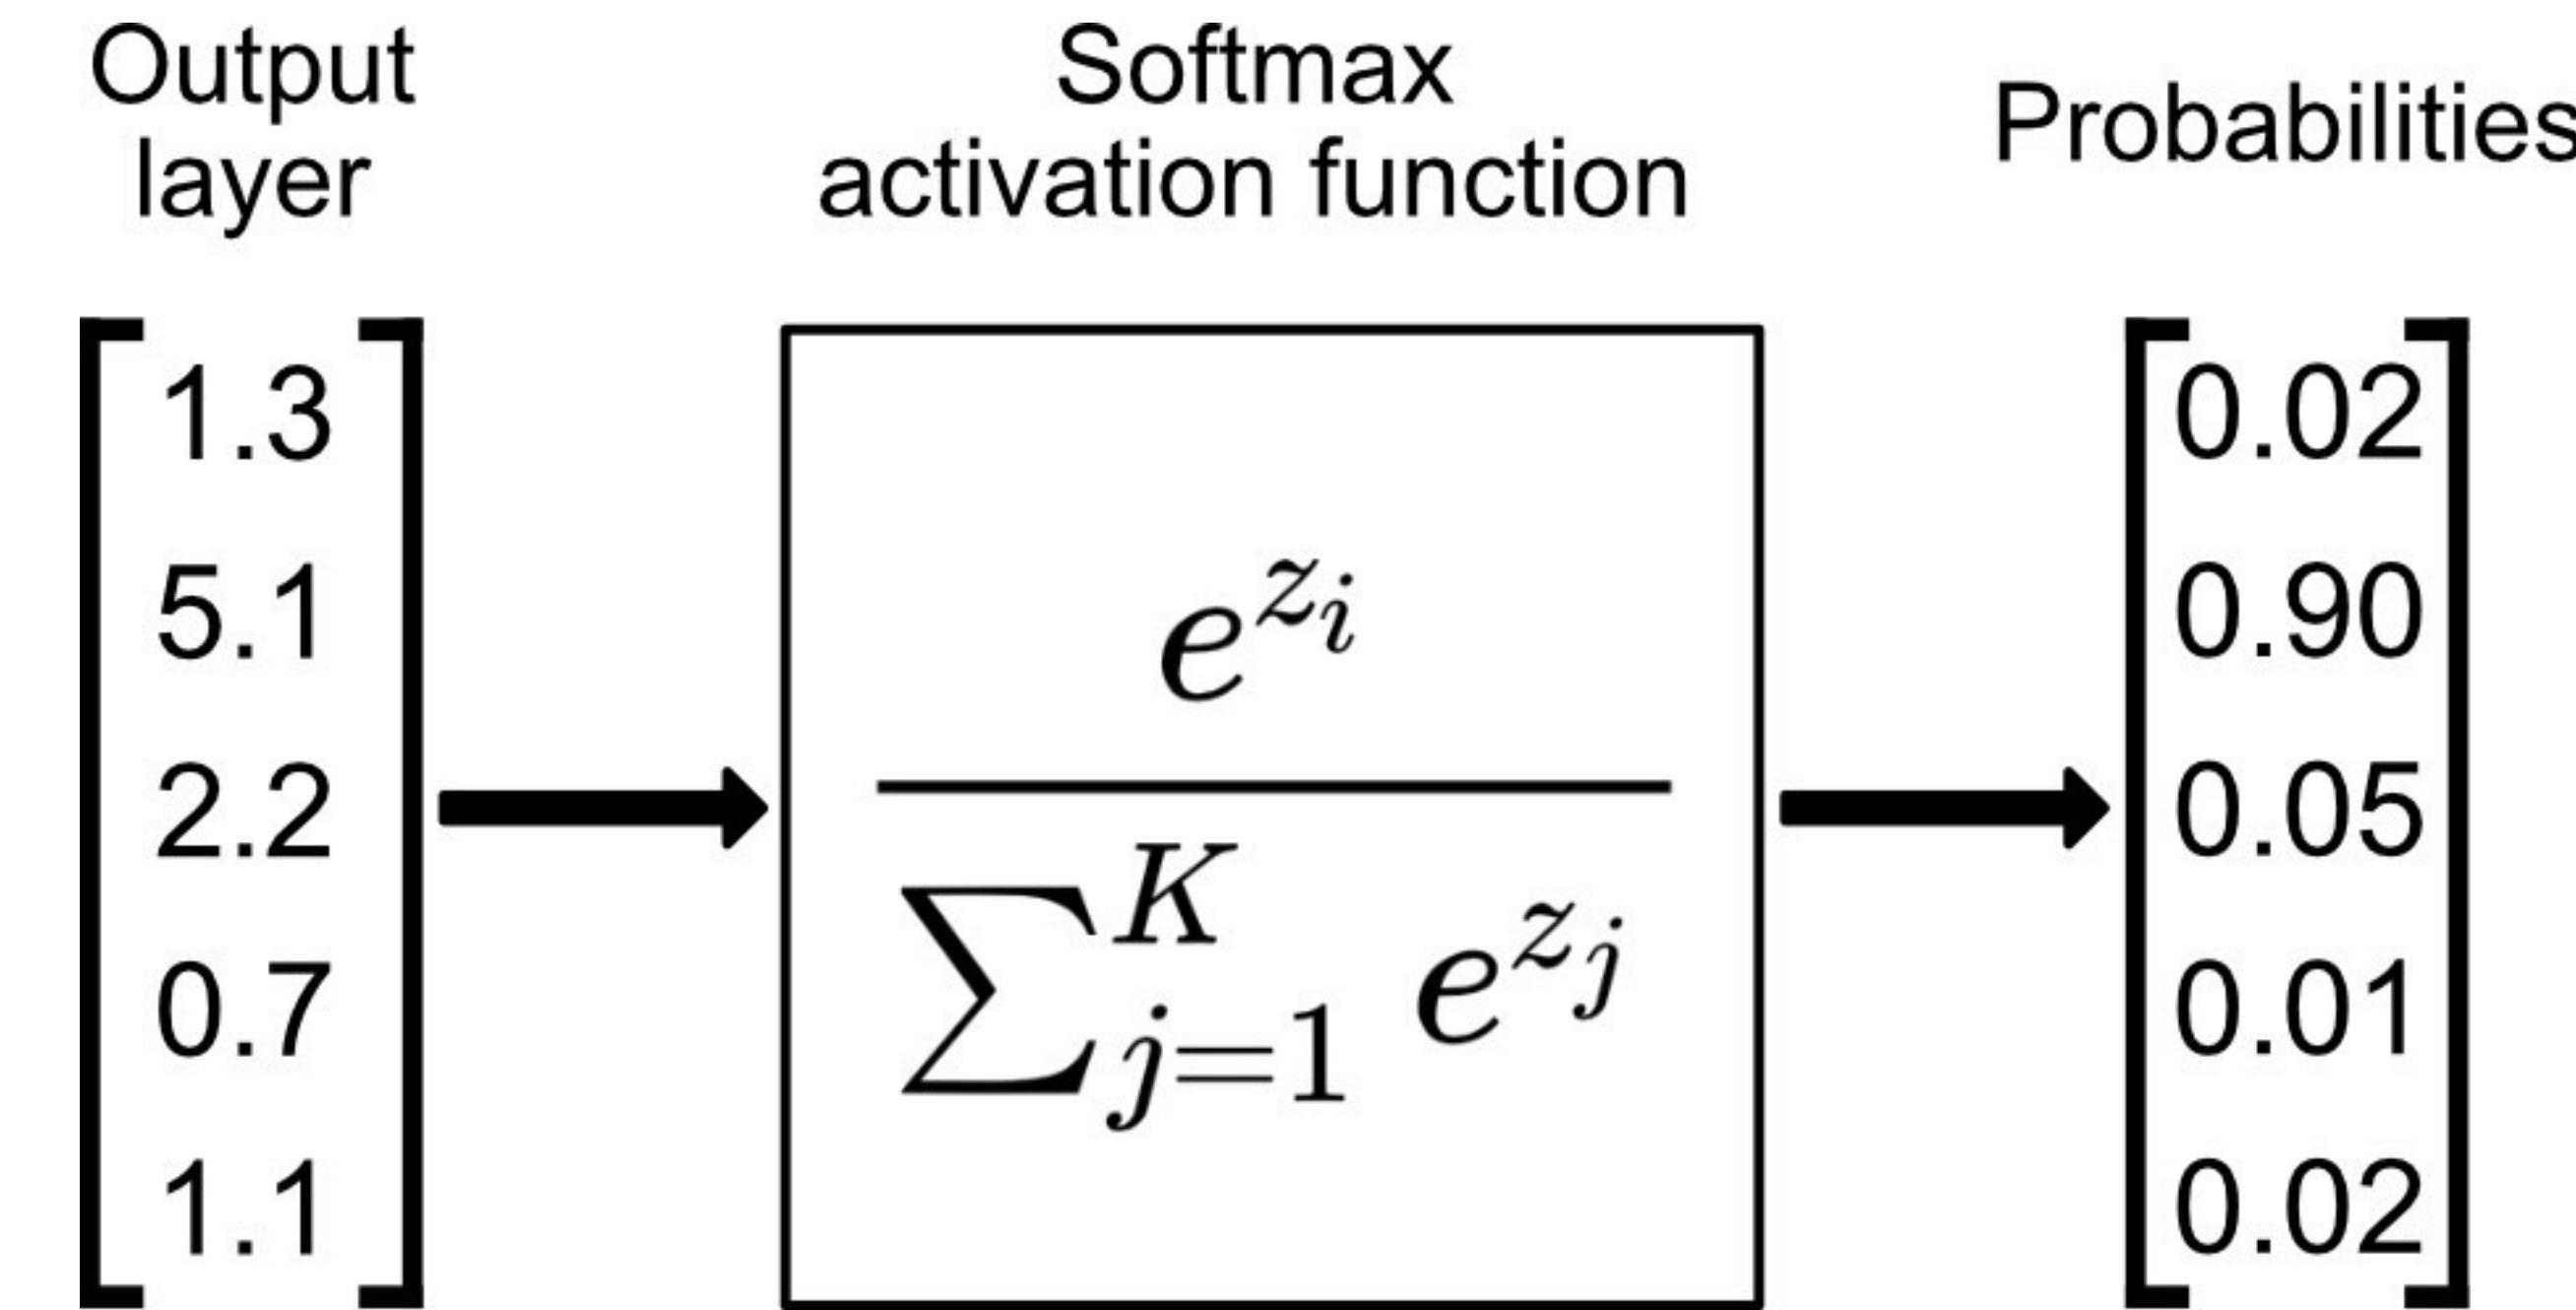
We define a simple neural network, with
One input layer
One hidden layer
Dropout is applied to the layer defined
previously, in our case the hidden layer
One output layer

Compile defines the **loss function**, the
optimiser and the . You **compile only when**
training!

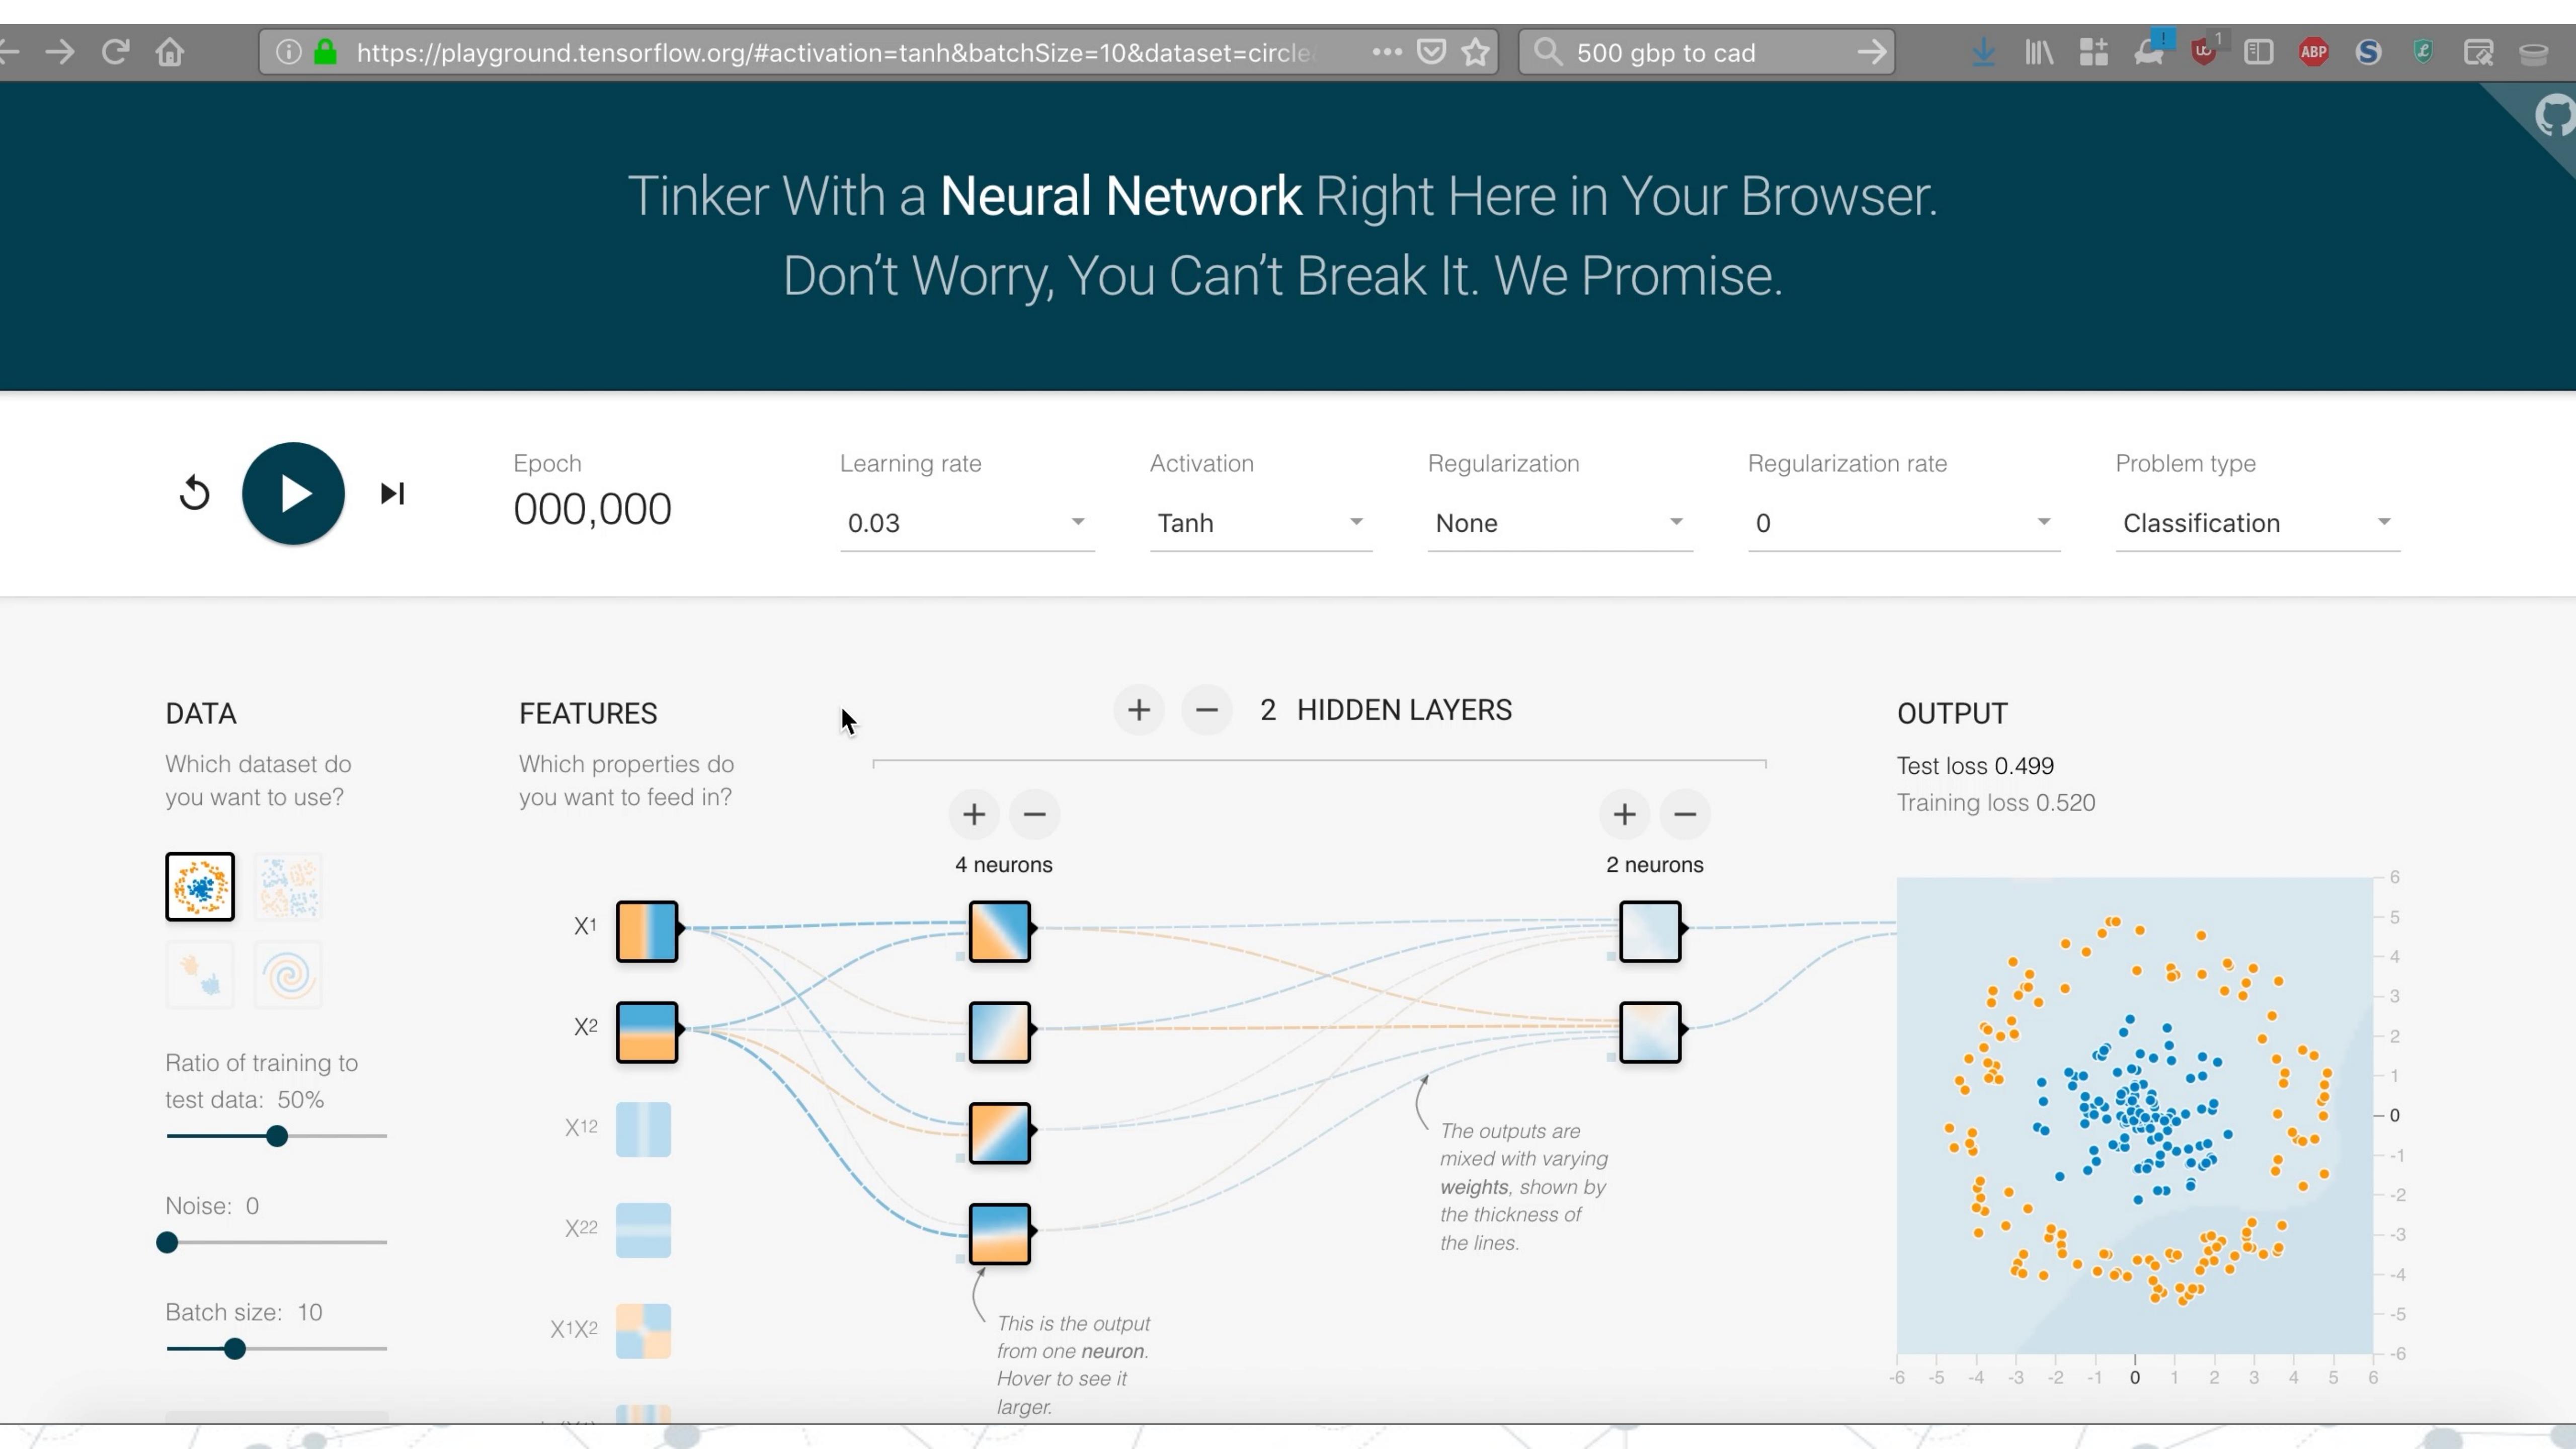
Epochs represents how many times we go
through the entire dataset in the training

Evaluate will return the accuracy on the test
set

TensorFlow V2.0



The lab





>HelloWorld-MNIST.ipynb

File Edit View Insert Runtime Tools Help

OPEN IN PLAYGROUND

Let's get started, import the TensorFlow library into your program, along with the library to plot graphs (matplotlib)

```
[ ] %matplotlib inline
import tensorflow as tf
import matplotlib.pyplot as plt
plt.rcParams['figure.figsize'] = (7,7) # Make the figures a bit bigger
import numpy as np
```

Load and prepare the MNIST dataset. Convert the samples from integers to floating-point numbers:

```
[ ] mnist = tf.keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0
```

Visualize some of the dataset images

```
[ ] for i in range(9):
    plt.subplot(3,3,i+1)
    plt.imshow(x_train[i], cmap='gray', interpolation='none')
    plt.title("Class {}".format(y_train[i]))
```





Thank you!